

Question #1 (Topic: Exam A)

Topic 1

The synthesis agent receives summarized findings from the web search and document analysis agents, then passes a consolidated summary to the report generator. During testing, you discover the generated reports make factual claims without proper citations – the report generator cannot attribute statements to their original sources because that metadata was lost during the summarization steps. What’s the most effective approach to ensure proper source attribution in the final reports?

- A. Have the report generator query the web search agent to re-locate sources for claims in the final report.
- B. Have each agent output structured data separating content summaries from source metadata (URLs, document names, page numbers).
Most Voted
- C. Skip summarization and pass full raw outputs from web search and document analysis directly to the report generator.
- D. Instruct the synthesis agent to embed source references inline within its summary text using a consistent citation format.

Correct Answer: B

Community vote currently supports Answer B

Community vote distribution

B**100%**

Question #2 (Topic: Exam A)

Topic 1

After the web search agent finds 25 sources (120K tokens of raw content), the document analysis agent extracts key insights (15K tokens), and the synthesis agent produces a coherent narrative draft (3K tokens), the coordinator must pass context to the report generation agent for the final output with proper source citations. What context-passing strategy provides the best balance of completeness and efficiency?

- A. Pass the full accumulated context from all prior agents.
- B. Pass the synthesis draft along with a structured source index that maps key claims to their source URLs and relevant excerpts.
Most Voted
- C. Pass only the synthesis draft and have a separate post-processing pipeline match claims to sources and insert citations after the report is generated.
- D. Pass a condensed summary of all prior stages that preserves the main findings and attributes them to sources by name only.

Correct Answer: B

Community vote currently supports Answer B

Community vote distribution

B**100%**

Your multi-agent research pipeline crashed after processing 12 of 28 documents. The web search agent had identified relevant sources, the document analyzer had partially completed extraction, and the synthesizer had begun pattern identification. You need to resume processing without repeating work or losing fidelity of prior findings. What state management approach best balances information fidelity with context efficiency when restoring agent state?

- A. Have each agent persist a structured export to a known location. On resume, the coordinator loads the manifest and injects relevant state into agent prompts.
- B. Index all agent outputs in a shared vector store. When resuming, each agent queries the store using semantic search to retrieve relevant prior findings.
- C. Have each agent maintain its own persistent state file and reload it independently at the start of each session.
- D. Persist the coordinator's conversation log containing all task delegations and responses, providing this to agents when resuming.

Correct Answer: A

Community vote currently supports Answer

Community vote distribution**A****100%**

You've configured the system so that all four subagents have access to the complete set of 18 tools. During testing, agents frequently call tools outside their specialization – the synthesis agent attempts web searches, and the report generator tries to analyze documents. What is the primary cause of this poor tool selection behavior?

- A. Choosing from 18 tools instead of 4-5 relevant ones increases decision complexity beyond reliable selection thresholds.
- B. The tool definitions consume too much context window space, leaving insufficient room for task content.
- C. The agents' role descriptions in their system prompts conflict with having access to tools outside that role.
- D. The coordinator cannot track which capabilities each subagent has, leading to misrouted tasks.

Correct Answer: A

Community vote currently supports Answer

Community vote distribution**A****100%**

The coordinator provides detailed step-by-step instructions to the web search subagent, specifying exact search queries, source priorities, and date filters. Production monitoring reveals three issues: (1) the subagent reports “insufficient results” rather than trying alternative approaches when pre-specified searches fail, (2) research quality drops for emerging topics that don’t match expected patterns, and (3) the subagent rarely surfaces valuable tangential sources. What’s the most effective way to improve subagent adaptability?

- A. Implement a topic classification step where the coordinator categorizes requests as “well-defined” or “exploratory” and uses different instruction styles for each category.
- B. Add explicit fallback directives to the detailed instructions: “If specified searches yield fewer than N results, attempt alternative query formulations before reporting failure.”
- C. Remove procedural details entirely, delegating with simple goals like “research X thoroughly” and relying on the subagent’s general capabilities.
- D. Specify research goals and quality criteria (coverage breadth, source diversity, recency) rather than procedural steps, letting the subagent determine its search strategy.

Correct Answer: D

The synthesis agent completes its initial pass but flags that three key research questions remain unanswered because the web search and document analysis agents didn’t find relevant information on those specific subtopics. The coordinator currently proceeds directly to report generation, producing reports with incomplete coverage. What change would most effectively improve research completeness?

- A. Have the report generation agent note which research questions couldn’t be answered, so users understand the limitations of the final output.
- B. Increase the initial breadth of queries sent to web search and document analysis to reduce the probability of missing relevant information.
- C. Have the coordinator evaluate synthesis output for gaps, then re-delegate to web search and document analysis with targeted queries before invoking synthesis again. Most Voted
- D. Give the synthesis agent direct access to web search tools so it can autonomously fill knowledge gaps without returning control to the coordinator.

Correct Answer: C

Community vote currently supports Answer C

Community vote distribution

C

100%

- The web search agent has gathered several relevant sources for a research topic. The document analysis agent now needs to examine these sources. How does information typically flow between these two specialized subagents?
- A. The agents communicate through an event-driven message queue, with the document analysis agent subscribing to web search completion events.
 - B. The web search agent directly invokes the document analysis agent, passing the discovered sources as parameters.
 - C. The coordinator agent receives the web search agent’s output and includes relevant findings in the prompt when invoking the document analysis agent. Most Voted
 - D. Both agents access a shared memory store where the web search agent writes findings and the document analysis agent reads them.

Correct Answer: C

Community vote currently supports Answer C

Community vote distribution

C

100%

- Production reviews reveal inconsistent handling of uncertainty in final reports. Sometimes conflicting subagent findings are synthesized into a single confident statement (losing nuance), while other times reports over-hedge with excessive qualifications (becoming unhelpful). When the web search agent returns “industry analysts estimate \$50B market size (methodology varies)” and the document analysis agent returns “peer-reviewed study estimates \$35B (±\$7B, 95% CI),” the coordinator either picks one arbitrarily or produces vague statements like “the market may be \$35B-\$50B depending on factors.” What systematic approach best addresses this?
- A. Configure subagents to only report findings meeting a high-confidence threshold, filtering uncertain information before it reaches the coordinator.
 - B. Add a verification subagent that cross-references findings across sources, only passing claims to synthesis that are corroborated by at least two independent sources.
 - C. Instruct the synthesis agent to structure reports with explicit sections distinguishing well-established findings from contested ones, preserving original source characterizations and methodological context. Most Voted
 - D. Implement a confidence calibration layer that normalizes subagent uncertainty expressions to standardized probability scores (0.0-1.0), then weight-average findings by their calibrated confidence.

Correct Answer: C

Community vote currently supports Answer C

Community vote distribution

B

14%

C

71%

D

14%

In production, final reports frequently contain claims without proper source attribution. Investigation shows that while the web search and document analysis agents correctly attach citations to their outputs, the synthesis agent loses track of which sources support which conclusions when combining findings. What’s the most effective architectural change?

- A. Require all subagents to output structured claim-source mappings that the synthesis agent must preserve and merge when combining findings from multiple sources.
- B. Maintain complete transcripts of all subagent interactions and add a citation-resolution agent to analyze logs and determine attributions before report generation.
- C. Add a verification step where the report generator uses semantic similarity matching against original sources to reconstruct which claims came from which documents.
- D. Have the coordinator inject source identifier prefixes into text before each handoff, then parse these prefixes at report generation to reconstruct citations.

Correct Answer: A

Community vote currently supports Answer

Community vote distribution

A

43%

D

57%

After the web search agent and document analysis agent complete their tasks, the coordinator invokes the synthesis agent. However, the synthesis agent responds that it cannot complete the task because no research findings were provided. What is the most likely cause of this issue?

- A. The subagents need to share a single API connection to enable automatic context sharing between invocations.
- B. The synthesis agent needs tools that can fetch results directly from the other agents’ conversation histories.
- C. The coordinator did not include the outputs from the previous agents in the synthesis agent’s prompt.

Most Voted
- D. The synthesis agent’s context window is not large enough to hold the combined outputs from both previous agents.

Correct Answer: C

Community vote currently supports Answer C

Community vote distribution

C

100%