

CLAUDE CERTIFIED ARCHITECT - PROFESSIONAL · EXAM CODE CCAR-P

Predict your score. Then prove it.

Many people rate themselves higher than they score. Write your practice-score guess below before you read on. Then answer seven scenarios, sit a full timed mock, and compare the result with your guess.

[Sit a full timed mock →](#)[Or the 10-question diagnostic, no account needed → /practice/diagnostic?cert=ccar-p](/practice/diagnostic?cert=ccar-p)

Full mocks need a plan. Page 16 lists the four. The diagnostic needs no plan or account.

63

SCORED ITEMS

720

TO PASS, OF 1000

7

DOMAINS

120 min

EXAM TIME

12 mo

VALIDITY

[How every fact here is sourced → /evidence](/evidence)

BEFORE YOU READ ON · YOUR FIRST PREDICTION

IF I SAT IT TODAY / 1000

MY TARGET / 1000

MY EXAM DATE

INSIDE THIS LEDGER

- 1 Your first prediction
- 2 Every core topic on one sheet
- 3 Find the constraint first
- 4 Tag every miss with one letter
- 5-11 Seven scenarios, one per area
- 12-13 The answers, and your first gap
- 14 Ten-minute reviews
- 15 Review two mocks before you book
- 16 Pick how many mocks you need

[NEXT · PAGE 2 Every core topic on one sheet →](#)

The CCA-P exam on one sheet

The core topics, by area and weight. Tick the first box after you study a topic. Tick the second box only after you get its questions right on a timed test. The empty second boxes are what to study next.

Open this sheet online, every concept linked → </exam-guide/ccar-p>

D1 Solution Design & Architecture PAGE 5 17% · ≈11 ITEMS STUDIEDTESTED Translating Business Requirements into Claude Solutions → ☐ ☐ End-to-End Architecture Design (Input to Feedback Loop) → ☐ ☐ Workflow vs Agent: Choosing the Right Architecture → ☐ ☐ Aligning Solutions to Business Value Pillars → ☐ ☐	D5 Governance, Safety & Risk PAGE 9 14% · ≈9 ITEMS STUDIEDTESTED Guardrails & Safe Deployment Practices → ☐ ☐ LLM System Risks & Failure Modes → ☐ ☐ Regulatory Compliance for AI Systems (GDPR, HIPAA, FedRAMP) → ☐ ☐ Ethical AI Considerations (Bias, Fairness, Transparency) → ☐ ☐
D2 Claude Models, Prompting & Context Engineering PAGE 6 13% · ≈8 ITEMS Related: Claude Model Selection & Tradeoffs →	D6 Stakeholder Communication & Lifecycle PAGE 10 14% · ≈9 ITEMS STUDIEDTESTED Structured Stakeholder Discovery → ☐ ☐ Communicating AI Decisions & Value to Stakeholders → ☐ ☐ Stakeholder Feedback Loops & SLA Alignment → ☐ ☐ Architecture Documentation & Implementation Handoff → ☐ ☐ AI Solution Lifecycle Management → ☐ ☐
D3 Integration PAGE 7 19% · ≈12 ITEMS STUDIEDTESTED Authentication & Access Control for Agent Systems → ☐ ☐ Accuracy-Latency-Cost Tradeoff Analysis → ☐ ☐ Observability & Monitoring for Agents at Scale → ☐ ☐ Progressive Discovery vs Monolithic Context Strategy → ☐ ☐	D7 Developer Productivity & Operational Enablement PAGE 11 7% · ≈4 ITEMS Related: Skills →
D4 Evaluation, Testing & Optimization PAGE 8 16% · ≈10 ITEMS Related: Tool Evaluation & Testing →	

Find the constraint first

Most questions follow the same shape. Knowing the shape is worth a few marks on its own.

1 • The situation	A short description of a real setup and what is going wrong.
2 • The constraint	The phrase that rules options out: "without adding infrastructure", "lowest latency", "must be auditable". Find it first.
3 • The question	"Which is most effective..." asks for the best fit, not the only correct one.
4 • The options	Usually one right, one that does too much, one that patches a symptom, one that almost fits.

TRY IT · UNDERLINE THE CONSTRAINT, THEN PICK

A support agent must answer from the policy manual only, with no new services added. Answers sometimes cite rules that do not exist. What is the most effective fix?

- A • Fine-tune a model on the manual
- B • Put the manual in context and require a cited section for every answer
- C • Add a second model to check every answer
- D • Lower the temperature to 0

Answer: B. A adds work the constraint rules out. C is the bigger hammer. D patches a symptom.

More worked scenarios → /scenarios

Tag every miss with one letter

Tag every wrong answer with one letter. Pick the first reason that fits. Only K means you need new knowledge. The other four mean you knew it and lost the point anyway.

K

Didn't know it

"I had never seen this idea."

FIX

Read the concept page. Then explain the idea back in one sentence.

Concept pages → </concepts>

C

Missed the constraint

"I knew it, but I skipped a word like only, first or without."

FIX

Underline every constraint in the scenario before you read the options.

Practise reading scenarios → </scenarios>

H

Picked the bigger hammer

"I chose the strongest tool, not the one the scenario needs."

FIX

Ask: what is the smallest change that meets the need?

Common anti-patterns → </exam-guide/ccar-p/anti-patterns>

S

Patched a symptom

"My answer hides the problem but leaves the cause."

FIX

Name the root cause in one line before you choose.

Common anti-patterns → </exam-guide/ccar-p/anti-patterns>

T

Ran out of time

"I rushed or guessed near the end."

FIX

Practise to a clock. Our full mock gives 90 seconds per question.

Timed full mock → </pricing?cert=ccar-p>

Solution Design & Architecture

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION An architect gets a one-line ask: 'automate contract review.' She sketches a multi-agent diagram that same afternoon.

CONSTRAINT Pin down requirements before any architecture pattern gets chosen.

QUESTION What should the architect do first?

- ☐ A Design a multi-agent orchestrator now, since review clearly needs dynamic judgment.
- ☐ B Build a quick prototype first, then write requirements once stakeholders see it.
- ☐ C Write the functional requirements, infrastructure needs, constraints, and a success metric.
- ☐ D Get the stakeholder to sign off on the diagram, then start building.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

Claude Models, Prompting & Context Engineering

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A team runs a simple, high-volume classification job on Opus. Their reasoning: bigger is always safer.

CONSTRAINT Match the model tier to the task, not to habit.

QUESTION What should the team check first?

- ☐ **A** Check whether Haiku handles this high-volume, straightforward task at lower cost.
- ☐ **B** Move directly to Fable for its higher capability.
- ☐ **C** Hide per-request costs from the dashboard while keeping Opus.
- ☐ **D** Add more few-shot examples to Opus so the extra cost is justified.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 7 **Integration** →

Integration

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION An architect wires an agent to an internal MCP server over STDIO. She assumes the transport handles authorization on its own.

CONSTRAINT **MCP authorization is optional by spec, not guaranteed.**

QUESTION **What should she do first?**

- ☐ **A** Confirm the server's documentation at least mentions security somewhere.
- ☐ **B** Deploy a separate authorization gateway for each tool before auditing access.
- ☐ **C** Add a warning in the agent's system prompt about unsafe actions.
- ☐ **D** Check what credentials and permissions the STDIO process actually inherits from its environment.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 8 **Evaluation, Testing & Optimization** →

Evaluation, Testing & Optimization

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A team finishes a new prompt version. Three teammates read the outputs and like them better.

CONSTRAINT A decision needs a measured comparison against ground truth.

QUESTION What should the team do before shipping it?

- ☐ A Add five more reviewers and take a majority vote on the outputs.
- ☐ B Compare both versions on held-out cases with independently validated expected answers.
- ☐ C Ship it today, since three independent readers already prefer it.
- ☐ D Let it run in production for a week and watch for complaints.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 9 **Governance, Safety & Risk** →

Governance, Safety & Risk

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A team ships a customer-facing agent with tool access. They rely only on a carefully written system prompt to block harm.

CONSTRAINT Tool access needs controls beyond system-prompt instructions.

QUESTION What should the team do next?

- ☐ A Replace the system prompt with a much longer, stricter one.
- ☐ B Log every tool call so issues can be reviewed later.
- ☐ C Keep the prompt as is; it already states clear refusal instructions.
- ☐ D Add permission policies and require human approval for high-risk tool calls.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 10 Stakeholder Communication & Lifecycle →

Stakeholder Communication & Lifecycle

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION An architect runs a one-hour kickoff for a ticket-triage system. The whole hour goes to which APIs to integrate.

CONSTRAINT **Discovery must cover success criteria, not only integrations.**

QUESTION **What should the architect do next?**

- ☐ **A** Book a second kickoff to cover even more integration options.
- ☐ **B** Move straight to build, since the integration list is already locked.
- ☐ **C** Define acceptable output quality and a measurable accuracy threshold.
- ☐ **D** Schedule a follow-up to document the current APIs in more detail.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 11 **Developer Productivity & Operational Enablement** →

Developer Productivity & Operational Enablement

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A team keeps re-typing the same five-step review checklist into Claude Code. It happens at the start of every PR review.

CONSTRAINT Share the checklist through the repo for use when teammates request PR reviews.

QUESTION What is the most effective fix?

- ☐ A Package the steps as a project skill shipped with the repo.
- ☐ B Keep a pinned message with the steps to copy and paste.
- ☐ C Write a custom MCP server that serves the steps on demand.
- ☐ D Add the steps to each engineer's own personal CLAUDE.md file.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 12 The answers, and your first gap →

The answers, and your first gap

Mark each area right or wrong on its page. If you were sure and still wrong, that is your first gap. It is the mistake you would not have caught alone. None like that? Start with an unsure one. All seven clear? Take the full mock.

RIGHT

____ of 7

SURE + WRONG

____ of 7

MY FIRST GAP IS

D ____

A score for every area: sit a full timed mock →

Needs a plan (page 16). Scored by area.

D1
Solution Design
RIGHT: C

A written spec is what any pattern choice gets checked against afterward.

A · C

The requirements must be defined before choosing an architecture.

B · C

The stated process requires requirements before prototyping.

D · C

Sign-off on a diagram isn't the four-part requirements spec the constraint needs.

Read next: Translating Business Requirements into Claude Solutions →

/concepts/requirements-to-solution-design

D2
Claude Models, Prompting
RIGHT: A

Haiku fits high-volume, straightforward work at the lowest cost and latency.

B · H

Fable suits work no other tier covers; this routine task doesn't need it.

C · S

Hiding costs leaves model suitability untested.

D · C

More examples keep the habit going instead of matching tier to task.

Read next: Claude Model Selection & Tradeoffs → /concepts/model-selection-tradeoffs

D3
Integration
RIGHT: D

STDIO pulls credentials from the local environment -- that's the real surface to audit.

A · C

A documentation mention isn't enforcement; MCP authorization is optional by spec regardless.

B · H

Separate gateways add complexity before the access requirements are known.

C · C

A warning still skips verifying the credentials and permissions the constraint requires checking.

Read next: Authentication & Access Control for Agent Systems → /concepts/auth-and-access-control-for-agents

D4
Evaluation, Testing
RIGHT: B

Held-out cases with validated answers give a trustworthy ground-truth comparison.

A · H

More reviewers is still opinion, not a measured comparison against ground truth.

C · C

Opinions aren't a measured comparison; the constraint requires testing against ground truth.

D · C

It skips the constraint's pre-release comparison and checks only after shipping.

Read next: Tool Evaluation & Testing → /concepts/evaluation

The answers, continued

D5 Governance, Safety
RIGHT: D
Guardrails are a layered system; permission policies and sandboxing matter alongside the prompt.

A · H
 A longer prompt is still just a prompt, the same insufficient layer.

B · C
 Logging records actions but does not prevent unauthorized tool use.

C · C
 Refusal wording alone isn't a control beyond the prompt the constraint requires.

Read next: Guardrails & Safe Deployment Practices → /concepts/guardrails-and-safety-controls

D6 Stakeholder
RIGHT: C
Discovery must establish acceptable quality and measurable success criteria.

A · H
 More integration talk repeats the same gap instead of closing it.

B · C
 A locked integration list still skips the success criteria the constraint requires.

D · C
 This still skips the constraint's success-criteria requirement, only adding more integration detail.

Read next: Structured Stakeholder Discovery → /concepts/stakeholder-discovery-and-requirements

D7 Developer Productivity
RIGHT: A
A repeated task Claude is told every time is exactly what a skill is for.

B · C
 Copy-pasting still skips sharing the checklist through the repo the constraint requires.

C · H
 A new MCP server is far more than this repeated task needs.

D · C
 Personal files are per-user, so they never reach the shared repo this constraint needs.

Read next: Skills → /concepts/skills

SEVEN QUESTIONS ARE A SAMPLE · A FULL MOCK GIVES YOU

63 questions at 90 seconds each, a tighter clock than the real exam. You get a score out of 1000, a score for each area, and every miss linked to its concept page.

Ten minutes a day, from memory

Pick one area a day. Close your notes, answer the prompt, then check the concept page. A wrong answer here costs nothing. It tells you what to read before the mock.

Check all seven areas in one full mock →

Needs a plan (page 16). Sit it after the seven reviews.

ONE PROMPT PER AREA · ANSWER FROM MEMORY, THEN CHECK

D1	Solution Design When do you pick a workflow over a full agent? Is 'agent' always the more advanced choice? Check: Workflow vs Agent: Choosing the Right Architecture →	☐ GOT IT ☐ SHAKY ☐ MISSED
D2	Claude Models, Prompting What does prompt caching cut costs on? Which Claude tier fits high-volume, straightforward work best? Check: Prompt Caching →	☐ GOT IT ☐ SHAKY ☐ MISSED
D3	Integration Does any single dashboard cover agent observability end to end? What does MCP authorization depend on by default? Check: Observability & Monitoring for Agents at Scale →	☐ GOT IT ☐ SHAKY ☐ MISSED
D4	Evaluation, Testing Is a 200 response with a bad stop_reason an HTTP error? What makes an A/B comparison trustworthy? Check: Debugging LLM System Failures →	☐ GOT IT ☐ SHAKY ☐ MISSED
D5	Governance, Safety Name two risk categories unique to agentic systems with tool access, beyond plain hallucination. Check: LLM System Risks & Failure Modes →	☐ GOT IT ☐ SHAKY ☐ MISSED
D6	Stakeholder Which three levers does Anthropic name for presenting a model choice to a non-technical sponsor? Check: Communicating AI Decisions & Value to Stakeholders →	☐ GOT IT ☐ SHAKY ☐ MISSED
D7	Developer Productivity When should Plan Mode run before execution? What is the difference between a skill and a CLAUDE.md rule? Check: Plan Mode →	☐ GOT IT ☐ SHAKY ☐ MISSED

Review two mocks before you book

This is the only score sheet. Fill one column at a time: your guess, the scenarios, then two full mocks. Read each row across. The weakest row is what to study next.

THE LEDGER · ONE COLUMN AT A TIME

AREA	WEIGHT	MY GUESS	SCENARIO	MOCK 1	MOCK 2	LETTER
D1 Solution Design	17%					
D2 Claude Models, Prompting	13%					
D3 Integration	19%					
D4 Evaluation, Testing	16%					
D5 Governance, Safety	14%					
D6 Stakeholder	14%					
D7 Developer Productivity	7%					
Overall	/ 1000					

THE BOOKING RULE · USE THE TARGET FROM PAGE 1

IF MOCK 2 IS...	THEN
● At or above your target, and no area is weak	Book by your own rule. A mock score is practice, not a prediction.
● Up from mock 1, but below your target	Add a week. Fix your new top three, then sit another mock.
● Flat or down from mock 1	Go back to page 4. Fix the letter you repeat most.

720 is the real exam's pass mark. Set your target with some room above it.

WRITE DOWN YOUR PRACTICE DATES

MOCK 1 · MY DATE

MOCK 2 · MY DATE

MY EXAM · BOOKED FOR

Pick how many mocks you need

The booking rule on page 15 needs two full mocks. A mock is a safe place to be wrong. It never counts toward the real exam.

[Start mock 1 with Pro →](#)
[Compare plans → /pricing](#)

Pro gives three full mocks per certification. That covers mock 1, mock 2, and a spare if the rule asks for an extra week.

THE FOUR PLANS · WHAT EACH ONE COVERS

Single	One full mock, on one certification you pick. See the Single mock → <code>/checkout?plan=single&cert=ccar-p</code>	FOR THIS LEDGER Mock 1 only
Base	One full mock for each of the four certifications. See Base → <code>/checkout?plan=base&cert=ccar-p</code>	FOR THIS LEDGER Mock 1 only
Pro	Three full mocks per certification, plus every study tool: warm-ups, an adaptive study plan, a readiness meter, progress emails and flashcards. See Pro → <code>/checkout?plan=pro&cert=ccar-p</code>	FOR THIS LEDGER Mock 1, mock 2 and a spare
All-Access	Unlimited full mocks on all four certifications, plus every study tool. See All-Access → <code>/checkout?plan=all-access&cert=ccar-p</code>	FOR THIS LEDGER As many as you need

EVERY PLAN

90 DAYS

Access lasts 90 days.

ONE PAYMENT

No auto-renew. Nothing is charged again.

MOVE UP LATER

While a plan is active, pay only the difference. The end date stays the same.

[Exam day checklist →](#) `/exam-guide/ccar-p/day-of`

[Registration and cost →](#) `/exam-guide/ccar-p/registration`

Good luck with the exam.

CAC Editorial