Question #1 (Topic: Exam A) *Topic 1*

HOTSPOT

You are a solution architect evaluating candidate use cases for a Claude-based program.

For each scenario, select yes if Claude is appropriate as the primary solution at the architectural level Otherwise, select no.

Answer Area

Scenarios:	Select Yes or No:
drafting first-pass investigative reports from semi-structured incident logs for analyst review	Yes / No
computing real-time fraud scores at sub-50-millisecond latency on a streaming pipeline	Yes / No
long-context contract review with structured clause extraction and deviation flagging	Yes / No
replacing a vector index for semantic retrieval over a multi-million-document corpus	Yes / No
routing inbound support tickets into 30 categories with reasoning for the chosen routing	Yes / No

Correct Answer:

Scenarios:	Select Yes or No:
drafting first-pass investigative reports from semi-structured incident logs for analyst review	Yes
computing real-time fraud scores at sub-50-millisecond latency on a streaming pipeline	No
long-context contract review with structured clause extraction and deviation flagging	Yes
replacing a vector index for semantic retrieval over a multi-million-document corpus	No
routing inbound support tickets into 30 categories with reasoning for the chosen routing	Yes

Question #2 (Topic: Exam A) *Topic 1*

A managed agent deployment for claims triage has grown from 6 tools to 34 tools over 18 months as product teams added capabilities. Triage accuracy has declined from 91 percent to 78 percent, and average tool selection latency has increased by 2.3 seconds. A junior engineer has proposed adding a tool router agent in front of the current agent to filter the tool list per request.
Which two findings should you present to justify capability decomposition before adding the router? (Choose two.)
Each correct answer presents part of the solution.

A. Several tools have not been invoked across the most recent 90 days of traffic.

B. The proposed router introduces an additional model call on every incoming request.

C. Tool descriptions overlap across multiple claim categories within the agent's tool set.

D. The 34 tools serve four distinct claim workflow domains within the triage scope.

E. The router pattern is well-documented across publicly available agent literature.

Correct Answer: CD

Community vote currently supports Answer CD

Community vote distribution

Question #3 (Topic: Exam A) *Topic 1*

HOTSPOT

You are reviewing a peer's end-to-end design for a Claude-based platform expected to scale to thousands of concurrent users.

For each statement, indicate yes if the statement reflects a sound architectural practice. Otherwise, select no.

Answer Area

Statements:	Select Yes or No:
Authentication and authorization run before retrieval so retrieval can filter by identity.	Yes / No
An asynchronous queue absorbs bursty traffic between intake and the model invocation layer.	Yes / No
Tax computation is encoded directly in the system prompt rather than in code.	Yes / No
Conversation logs include unredacted government identifiers to maximize signal for tuning.	Yes / No

Correct Answer:

Statements:	Select Yes or No:
Authentication and authorization run before retrieval so retrieval can filter by identity.	Yes
An asynchronous queue absorbs bursty traffic between intake and the model invocation layer.	Yes
Tax computation is encoded directly in the system prompt rather than in code.	No
Conversation logs include unredacted government identifiers to maximize signal for tuning.	No

Question #4 (Topic: Exam A) *Topic 1*

You are comparing patterns for a batch document-classification job that follows fixed steps: extract metadata, classify, summarize, and persist. Which pattern is the best fit and why?

A. A workflow pattern, because the steps are well-defined and predictable per-request token cost is preferred over flexibility. **Most Voted**

B. An agentic pattern, because agents are categorically more accurate than workflows and therefore always preferable regardless of step predictability.

C. An agentic pattern, because the open-ended planning capability of agents produces more consistent structured outputs than a fixed workflow graph on classification tasks.

D. A workflow pattern, because agentic patterns do not support tool invocations and therefore cannot execute the persist step that writes results to the downstream store.

Correct Answer: A

Community vote currently supports Answer A

Community vote distribution

Question #5 (Topic: Exam A) *Topic 1*

A platform team operates a self-hosted multi-agent system on Kubernetes that orchestrates seven specialized agents for invoice processing. The team spends approximately 40 percent of engineering capacity on infrastructure maintenance, message bus reliability, and agent state recovery. The CFO has asked you to evaluate moving to managed agent infrastructure to reclaim engineering capacity. The security officer requires that all customer financial data remain within an approved network boundary.
Which factor should most heavily influence your recommendation?

A. whether the current seven agents map cleanly to the patterns supported by managed agents

B. whether managed agents support the current bus topology used by the platform team

C. whether managed agents reduce per-invoice token costs across the existing processing volume

D. whether managed agent data handling satisfies the network boundary required by security **Most Voted**

Correct Answer: D

Community vote currently supports Answer D

Community vote distribution

Question #6 (Topic: Exam A) *Topic 1*

You are selecting a pattern for a compliance Q&A assistant that must answer policy questions with citations to the authoritative internal source set. Latency, cost, and audit predictability are prioritized.
Which pattern is the best fit?

A. multi-agent orchestration with a planner, researcher, and writer agent for every query

B. a static prompt with the entire policy corpus concatenated into every request

C. augmented LLM with retrieval-augmented generation over the indexed authoritative corpus and citation rendering on each answer

D. a pure agent loop with open web-browsing tools to surface the most current policy information, without constraining retrieval to the authoritative internal corpus

Correct Answer: C

Question #7 (Topic: Exam A) *Topic 1*

You are a solution architect designing a Claude-based assistant with access to 60 internal tools across multiple business domains. Loading every tool definition on every request increases token usage and time to first response.
Which design pattern best addresses this issue without sacrificing capability breadth?

A. Use a separate model call to summarize all 60 tool definitions before each user turn.

B. Apply progressive tool discovery so a curated initial subset is exposed and additional tools are loaded on demand based on the task.

C. Hard-code a fixed set of five tools per request to reduce token usage, regardless of whether those tools are relevant to the current task.

D. Increase the maximum context length and load all 60 tool definitions on every request, accepting the higher token cost and latency as necessary for full capability.

Correct Answer: B

Question #8 (Topic: Exam A) *Topic 1*

HOTSPOT

You are sequencing decomposed components in an invoice-processing pipeline.

For each of the decomposed components, select the execution layer it belongs to: “Pre-Processing”, “Model Stage”, or “Post-Processing”.

Answer Area

Decomposed Components:	Select the Correct Execution Layer:
Claude-based classification of invoice type	Pre-Processing / Model Stage / Post-Processing
persistence of validated records to the system of record	Pre-Processing / Model Stage / Post-Processing
schema validation of the extracted fields	Pre-Processing / Model Stage / Post-Processing
Claude-based extraction of structured invoice fields	Pre-Processing / Model Stage / Post-Processing
PII redaction on extracted text prior to model invocation	Pre-Processing / Model Stage / Post-Processing
optical character recognition on scanned invoice images	Pre-Processing / Model Stage / Post-Processing

Correct Answer:

Decomposed Components:	Select the Correct Execution Layer:
Claude-based classification of invoice type	Model Stage
persistence of validated records to the system of record	Post-Processing
schema validation of the extracted fields	Post-Processing
Claude-based extraction of structured invoice fields	Model Stage
PII redaction on extracted text prior to model invocation	Pre-Processing
optical character recognition on scanned invoice images	Pre-Processing

Question #9 (Topic: Exam A) *Topic 1*

You are reviewing a customer-support agent's configuration. Each candidate tool falls into one of four categories:
(1) required to complete defined tasks
(2) frequently used and reduces hand-offs
(3) occasionally useful for unrelated work
(4) speculative future utility
Which categories should typically remain in the agent configuration?

A. Only category 3, because occasionally useful tools for unrelated work provide broader coverage and should take priority over tools required for the agent's defined tasks.

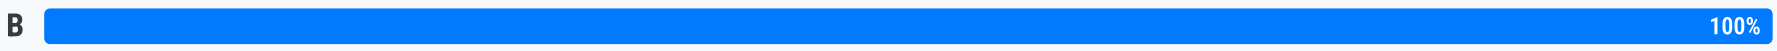
B. Categories 1 and 2 only, because they map to defined tasks and the agent's regular hand-offs.

C. All four categories, because broader tool access is categorically better for agent performance regardless of whether the tools map to defined tasks or regular hand-offs.

D. Only category 4, because speculative future-utility tools provide the most flexibility and should be configured even when no defined task currently requires them.

Correct Answer: B Community vote currently supports Answer

Community vote distribution



Question #10 (Topic: Exam A) *Topic 1*

HOTSPOT

You are classifying token-management tactics by where in the request lifecycle each tactic applies.

For each of the tactics, select where in the request lifecycle each tactic applies: “Input Preparation”, “Prompt Construction”, or “Output Handling”.

Answer Area

Tactic:	Select the Correct Request Lifecycle:
Persist the validated output for downstream consumption and audit.	Input Preparation / Prompt Construction / Output Handling
Order the prompt sections so cacheable content sits before per-request content.	Input Preparation / Prompt Construction / Output Handling
Summarize prior conversation history when full history is no longer needed.	Input Preparation / Prompt Construction / Output Handling
Validate the model's structured output against the expected schema.	Input Preparation / Prompt Construction / Output Handling
Move stable repeated content into a cacheable prefix at the head of the prompt.	Input Preparation / Prompt Construction / Output Handling
Trim retrieved passages to spans relevant to the user's question.	Input Preparation / Prompt Construction / Output Handling

Correct Answer:

Tactic:	Select the Correct Request Lifecycle:
Persist the validated output for downstream consumption and audit.	Output Handling
Order the prompt sections so cacheable content sits before per-request content.	Prompt Construction
Summarize prior conversation history when full history is no longer needed.	Input Preparation
Validate the model's structured output against the expected schema.	Output Handling
Move stable repeated content into a cacheable prefix at the head of the prompt.	Prompt Construction
Trim retrieved passages to spans relevant to the user's question.	Input Preparation

Next Questions →

Browse atleast **50%** to increase passing rate

Viewing page 1 out of 13 pages.

Viewing questions **1-10** out of 125 questions