

CLAUDE CERTIFIED ASSOCIATE - FOUNDATIONS · EXAM CODE CCAO-F

Predict your score. Then prove it.

Many people rate themselves higher than they score. Write your practice-score guess below before you read on. Then answer seven scenarios, sit a full timed mock, and compare the result with your guess.

[Sit a full timed mock →](#)[Or the 10-question diagnostic, no account needed → /practice/diagnostic?cert=ccao-f](/practice/diagnostic?cert=ccao-f)

Full mocks need a plan. Page 16 lists the four. The diagnostic needs no plan or account.

60

SCORED ITEMS

720

TO PASS, OF 1000

7

DOMAINS

120 min

EXAM TIME

12 mo

VALIDITY

[How every fact here is sourced → /evidence](/evidence)

BEFORE YOU READ ON · YOUR FIRST PREDICTION

IF I SAT IT TODAY / 1000

MY TARGET / 1000

MY EXAM DATE

INSIDE THIS LEDGER

- 1 Your first prediction
- 2 Every core topic on one sheet
- 3 Find the constraint first
- 4 Tag every miss with one letter
- 5-11 Seven scenarios, one per area
- 12-13 The answers, and your first gap
- 14 Ten-minute reviews
- 15 Review two mocks before you book
- 16 Pick how many mocks you need

[NEXT · PAGE 2 Every core topic on one sheet →](#)

The CCA-A exam on one sheet

The core topics, by area and weight. Tick the first box after you study a topic. Tick the second box only after you get its questions right on a timed test. The empty second boxes are what to study next.

Open this sheet online, every concept linked → </exam-guide/ccao-f>

D1 Prompting & Task Execution PAGE 5 14% · ≈8 ITEMS STUDIED TESTED Prompting Strategy by Task Type: Analysis, Research, Drafting, Brainstorming → ☐ ☐	D4 Workflow Integration & Solution Design PAGE 8 16% · ≈10 ITEMS STUDIED TESTED Using Claude for Research, Planning & Solution Development → ☐ ☐ Integrating Claude into Existing Workflows → ☐ ☐
D2 Output Evaluation & Validation PAGE 6 21% · ≈13 ITEMS STUDIED TESTED Evaluating Claude Output for Accuracy & Completeness → ☐ ☐ Identifying Hallucinations, Inconsistencies & Bias in AI Output → ☐ ☐ Editing & Adapting Outputs for the Intended Audience → ☐ ☐ Choosing Output Formats: Artifacts vs Inline vs Structured Data → ☐ ☐	D5 Configuration & Knowledge Management PAGE 9 12% · ≈7 ITEMS STUDIED TESTED Configuring & Maintaining Claude Projects → ☐ ☐ Managing Uploaded Knowledge & Connectors → ☐ ☐
D3 Product & Model Selection PAGE 7 12% · ≈7 ITEMS STUDIED TESTED Selecting the Right Claude Product Feature → ☐ ☐	D6 Governance, Risk & Responsible Use PAGE 10 15% · ≈9 ITEMS STUDIED TESTED Identifying Appropriate vs Inappropriate AI Use Cases → ☐ ☐ Following Organizational AI Governance Policy → ☐ ☐
	D7 Troubleshooting & Optimization PAGE 11 10% · ≈6 ITEMS STUDIED TESTED Adjusting Approach Based on Feedback & Results → ☐ ☐ Optimizing AI Workflows for Efficiency → ☐ ☐

Find the constraint first

Most questions follow the same shape. Knowing the shape is worth a few marks on its own.

1 • The situation	A short description of a real setup and what is going wrong.
2 • The constraint	The phrase that rules options out: "without adding infrastructure", "lowest latency", "must be auditable". Find it first.
3 • The question	"Which is most effective..." asks for the best fit, not the only correct one.
4 • The options	Usually one right, one that does too much, one that patches a symptom, one that almost fits.

TRY IT · UNDERLINE THE CONSTRAINT, THEN PICK

A support agent must answer from the policy manual only, with no new services added. Answers sometimes cite rules that do not exist. What is the most effective fix?

- A • Fine-tune a model on the manual
- B • Put the manual in context and require a cited section for every answer
- C • Add a second model to check every answer
- D • Lower the temperature to 0

Answer: B. A adds work the constraint rules out. C is the bigger hammer. D patches a symptom.

More worked scenarios → /scenarios

Tag every miss with one letter

Tag every wrong answer with one letter. Pick the first reason that fits. Only K means you need new knowledge. The other four mean you knew it and lost the point anyway.

K

Didn't know it

"I had never seen this idea."

FIX

Read the concept page. Then explain the idea back in one sentence.

Concept pages → /concepts

C

Missed the constraint

"I knew it, but I skipped a word like only, first or without."

FIX

Underline every constraint in the scenario before you read the options.

Practise reading scenarios →
/scenarios

H

Picked the bigger hammer

"I chose the strongest tool, not the one the scenario needs."

FIX

Ask: what is the smallest change that meets the need?

Common anti-patterns →
/exam-guide/ccao-f/anti-patterns

S

Patched a symptom

"My answer hides the problem but leaves the cause."

FIX

Name the root cause in one line before you choose.

Common anti-patterns →
/exam-guide/ccao-f/anti-patterns

T

Ran out of time

"I rushed or guessed near the end."

FIX

Practise to a clock. Our full mock gives 90 seconds per question.

Timed full mock →
/pricing?cert=ccao-f

Prompting & Task Execution

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A consultant reuses one detailed template for every task: structure first, then conclusions. This time the task is brainstorming new product names.

CONSTRAINT Match the prompt's shape to the task type, not one template.

QUESTION What should change for this request?

- ☐ A Ask Claude to generate fifty extra names to compensate.
- ☐ B Ask for varied naming directions before narrowing the shortlist.
- ☐ C Keep the same template, but add 'be creative' to the wording.
- ☐ D Use the same template again; a well-crafted one fits every task type.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 6 **Output Evaluation & Validation** →

Output Evaluation & Validation

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A report from Claude reads fluently and confidently. A reviewer skims just the opening paragraph and approves it.

CONSTRAINT **Fluent, confident text is not evidence of accuracy or completeness.**

QUESTION **What should the reviewer do next?**

- ☐ **A** Ask Claude to regenerate the report three more times and compare tone.
- ☐ **B** Reformat the report with headings so it looks more organized.
- ☐ **C** Approve it as written; confident, well-formatted text is a sign it's solid.
- ☐ **D** Verify claims against sources and check coverage against the request.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 7 **Product & Model Selection** →

Product & Model Selection

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A user on the base, non-paying plan asks how to turn on 'Research mode.' Support just tells them to phrase the question differently.

CONSTRAINT Research needs a paid plan; phrasing a chat prompt doesn't add it.

QUESTION What should this user actually be told?

- ☐ A That Research requires upgrading to a paid plan to use it.
- ☐ B That chat already runs the same decomposed search Research does.
- ☐ C To instead build a Project and upload dozens of market reports.
- ☐ D To ask the same question five different ways in chat.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 8 **Workflow Integration & Solution Design** →

Workflow Integration & Solution Design

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A department wants to skip straight to redesigning its intake process around Claude. No pilot has run yet.

CONSTRAINT Anthropic's rollout guidance favors a validated pilot before redesign.

QUESTION What should this department do first?

- ☐ A Redesign the intake process and three adjacent processes all at once.
- ☐ B Rename the current process steps to sound more 'AI-native.'
- ☐ C Run a narrow, validated pilot that augments one step first.
- ☐ D Move straight to the redesign; it delivers the most value fastest.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 9 Configuration & Knowledge Management →

Configuration & Knowledge Management

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A consultant discusses a tricky client-pricing question at length inside a Project chat. He opens a new chat and expects it to already know.

CONSTRAINT Make this context reusable without pasting it into each chat.

QUESTION What should the consultant do?

- ☐ A Rebuild the Project and repeat the discussion there.
- ☐ B Paste a summary of the discussion into the new chat each time.
- ☐ C Carry on; every chat in a Project already shares the others' history.
- ☐ D Save a reusable summary in project knowledge.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

Governance, Risk & Responsible Use

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A staffer argues Anthropic's published safety policy already covers their team's data rules. They've been pasting HR records into Claude on that basis.

CONSTRAINT Vendor-level model safety and org-level data compliance are separate layers.

QUESTION What should the team do now?

- ☐ A Scale the same practice to every other team right away.
- ☐ B Check the company's own data-handling rules before pasting any more records.
- ☐ C Delete the records from the chat right after pasting them.
- ☐ D Carry on; a published safety policy covers any downstream use of the data.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 11 **Troubleshooting & Optimization** →

Troubleshooting & Optimization

1 · THE SCENARIO · FIND THE CONSTRAINT FIRST (PAGE 3)

SITUATION A coordinator reruns an unchanged weekly-summary prompt three times. She picks whichever draft reads best that day.

CONSTRAINT Change the specific input causing the gap, not just re-run it.

QUESTION What should she do instead?

- ☐ **A** Identify the missing requirement, change its prompt instruction, then compare results.
- ☐ **B** Rewrite the entire prompt from scratch, just to be safe.
- ☐ **C** Ask a colleague to pick their favorite of the three drafts.
- ☐ **D** Keep doing this; running it multiple times and picking the best works fine.

2 · MY ANSWER

☐ A ☐ B ☐ C ☐ D

3 · HOW SURE AM I?

☐ Sure ☐ Unsure

THE ANSWER

On page 12, after all seven

4 · THE CONSTRAINT, IN MY OWN WORDS

5 · THE OPTION I RULED OUT FIRST, AND WHY

NEXT · PAGE 12 The answers, and your first gap →

The answers, and your first gap

Mark each area right or wrong on its page. If you were sure and still wrong, that is your first gap. It is the mistake you would not have caught alone. None like that? Start with an unsure one. All seven clear? Take the full mock.

RIGHT

____ of 7

SURE + WRONG

____ of 7

MY FIRST GAP IS

D ____

A score for every area: sit a full timed mock →

Needs a plan (page 16). Scored by area.

D1 Prompting & Task

RIGHT: B

Varied directions give room to diverge, which brainstorming needs more than a fixed structure.

A · H

More output volume doesn't fix a prompt shaped for the wrong task type.

C · C

A phrase addition still keeps the same template shape the constraint says must change.

D · C

The exam trap is exactly this. One template doesn't fit every task.

 Read next: [Prompting Strategy by Task Type: Analysis, Research, Drafting, Brainstorming →](#)
</concepts/task-type-prompting-strategies>

D2 Output Evaluation

RIGHT: D

Accuracy and completeness are independent checks; fluency is evidence of neither.

A · H

Comparing tone does not verify accuracy or completeness.

B · S

Better formatting doesn't verify a single claim or catch a dropped item.

C · C

Fluent, well-formatted text is exactly what looks right while still being wrong.

 Read next: [Evaluating Claude Output for Accuracy & Completeness →](#)
</concepts/evaluating-output-accuracy-completeness>

D3 Product & Model

RIGHT: A

Research is a paid-tier feature; no account workaround changes that fact.

B · C

Chat doesn't substitute for Research, which the constraint says needs a paid plan.

C · H

A Project's fixed knowledge base isn't a substitute for Research's live search.

D · C

Rephrasing a chat prompt doesn't add a feature the plan lacks.

 Read next: [Selecting the Right Claude Product Feature →](#) </concepts/claude-product-feature-selection>

D4 Workflow Integration

RIGHT: C

Augmentation proves the use case before the workflow's shape actually changes.

A · H

A wider redesign skips the validation step the guidance calls for.

B · S

A new name doesn't validate anything; the process itself is unchanged.

D · C

Skipping the validated pilot is exactly what the constraint says to avoid.

 Read next: [Integrating Claude into Existing Workflows →](#) </concepts/integrating-claude-into-existing-workflows>

 NEXT · PAGE 13 **The answers, continued →**

The answers, continued

D5 Configuration **RIGHT: D** Project knowledge persists across chats, so this avoids re-pasting it every time.

A · H

Nothing is broken, and rebuilding doesn't make the discussion reusable without re-pasting it.

B · C

Pasting it each time is exactly the repeated work the constraint asks to avoid.

C · C

Chats don't share history automatically, so the constraint's reuse problem still needs solving.

Read next: Configuring & Maintaining Claude Projects → </concepts/configuring-claude-projects>

D6 Governance, Risk **RIGHT: B** That policy is vendor-level safety; it says nothing about this company's own compliance.

A · C

This ignores the constraint that vendor safety and org compliance are separate layers.

C · S

Deleting it after the fact doesn't undo an already-made compliance decision.

D · C

Vendor safety and this company's own data rules are separate layers, per the constraint.

Read next: Following Organizational AI Governance Policy → </concepts/organizational-ai-governance-policy>

D7 Troubleshooting **RIGHT: A** Diagnosing the actual gap first is what makes the input change meaningful, not random edits.

B · H

A full rewrite risks breaking what already worked, and isn't targeted.

C · S

A second opinion still only picks among random variants, not a fix.

D · C

The constraint needs the input changed; repeating it doesn't change anything.

Read next: Adjusting Approach Based on Feedback & Results → </concepts/adjusting-approach-from-feedback>

SEVEN QUESTIONS ARE A SAMPLE · A FULL MOCK GIVES YOU

60 questions at 90 seconds each, a tighter clock than the real exam. You get a score out of 1000, a score for each area, and every miss linked to its concept page.

Ten minutes a day, from memory

Pick one area a day. Close your notes, answer the prompt, then check the concept page. A wrong answer here costs nothing. It tells you what to read before the mock.

Check all seven areas in one full mock →

Needs a plan (page 16). Sit it after the seven reviews.

ONE PROMPT PER AREA · ANSWER FROM MEMORY, THEN CHECK

D1	Prompting & Task Name the four task types this objective tests. Which one needs source-scoping rather than a rigid structure? Check: Prompting Strategy by Task Type: Analysis, Research, Drafting, Brainstorming →	☐ GOT IT ☐ SHAKY ☐ MISSED
D2	Output Evaluation Does hallucinated content always look obviously wrong? Is bias a claim you can spot-check like a single fact? Check: Identifying Hallucinations, Inconsistencies & Bias in AI Output →	☐ GOT IT ☐ SHAKY ☐ MISSED
D3	Product & Model Are Projects, Research, chat, and Artifacts mutually exclusive? Can a Project trigger Research and still produce an Artifact? Check: Selecting the Right Claude Product Feature →	☐ GOT IT ☐ SHAKY ☐ MISSED
D4	Workflow Integration What is the actual value of using Claude for research, beyond one chat answer? Is augment the same as redesign? Check: Using Claude for Research, Planning & Solution Development →	☐ GOT IT ☐ SHAKY ☐ MISSED
D5	Configuration Can the Gmail connector send an email on its own? What can the Calendar connector delete? Check: Managing Uploaded Knowledge & Connectors →	☐ GOT IT ☐ SHAKY ☐ MISSED
D6	Governance, Risk Does 'Claude can produce this output' mean the use case is appropriate? Name one category that's prohibited outright. Check: Identifying Appropriate vs Inappropriate AI Use Cases →	☐ GOT IT ☐ SHAKY ☐ MISSED
D7	Troubleshooting Does fixing latency or cost always mean switching to a different model? What should you check first? Check: Optimizing AI Workflows for Efficiency →	☐ GOT IT ☐ SHAKY ☐ MISSED

Review two mocks before you book

This is the only score sheet. Fill one column at a time: your guess, the scenarios, then two full mocks. Read each row across. The weakest row is what to study next.

THE LEDGER · ONE COLUMN AT A TIME

AREA	WEIGHT	MY GUESS	SCENARIO	MOCK 1	MOCK 2	LETTER
D1 Prompting & Task	14%					
D2 Output Evaluation	21%					
D3 Product & Model	12%					
D4 Workflow Integration	16%					
D5 Configuration	12%					
D6 Governance, Risk	15%					
D7 Troubleshooting	10%					
Overall	/ 1000					

THE BOOKING RULE · USE THE TARGET FROM PAGE 1

IF MOCK 2 IS...	THEN
● At or above your target, and no area is weak	Book by your own rule. A mock score is practice, not a prediction.
● Up from mock 1, but below your target	Add a week. Fix your new top three, then sit another mock.
● Flat or down from mock 1	Go back to page 4. Fix the letter you repeat most.

720 is the real exam's pass mark. Set your target with some room above it.

WRITE DOWN YOUR PRACTICE DATES

MOCK 1 · MY DATE

MOCK 2 · MY DATE

MY EXAM · BOOKED FOR

Pick how many mocks you need

The booking rule on page 15 needs two full mocks. A mock is a safe place to be wrong. It never counts toward the real exam.

[Start mock 1 with Pro →](#)
[Compare plans → /pricing](#)

Pro gives three full mocks per certification. That covers mock 1, mock 2, and a spare if the rule asks for an extra week.

THE FOUR PLANS · WHAT EACH ONE COVERS

Single

One full mock, on one certification you pick.

[See the Single mock →](#)

</checkout?plan=single&cert=ccao-f>

FOR THIS LEDGER

Mock 1 only

Base

One full mock for each of the four certifications.

[See Base →](#)

</checkout?plan=base&cert=ccao-f>

FOR THIS LEDGER

Mock 1 only

Pro

Three full mocks per certification, plus every study tool: warm-ups, an adaptive study plan, a readiness meter, progress emails and flashcards.

[See Pro →](#)

</checkout?plan=pro&cert=ccao-f>

FOR THIS LEDGER

Mock 1, mock 2 and a spare

All-Access

Unlimited full mocks on all four certifications, plus every study tool.

[See All-Access →](#)

</checkout?plan=all-access&cert=ccao-f>

FOR THIS LEDGER

As many as you need

EVERY PLAN

90 DAYS

Access lasts 90 days.

ONE PAYMENT

No auto-renew. Nothing is charged again.

MOVE UP LATER

While a plan is active, pay only the difference. The end date stays the same.

[Exam day checklist → /exam-guide/ccao-f/day-of](#)

[Registration and cost → /exam-guide/ccao-f/registration](#)

Good luck with the exam.

CAC Editorial