

1. Claude's answers are inconsistent across similar requests. What optimization helps most?

- A. Attach broader background files so every request shares the maximum context.
- B. Increase the level of detail requested so the answers converge on completeness.
- C. Standardize prompts/templates and move durable rules into project instructions.**
- D. Ask each question in multiple chats and keep whichever answer recurs the most.

Inconsistency often comes from prompt drift. Templates and durable project instructions stabilize behavior; detail inflation, majority voting, and bulk context treat the symptom instead of the drift.

2. Which signal most strongly suggests a Claude summary may be incomplete?

- A. Key sections of the source document are never mentioned despite being relevant to the question.**
- B. The summary reorders the topics instead of following the source document's own section structure.
- C. The summary is noticeably shorter than the original source document.
- D. The summary paraphrases the source rather than quoting it directly.

Relevant omissions are a primary incompleteness signal. Summaries are expected to be shorter, may reorder for the reader, and normally paraphrase; none of those traits alone means content is missing.

3. Claude cites "a recent industry study" without title, author, date, or link. How should you treat that claim?

- A. As reliable if it matches your general industry impressions.
- B. As unverified until a real source is identified and checked.**
- C. As usable once Claude confirms the study exists when asked.
- D. As acceptable background as long as it is not quoted verbatim.

Vague study references are classic unsupported claims. Validation requires identifiable, checkable sources; neither topical plausibility nor the model's own confirmation substitutes for locating the actual study.

4. You ask Claude for a competitive analysis your team will edit and reuse over several weeks. Which output format request fits best?

- A. A series of short chat replies, one per competitor, to keep each answer tightly focused.
- B. An artifact containing the analysis document, so it can be revised and reused iteratively.**
- C. An inline chat answer, so the analysis stays attached to the conversation flow.
- D. A JSON structure, so every finding is machine-readable for future automated parsing.

A substantial document meant for iteration and reuse is the classic artifact case: it renders in a dedicated panel, supports revision, and can be returned to later. Inline replies suit quick answers and structured data suits machine processing, not team editing.

5. When is enhanced project knowledge with RAG most relevant?

- A. When you want project chats to run faster by skipping the knowledge base entirely
- B. When paid-plan project knowledge grows large enough that retrieval is needed to scale beyond raw context limits.**
- C. When a project holds only a handful of short documents that already fit comfortably within the regular context window
- D. When free-plan users need a way to raise the per-chat context ceiling

On paid plans, Claude automatically enables RAG mode when project knowledge approaches the context window limit, expanding capacity while maintaining quality. It is for large knowledge bases, not small ones, and it is not a knowledge bypass or a free-tier feature.

6. What is the best way to validate Claude's extraction of action items from meeting notes?

- A. Spot-check each action item against the transcript for owner, task, and deadline accuracy.**
- B. Have Claude re-extract the list a second time and keep the items that appear in both runs.
- C. Confirm the list covers every agenda topic that was scheduled for the meeting.
- D. Ask attendees whether the list feels complete without consulting the transcript.

Extraction quality is judged by fidelity to source: correct tasks, owners, and dates. Agenda coverage, run-to-run agreement, or attendee recall can all look reassuring while the transcript says otherwise.

7. When is Sonnet typically the most practical default for everyday business writing and analysis in Claude?

- A. When you need a strong balance of quality, speed, and cost for common knowledge-work tasks.**
- B. Only for coding tasks, because balanced tires are not really tuned for prose.
- C. When the work is high-volume and trivially simple, so per-task cost matters most.
- D. Only when a task has already failed on both the fastest tier and the most capable tier available.

Sonnet-class models are commonly positioned as a balanced default for mainstream productive work-writing, analysis, and tool-assisted tasks-when maximum depth or maximum speed is not the sole priority.

8. A project's knowledge should track a policy folder your team maintains in Google Drive. What is the most maintainable setup?

- A. Add the folder's URL to a knowledge file so Claude can fetch the documents on demand.
- B. Paste the text of the single most-used policy into the project instructions and skip the rest.
- C. Download the folder every month and re-upload each of the files to the project by hand.
- D. Connect the Drive source so project knowledge stays aligned with the maintained folder.**

Connecting the maintained source, such as Google Drive integration, keeps project knowledge aligned with the folder the team already updates. Manual re-upload cycles drift stale, instructions are for behavior rather than document storage, and a pasted link is not ingested content.

9. You are iterating on a live interactive calculator that colleagues will reuse. Which Claude product surface is most appropriate?

- A. A long chat message that walks colleagues through performing the calculation manually
- B. Artifacts, which render standalone interactive content beside the chat for iteration and reuse.**
- C. A project knowledge file that stores the calculator logic for each new chat to reference
- D. A regular chat reply containing the calculator's code for colleagues to copy out themselves.

Artifacts are designed for substantial standalone content such as interactive tools, documents, and apps in a dedicated panel for iteration and sharing. Ordinary ephemeral chat is weaker for reusable interactive calculators.

10. A colleague wants Claude to write a performance review using private HR notes about another employee. What is the responsible approach?

- A. Proceed anyway, because performance feedback is a routine and comparatively low-sensitivity workplace document.
- B. Proceed after replacing the employee's name with initials in the pasted HR notes.
- C. Follow company policy and use others' sensitive HR data only in explicitly approved tooling with controls.**
- D. Proceed in a personal Claude account so the data stays outside company systems.

Employee HR data is sensitive. Responsible use follows policy and approved tooling with controls; superficial masking and personal-account workarounds do not make the processing compliant.

11. How should your prompting differ between a brainstorming task and a compliance-analysis task?

- A. Use identical prompts for both tasks so that the two sets of results stay directly comparable.
- B. Make both prompts as short as possible, since brevity outperforms structure in either case.
- C. Add strict formatting constraints to the brainstorm and keep the compliance prompt open-ended.
- D. Encourage breadth for brainstorming; require grounded, cited, tightly constrained reasoning for compliance.**

Prompting adapts to task type: divergent tasks benefit from breadth and deferred judgment, while compliance analysis needs grounding, citations, and hard constraints. Reversing these or flattening both wastes each mode's strengths.

12. A workflow needs both rapid brainstorming and a final careful risk assessment. What selection pattern is sound?

- A. Run both stages on the fastest model and compensate with one extra proofreading pass.
- B. Use one mid-tier model throughout so that the results stay stylistically consistent.
- C. Use the most capable model for open-ended brainstorming, then a faster model to write up the final risk assessment
- D. Use a faster model for ideation drafts, then a more capable model and stricter review for the final risk assessment**

Staging model choice by subtask-fast ideation, higher capability for high-stakes assessment-is a practical selection pattern. Putting the weakest model on the riskiest step inverts the logic.

13. Client data is covered by an NDA. Your Claude Team project is shared too broadly. What should you do?

- A. Leave the membership as it is but ask current members not to open the sensitive files.
- B. Export the sensitive files to an email thread and simply continue the analysis outside the project entirely.
- C. Restrict project membership to need-to-know users and remove sensitive files if access was excessive.**
- D. Wait for the client to raise concerns before adjusting who has project access.

Need-to-know access is required under NDAs. Oversharing should be corrected immediately by tightening membership and removing sensitive content as needed; honor-system requests, email export, and waiting all fail that duty.

14. On Team/Enterprise plans, what does the "Can view" project permission generally allow?

- A. See project contents/knowledge/instructions and chat in the project, without editing project configuration.**
- B. See only the project's name and description while its knowledge stays hidden until upgraded.
- C. Chat in the project and make small edits to the instructions but never delete files.
- D. View and chat in the project, plus add new knowledge files as long as none of the existing ones are removed.

"Can view" allows using the project (seeing contents, knowledge, and instructions, and chatting) without any edit rights. Anthropic's help with documentation reserves editing instructions, knowledge, and membership for "Can edit."

15. Claude produces a market-size figure that is not present on your uploaded deck. What should you do first?

- A. Use the figure and attribute it to the uploaded deck so that reviewers can trace the number later.
- B. Treat the figure as unverified until you find a primary source or confirm it is not supported.**
- C. Ask Claude to double-check its own figure and proceed once it confirms the number.
- D. Keep the figure but round it so that any error would be less material to the story.

Numbers not grounded in provided sources are high-risk hallucinations or external inventions. Validation means verifying against primary sources before use. Shipping unverified metrics is a classic evaluation failure.

16. An analyst needs figures from industry reports published this week, beyond Claude's training data. Which capability addresses this directly?

- A. A more capable model tier, which is trained on more recent material as a matter of course.
- B. Claude's web search, which retrieves current information beyond the model's training cutoff.**
- C. Longer prompts that describe the missing reports in enough detail to reconstruct their figures.
- D. A larger context window, which extends how recent the model's built-in knowledge can be.

Web search lets Claude retrieve current sources published after its training data. Context window size governs how much you can provide, not knowledge recency; tier does not guarantee recency; and reconstruction from description invents data.

17. You want Claude to critique your draft before rewriting it. Which instruction pattern is strongest?

- A. "Tell me what you like most about the draft, then lightly polish the weaker wording."
- B. "Rewrite the draft right away and note down afterwards anything that you happened to change along the way."
- C. "Rewrite it three different ways so I can simply pick whichever version I prefer."
- D. "First list strengths and issues against these criteria; do not rewrite until I approve the critique."**

Separating evaluation from rewriting surfaces issues you can accept or reject before content changes. Immediate rewrites hide the diagnosis. Praise-only or random edits do not improve quality control.

18. What is a core principle when using Claude for regulated or compliance-sensitive content?

- A. A senior manager's sign-off substitutes for qualified regulatory review of the content
- B. Compliance review is needed only for documents that actually leave the company.
- C. Human experts remain accountable; Claude drafts require qualified review against official requirements.**
- D. Claude may finalize regulated text whenever its citations all point to official government source documents.

Accountability stays with humans and organizations. Citation-shaped references do not certify compliance, internal regulated content also carries obligations, and seniority is not regulatory qualification.

19. Claude summarizes customer feedback and claims "most users hate feature X," but the sample has 8 comments with mixed views. Best evaluation response?

- A. Reject overgeneralization and require counts or proportions tied to the actual sample.**
- B. Accept the phrasing, since "most" is an appropriately hedged and non-absolute word.
- C. Ask Claude to add more supporting detail that justify the existing claim.
- D. Rerun the analysis repeatedly until a milder version of the summary statement happens to appear.

Overgeneralization from tiny mixed samples is a common evaluation failure. Demand counts or proportions tied to the actual sample; hedge words, added justification, and rerolls do not fix an unsupported majority claim.

20. Which is the strongest method to validate a Claude-generated process checklist against an SOP PDF?

- A. Check that the steps appear in the same order as the SOP's table of contents.
- B. Confirm the checklist contains the same number of steps as the SOP has sections.
- C. Walk through the checklist together with a colleague who knows the procedure well from memory.
- D. Map each checklist step to a specific SOP section and note any extra or missing steps.**

Traceability from checklist steps to SOP sections verifies coverage and detects invented or omitted steps. Count parity, colleague recall, and order alone are weaker proxies that miss content-level drift.

21. Team members need shared access to the same project knowledge with edit controls. Which plan/feature context matters?

- A. A Personal Pro accounts exchanging exported chat transcripts by email.
- B. Any Claude plan, since sharing controls are identical across all of them.
- C. Team or Enterprise project sharing with view/edit permissions.**
- D. Any plan provided by every teammate signs in with the same login credentials.

Project collaboration with view/edit permissions is a Team/Enterprise capability for shared organizational work. Transcript passing, plan-agnostic assumptions, and credential sharing do not provide managed collaboration.

22. A finance team runs the same month-end variance commentary with Claude every month. What turns this from one-off chats into a dependable workflow?

- A. A calendar reminder assigning whoever is free that month to ask Claude for the commentary.
- B. Saving last month's chat link so the next person can read through how it was done before.
- C. A Project with standing instructions, a fixed input checklist, and a named reviewer before distribution**
- D. Asking Claude at each month-end to repeat whatever it did for the previous month's version.

Repeatability comes from designed process: standing instructions and knowledge in a Project, standardized inputs, and a named reviewer. Informal reminders, chat archaeology, and asking the model to self-repeat leave the process undefined.

23. How should you structure project instructions for clarity?

- A. Role, purpose, tone/format rules, must-do constraints, and must-not-do boundaries**
- B. A set of prior outputs to imitate, with no stated rules about why they worked.
- C. A short mission statement only, trusting Claude to infer the operating details.
- D. A copy of every email thread where formatting preferences were ever discussed.

Clear instructions cover role, purpose, tone and format rules, and hard constraints and boundaries. Mission statements alone, raw archives, and unexplained exemplars fail to steer outcomes reliably.

24. A Claude answer mixes accurate points with one critical incorrect claim. What is the best evaluation outcome?

- A. Publish with a blanket disclaimer telling readers to verify all claims independently.
- B. Approve the accurate sections for use now and schedule the critical claim for correction in a later revision.
- C. Approve it with the incorrect claim struck through so that readers know to skip it.
- D. Reject or revise the whole deliverable until the critical claim is fixed; partial accuracy is not enough.**

Critical errors can dominate outcomes even if most text is fine. Evaluation should block or fix material mistakes before circulation; strikethroughs, deferred fixes, and blanket disclaimers all leave the error doing damage.

25. What should you typically put in project knowledge for a brand-writing project?

- A. Style guide, messaging pillars, approved product facts, and audience profiles**
- B. A list of links to external style websites that Claude should consult as needed.
- C. Every historical draft the team has ever produced, for maximum possible context.
- D. The company org chart and the meeting calendar for background team awareness

Brand-writing projects benefit from style, messaging, product facts, and audience context that is curated and current. Bulk archives, unrelated logistics, and bare link lists add noise without brand guidance.

26. Claude's first draft of the proposal section is 80% right but misses the client's stated budget constraint. What is the best iteration?

- A. Rewrite the section yourself by hand, since iterating on drafts rarely improves them.
- B. Regenerate the response without changes until the budget constraint eventually appears.
- C. Point out the specific gap, restate the budget constraint, and ask for a targeted revision.**
- D. Start over with a new chat and hope the next draft lands closer to what the client needs.

Effective iteration names the specific gap and supplies the missing constraint for a targeted revision. Restarting or regenerating without new information discards the 80% that worked, and abandoning iteration forfeits the tool's value.

27. Which metric best indicates a Claude workflow is adding business value?

- A. Total number of Claude messages sent by the team each month.
- B. Average length of the outputs Claude produces per request.
- C. The share of employees who have opened Claude at least once.
- D. Time-to-quality-draft and acceptance rate after human review**

Value shows up faster high-quality drafts and higher acceptance after review. Activity metrics such as message counts, adoption breadth, and output length do not measure business value.

28. You must produce a one-page HTML prototype landing page from a brief. Best product path?

- A. Ask Claude to create an Artifact (HTML) and iterate it in the artifact panel.**
- B. Ask for a slide-deck outline of the page and then design it manually from that.
- C. Ask Claude to describe the page section by section for a developer to build later.
- D. Ask for the HTML in a plain chat reply and paste it into a separate editor to view.

Artifacts support single-page HTML and iterative visual development beside the chat. Copy-out workflows, verbal specs, and slide outlines all postpone or bypass the working prototype the brief asks for.

29. Chats in a project do not reliably share one another's full content; only project knowledge and instructions are guaranteed context for every chat. What operational practice follows?

- A. Put durable reference material into project knowledge rather than relying on prior chat memory alone.**
- B. Manually paste a running recap of every prior chat at the start of each new conversation in the project
- C. Keep one single continuous chat running forever so nothing has to be written down.
- D. Rely on Claude to recall the earlier project chats whenever a new chat needs those details.

Only project knowledge and instructions are guaranteed to reach every chat in a project; memory features surface summaries of past chats rather than their full content and vary by plan and settings. Durable facts belong in knowledge files or instructions; recall assumptions, endless threads, and manual recaps are all fragile substitutes.

30. What is good practice when uploading knowledge files?

- A. Prefer current, relevant, well-named documents and avoid dumping unrelated or sensitive files.**
- B. Upload the complete department archive so no potentially useful file is ever missing.
- C. Keep the superseded policy versions alongside the current ones to preserve valuable historical context.
- D. Prefer scanned image PDFs because they best preserve the original document formatting.

Knowledge quality depends on relevance, currency, and safe content selection: current, well-named, text-readable documents. Archive dumps, image-only scans, and superseded versions degrade both security and answer quality.

31. A manager asks Claude to produce both a client-facing summary and an internal risk list from the same notes. What prompting practice best prevents mixed audiences?

- A. Produce the client summary first, then let the risk list reuse the same tone and wording.
- B. Let Claude infer from the notes which audience each point belongs to as it writes.
- C. Request two clearly labeled outputs with separate tone and detail rules for each audience.**
- D. Ask for one carefully worded document that serves the client and the internal team at once.

Separating deliverables by audience keeps tone, disclosure, and detail appropriate. Blending audiences often leaks internal risk language to clients or oversimplifies internal needs. Explicit dual outputs with labels prevent mixing.

32. For a multi-constraint scheduling problem where you want Claude to reason carefully before answering, which feature is designed to help?

- A. A concise response style, which strips the answer down to only its final conclusions.
- B. Artifacts, which move the model's reasoning into a separate panel for easier review.
- C. A longer chat history, which gradually trains Claude on your reasoning preferences.
- D. Extended thinking, which lets Claude work through the problem step by step before responding.**

Extended thinking gives Claude room to reason through complex problems before answering, which suits multi-constraint tasks. Response styles shape presentation, artifacts render outputs, and chat history provides context rather than training.

33. A user needs Claude on the desktop with local connectors for workplace tools. What product consideration matters?

- A. Desktop Claude is a separate product line whose chats cannot sync with the web app account.
- B. Local integrations require the business user to build custom API code before any connector can be used at all.
- C. Claude Desktop with connector/MCP-style integrations extends Claude beyond browser chat into local tools.**
- D. Connectors let Claude read workplace tools, so responsible-use review is no longer necessary.

Claude Desktop and connectors extend chat with tool access for workplace workflows without custom development, and responsible-use obligations still apply to everything the tools can reach.

34. A strategy lead needs deep multi-document synthesis with careful trade-off reasoning for an executive decision. Which model direction fits best?

- A. A fast model runs several times, merging the answers into a single synthesis.
- B. A higher-capability model tier such as Opus for complex reasoning and synthesis**
- C. The fastest tier, since executive deadlines make latency the deciding factor.
- D. Whichever model is currently set as the default, to keep the workflow simple.

Complex, high-stakes synthesis benefits from a higher-capability model plus human review. Speed-first or default-first selection under-serves the task, and stitching fast passes together is not depth.

35. A manager wants productivity metrics from Claude usage. Which approach is most responsible?

- A. Read individual employees' chats directly, since work accounts imply monitoring consent.
- B. Rank employees on a public leaderboard by their weekly Claude message counts.
- C. Measure outcomes and quality through privacy-aware analytics aligned with policy.**
- D. Infer productivity from who appears latest at night in the Claude usage logs.

Responsible analytics focus on outcomes and quality under policy and privacy constraints, not on reading private content, leaderboards, or surveillance-style proxies.

36. A user repeatedly gets off-topic answers. What prompt adjustment most helps?

- A. Move to a brand-new chat and ask the identical question again from scratch.
- B. Add more background documents so Claude has broader material to draw from
- C. Restate the primary objective and add an out-of-scope list of topics to avoid.**
- D. Ask the question several different ways within the same message for coverage.

Restating the objective and declaring out-of-scope topics re-centers the model on the job to be done. Adding volume, multiple phrasings, or a chat reset without a sharper objective leaves the underlying ambiguity in place.

37. Outputs became worse after you uploaded dozens of semi-related PDFs. Likely cause and fix?

- A. Model drift; switch model tiers and re-ask the same question over the same files.
- B. Context overload/noise; prune to the most relevant current sources and restate the task.**
- C. File-format problems; convert every PDF to plain text and re-upload all of them.
- D. Insufficient context: add the remaining PDFs so Claude finally sees the full picture.

Too much semi-related material can dilute attention and produce muddled answers. Prune to relevant current sources and restate the goal; adding more volume, switching tiers, or reformatting everything keeps the noise in place.

38. You hit usage limits during a busy week. What is practical optimization?

- A. Keep one very long chat open all week long so the accumulated context never has to be re-sent to Claude.
- B. Batch related tasks, reuse projects and templates, and route simple work to faster, cheaper models.**
- C. Run every task on the most capable model so that fewer retries will be needed.
- D. Duplicate the important requests across chats to guarantee at least one completes.

Under limits, optimize throughput: batch related work, reuse project context and templates, match simpler models to simple tasks, and prioritize. Top-tier defaults, ever-growing threads, and duplicated requests all waste quota.

39. What is a good pattern for Claude-assisted brainstorming that still leads to decisions?

- A. Have Claude score all of the ideas and adopt its top-ranked option automatically.
- B. Limit each session to the first three ideas Claude generates to force efficiency.
- C. Keep the sessions purely divergent so that evaluation criteria never constrain the creative flow of ideas.
- D. Diverge with Claude for options, then converge with explicit scoring criteria and an owner decision.**

Effective ideation separates divergence from convergence and ends with owner accountability. Endless divergence never yields a decision, automatic adoption removes accountability, and arbitrary caps discard breadth.

40. A research workflow uses Claude to synthesize interviews. Where should human judgment concentrate?

- A. On formatting the report so themes are visually distinct and easy to skim
- B. On standardizing interviewee wording the quotes are read consistently
- C. On expanding every theme with additional supporting language and detail
- D. On theme validity, quote fidelity, and decision implications.**

In qualitative synthesis, humans should validate themes, quote accuracy, and decision impact. Formatting and stylistic smoothing are secondary, and rewording participants' quotes actively harms fidelity.

41. In a Claude Project, where should the brand tone rules and the 40-page product manual each live?

- A. Both were uploaded as knowledge files, leaving the project instructions empty.
- B. Both pasted into the project instructions so that nothing depends on retrieval.
- C. The manual in the instructions and the tone rules as a knowledge file
- D. Tone rules in the project instructions; the manual as a project knowledge file**

Instructions hold durable behavior rules such as tone, format, and boundaries, while large reference documents belong in knowledge files. Reversing them or piling everything into one slot degrades steering and retrieval.

42. What is a practical human-in-the-loop checkpoint for Claude-assisted research memos?

- A. Have a reviewer check the formatting and tone while trusting the factual content.
- B. Require a reviewer to verify key claims, sources, and recommendations before distribution.**
- C. Ask Claude to self-review the memo and confirm that it is ready for distribution.
- D. Circulate the memo to a small group first and correct issues if anyone happens to reply.

Human review of claims, sources, and recommendations is the core HITL control for research memos. Self-review, send-then-fix, and cosmetics-only checks miss material risk.

43. Which practice best protects intellectual property when using Claude?

- A. Treat anything found on the open internet as clear for upload and unrestricted reuse.
- B. Rely on Claude to automatically detect and flag any uploaded content that might be under copyright before you make use of it.
- C. Do not upload third-party proprietary materials you are not allowed to share, respect licenses and client confidentiality.**
- D. Upload licensed third-party reports in full, since a license always covers internal AI tools.

IP-responsible use respects licenses, NDAs, and permission boundaries. License scope varies, public availability is not permission, and copyright clearance is a human and legal responsibility rather than something detectable from text.

44. Your team complains Claude workflows feel slower than manual work. How do you diagnose?

- A. Measure time spent prompting, waiting, reviewing, and reworking, then fix the slowest stage.**
- B. Deploy a faster model across the board on the assumption that generation is the bottleneck.
- C. Survey the team on which stage feels slowest and optimize based on the consensus vote.
- D. Compare only the total end-to-end time for Claude versus manual work on one sample task.

Speed complaints need stage-level measurement: prompt time, generation wait, review, and rework. Unclear prompts and vague review standards are common bottlenecks; feelings, single-number comparisons, and assumptions do not localize the problem.

45. You are designing an end-to-end content refresh workflow using Claude. What is the correct sequencing?

- A. Audit live content → Prioritize pages → Claude redrafts against brief → SME review → Publish → Measure.**
- B. Publish Claude's redrafts on a rolling basis and audit only whichever pages later draw reader complaints or corrections.
- C. Redraft every page with Claude first, then decide which of the redrafts merit review.
- D. Start with SME review of existing pages, then measure, then define the refreshing goal.

Content refresh should audit and prioritize, draft against a brief, review, publish, then measure. Drafting before prioritizing wastes effort, complaint-driven auditing lets errors go live, and measuring before goals produces unanchored data.

46. Which design best supports handoff between a Claude-assisted analyst and a manager?

- A. Deliver a quick one-line verbal summary and offer to answer any questions that happen to come up later.
- B. Send the final numbers only, keeping the assumptions internal to avoid confusing anyone.
- C. Deliver a structured brief with sources, assumptions, options, recommendation, and open questions.**
- D. Forward the complete chat history so the manager can trace how the conclusions emerged.

Handoffs need decision-oriented structure: evidence, assumptions, options, recommendation, and unknowns. Raw history, one-liners, and numbers without assumptions force the manager to reconstruct meaning.

47. Why should you disclose AI assistance when required by stakeholders or policy?

- A. Transparency supports trust, accountability, and informed review of AI-influenced work.**
- B. Disclosure transfers responsibility for any errors from the author to the tool
- C. Disclosure only becomes necessary when the work product later turns out to contain errors.
- D. Disclosed AI content becomes exempt from the ordinary review requirements.

Transparency about AI assistance enables proper review, trust, and accountability. It does not shift responsibility to the tool, waive review, or wait until something goes wrong.

48. A PMO wants Claude to generate RAID logs from meeting transcripts. Which design is correct?

- A. Have Claude write its extracted entries straight into the official risk register so no manual validation delays the process.
- B. Extract only the risks, since actions, issues, and dependencies rarely matter to a PMO.
- C. Keep the RAID log in the project chat where it was generated for convenient reference.
- D. Extract candidate risks/actions with fields, then a PM validates and enters approved items into the system of record.**

Claude can propose RAID items, but a PM should validate them and commit them to the system of record with owners and dates. Auto-committed entries, chat-only storage, and risk-only extraction all weaken PM control.

49. Claude keeps missing a required section in deliverables. What is the first troubleshooting step?

- A. Escalate the deliverable to a more capable model tier, on the theory that section omissions signal a capability limit.
- B. Regenerate the response until a version happens to include the missing section.
- C. Make the required sections explicit in the prompt or project instructions and provide a checklist to verify.**
- D. Split the deliverable into separate chats, one per section, and merge them manually.

Missing sections usually mean the requirements were implicit. Explicit section lists and verification checklists fix the process; capability upgrades, fragmentation, and lucky regeneration do not.

50. Which prompt is best for extracting structured data from messy notes?

- A. "Return JSON with keys issue, owner, due_date, severity. Use null when a field is missing. Do not invent values."**
- B. "Summarize the notes in a few clear paragraphs, mentioning owners and dates where they are visible."
- C. "List the issues in whatever structure feels most natural for this particular material."
- D. "Return JSON for each issue and fill in your best estimate whenever a field is not explicitly stated in the notes."

Explicit schema, nulls for missing fields, and a ban on invention produce machine-usable and auditable extraction. Prose or free-form structure is harder to validate and reuse, and estimated values undermine data quality.

51. You want Claude to help prepare for meetings. Which workflow is most effective?

- A. Ask Claude for general advice on running effective meetings of this particular type.
- B. Provide the agenda, attendee goals, and prior notes, then ask for talking points, risks, and questions to refine.**
- C. Prepare entirely by hand and use Claude only to format the final notes afterwards.
- D. Ask Claude to produce a full predicted transcript of how the meeting is likely to unfold, speaker by speaker and minute by minute.

Meeting prep works when Claude is grounded in the agenda, goals, and prior notes, then refined by a human. Generic advice, predicted transcripts, and after-the-fact formatting miss the preparation value.

52. You need a fast triage of hundreds of short support tickets into a few categories. Which model tier is generally the best first choice?

- A. Haiku, optimized for speed and high-volume simpler tasks.**
- B. Sonnet with extended thinking enabled for every individual ticket.
- C. A different model tier for each ticket category to diversify the results.
- D. Opus, because classification accuracy always justifies the most capable tier.

High-volume, relatively structured classification benefits from a fast, cost-efficient tier such as Haiku. Reserving the most capable tiers for harder reasoning keeps cost and latency proportionate to the task.

53. How should Claude fit into a quarterly planning cycle?

- A. Use Claude to pick the quarter's goals so planning starts from a neutral outside view.
- B. Restrict Claude to formatting the final plan document once leaders finish the thinking.
- C. Use Claude to synthesize inputs and draft plan options, while leaders set goals and approve trade-offs.**
- D. Have Claude grade last quarter's performance and mechanically set the new quarter's targets from that grade.

Claude can accelerate synthesis and drafting, but leaders own goals and trade-off approval. Delegating goal setting, format-only use, and grade-driven targets each misplace the division of labor.

54. Claude refuses a request you believe is legitimate business analysis. Best troubleshooting approach?

- A. Rephrase to clarify lawful business intent, remove ambiguous harmful interpretations, and split the task if needed.**
- B. Prepend a statement that you accept full responsibility for whatever is produced.
- C. Submit the identical request again several times in a row, on the theory that the refusal will eventually stop appearing.
- D. Ask for the same analysis reframed as a purely hypothetical fiction exercise.

Legitimate work sometimes triggers refusals due to ambiguous framing. Clarify benign intent and scope; verbatim retries change nothing, disclaimers do not alter the request, and fiction-framing to slip past safeguards is a jailbreak pattern.

55. Free Claude accounts can create projects with limits. Which statement is accurate based on Claude Projects' guidance?

- A. Projects are available to free users, with a maximum number of projects (currently five)
- B. Free accounts get unlimited projects but cannot upload knowledge files into them.
- C. Projects require at least a Pro subscription; free accounts remain chat-only.
- D. Free accounts may use projects that others share with them but can never create their own projects.**

Claude Help states projects are available to all users, including free accounts, with free users limited to a maximum of five projects. The limit is on count, not on whether free users can create projects or add knowledge.

56. When should you prefer a new standard chat over creating a Project?

- A. For a one-off question with no reusable files, instructions, or ongoing thread of work**
- B. Whenever a task involves any uploaded file, since files always require project knowledge.
- C. Whenever the work is confidential, since standard chats have stronger privacy settings.
- D. When you expect several related follow-up chats that need the same background

Projects shine for ongoing work with reusable knowledge and instructions; one-off questions are fine in standard chat. File attachment and confidentiality do not force either choice, or repeated related chats are exactly when a Project pays off.

57. Which data should you generally avoid pasting into Claude without explicit authorization and safeguards?

- A. Publicly filed financial statements of listed competitor companies
- B. Anonymized and aggregated survey statistics that have already been cleared for internal circulation.
- C. Marketing copy that is already scheduled for publication later this week
- D. Secrets such as passwords, API keys, regulated personal data, and highly confidential deal terms.**

Responsible use minimizes exposure of secrets, regulated personal data, and highly confidential material unless policy and safeguards allow it. Cleared aggregates, public filings, and about-to-publish marketing copy are comparatively low-risk classes.

58. How should multi-person content operations use Claude without version chaos?

- A. Define owners, a single source brief, review stages, and a final approver before publishing.**
- B. Route all edits through Claude chats so the chat history doubles as version control.
- C. Freeze every draft after first review so later changes cannot introduce conflicts.
- D. Let each contributor keep a personal draft and merge everything just before publication.

Content ops need ownership, a shared brief, staged review, and a final approver. Late merges and chat-as-version-control invite chaos, while freeze block needed corrections.

59. An Artifact shows an error after a code change. What is a good next step in Claude?

- A. Rebuild the whole artifact from scratch in a brand-new chat on the theory that hidden state is causing the error.
- B. Use the error details with Claude (for example, fix-with-Claude flows) to diagnose and iterate, then re-test**
- C. Roll back to an earlier version of the artifact and abandon the attempted change.
- D. Post the artifact link to the team channel and ask whether the error matters to anyone.

Artifact errors should be diagnosed with the error details and iterative fixes, then re-tested. Rollback abandons the improvement, scratch rebuilding cost far more than one targeted fix, and shipping a broken tool outsources debugging to colleagues.

60. How do project instructions differ from ordinary chat messages?

- A. Project instructions override organization-level policies whenever the two conflict.
- B. Project instructions persist and guide all chats in that project until changed.**
- C. Project instructions apply account-wide, including chats outside the project.
- D. Project instructions are visible to Claude only in the first chat created after saving.

Project instructions are durable settings for every chat inside that project: scoped to it, persistent until changed, and not an override of organizational controls.

61. A team will repeatedly ask questions against the same policy pack and brand voice rules. Which product choice fits best?

- A. A Claude Project with knowledge files and project instructions**
- B. A shared document of past Claude answers for teammates to search before asking.
- C. A saved chat that teammates reopen and continue whenever questions come up.
- D. Individual chats where each teammate re-uploads the policy pack as needed.

Projects provide persistent knowledge and instructions across chats for a body of work. Single long chats, answer archives, and per-person re-uploads all recreate the context problem Projects solve.

62. Which solution design best reduces rework when many employees use Claude for similar tasks?

- A. Let best practices spread through hallway conversation over time.
- B. Ask the most experienced user to handle all Claude requests for the whole team.
- C. Publish approved prompt playbooks and project templates for common tasks.**
- D. Rotate the tools frequently so employees do not settle into fixed habits.

Shared playbooks and templates spread effective patterns and reduce rework. Single-operator bottlenecks, informal diffusion, and tool churn either fail to scale or destroy accumulated skill.

63. You asked for a neutral competitive comparison. Claude's draft heavily favors one vendor with emotional language. What should you do?

- A. Request a rebalance with explicit criteria and comparable evidence for each vendor.**
- B. Keep the conclusion but ask Claude to soften the emotional wording slightly.
- C. Balance it by adding equally emotional language in favor of the other vendors.
- D. Accept the draft as written if the favored vendor really is the current overall market leader.

Evaluation includes fairness and criteria balance. Requiring comparable evidence and striving unsupported preference language restores neutrality, softening or offsetting the bias does not.

64. Which factor should most influence choosing a more capable model over a faster one?

- A. Whether the task's outputs will eventually be shared outside the immediate team
- B. How many separate steps the overall workflow contains from end to end.
- C. The total number of tokens the task is expected to consume from the very start through to the finish.
- D. Task difficulty, error cost, and need for deeper reasoning outweigh latency and cost constraints.**

Model selection is a trade-off among capability, speed, and cost. Higher capability is justified when task difficulty, error cost, and reasoning depth demand it, not by token volume, audience, or step count alone.

65. You need Claude to rewrite a dense technical paragraph for non-technical executives. Which role framing is most effective?

- A. "Act as a technical translator who explains complex ideas in plain business language for busy executives."**
- B. "Act as a PhD researcher who carefully preserves every technical term, caveat, and citation in full detail."
- C. "Act as a thought leader who showcases an advanced vocabulary to impress the audience."
- D. "Act as a creative fiction writer who dramatizes the material, so it entertains readers."

Role framing works when it points Claude at a useful skill and audience. A technical translator for executives signals plain language, business relevance, and brevity. Roles that optimize for technical completeness, entertainment, or vocabulary display pull the rewrite away from what executives need.

66. Multiple teammates edit the same project knowledge without coordination. What risk should you manage?

- A. Claude declining to answer questions while more than one editor is active.
- B. Automatic deletion of the oldest knowledge files whenever new ones is added.
- C. Conflicting instructions/files that cause inconsistent answers; use ownership and change notes.**
- D. The project knowledge base locked itself permanently after registering too many concurrent edits.

Uncoordinated edits create conflicting guidance, which quietly degrades answer consistency. Ownership and lightweight change notes reduce the inconsistency; the invented lockout, refusal, and auto-deletion behaviors do not exist.

67. Your task is not on the company's approved Claude case list, but it is not prohibited either. What is the next step responsible?

- A. Proceed but keep the chat private so that the task does not set any precedent.
- B. Proceed, since anything that is not expressly prohibited is by definition permitted.
- C. Route the task through a personal account where the company's list does not apply.
- D. Ask the policy owner or governance channel to classify the task before proceeding.**

Gray-zone tasks go to the policy owner for classification; that is what governance channels exist for. Not-prohibited is not the same as approved, secrecy is not compliance, and personal-account routing evades controls entirely.

68. Under a typical corporate AI policy, which task could proceed in Claude without extra approval?

- A. Summarizing a publicly published industry report for an internal briefing**
- B. Reviewing the details of an employee's medical accommodation request
- C. Drafting a response to a regulator using unfiled internal investigation records
- D. Analyzing a spreadsheet of customer names, emails, and purchase histories

Published public-source material is the lowest-risk data class and typically needs no special approval. Identifiable customer data, unfiled investigation records, and employee health information are sensitive classes that normally require policy checks, approvals, or approved tooling.

69. Quality is good in short chats but degrades in very long threads. What is a sensible fix?

- A. Push the thread further, since models warm up as their conversations grow longer.
- B. Switch to a more capable model within the same thread and simply continue.
- C. Start a fresh chat in the same Project, restate the goal, and keep only necessary recent context.**
- D. Paste the full conversation history into an even longer master prompt so continuity is preserved.

Long threads can accumulate noise and drift. A fresh chat inside the Project with a restated goal and minimal needed context often restores quality; more length, duplicated history, or a tier change within the overloaded thread do not.

70. Which Claude-drafted item most clearly requires qualified human review before it is used?

- A. A casual reminder message about a rescheduled daily stand-up
- B. A regulatory disclosure paragraph for investor communication**
- C. A brainstorm list of possible names for an internal team wiki
- D. A first-draft agenda for a routine internal team meeting

Review depth should scale with stakes. Regulated, external, high-consequence content such as investor disclosures requires qualified review; low-stakes internal drafts can ship after a quick self-check.

71. A project keeps giving outdated answers after a policy update. What configuration fix is best?

- A. Add the new policy document alongside the old ones so that Claude can weigh both versions and decide which applies.
- B. Tell Claude at the start of each chat to ignore anything that sounds out of date.
- C. Rename the project so its title signals that it now covers the updated policy.
- D. Update or replace the knowledge files and instructions to reflect the new policy and remove obsolete docs.**

Knowledge management requires retiring stale sources and updating instructions when policies change. Keeping both versions invites mixed answers, per-chat disclaimers are unreliable, and renaming changes nothing Claude retrieves.

72. Which governance artifact helps teams use Claude consistently and safely?

- A. A one-time launch email announcing that Claude is now available to all staff.
- B. An informal norm that people simply ask whichever manager is available whenever they are unsure about a use case.
- C. An internal acceptable-use guide covering data classes, review requirements, disclosure, and banned use cases.**
- D. A folder of screenshots showing examples of good Claude conversations.

Acceptable-use guides codify data handling, reviews, disclosure, and prohibitions, which is the governance teams need for consistent safe adoption. Announcements, for example screenshots, and ask-a-manager norms do not create standards.

73. What is a sound default workflow for Claude-assisted client deliverables?

- A. Claude drafts → Automated grammar check → Send once no errors are flagged.
- B. Intake requirements → Claude draft → Human expert review → Finalize and send.**
- C. Claude draft → Client review of the draft → Internal expert corrections afterwards
- D. Intake requirements → Human outline → Claude writes final → Send without review.

Reliable business workflows put human expert review before anything external and start from clear intake. Client-first review, review-free sends, and grammar-only gates all leave the riskiest step uncontrolled.

74. How should Claude be integrated into weekly business reporting?

- A. Automate direct distribution to stakeholders and quietly correct any metric issues in the following week's edition.
- B. Let each analyst prompt in their own personal style so reports keep an individual voice.
- C. Have Claude pull last week's report and update whichever numbers it remembers changing.
- D. Standardize inputs and templates, generate a draft, then have an owner validate metrics before distribution.**

Repeatable reporting needs standard inputs, templates, and owner validation of metrics before distribution. Memory-based updates, ad-hoc prompting, and fix-it-next-week distribution all put unreliable numbers in front of stakeholders.

75. A Claude answer includes a URL that returns 404. What does this imply for validation?

- A. The reference is not currently usable, and the related claim needs another source or removal.**
- B. The claim is fine if searching the page title returns similar-looking articles.
- C. The page probably just moved, so the citation can stand with a note about link rot.
- D. The claim is still adequately supported because the URL presumably existed when Claude learned it.

A dead link means the citation cannot support the claim as presented. Find and verify a working source (or the moved page) or drop and qualify the claim; assumptions about link rot are not verification.

76. When using Claude for external communications, which governance check matters most?

- A. Whether the statement can be published before competitors have commented
- B. Whether the draft matches the CEO's personal writing style closely enough
- C. Whether the wording is kept vague enough that no specific claim can ever be pinned down
- D. Accuracy, brand voice, required disclosures, and approval path for public statements.**

External comms governance centers on truth, brand, disclosures, and approvals. Style mimicry is cosmetic, timing pressure does not excuse skipped approvals, and strategic vagueness undermines required accuracy.

77. A user asks Claude how to build malware to harm others. What is the product expectation responsible?

- A. Claude should refuse harmful misuse and not provide assistance for criminal attacks.**
- B. Claude should comply whenever the requester claims to be a security researcher.
- C. Claude should provide partial guidance while omitting only the most dangerous step.
- D. Claude should answer fully, since responsibility rests solely with the person asking

Responsible AI systems refuse assistance for harmful criminal misuse such as malware for attacks. Unverified researcher claims, partial harm, and blame-shifting to the user do not satisfy that expectation.

78. How can Claude improve an RFP response workflow?

- A. Draft every answer from scratch in one long chat so the response voice stays uniform.
- B. Map RFP questions to owners, draft from approved content libraries, then run compliance and win-theme reviews.**
- C. Answer only the scored sections in depth and paste boilerplate into the remaining ones.
- D. Assign the entire response to Claude end to end and review everything in a single pass right at the submission deadline.

Strong RFP workflows combine requirement mapping, grounded drafting from approved content, and human compliance and win-theme review, with time left to fix what the review finds.

79. You compare Claude's analysis to a spreadsheet source and find a transposed percentage. What is the right operational response?

- A. Add a footnote that the numbers are approximate and distribute the analysis as it stands.
- B. Correct the single figure and move on, since the rest of the analysis is read correctly.
- C. Ask Claude to regenerate the analysis and go with whichever version looks cleaner.
- D. Correct the figure, note the error pattern, and re-check other numbers for similar mistakes.**

A found transposition should trigger correction and a broader numeric QA pass. Single errors can indicate a pattern (row shifts, unit mix-ups). Publishing unchanged is unacceptable; banning all numbers is overreaction.

80. Which statement best reflects Anthropic's responsible AI direction for Claude's behavior?

- A. Claude's safety rules apply only to enterprise customers rather than individual users.
- B. Claude is trained to be helpful while avoiding harmful, dishonest, or disallowed assistance under safety policies.**
- C. Claude defends all safety judgments against whatever acceptable-use policy each individual customer organization has adopted.
- D. Claude optimizes purely user satisfaction scores in every single interaction.

Claude is built to be helpful within safety boundaries that apply across the product; those boundaries are not tire-specific and are not delegated entirely to customers.

81. Which practice best reduces silent errors when Claude reformats a table?

- A. Compare row counts, totals, and a sample of cell values to the source table after reformatting.**
- B. Verify that the column headers and the row labels transferred across correctly.
- C. Read the reformatted table from top to bottom to see whether anything looks out of place.
- D. Have Claude reformat the table a second time and then confirm that the two outputs agree with each other.

Structural QA (counts, totals, sampled cells) catches silent reformatting errors. Visual inspection and header-only checks miss body errors, and agreement between two model runs is not a comparison against the source.

82. You paste a long transcript and ask Claude to "find insights." Results are shallow. What is a stronger task formulation?

- A. "List everything interesting that anyone said at any point during the call."
- B. "Read the transcript again carefully and give me deeper, more meaningful insights this time."
- C. "Extract the top 5 customer objections with supporting quotes and one suggested response for each."**
- D. "Summarize the whole conversation for me so that I can hunt through it for the interesting parts myself."

Specific extraction criteria (top 5 objections, quotes, suggested responses) create a clear, evaluable deliverable. Vague or unbounded requests produce generic commentary; a defined deliverable with quotes keeps the analysis grounded in the actual conversation.

83. You asked Claude for recommendations based only on an attached policy PDF. The answer cites a regulation not in the PDF. Best evaluation response?

- A. Ask Claude to expand the answer with several additional supporting citations.
- B. Flag the unsupported citation and require grounding in the provided policy or an explicit external source label.**
- C. Accept the citation as long as the referenced regulation sounds consistent with the policy's general subject matter.
- D. Reward the recommendation so the outside citation is implied rather than stated.

When you constrain Claude to a source set, new legal citations are a red flag unless labeled as outside knowledge and then verified. Credibility theater via extra unsourced cites makes evaluation worse.

84. A sales ops analyst wants Claude to rewrite email drafts. Past drafts are too salesy. Which constraint is most useful?

- A. "Use a consultative tone, max 120 words, no superlatives, end with one clear next step."**
- B. "Add urgency and energy so the emails feel excited and drive immediate replies."
- C. "Improve these drafts and make them feel more compelling to prospective buyers."
- D. "Rewrite these in a bold promotional voice that showcases our product's superiority."

Specific tone, length, banned patterns, and CTA constraints give Claude measurable targets. Vague improvement requests leave the failure mode unchanged, and high-energy promotional styles amplify the salesy tone the analyst is trying to remove.

85. What is the best way to integrate Claude into a document translation-plus-localization workflow?

- A. Require human translators to start from scratch, using Claude solely to count words.
- B. Machine draft with Claude, then bilingual reviewer adapts tone, idioms, and regulated terms.**
- C. Have Claude translate and also self-certify the idioms and regulated terms per locale.
- D. Translate with Claude and send it to native speaker review only if readers actually complain.

Localization workflows pair model drafts with qualified human review for tone, idioms, and regulated terminology. Self-certification and complaint-driven review let errors reach the audience; scratch-only work forfeits legitimate drafting speed.

86. Claude's summary says revenue grew 12%, but the source spreadsheet's summary row shows 8% growth. What is the soundest verification sequence?

- A. Average the two figures, since the true value most likely falls somewhere in between them.
- B. Check whether the figures use different scopes or periods, then correct the summary against the confirmed source definition.**
- C. Prefer whichever of the two figures appeared most recently in conversation history.
- D. Prefer the summary's figure, since Claude weighed the entire spreadsheet rather than a single row.

Discrepancies often trace to scope or period mismatches, such as total versus segment or fiscal versus calendar year. Resolve the definition first, then correct the summary against the confirmed source. Averaging and recency have no evidentiary basis.

87. Which statement best reflects responsible product selection for confidential M&A analysis?

- A. Upload the full data room up front so the analysis never lacks context, then restrict the project sharing afterwards.
- B. Keep all of the deal work in personal accounts so it stays outside company oversight.
- C. Use company-approved Claude surfaces, minimize sensitive uploads, and keep project access tightly controlled.**
- D. Use whichever surface the deal team finds fastest and rely on the NDA to cover the risk.

For confidential work, product choice includes policy compliance, data minimization, and controlled access. NDAs, retroactive restriction, and personal-account workarounds do not substitute for those controls.

88. Responses are too verbose for busy executives. What optimization should you apply?

- A. Adopt a norm that executives read only the first paragraph of each response.
- B. Delete most of the background context so there is simply less source material to talk about
- C. Set strict length limits, request an executive summary first, and specify bullet structure.**
- D. Ask Claude to write more formally, since formal prose tends to run shorter.

Verbosity is controlled with length caps, summary-first structure, and bullets for scanability. Formality is independent of length, stripping context degrades quality, and skim-only norms hide content instead of shaping it.

89. What is a practical data minimization habit with Claude?

- A. Paste the full records but instruct Claude to read only the fields that are relevant.
- B. Share the complete datasets on a quarterly schedule rather than sending scoped excerpts per task.
- C. Share only the fields and excerpts needed for the task, redacting unnecessary personal data.**
- D. Rely on deleting the chats afterwards instead of limiting what goes into them.

Minimization means providing only necessary data and redacting extra personal fields before sharing. Read-only instructions, after-the-fact deletion, and batched full-dataset sharing all expose data that never needed to leave the source.

90. What does "decision-ready" mean when evaluating a Claude deliverable?

- A. It has been carefully reviewed for grammar, formatting, and consistent brand voice.
- B. It presents a single firm recommendation while leaving no open questions on the record.
- C. A human could act on it with acceptable risk after reviewing evidence, assumptions, and residual unknowns.**
- D. It includes every piece of the underlying source data so readers can redo the entire analysis for themselves.

Decision-ready outputs are clear enough to act on, with known evidence, assumptions, and remaining risks. Polished surfaces, exhaustive raw data, or manufactured certainty do not equal readiness.

91. What is the right response if Claude generates medical advice that could be taken as a diagnosis?

- A. Forward it to the affected person as a convenient and free second medical opinion.
- B. Keep it for your own personal decisions while avoiding sharing it with anyone else.
- C. Treat it as reliable whenever a search engine independently surfaces the same suggestion in its top results.
- D. Treat it as non-authoritative information and direct real medical questions to qualified professionals.**

Responsible use keeps Claude from being positioned as a licensed clinician. Users should seek qualified professionals for diagnosis and treatment decisions; forwarding, search-engine corroboration, and private reliance all still treat the output as a diagnosis.

92. Claude provides two alternative strategies with trade-offs. How do you evaluate which is better for your team?

- A. Adopt the option Claude presents first, since models usually lead with the stronger case.
- B. Choose the option with fewer listed risks, because fewer risks mean a safer plan.
- C. Ask Claude which option it would choose and simply follow that recommendation.
- D. Score each option against your stated goals, constraints, risks, and success metrics.**

Comparative evaluation should map options to goals, constraints, risks, and metrics. Presentation order, raw risk counts, and delegating the pick to the model are not decision criteria.

93. How should you handle Claude outputs that could be mistaken for official company policy?

- A. Let individual team leaders adopt them locally without any central policy review.
- B. Share them verbally rather than in writing so no official-looking record gets created.
- C. Label them as drafts, route through policy owners, and avoid presenting unapproved text as policy.**
- D. Publish them quickly on the intranet so employees have some interim guidance while the owners catch up.

Unapproved AI text must not be treated as policy. Label drafts, use owners, and preserve auditability; speed, verbal circulation, and fragmented local adoption all create governance failures.

94. A sales team wants Claude to draft proposals. Which integration design is safest and most effective?

- A. Reuse the most recent winning proposal wholesale, updating only the customer's name.
- B. Ground drafts in approved pricing and product sources, with seller and deal-desk review before send.**
- C. Have Claude personalize each proposal from whatever it can infer about the customer online.
- D. Let sellers send Claude drafts directly to customers for small deals and review only the larger ones.

Proposal workflows ground on approved pricing and product facts with human commercial review before sending. Inference-based personalization and wholesale reuse introduce factual, pricing, and confidentiality risk.

95. You have a multi-step research task: summarize three PDFs, extract risks, then draft recommendations. What is the best prompting approach?

- A. Break the work into sequenced steps with intermediate checks before the final recommendation.**
- B. Ask for the recommendations first, then work backwards to fill in the supporting analysis.
- C. Paste all three PDFs and every instruction into one message and review only the final output.
- D. Run each PDF in a separate unrelated chat so the summaries cannot influence each other.

Complex work is more reliable when decomposed: summarize, extract risks, then recommend, with human review of intermediates. A single mega prompt hides errors and makes quality control harder. Sequencing improves accuracy and auditability.

96. A complex analysis request fails when asked all at once. Which prompting strategy is most appropriate?

- A. Resubmit the full request with a stronger instruction to think much more carefully this time.
- B. Ask for a much shorter final answer so the task fits within one manageable response.
- C. Use progressive disclosure: first outline the approach, then analyze section by section with checkpoints.**
- D. Split the work across several entirely unrelated chats so each part starts from a completely clean context.

Progressive disclosure and checkpoints let you correct course early and reduce cascading errors. Effort exhortations, fragmented contexts, or compressed answers leave the underlying task structure unmanaged.

97. Claude produces a hiring screen summary that stereotypes candidates by name origin. What should you do?

- A. Treat the output as one neutral data point, since models have no discriminatory intent.
- B. Keep the summary but delete the candidate names so the pattern is less visible.
- C. Keep using the summary for screening decisions while HR investigates whether the bias actually influenced any scores.
- D. Reject the biased framing, re-run with job-related criteria only, and review for fairness before any decision use.**

Responsible use requires removing non-job-related biased criteria and enforcing fairness review before any decision use. Masking names, intent arguments, and use-during investigation all leave biased output in the decision path.

98. A project lead wants Claude to produce a project status update in a fixed format every week. What should the prompt include to make the output consistent?

- A. A different example updates every week, so the format keeps evolving alongside the work.
- B. A reusable template with required sections, field definitions, and an example of a good status update**
- C. A rule that the update must stay under fifty words, so the structure remains simple.
- D. A request for a status update plus a note that Claude may pick whatever layout best fits each new week.

Consistency comes from explicit structure plus examples. A template with required sections and a model sample teaches Claude the expected shape and level of detail. Open-ended requests produce format drift week to week.

99. You are rolling out a Claude-assisted contract-summary workflow to 40 paralegals. Which rollout design most reduces risk?

- A. Restrict the workflow permanently to the two most senior paralegals on the team.
- B. Enable it for all 40 users at once so that every user's feedback arrives in one single wave.
- C. Roll it out silently with no training so that natural usage patterns can emerge on their own.
- D. Pilot with a few users against a quality bar, fix failure patterns, then scale with training and a feedback channel**

Staged rollout catches systematic errors while the blast radius is small: pilot, measure against a quality bar, fix failure modes, then scale with training and feedback. Big-bang, silent, and permanently restricted rollouts either multiply risk or forfeit the value.

100. When should you provide few-shot examples in a Claude prompt?

- A. Only when project instructions are unavailable to hold the formatting rules instead.
- B. Mainly when the prompt is already long, because examples work best with maximum context.
- C. Only when the task has a single objectively correct answer Claude might otherwise miss.
- D. When output format, style, or edge-case handling is hard to describe in rules alone.**

Examples teach pattern recognition for format, tone, and tricky cases that pure rules under-specify. They are especially useful for classification labels, writing style, and structured extraction, regardless of prompt length or where durable rules are stored.

101. What are Artifacts best used for in configuration terms?

- A. Creating and iterating substantial standalone outputs like docs, code, and interactive tools in a dedicated panel**
- B. Automatically archiving every chat response, the account produces so that compliance retention requirements are satisfied.
- C. Scheduling recurring Claude tasks that run in the background without an open chat
- D. Storing project knowledge files in a compressed form that Claude can search faster.

Artifacts surface substantial standalone content in a dedicated window for editing, iteration, and reuse. They are an output workspace, not storage, archiving, or scheduling infrastructure.

102. Which prompt element most reduces Claude inventing missing details in a business write-up?

- A. Explicit instructions to use only provided facts and to mark unknowns as unknown.**
- B. An instruction to produce a longer and more detailed draft of the same content.
- C. A reminder that the deliverable is high-stakes and therefore must be flawless.
- D. A request for a confident, authoritative tone throughout the entire write-up

Grounding instructions tell Claude to stick to supplied facts and surface gaps instead of filling them inventively. Confidence, length, or perfection demands can increase fluent but unsupported claims.

103. Claude chats for your whole team have failed to load files all morning, across projects and browsers. What is the right troubleshooting posture?

- A. Treat it as a platform or IT issue: check the status page and escalate through support channels.**
- B. Treat it as a model-selection problem and switch tiers until the uploads work.
- C. Treat it as a prompting problem and iterate on clearer file-handling instructions.
- D. Treat it as a knowledge-based problem and rebuild each affected project's files.

Team-wide, cross-project, cross-browser failures point to platform or environment issues, which call for status checks and support escalation rather than prompt or configuration churn. Recognizing when a problem is technical rather than usage-level is part of troubleshooting.

104. Which failure mode is most dangerous when Claude drafts financial commentary from a spreadsheet?

- A. An obviously garbled table that fails to render and stops the draft from being used.
- B. Minor rounding differences that are visible when compared against the source data.
- C. Correct-looking narrative that misstates trends because it swapped columns or periods.**
- D. A narrative that hedges heavily and repeatedly asks the analyst to verify the figures.

Fluent but factually inverted trend narratives can drive wrong financial decisions precisely because nothing looks broken. Visible breakage, hedged text, and small rounding gaps stand out and get corrected; silent column or period swaps do not.

105. How should you validate Claude's paraphrase of a legal clause for an internal explainer?

- A. Compare the paraphrase to the original clause for meaning drift, then have a qualified reviewer approve if stakes are high.**
- B. Ask Claude to rate his own confidence in the paraphrase's legal accuracy and record it.
- C. Read the paraphrase on its own to confirm it reads clearly, flows well, and is entirely free of dense legal jargon for employees.
- D. Accept it if each key term from the clause also appears somewhere in the paraphrase.

Paraphrase validation checks semantic fidelity against the original clause, with qualified review when legal stakes are high. Clarity-only reads, self-rated confidence, and keyword overlap all silent changes to obligations and exceptions.

106. You need Claude to follow a company style guide. Where should the enduring style rules live for best ongoing results?

- A. In the first chat of the project only, since later chats inherit earlier conversations.
- B. Pasted manually into each chat message whenever someone remembers the rules apply.
- C. In a shared document outside Claude that writers consult after drafts are generated
- D. In project instructions or a reusable prompt block, not only in one-off chat messages**

Durable rules belong in project instructions or a standard prompt template, so every related chat inherits them. One-off messages are easy to forget. Rules stored outside the conversation cannot guide Claude.

107. You receive a Claude draft that is well written but does not answer the actual question asked. What is the evaluation verdict?

- A. Pass with a note, since strong writing suggests the analysis behind it is sound.
- B. Fail on style and request shorter paragraphs before re-reviewing the content.
- C. Failure on task relevance; require a rewrite that addresses the question asked first.**
- D. Pass if the draft would be genuinely useful for a different but related purpose.

Relevance to the requested question is a primary evaluation dimension. Eloquence or incidental usefulness cannot compensate for answering the wrong job; fix relevance before style.

108. A Project answers with outdated product names after a rebrand. What do you check first?

- A. Whether the model version needs upgrading to one trained after the rebrand
- B. Whether the chat history length has hit some limit that forces older cached completions
- C. Whether teammates have been deliberately editing prompts to use the old names
- D. Whether knowledge files and instructions still contain old names and need updating**

Outdated branding in a Project almost always comes from stale knowledge and instructions, so those sources are the first check. Training vintage is overridden by project knowledge for names it defines, and the other mechanisms do not exist.

109. Claude claims a competitor "always" fails on reliability, based on one anecdote in your notes. Best evaluation action?

- A. Balance it by adding an equally strong positive claim about the same competitor.
- B. Keep the claim but attribute it to the customer who reported the original incident.
- C. Downgrade the claim to match the evidence strength and separate anecdote from pattern.**
- D. Move the claim into a footnote so that it carries less weight in the document.

Evidence strength must match claim strength: one anecdote cannot support "always." Attribution, placement, and counterweight claims do not fix scope; qualified, evidence-tied language does.

110. Claude returns a confident answer with no uncertainty language on a question where your data is incomplete. What should you do?

- A. Probe for what is unknown, request assumptions explicitly, and avoid treating confidence as evidence.**
- B. Accept the answer, since confident phrasing usually reflects strong underlying evidence.
- C. Ask Claude to restate the same answer more cautiously so the tone matches the data.
- D. Present the answer to stakeholders alongside a general disclaimer that AI systems can sometimes be wrong.

Fluent confidence is not epistemic certainty. Good evaluation surfaces assumptions and unknowns, especially under incomplete data. Hiding uncertainty or equating tone with truth degrades decision quality.

111. A support team wants Claude to draft replies. What design element is essential?

- A. A blanket policy that agents rewrite every single Claude draft from scratch to guarantee a human touch.
- B. Full automation for billing and legal tickets first, since those follow the strictest rules.
- C. Macros/prompts tied to issue types, escalation rules, and mandatory human send for sensitive cases.**
- D. A single universal reply prompt used for every ticket type to keep the training simple.

Support workflows need issue-typed prompts, escalation rules, and human control over sensitive commitments. Automating the riskiest categories first or mandating full rewrites both misfires.

112. A marketing manager asks Claude: "Write something about our product." The first draft is generic and unusable. Which change most improves the prompt?

- A. Add a clear goal, audience, product facts, tone, length, and success criteria.**
- B. Switch to a more capable model tier before clarifying what the task requires.
- C. Send the same request again in a fresh chat so Claude tries a different angle.
- D. Ask Claude to make the next draft more creative without adding any other detail.

Effective prompts specify the goal, audience, key facts, tone, length, and what good looks like. Vague requests produce generic drafts because Claude must invent missing requirements. Adding concrete constraints is the highest-leverage fix before model switching or re-asking.

113. What is the primary purpose of a Claude Project?

- A. To create a self-contained workspace with its own chats, instructions, and knowledge base for a body of work**
- B. To provide a shared inbox where teammates route their Claude conversations for triage.
- C. To permanently increase the usage limits available to every chat created inside it.
- D. To train a custom, organization-specific Claude model on the documents your company chooses to upload into the workspace.

Projects are self-contained workspaces with chat histories, project instructions, and knowledge files for focused work. They organize context; they do not raise limits, train models, or route conversations.

114. Which evaluation criterion is most important before sending a Claude-drafted customer email?

- A. Factual correctness, appropriate tone, correct recipient context, and absence of confidential leakage**
- B. Whether the draft was produced by the most capable model tier available
- C. Whether the draft closely matches the length and structure of previous emails sent to that same customer
- D. Whether the writing sounds polished and is free of grammatical or spelling errors

Customer-facing email quality hinges on truth, tone, recipient context, and confidentiality. Formatting consistency, surface polish, and model pedigree are secondary to those checks.