Status	Started	Completed
Finished	Wednesday, 11 November 2026, 2:39 PM	Wednesday, 11 November 2026, 2:44 PM

Duration	Grade
4 mins 45 secs	**19.00** out of 20.00 (**95**%)

Question 1

Flag question

Correct Mark 1.00 out of 1.00

Canary deployments for prompt changes work by:

a. Testing changes on synthetic data only
b. Deploying changes to a separate region
c. Requiring approval from all stakeholders
d. Routing a small percentage of traffic to the new version first

The correct answer is: Routing a small percentage of traffic to the new version first

Question 2

Flag question

Correct Mark 1.00 out of 1.00

What is metadata filtering in a RAG system?

a. Cleaning HTML and special characters from documents
b. Compressing metadata to reduce storage
c. Removing sensitive data before indexing
d. Restricting retrieval to documents matching specific attributes like date or category

The correct answer is: Restricting retrieval to documents matching specific attributes like date or category

Question 3

Flag question

Correct Mark 1.00 out of 1.00

Audit logging in a production Claude deployment must capture:

a. Only outputs that triggered safety filters
b. Inputs, outputs, model versions, and timestamps for every request ✓
c. Only requests that returned errors
d. User feedback ratings for quality review

The correct answer is: Inputs, outputs, model versions, and timestamps for every request

Question 4

Flag question

Incorrect Mark 0.00 out of 1.00

Environment variables are the standard way to manage API keys because:

a. They automatically rotate keys on schedule ✗
b. They are required by Anthropic's API specification
c. They encrypt keys at the OS level always
d. They keep secrets out of source code and version control

The correct answer is: They keep secrets out of source code and version control

Question 5

Flag question

Correct Mark 1.00 out of 1.00

Immutable audit logs differ from standard logs by:

a. Being stored in a separate cloud region
b. Compressing automatically after 30 days
c. Requiring admin approval to write entries
d. Not being alterable after writing ✓

The correct answer is: Not being alterable after writing

Question 6

Flag question

Correct Mark 1.00 out of 1.00

Storing an API key as a hardcoded string in source code is problematic because:

a. It slows API response time
b. It exposes the key in version control repositories ✓
c. It prevents multiple developers from using Claude
d. It causes authentication failures in production

The correct answer is: It exposes the key in version control repositories

Question 7

Flag question

Correct Mark 1.00 out of 1.00

API key rotation should occur:

a. Regularly on a defined schedule and immediately after suspected compromise ✓
b. Only after a confirmed security incident
c. Only when a developer leaves the team
d. Only when the API plan is upgraded

The correct answer is: Regularly on a defined schedule and immediately after suspected compromise

Question 8

Flag question

Correct Mark 1.00 out of 1.00

The stop_reason field in an API response indicates:

a. The model's confidence level in the response
b. Whether the response used tool calls
c. Why generation stopped — end_turn, max_tokens, or tool_use ✓
d. A safety filter that triggered on output

The correct answer is: Why generation stopped — end_turn, max_tokens, or tool_use

Question 9

Flag question

Correct Mark 1.00 out of 1.00

The recommended retry strategy for 429 and 529 errors is:

a. Immediate retry three times
b. Wait exactly 30 seconds then retry
c. Exponential backoff with added jitter ✓
d. Switch to a different model tier

The correct answer is: Exponential backoff with added jitter

Question 10

Flag question

Correct Mark 1.00 out of 1.00

The /v1/messages/count_tokens endpoint is used to:

a. Count tokens before submitting a full request ✓
b. Calculate cost estimates for billing
c. Monitor token usage across deployments
d. Validate system prompt token limits

The correct answer is: Count tokens before submitting a full request

Question 11

Flag question

Correct Mark 1.00 out of 1.00

Re-ranking retrieved chunks before injection improves quality by:

a. Reordering chunks by relevance to the specific query ✓
b. Reducing the total number of chunks used
c. Compressing chunk content to save tokens
d. Removing duplicate content across chunks

The correct answer is: Reordering chunks by relevance to the specific query

Question 12

Flag question

Correct Mark 1.00 out of 1.00

The primary risk of injecting too many chunks per query is:

a. Retrieval latency becomes unacceptable
b. The embedding model overloads
c. The knowledge base index becomes corrupted
d. Relevant content gets diluted by less relevant chunks in context ✓

The correct answer is: Relevant content gets diluted by less relevant chunks in context

Question 13

Flag question

Correct Mark 1.00 out of 1.00

Model version pinning in a production deployment means:

a. Specifying an exact model version to prevent unexpected behaviour changes ✓
b. Preventing any model updates for 12 months
c. Locking the deployment to one model tier permanently
d. Requiring Anthropic approval before any version change

The correct answer is: Specifying an exact model version to prevent unexpected behaviour changes

Question 14

Flag question

Correct Mark 1.00 out of 1.00

The usage object in an API response contains:

a. input_tokens and output_tokens counts ✓
b. Billing amount for the request
c. Cache hit and miss statistics only
d. API latency metrics for the call

The correct answer is: input_tokens and output_tokens counts

Question 15

Flag question

Correct Mark 1.00 out of 1.00

The most important operational document for a production Claude deployment is:

a. A runbook covering monitoring, alerting, escalation, and update procedures
b. The API cost breakdown for the billing team
c. The architectural diagram from the design phase
d. The original project requirements document

The correct answer is: A runbook covering monitoring, alerting, escalation, and update procedures

Question 16

Flag question

Correct Mark 1.00 out of 1.00

System prompt version tagging in logs allows:

a. Sharing prompts across deployments securely
b. Correlating quality metrics to specific prompt versions
c. Automatically reverting poor-performing prompts
d. Encrypting prompt content in transit

The correct answer is: Correlating quality metrics to specific prompt versions

Question 17

Flag question

Correct Mark 1.00 out of 1.00

When a user query is ambiguous for retrieval, the best practice is:

a. Rewrite the query to improve retrieval precision before searching
b. Ask Claude to select the best chunk itself
c. Retrieve the maximum number of chunks available
d. Use the raw user query without modification

The correct answer is: Rewrite the query to improve retrieval precision before searching

Question 18

Flag question

Correct Mark 1.00 out of 1.00

A/B testing two system prompts in production requires:

a. Running both versions on the same user simultaneously
b. Deploying both versions to different regions
c. Testing on synthetic data before any live traffic
d. Routing a defined traffic split to each version and measuring outcomes ✓

The correct answer is: Routing a defined traffic split to each version and measuring outcomes

Question 19

Flag question

Correct Mark 1.00 out of 1.00

HTTP 529 from the Anthropic API means:

a. API is overloaded — retry with backoff ✓
b. Request format is invalid
c. Context window limit exceeded
d. Authentication token expired

The correct answer is: API is overloaded — retry with backoff

Question 20

Flag question

Correct Mark 1.00 out of 1.00

An enterprise knowledge assistant must not show one department's documents to another. The architectural solution is:

a. Requiring users to specify their department in each query
b. Using separate Claude deployments per department
c. Encrypting documents before indexing them
d. Applying access control filters during retrieval based on user permissions ✓

The correct answer is: Applying access control filters during retrieval based on user permissions

Previous activity
Domain 5 - Product and Model Selection

Jump to...

Next activity
Domain 7 - Troubleshooting and Optimization