Status	Started	Completed
Finished	Wednesday, 11 November 2026, 7:20 PM	Wednesday, 11 November 2026, 8:01 PM

Duration	Grade
41 mins 38 secs	**54.00** out of 65.00 (**83.08**%)

Question 1

Flag question

Correct Mark 1.00 out of 1.00

Systematic test coverage for a Claude deployment must include:

a. Only adversarial inputs for security purposes
b. Only happy-path tests with valid in-scope inputs
c. Common cases, edge cases, out-of-scope queries, and adversarial inputs
d. Only edge cases since common cases rarely fail

The correct answer is: Common cases, edge cases, out-of-scope queries, and adversarial inputs

Question 2

Flag question

Correct Mark 1.00 out of 1.00

Summarising completed tool interactions in a long agentic session is done to:

a. Reduce context size while preserving semantic information
b. Reset Claude's tool selection preferences
c. Speed up future tool call execution
d. Log tool interactions for audit purposes

The correct answer is: Reduce context size while preserving semantic information

Question 3

Flag question

Correct Mark 1.00 out of 1.00

In streaming responses, content_block_delta events contain:

a. Incremental text chunks as they generate ✓
b. The final complete response text
c. Tool call results being returned
d. Metadata about generation progress only

The correct answer is: Incremental text chunks as they generate

Question 4

Flag question

Correct Mark 1.00 out of 1.00

The budget_tokens parameter in extended thinking controls:

a. Maximum output tokens for the response
b. Total tokens in the full API request
c. Tokens Claude can spend on internal reasoning ✓
d. Cache token budget for the request

The correct answer is: Tokens Claude can spend on internal reasoning

Question 5

Flag question

Correct Mark 1.00 out of 1.00

An MCP server's initialisation handshake establishes:

a. The user's authentication credentials
b. The default model tier to use
c. Protocol version and capability negotiation ✓
d. The maximum tool calls per session

The correct answer is: Protocol version and capability negotiation

Question 6

Flag question

Correct Mark 1.00 out of 1.00

HIPAA is most relevant to a Claude deployment that:

- a. Handles protected health information for a US covered entity ✓
- b. Processes any health-related search queries
- c. Operates in a healthcare-adjacent industry
- d. Uses health-related training data

The correct answer is: Handles protected health information for a US covered entity

Question 7

Flag question

Correct Mark 1.00 out of 1.00

When onboarding a developer to Claude Code, the most effective first step is:

- a. Providing the full Claude Code documentation to read
- b. Starting with a bounded real task from their actual work ✓
- c. Starting with the most complex available use case
- d. Having them complete all training modules first

The correct answer is: Starting with a bounded real task from their actual work

Question 8

Flag question

Correct Mark 1.00 out of 1.00

Metadata filtering in a RAG system restricts retrieval to:

- a. Documents from the most recent index build
- b. Documents matching specific attributes like date or category ✓
- c. Documents with the highest embedding scores
- d. Documents under a maximum token length

The correct answer is: Documents matching specific attributes like date or category

Question 9

Flag question

Incorrect Mark 0.00 out of 1.00

tool_choice set to {"type": "any"} forces Claude to:

a. Call at least one tool from the available list
b. Use the first tool alphabetically listed
c. Use all tools simultaneously in one call
d. Select tools randomly from the available list ⊗

The correct answer is: Call at least one tool from the available list

Question 10

Flag question

Correct Mark 1.00 out of 1.00

The anthropic-beta header in an API request is used to:

a. Pin the API contract version used
b. Enable verbose debug logging output
c. Opt into a specific pre-release feature ✓
d. Declare the SDK version making the call

The correct answer is: Opt into a specific pre-release feature

Question 11

Flag question

Correct Mark 1.00 out of 1.00

Claude Code handling production database credentials must:

a. Store them in CLAUDE.md for convenience
b. Use environment variables — never expose them in context ✓
c. Embed them in project configuration files
d. Pass them as command-line arguments instead

The correct answer is: Use environment variables — never expose them in context

Question 12

Flag question

Correct Mark 1.00 out of 1.00

"Galaxy-brained reasoning" in an agentic context refers to:

a. Claude misidentifying tool input parameters
b. Claude refusing valid tasks due to excessive caution
c. Claude constructing plausible but flawed reasoning toward a harmful conclusion ✓
d. Claude generating overly verbose multi-step responses

The correct answer is: Claude constructing plausible but flawed reasoning toward a harmful conclusion

Question 13

Flag question

Correct Mark 1.00 out of 1.00

Context window growth in a long agentic session is caused by:

a. Accumulating history and tool results with each iteration ✓
b. Tool definitions expanding with each call
c. The system prompt growing automatically
d. The API appending session metadata each turn

The correct answer is: Accumulating history and tool results with each iteration

Question 14

Flag question

Correct Mark 1.00 out of 1.00

An eval programme showing consistently high scores with persistently low user satisfaction indicates:

a. The eval metrics do not measure what users actually value ✓
b. The model tier is insufficient for the user base
c. The evaluation runs too infrequently currently
d. Users have unrealistic quality expectations

The correct answer is: The eval metrics do not measure what users actually value

Question 15

Flag question

Correct Mark 1.00 out of 1.00

Claude Sonnet is the recommended starting model because:

a. It has the largest context window available
b. It supports all beta features by default
c. It always outperforms Haiku on every task
d. It balances capability and cost for most tasks ✓

The correct answer is: It balances capability and cost for most tasks

Question 16

Flag question

Correct Mark 1.00 out of 1.00

Extended thinking content blocks returned by Claude have type:

a. reasoning
b. scratchpad
c. internal
d. thinking ✓

The correct answer is: thinking

Question 17

Flag question

Correct Mark 1.00 out of 1.00

Testing tool use integrations must cover:

a. Tests for error handling, invalid inputs, and server unavailability ✓
b. Only the most common tool call pattern
c. Only latency benchmarks for tool execution
d. Only valid inputs across happy-path scenarios

The correct answer is: Tests for error handling, invalid inputs, and server unavailability

Question 18

Flag question

Correct Mark 1.00 out of 1.00

For a customer-facing product, the API key must be:

a. Kept server-side, never exposed to clients ✓
b. Stored in localStorage for frontend access
c. Embedded in the mobile app for direct use
d. Distributed to trusted customers on request

The correct answer is: Kept server-side, never exposed to clients

Question 19

Flag question

Correct Mark 1.00 out of 1.00

MCP's open protocol standard benefits developers by:

a. Allowing any developer to build MCP-compatible servers and clients ✓
b. Requiring all MCP servers to be open-source
c. Mandating servers expose all capabilities
d. Restricting MCP client development to Anthropic

The correct answer is: Allowing any developer to build MCP-compatible servers and clients

Question 20

Flag question

Incorrect Mark 0.00 out of 1.00

Zero-shot prompting provides Claude with:

a. A system prompt with no user message
b. Instructions and multiple worked examples ✗
c. Only the task context and nothing else
d. Task instructions only with no examples

The correct answer is: Task instructions only with no examples

Question 21

Flag question

Correct Mark 1.00 out of 1.00

The recommended Python SDK method for a synchronous Claude API call is:

a. client.chat.complete()
b. client.generate()
c. client.messages.create() ✓
d. client.send_message()

The correct answer is: client.messages.create()

Question 22

Flag question

Correct Mark 1.00 out of 1.00

The most effective way to stop Claude from opening with "Certainly!" is:

a. Reducing the temperature parameter
b. Instructing Claude to be more direct always
c. Naming the exact phrase to prohibit explicitly ✓
d. Switching to a more capable model version

The correct answer is: Naming the exact phrase to prohibit explicitly

Question 23

Flag question

Correct Mark 1.00 out of 1.00

Demographic bias testing for a hiring tool tests whether:

a. Identical profiles with different demographic signals receive consistent assessments ✓
b. The system prompt references demographic groups
c. Response latency varies across user types
d. Output length is equal for all candidates

The correct answer is: Identical profiles with different demographic signals receive consistent assessments

Question 24

Flag question

Correct Mark 1.00 out of 1.00

Computer use tools in the API allow Claude to:

a. Access real-time internet data directly
b. Navigate GUIs by observing screenshots and taking actions ✓
c. Control hardware peripherals directly
d. Execute code in the user's local terminal

The correct answer is: Navigate GUIs by observing screenshots and taking actions

Question 25

Flag question

Correct Mark 1.00 out of 1.00

Self-consistency prompting selects the answer by:

a. Choosing the most frequent answer across multiple independent outputs ✓
b. Presenting all outputs to users for selection
c. Using the longest response as most complete
d. Averaging numerical values across all runs

The correct answer is: Choosing the most frequent answer across multiple independent outputs

Question 26

Flag question

Correct Mark 1.00 out of 1.00

Audience information in a prompt helps Claude calibrate:

a. API routing and model selection logic
b. Retrieval query construction in RAG
c. Token budget allocation per response
d. Vocabulary depth, tone, and detail level ✓

The correct answer is: Vocabulary depth, tone, and detail level

Question 27

Flag question

Correct Mark 1.00 out of 1.00

The /v1/messages/count_tokens endpoint is used to:

a. Monitor token usage across all deployments
b. Count tokens before submitting a full request ✓
c. Calculate billing estimates for the month
d. Validate system prompt compliance limits

The correct answer is: Count tokens before submitting a full request

Question 28

Flag question

Correct Mark 1.00 out of 1.00

The usage object in an API response reports:

a. Cache performance statistics only
b. Billing amount for the request made
c. API latency metrics for the call
d. input_tokens and output_tokens counts ✓

The correct answer is: input_tokens and output_tokens counts

Question 29

Flag question

Correct Mark 1.00 out of 1.00

Rate limiting for MCP tool calls in production should be implemented:

a. Via Claude's tool_choice configuration
b. Only inside each individual tool function
c. At the MCP server level to protect backend services ✓
d. Only at the Anthropic API layer

The correct answer is: At the MCP server level to protect backend services

Question 30

Flag question

Correct Mark 1.00 out of 1.00

For a complex codebase, the most effective context strategy for Claude Code is:

a. Asking Claude Code to search documentation online
b. Requiring developers to pre-select all files manually
c. CLAUDE.md for persistent context plus targeted file reading per task ✓
d. Loading the entire codebase on every session

The correct answer is: CLAUDE.md for persistent context plus targeted file reading per task

Question 31

Flag question

Correct Mark 1.00 out of 1.00

For a multi-turn conversation approaching the context limit, the correct strategy is:

a. Summarise older turns into a compact memory block ✓
b. Ask the user to start a new session
c. Delete the oldest messages abruptly
d. Move old messages to the system prompt

The correct answer is: Summarise older turns into a compact memory block

Question 32

Flag question

Correct Mark 1.00 out of 1.00

After Claude Code modifies multiple repository files, the developer should:

a. Commit immediately to preserve the changes
b. Review all changes before committing to version control ✓
c. Accept changes since Claude Code ran tests
d. Revert all changes and redo them manually

The correct answer is: Review all changes before committing to version control

Question 33

Flag question

Correct Mark 1.00 out of 1.00

cache_creation_input_tokens in the usage object represents:

a. Tokens paid to create a new cache entry
b. Tokens saved by prompt compression
c. Tokens retrieved from RAG at query time
d. Tokens served from an existing cache hit

The correct answer is: Tokens paid to create a new cache entry

Question 34

Flag question

Correct Mark 1.00 out of 1.00

For a deployment processing 200,000 short documents overnight with no real-time need:

a. Use the Message Batches API for async processing
b. Use streaming with asynchronous collection
c. Use parallel synchronous calls with retry
d. Use synchronous calls with high concurrency

The correct answer is: Use the Message Batches API for async processing

Question 35

Flag question

Correct Mark 1.00 out of 1.00

For a high-concurrency deployment receiving traffic spikes, the best architectural control is:

a. Switching to a faster model at peak times
b. A queue that absorbs spikes and smooths request flow to the API
c. Caching all possible responses in advance
d. Disabling streaming during high-load periods

The correct answer is: A queue that absorbs spikes and smooths request flow to the API

Question 36

Flag question

Correct Mark 1.00 out of 1.00

The Anthropic API Workbench is most useful for:

a. Monitoring production API usage in real time
b. Generating API keys for new team members
c. Testing prompts interactively without writing integration code
d. Running automated regression test suites

The correct answer is: Testing prompts interactively without writing integration code

Question 37

Flag question

Incorrect Mark 0.00 out of 1.00

In a multi-agent architecture, the orchestrator must treat sub-agent outputs as:

a. Authoritative since sub-agents share training
b. Cached responses from prior sessions
c. Pre-verified results ready for direct action
d. Untrusted content that may have been prompt- injected

The correct answer is: Untrusted content that may have been prompt- injected

Question 38

Flag question

Correct Mark 1.00 out of 1.00

The EU AI Act classifies an AI hiring tool as:

a. Prohibited under automated hiring rules
b. High risk requiring conformity assessment before deployment
c. Minimal risk with no specific obligations
d. Limited risk requiring only transparency notices

The correct answer is: High risk requiring conformity assessment before deployment

Question 39

Flag question

Incorrect Mark 0.00 out of 1.00

When Claude returns multiple tool_use blocks in one response, the developer must:

a. Execute all tools and return all results in one user message
b. Execute them sequentially one per API call ⊗
c. Ask Claude which tool to prioritise first
d. Execute only the first tool and discard the rest

The correct answer is: Execute all tools and return all results in one user message

Question 40

Flag question

Correct Mark 1.00 out of 1.00

Red-team testing for a Claude integration uses:

a. Inputs matching the highest-volume production category
b. Latency benchmarks under peak load conditions
c. Adversarial inputs designed to find safety and quality boundary failures ✓
d. Inputs from external domain experts only

The correct answer is: Adversarial inputs designed to find safety and quality boundary failures

Question 41

Flag question

Correct Mark 1.00 out of 1.00

Claude Code is installed via:

a. A standalone desktop application download
b. A browser extension for supported IDEs
c. npm install -g @anthropic-ai/claude-code ✓
d. The Anthropic web dashboard settings page

The correct answer is: npm install -g @anthropic-ai/claude-code

Question 42

Flag question

Correct Mark 1.00 out of 1.00

Human-in-the-loop gates are most critical before:

a. Making read-only calls to external APIs
b. Adding new turns to conversation history
c. Executing irreversible high-stakes external actions ✓
d. Generating text responses to users

The correct answer is: Executing irreversible high-stakes external actions

Question 43

Flag question

Incorrect Mark 0.00 out of 1.00

The anthropic-version header prevents:

a. Rate limiting from activating unexpectedly
b. Unauthorised access to beta features
c. Silent API contract changes breaking the integration
d. Model version changes between requests ✗

The correct answer is: Silent API contract changes breaking the integration

Question 44

Flag question

Correct Mark 1.00 out of 1.00

System prompts must be stored in version control because:

a. It enables audit history, rollback, and regression testing before deployment ✓
b. Version control compresses prompt tokens automatically
c. The API requires version-controlled prompts for production access
d. Anthropic audits all production system prompts annually

The correct answer is: It enables audit history, rollback, and regression testing before deployment

Question 45

Flag question

Incorrect Mark 0.00 out of 1.00

MCP resources are:

- a. Reusable prompt templates for common tasks ⊗
- b. Authentication tokens for external services
- c. Data sources Claude can read via a URI
- d. Callable actions that perform operations

The correct answer is: Data sources Claude can read via a URI

Question 46

Flag question

Correct Mark 1.00 out of 1.00

"Metatransparency" in Anthropic's framework means:

- a. Anthropic publishes its framework for what operators can instruct Claude to do ✓
- b. Operators must publish all system prompts
- c. Claude always reveals its system prompt on request
- d. Users receive a detailed audit of all interactions

The correct answer is: Anthropic publishes its framework for what operators can instruct Claude to do

Question 47

Flag question

Correct Mark 1.00 out of 1.00

Graceful degradation in a Claude integration means:

- a. Logging failures silently for later review
- b. Providing a useful fallback when Claude is unavailable ✓
- c. Reducing output length during high load
- d. Switching automatically to a smaller model

The correct answer is: Providing a useful fallback when Claude is unavailable

Question 48

Flag question

Correct Mark 1.00 out of 1.00

Prompt sensitivity testing submits:

a. Paraphrase variants of the same intent to detect output brittleness
b. Safety-probing inputs to check boundary adherence
c. Prompts at different temperature settings to compare stability
d. Only incorrectly formatted inputs to test resilience

The correct answer is: Paraphrase variants of the same intent to detect output brittleness

Question 49

Flag question

Correct Mark 1.00 out of 1.00

An operator instruction is illegitimate when it directs Claude to:

a. Deceive or actively harm the users it serves
b. Give Claude a specific persona name
c. Restrict Claude's topic scope
d. Require Claude to use formal language

The correct answer is: Deceive or actively harm the users it serves

Question 50

Flag question

Correct Mark 1.00 out of 1.00

A team engineering acceptable use policy for Claude Code must define:

a. How often API keys must be rotated
b. The maximum session length per developer
c. Which operations require mandatory human review before execution
d. Which Claude model tier is approved always

The correct answer is: Which operations require mandatory human review before execution

Question 51

Flag question

Correct Mark 1.00 out of 1.00

Claude Code's headless mode enables:

a. Non-interactive automated pipeline execution
b. Running without a display connected
c. Disabling all confirmation prompts always
d. Operating without internet connectivity

The correct answer is: Non-interactive automated pipeline execution

Question 52

Flag question

Incorrect Mark 0.00 out of 1.00

Valid values for the role field in the messages array are:

a. input and output
b. human and model
c. system and user
d. user and assistant

The correct answer is: user and assistant

Question 53

Flag question

Correct Mark 1.00 out of 1.00

Regression testing for a Claude integration runs:

a. Only on the highest-volume production queries
b. Adversarial inputs designed to find safety gaps
c. A fixed eval set after changes to confirm no passing cases now fail
d. Only on new user inputs not in the original eval set

The correct answer is: A fixed eval set after changes to confirm no passing cases now fail

Question 54

Flag question

Incorrect Mark 0.00 out of 1.00

When agentic tasks involve external document processing, retrieved content should be:

a. Automatically summarised before Claude reads
b. — wait, checking key. Q27:A per key.
c. Trusted fully if from a reputable source
d. Rejected unless from an approved allowlist ⊗

The correct answer is: Trusted fully if from a reputable source

Question 55

Flag question

Correct Mark 1.00 out of 1.00

A false negative in safety evaluation means:

a. Harmful content passed through when it should have been blocked ✓
b. The eval dataset contains incorrect ground truth labels
c. The system incorrectly refused a legitimate request
d. The system prompt contains an unseen safety gap

The correct answer is: Harmful content passed through when it should have been blocked

Question 56

Flag question

Correct Mark 1.00 out of 1.00

Input validation for a Claude integration must occur:

a. Before the input reaches Claude's context ✓
b. Only when the user is unauthenticated
c. Only for inputs exceeding 1,000 tokens
d. Only after Claude generates a response

The correct answer is: Before the input reaches Claude's context

Question 57

Flag question

Correct Mark 1.00 out of 1.00

Model routing by query complexity optimises:

a. Both cost and quality simultaneously ✓
b. Context window utilisation per request
c. API authentication overhead per call
d. Cache hit rates across all requests

The correct answer is: Both cost and quality simultaneously

Question 58

Flag question

Incorrect Mark 0.00 out of 1.00

When a tool returns a very large response, the developer should:

a. Ask Claude to summarise it in a separate call first ✗
b. Include the full response verbatim always
c. Filter or summarise before returning it to Claude
d. Truncate to exactly 1,000 tokens always

The correct answer is: Filter or summarise before returning it to Claude

Question 59

Flag question

Incorrect Mark 0.00 out of 1.00

The sampling capability in MCP allows servers to:

a. Request Claude completions on the server's behalf
b. Sample random subsets of available tools ✗
c. Measure token usage per tool call
d. Test tool definitions before deployment

The correct answer is: Request Claude completions on the server's behalf

Question 60

Flag question

Incorrect Mark 0.00 out of 1.00

The "reflection" agentic pattern involves:

a. Users providing feedback after each response
b. Claude critiquing its own output and producing an improved version
c. A second model validating the first model's output
d. Claude reviewing prior conversation history

The correct answer is: Claude critiquing its own output and producing an improved version

Question 61

Flag question

Correct Mark 1.00 out of 1.00

When a Claude integration fails in production but works locally, the most likely cause is:

a. The context window is larger in production
b. The temperature setting differs locally
c. The local machine has faster CPU
d. Different model version or system prompt is active in production

The correct answer is: Different model version or system prompt is active in production

Question 62

Flag question

Correct Mark 1.00 out of 1.00

Prompt chaining improves reliability by:

a. Running multiple prompts in parallel always
b. Reducing the number of API calls needed
c. Allowing users to review each step output
d. Decomposing complex tasks into focused sequential steps with validated handoffs

The correct answer is: Decomposing complex tasks into focused sequential steps with validated handoffs

Question 63

Flag question

Correct Mark 1.00 out of 1.00

When latency is the critical production requirement, the first model to test is:

a. Extended for larger context support
b. Opus for maximum accuracy first
c. Sonnet as the balanced starting point
d. Haiku for its speed advantages ✓

The correct answer is: Haiku for its speed advantages

Question 64

Flag question

Correct Mark 1.00 out of 1.00

The "plan-and-execute" agentic pattern improves reliability by:

a. Generating multiple plans and selecting the best one
b. Using a separate model for planning versus execution
c. Running the plan past a human before any execution
d. Separating task planning from step-by-step execution in distinct calls ✓

The correct answer is: Separating task planning from step-by-step execution in distinct calls

Question 65

Flag question

Correct Mark 1.00 out of 1.00

Chain-of-thought prompting instructs Claude to:

a. Call multiple tools in a defined sequence
b. Chain outputs from sequential API calls
c. Generate multiple candidate answers first
d. Reason step-by-step before stating the final answer ✓

The correct answer is: Reason step-by-step before stating the final answer

Previous activity
Domain 7 - Troubleshooting and Optimization

Jump to...