

Associate Data Practitioner Sample Questions

The Data Practitioner sample questions will familiarize you with the format of exam questions and example content that may be covered on the exam.

The sample questions do not represent the range of topics or level of difficulty of questions presented on the exam. Performance on the sample questions should not be used to predict your Data Practitioner exam result.

Registration

First Name *

Gopikrishnan

Last Name *

TS

Primary Email *

gopikrishnants31102k2@gmail.com

Recovery Email

.....

Organization (Employer or School) *

Accenture

.....

Organization email (an email associated with your current organization)

.....

Country *

Dropdown

India



Primary Relationship to Google *

Dropdown

Partner



Send me offers, updates and useful tips for getting the most out of Google Cloud training and certification products and services. *

Dropdown

No



Question 1

- ✓ You have a logistics and shipping analytics solution that uses BigQuery as its **data warehouse**. You want to enhance the solution with **real-time** location data from your truck fleet and hourly traffic data. You want to use the same **data ingestion process for both real-time and batch data**. What should you do?

- ☒ A. Use Dataflow. ✓
- ☐ B. Use Cloud Data Fusion.
- ☐ C. Use Storage Transfer Service.
- ☐ D. Use BigQuery Data Transfer Service.

Feedback

A is correct because Dataflow allows for developing real-time batch and stream data processing pipelines, with the same pipeline you can insert real-time streams as well as batch data.

B is not correct because Cloud Data Fusion requires two different pipelines, one for real-time and one for batch, though Data Fusion helps build and manage data pipelines using fully managed, code-free data integration with a graphical interface.

C is not correct because Storage Transfer Service transfers data between Cloud Storage services such as AWS S3 and Cloud Storage and is not designed for real-time ingestion.

D is not correct because BigQuery Data Transfer Service automates data movement into BigQuery on a scheduled, not real-time, basis.

<https://cloud.google.com/...>


<https://cloud.google.com/...>

Question 2

- ✓ You need to create a new data pipeline. You want a serverless solution that meets the following requirements:

- Data is streamed from Pub/Sub and is **processed in real-time**.
- Data is **transformed** before being stored.
- Data is stored in a location that will allow it to be **analyzed** with SQL using Looker.

Which Google Cloud services should you recommend for the pipeline?



- ☐ A. 1. BigQuery, 2. Analytics Hub
- ☒ B. 1. Dataflow, 2. BigQuery
- ☐ C. 1. Dataproc Serverless, 2. Bigtable
- ☐ D. 1. Cloud Composer, 2. Cloud SQL for MySQL



Feedback

A is not correct because we want to transform the data before storage. This option performs transformations after the data is stored in BigQuery.

B is correct because Dataflow is a no code/low code solution to processing data, and the processing happens before the data is loaded into BigQuery.

C is not correct because Dataproc Serverless processes in batches, and Bigtable doesn't support SQL which is required by Looker in this pipeline.

D is not correct because Cloud Composer is used to manage DAGs, and Cloud SQL for MySQL isn't optimized for analytics.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

Question 3

- ✓ You have a **Dataproc cluster** that performs batch processing on data stored in Cloud Storage. You need to schedule a daily Spark job to generate a report that will be emailed to stakeholders. You need a **fully-managed solution that is easy to implement and minimizes complexity**. What should you do?
- ☐ A. Use Cloud Composer to orchestrate the Spark job and email the report.
 - ☐ B. Use Cloud Run functions to trigger the Spark job and email the report.
 - ☒ C. Use Dataproc workflow templates to define and schedule the Spark job, and to email the report. ✓
 - ☐ D. Use Cloud Scheduler to trigger the Spark job, and use Cloud Run functions to email the report.

Feedback

A is not correct because while Cloud Composer can orchestrate tasks, it might be overkill for a simple daily Spark job on an existing Dataproc cluster. Cloud Composer is not fully-managed.

B is not correct because Cloud Run functions are better suited for event-driven tasks and have limitations on execution time that might not be suitable for long-running Spark jobs.

C is correct because Dataproc workflow templates are specifically designed for managing and scheduling jobs on Dataproc clusters. They provide built-in functionality for email notifications.

D is not correct because while Cloud Scheduler can trigger the job, using it in conjunction with Cloud Run functions adds unnecessary complexity compared to using Dataproc workflow templates' built-in email functionality.

<https://cloud.google.com/...>

Question 4

- ✓ Your retail company has an on-premises MySQL database that contains millions of rows of customer data. You need to **efficiently migrate the data**, in daily increments, into BigQuery for analysis. What should you do?
- ☐ A. Export the data from MySQL in CSV format and upload the CSV file to Cloud Storage. Use BigQuery Data Transfer Service to load the data into BigQuery.
 - ☐ B. Use Cloud Data Fusion to connect to the MySQL database, create a pipeline to extract the data, and load the data into BigQuery.
 - ☒ C. Use Database Migration Service to replicate the MySQL database to Cloud SQL for MySQL. Use BigQuery Data Transfer Service to load the data into BigQuery. ✓
 - ☐ D. Use Dataflow to connect to the MySQL database, extract the data, and load the data into BigQuery.

Feedback

A is not correct because it's not the most efficient solution for incremental daily loads of large datasets. Manually exporting and uploading CSV files can be time-consuming and error-prone.

B is not correct because Cloud Data Fusion is primarily designed for complex ETL tasks and might be overkill for this scenario.

C is correct because it is the most suitable solution as it leverages Database Migration Service to efficiently replicate the MySQL database to Cloud SQL, ensuring minimal downtime. BigQuery Data Transfer Service can then be used to automate the incremental loading of data from Cloud SQL into BigQuery.

D is not correct because Dataflow requires more custom coding and is less efficient than using Database Migration Service for initial replication and BigQuery Data Transfer Service for ongoing incremental loads.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

<https://cloud.google.com/...>

✓ You have a BigQuery dataset containing sales data. This data is actively queried for the first 6 months. After that, the data is not queried but needs to be retained for 3 years for compliance reasons. You need to implement a data management strategy that meets access and compliance requirements, while keeping cost and administrative overhead to a minimum. What should you do?

- ☐ A. Store all data in a single BigQuery table without partitioning or lifecycle policies.
- ☒ B. Partition a BigQuery table by month. After 6 months, export the data to Coldline storage. Implement a lifecycle policy to delete the data from Cloud Storage after 3 years. ✓
- ☐ C. Set up a scheduled query to export the data to Cloud Storage after 6 months. Write a stored procedure to delete the data from BigQuery after 3 years.
- ☐ D. Use BigQuery long-term storage for the entire dataset. Set up a Cloud Run function to delete the data from BigQuery after 3 years.

Feedback

A is not correct because it doesn't leverage BigQuery's lifecycle management capabilities to optimize costs.

B is correct because partitioning allows for efficient querying of recent data, and the lifecycle policy automates the transition to cheaper storage and eventual deletion, meeting both cost and compliance needs.

C is not correct because exporting data and writing stored procedures adds administrative overhead.

D is not correct because while long-term storage is less expensive than active storage, it is still more expensive than Coldline for infrequently accessed data. Plus this option requires additional setup (overhead) for data deletion.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

✓ Your company uses BigQuery to store your company data for internal users. The users are all members of the “Internal-Reporting” Google group. You need to allow these users to **SELECT and JOIN** the data within BigQuery, but **not UPDATE or DELETE** any of the data. You need to follow the principle of least privilege. What should you do?

- ☐ A. Grant the BigQuery Admin IAM role to the “Internal-Reporting” Google group.
- ☐ B. Grant the BigQuery Data Owner IAM role to the “Internal-Reporting” Google group.
- ☒ C. Grant the BigQuery Data **Viewer** and BigQuery **Job User** IAM roles to the “Internal-Reporting” Google group. ✓
- ☐ D. Grant the BigQuery Metadata Viewer IAM role to the “Internal-Reporting” Google group.

Feedback

A is not correct because the BigQuery Admin IAM role is over-provisioned (violating principle of least privilege) for only allowing the ability to read data within BigQuery.
B is not correct because the BigQuery Data Owner IAM role allows not only for the reading, but also for the deletion, updating, and sharing of the data.
C is correct, as it allows just for the execution of BigQuery based jobs, including queries, and only for reading of data within BigQuery. It does not allow for the editing or deletion of any data.
D is not correct because users cannot run queries with this access.

<https://cloud.google.com/...>

✓ You are migrating data from a legacy on-premises MySQL database to Google Cloud. The database contains various tables with different data types and sizes, including large tables with millions of rows and transactional data. You need to **migrate this data while maintaining data integrity, and minimizing downtime and cost**. What should you do?

- ☒ A. Use Database Migration Service to replicate the MySQL database to a Cloud SQL for MySQL instance. ✓
- ☐ B. Use Cloud Data Fusion to migrate the MySQL database to MySQL on Compute Engine.
- ☐ C. Export the MySQL database to CSV files, transfer the files to Cloud Storage by using Storage Transfer Service, and load the files into a Cloud SQL for MySQL instance.
- ☐ D. Set up a Cloud Composer environment to orchestrate a custom data pipeline. Use a Python script to extract data from the MySQL database and load it to MySQL on Compute Engine.

Feedback

A is correct because this approach combines the strengths of Database Migration Service and BigQuery. Database Migration Service enables efficient replication of the MySQL database to Cloud SQL for MySQL with minimal downtime, ensuring continuous operation during the migration. This approach provides a balance between minimal downtime and efficient data transfer.

B is not correct because Cloud Data Fusion is not suitable for a migration scenario requiring minimal downtime and would not be the most cost-effective solution.

C is not correct because exporting the entire database to CSV files can be inefficient and time-consuming, especially for large tables. It also introduces potential data integrity issues during the export and import process.

D is not correct because setting up a Cloud Composer environment and a custom data pipeline adds significant complexity and overhead for this migration scenario.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

✓ Your organization has decided to move their on-premises Apache Spark-based workload to Google Cloud. You want to be able to manage the code without needing to provision and manage your own cluster. What should you do?

- ☒ A. Migrate the Spark jobs to Dataproc Serverless. ✓
- ☐ B. Migrate the Spark jobs to Dataproc on Google Kubernetes Engine.
- ☐ C. Migrate the Spark jobs to Dataproc on Compute Engine.
- ☐ D. Configure a Google Kubernetes Engine cluster with Spark operators, and deploy the Spark jobs.

Feedback

A is correct because Dataproc Serverless lets you deploy Spark applications without managing infrastructure.

B is not correct because Dataproc on Google Kubernetes Engine (GKE) requires you to manage your own GKE clusters.

C is not correct because Dataproc on Compute Engine requires you to manage your compute environment .

D is not correct because the GKE cluster requires you to manage both compute and GKE clusters.

<https://cloud.google.com/...>

✗ You are working on a data pipeline that will **validate and clean incoming data** before loading it into BigQuery for real-time analysis. You want to ensure that the **data validation and cleaning is performed efficiently and can handle high volumes of data**. What should you do?

- ☐ A. Use Dataflow to create a streaming pipeline that includes validation and transformation steps.
- ☐ B. Load the raw data into BigQuery using Cloud Storage as a staging area, and use SQL queries in BigQuery to validate and clean the data.
- ☐ C. Write custom scripts in Python to validate and clean the data outside of Google Cloud. Load the cleaned data into BigQuery.
- ☒ D. Use Cloud Run functions to trigger data validation and cleaning routines when new data arrives in Cloud Storage. ✗

Correct answer

- ☒ A. Use Dataflow to create a streaming pipeline that includes validation and transformation steps.

Feedback

A is correct because Dataflow allows for large-scale, real-time processing of data streams with built-in capabilities for validation and cleaning. Dataflow excels at handling high volumes of streaming data with minimal latency. Its distributed processing capabilities ensure efficient validation and cleaning, while the BigQueryIO connector enables seamless integration with BigQuery for real-time analysis.

B is not correct because validation should occur before the data is stored to ensure data integrity. Loading raw data into BigQuery before validation and cleaning introduces a delay in the pipeline. While BigQuery can handle large datasets and SQL queries are powerful, this approach is not optimized for continuous, high-velocity data streams.

C is not correct as external scripts may complicate the process and reduce efficiency compared to using Dataflow. It also lacks the seamless integration and scalability offered by managed services like Dataflow.

D is not correct, while Cloud Run functions can be effective for event-driven processing, they are not ideal for this scenario. Cloud Run functions are better suited for smaller-scale tasks. For large volumes of real-time data, Dataflow offers superior performance and scalability.

<https://cloud.google.com/...>

<https://beam.apache.org/r...>

<https://cloud.google.com/...>

Question 10

- ✓ You created a new dashboard in Looker that displays live, up-to-the-minute sales data for your company. You want to provide **access to your new dashboard** to your organization's sales team so they can see **live data whenever needed**. What should you do?
- ☐ A. Use the Add Filter modal and add a user attribute filter to the dashboard based on the users' email attribute.
 - ☒ B. Use the Manage Access panel for the folder where the dashboard is saved, and grant view access to the sales user group. ✓
 - ☐ C. Use the Dashboard actions menu at the top of the dashboard to get the link to the dashboard. Send the dashboard's URL to the sales team members.
 - ☐ D. Use the Dashboard actions menu at the top of the dashboard, and schedule a daily dashboard delivery to the sales team members' email addresses.

Feedback

A is not correct because filters affect the data returned from queries, but doesn't provide access to the content itself.

B is correct because users need access to the folder where content is saved in order to view it.

C is not correct because sharing a link to the dashboard doesn't provide access to the dashboard. A user that accesses the URL without access to the dashboard will see an error message.

D is not correct because while users would see data for the dashboard once per day, the question specifically asks for sharing access so they can see live data when they need it. Scheduling daily would show what the data in the dashboard looked like at the specific time the schedule was sent.

<https://support.google.co...>

<https://cloud.google.com/...>

<https://cloud.google.com/...>

<https://cloud.google.com/...>

Question 11

- ✓ You work for an ecommerce company that stores product image files in a Cloud Storage bucket. New product images are frequently uploaded as either new images or to replace existing images. The older images are accessed less frequently over time. You need to configure a cost-effective lifecycle policy for managing image files based on age and versions that meets the following requirements:

Age: Older images are automatically moved to a lower-cost storage class while maintaining the ability to access images quickly.

Versions: The latest three versions of each image are stored in the bucket with versioning enabled.

What should you do?

- ☒ A. Age: Transition objects to Nearline storage after 30 days, and to Coldline storage after 90 days. Versions: Store only the latest three versions and delete noncurrent versions after 30 days. ✓
- ☐ B. Age: Transition objects to Nearline storage after 30 days, and to Archive storage after 90 days. Versions: Store only the latest three versions and delete noncurrent versions after 30 days.
- ☐ C. Age: Transition objects to Nearline storage after 30 days, and to Coldline storage after 90 days. Versions: Delete noncurrent versions after 30 days.
- ☐ D. Age: Transition objects to Coldline storage after 30 days, and to Archive storage after 90 days. Versions: Delete noncurrent versions after 30 days.

Feedback

A is correct because it satisfies both requirements. It transitions older objects to lower-cost storage classes (Nearline and Coldline) based on age, ensuring cost-efficiency while maintaining access to recent images. It also enforces a three version limit for both live and noncurrent versions, optimizing storage usage.

B is incorrect because transitioning to Archive storage after 90 days would make the images too expensive to retrieve if needed, which doesn't meet the cost optimization requirement.

C is incorrect because it doesn't address the requirement for keeping three versions for objects.

D is incorrect because transitioning to Archive storage after 30 days would make the images too expensive to retrieve if needed, and it doesn't manage noncurrent versions.

Question 12

✗ Your BigQuery dataset contains a table with sensitive columns with personally identifiable information (PII). You want to ensure that authorized users and applications can query the data while maintaining strong encryption at the column level and the ability to decrypt with specific keys. What should you do?

- ☐ A. Apply a dynamic data masking rule to the PII columns.
- ☒ B. Use AEAD encryption functions with a Cloud Key Management System (Cloud KMS) keyset.
- ☐ C. Apply encryption with a customer-managed encryption key (CMEK) on the dataset that stores the table.
- ☒ D. Create BigQuery views that replace the PII columns with NULL values, and only grant read access to the BigQuery views. ✗

Correct answer

- ☒ B. Use AEAD encryption functions with a Cloud Key Management System (Cloud KMS) keyset.

Feedback

A is not correct because dynamic data masking does not encrypt columns, merely obfuscates their returned values.

B is correct because AEAD allows the user to apply encryption at the column or row level and provides a secondary level of encryption.

C is not correct because CMEK encryption for the dataset would not specifically encrypt the PII columns.

D is not correct because this option would not add an additional layer of encryption.

Question 13

✗ You are a Looker analyst. You need to add a new field to your Looker report that generates SQL that will run against your company's database. You do not have the Develop permission. What should you do?

- ☒ A. Create a new field in the LookML layer, refresh your report, and select your new field from the field picker. ✗
- ☐ B. Create a table calculation from the field picker in Looker, and add it to your report.
- ☐ C. Create a custom field from the field picker in Looker, and add it to your report.
- ☐ D. Create a calculated field using the Add a field option in Looker Studio, and add it to your report.

Correct answer

- ☒ C. Create a custom field from the field picker in Looker, and add it to your report.

Feedback

A is not correct because adding a new field in the LookML requires the Develop permission.

B is not correct because table calculations do not generate SQL; they only process logic on query result sets.

C is correct because custom fields generate SQL without needing to modify the LookML layer.

D is not correct because that is how to add a field to Looker Studio, not Looker.

<https://support.google.co...>

<https://cloud.google.com/...>

- ✓ You work for a home insurance company. You are frequently asked to create and save risk reports with charts for specific areas using a publicly available storm event dataset. You want to be able to quickly create and re-run risk reports when new data becomes available. What should you do?
- ☐ A. Copy the storm event dataset into your BigQuery project. Use BigQuery Studio to query and visualize the data in Looker Studio.
 - ☐ B. Export the storm event dataset as a CSV file. Import the file to Google Sheets, and use cell data in the worksheets to create charts.
 - ☒ C. Reference and query the storm event dataset using SQL in a Colab Enterprise notebook. Display the table results and document with Markdown, and use Matplotlib to create charts. ✓
 - ☐ D. Reference and query the storm event dataset using SQL in BigQuery Studio. Export the results to Google Sheets, and use cell data in the worksheets to create charts.

Feedback

A is not correct because we are copying public data into our dataset rather than merely referencing it. This increases our operational cost for no obvious benefit. While we can query from BigQuery Studio manually and also visualize in Looker Studio, none of that is easily repeatable for future data without manually re-executing.

B is not correct because the export of the CSV data is a manual step and is more complex to re-run in the future. It requires skills to export and re-import. In addition, complexity increases because now we don't have the power of SQL at our disposal. It requires skills to export and re-import. This conflicts with the desire to "quickly re-run a report as new data becomes available".

C is correct because the public datasets are continually updated and allow us to re-run immediately and see the latest information. Colab Enterprise notebooks are saved and hence we have the ability to re-run as needed. The markdown produces annotations and report text and Matplotlib provides charts and graphs.

D is not correct because the export of the SQL query results into Google Sheets is a manual step and is more complex to re-run in the future. It requires skills to export and re-import. This conflicts with the desire to "quickly re-run a report as new data becomes available".

<https://cloud.google.com/...>

<https://www.geeksforgeek...>

<https://cloud.google.com/...>

Question 15

- ✓ You have a large dataset stored in Cloud Storage in CSV format. The data **needs to be transformed and cleaned**, which involves transforming the data, filtering records, and converting data types. After the data is prepared, **it needs to be loaded into a Cloud SQL database**. You want to use the **Google-recommended approach**. What should you do?
- ☐ A. Use Dataproc with Spark to perform the transformations and load the transformed data into Cloud SQL.
 - ☒ B. Use Dataflow to process the data in batches, perform the transformations, and write the transformed data to Cloud SQL. ✓
 - ☐ C. Use Cloud Data Fusion to design the data pipeline, perform the transformations using built-in connectors, and write the transformed data to Cloud SQL.
 - ☐ D. Use a Cloud Run function to trigger a transformation job and load the transformed data into Cloud SQL.

Feedback

A is not correct because while Dataproc can perform transformations, it is typically less efficient for straightforward data preparation tasks compared to Dataflow, which is optimized for such scenarios.

B is correct because Dataflow is specifically designed for batch processing of large datasets and integrates seamlessly with both Cloud Storage and Cloud SQL, making it the most efficient option for this task.

C is not correct because while Cloud Data Fusion is capable of visually designing data pipelines, it may be more complex and resource-intensive for a simple transformation, where Dataflow would be more appropriate.

D is not correct because the Cloud Run function can handle data transformations, but they are primarily designed for event-driven tasks and are not the most efficient choice for processing large datasets in batch mode.

<https://cloud.google.com/...>

Question 16

✗ You manage a data pipeline that streams sales transactions from Pub/Sub to BigQuery. Consumers of the data are reporting that new data has stopped arriving in the table despite messages still being successfully published. You need to diagnose the issue. What should you do?

- ☒ A. Open the Logs Explorer and look for errors filtered by Pub/Sub or BigQuery. ✗
- ☐ B. Create a dead-letter topic associated with the subscription. Look for errors in the messages arriving in that topic.
- ☐ C. Change the subscription definition's retry policy from "Retry after exponential backoff delay" to "Retry immediately".
- ☐ D. Open the Pub/Sub subscription in the Google Cloud console, navigate to the Messages tab, and look for errors in the messages waiting to be delivered.

Correct answer

- ☒ B. Create a dead-letter topic associated with the subscription. Look for errors in the messages arriving in that topic.

Feedback

A is not correct because errors encountered as part of the delivery of Pub/Sub messages to BigQuery do not cause Cloud Logging records to be written. There would be nothing of value to be found in Cloud Logging.

B is correct because when a message can not be delivered through Pub/Sub its final disposition is a dead-letter topic. The message written to the dead-letter topic contains a description of the reason the message failed to be delivered. Specifically, the error is recorded in the attribute called CloudPubSubDeadLetterSourceDeliveryErrorMessage.

C is not correct because changing the retry policy changes how quickly (and when) retry attempts are made but does not give any indication on the underlying error or its resolution.

D is not correct because messages arriving on a topic that has a BigQuery subscription aren't modified/changed/updated if they fail to be delivered. Think of a message on a topic as being immutable. When the original message is placed in the topic, that message will not be updated to include any error indications.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

- ✓ You need to design a data pipeline that **ingests data** from CSV, Avro, and Parquet files into Cloud Storage. The data includes raw user input. You need to **remove all malicious SQL injections** before storing the data in BigQuery. Which data manipulation methodology should you choose?

- ☐ A. EL
- ☐ B. ELT
- ☒ C. ETL
- ☐ D. ETLT



Feedback

A is not correct because extracting and loading the data will not remove malicious code from your database.

B is not correct because the stem states that data needs to be transformed before loading.

C is correct because this is the only answer that transforms the data before loading it into BigQuery.

D is not correct because the situation doesn't require a second round of transformations after loading.

<https://cloud.google.com/...>

✓ You need to create a new encryption key to secure sensitive data in a BigQuery table. Local regulations require that your organization maintains control of the key and its location, protection level, and rotation cycle. The key needs to be stored outside of Google Cloud. What should you do?

- ☐ A. Create a Google-managed encryption key (GMEK) in a global key ring.
- ☐ B. Create a Google-managed encryption key (GMEK) in a regional key ring.
- ☐ C. Create a customer-supplied encryption key (CSEK).
- ☒ D. Create a customer-managed encryption key (CMEK). ✓

Feedback

A is not correct because GMEK involves a Google-managed key.

B is not correct because GMEK involves a Google-managed key.

C is not correct because CSEK keys are stored within GCP and the question explicitly says we need to provide it from an external source. CSEKs can also only be used to encrypt GCS and GCE.

D is correct because CMEKs can be externally provided and allow for the level of control required in the question.

<https://cloud.google.com/...>

✗ You have a large CSV file that is uploaded to Cloud Storage every morning. The file contains several nested and repeated fields. You need to load the data into BigQuery daily for analysis, while preserving the nested structure. You want to use the most efficient approach. What should you do?

- ☐ A. Use the bq load command with the --autodetect flag to load the CSV file.
- ☐ B. Use BigQuery Data Transfer Service to schedule daily data loads of the CSV file.
- ☐ C. Use Cloud Data Fusion to schedule the pipeline to parse the CSV file and create the BigQuery schema.
- ☒ D. Use Cloud Composer to schedule a Python script to flatten the data and execute the bq load command to load the CSV file into BigQuery. ✗

Correct answer

- ☒ C. Use Cloud Data Fusion to schedule the pipeline to parse the CSV file and create the BigQuery schema.

Feedback

A is not correct because while the bq load command is a common way to load data into BigQuery, the --autodetect flag may not accurately handle nested and repeated data structures in CSV files. It might flatten the data or create a schema that doesn't reflect the original nested structure, leading to data loss or difficulties in querying the data later.

B is not correct because BigQuery Data Transfer Service is not designed for complex data transformations like handling nested and repeated data in CSV files. It's more suitable for straightforward data loads from supported sources. A preloaded schema is required for loading CSV files.

C is correct because Cloud Data Fusion can handle complex data transformations, including parsing nested and repeated data in CSV files, and automatically generate a corresponding BigQuery schema, making it the most efficient way to load this type of data into BigQuery while preserving its structure.

D is not correct because this approach can be time-consuming and error-prone, especially for large datasets with complex nested structures. It also loses the original nested structure, which might be crucial for analysis and querying the data later.

<https://cloud.google.com/...>

<https://cloud.google.com/...>

<https://codelabs.develope...>

<https://cloud.google.com/...>

✓ Your company currently uses an on-premises network file system (NFS) and is migrating data to Google Cloud. You want to be able to control how much bandwidth is used by the data migration while capturing detailed reporting on the migration status. What should you do?

- ☐ A. Use gcloud storage commands.
- ☐ B. Use Cloud Storage FUSE.
- ☐ C. Use a Transfer Appliance.
- ☒ D. Use Storage Transfer Service.



Feedback

A is not correct because gcloud commands won't give you control over bandwidth or detailed reporting.

B is not correct because FUSE lets you interact with data similar to a NFS, but is not actually used for data migration.

C is not correct because this is an offline transfer and bandwidth control does not apply.

D is correct because Storage Transfer Service allows you to decide the bandwidth used and provides detailed reporting.

 <https://cloud.google.com/...>

 <https://cloud.google.com/...>

This form was created inside Google.com. [Privacy & Terms](#)

Google Forms

