



- Expert Verified, Online, **Free**.

⚙ Custom View Settings

You are migrating workloads to the cloud. The goal of the migration is to serve customers worldwide as quickly as possible According to local regulations, certain data is required to be stored in a specific geographic area, and it can be served worldwide. You need to design the architecture and deployment for your workloads.

What should you do?

- A. Select a public cloud provider that is only active in the required geographic area
- B. Select a private cloud provider that globally replicates data storage for fast data access
- C. Select a public cloud provider that guarantees data location in the required geographic area
- D. Select a private cloud provider that is only active in the required geographic area

Correct Answer: C

Community vote distribution

C (100%)

- NU_1234

Highly Voted

2 years, 2 months ago

Selected Answer: C

Key words - as quickly as possible -> should be public cloud with low latency and also guaranteeing data storage in the required location
upvoted 13 times
- ahsangh

Highly Voted

2 years, 5 months ago

Selected Answer: C

C. Select a public cloud provider that guarantees data location in the required geographic area
upvoted 11 times
- Certslab

Most Recent

1 day, 10 hours ago

Selected Answer: C

Certslab.com (no Discussion C ;)
upvoted 1 times
- YoExam

1 month, 1 week ago

If someone has the pdf with all the questions please send me to mansiau1019@gmail.com Thanks in advance :)
upvoted 1 times
- ES_Capgemini_landD

2 months, 1 week ago

Selected Answer: C

Simple, es la C
upvoted 1 times
- Surek

5 months, 3 weeks ago

Answer is c
upvoted 1 times
- Lakshvenkat

6 months ago

Answer is C
upvoted 1 times
- Guru4Cloud

6 months, 1 week ago

Selected Answer: C

C. Select a public cloud provider that guarantees data location in the required geographic area
upvoted 1 times
- MR_max

6 months, 3 weeks ago

Anyone having access to the full question set for this certification. I'm also in the learning phase for the same certification. Could you pls help me with the dumps available by sending them in mail aravindhanmax777@gmail.com. It would be very very helpful. Many thanks in advance.
upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

  **__rajan__** 8 months, 2 weeks ago

Selected Answer: C

C is correct as they require their data to be location specific.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

Select a public cloud provider that guarantees data location in the required geographic area

upvoted 1 times

  **dienm** 9 months, 2 weeks ago

Selected Answer: C

can be accessed worldwide: public cloud

data regulation to be stored on specific geo location: guarantees data location

upvoted 3 times

  **Pauloafreitas** 9 months, 2 weeks ago

If someone has the pdf with all the questions please send me to paulofsax@gmail.com. Thanks in advance!

upvoted 1 times

  **firdaus27** 9 months, 4 weeks ago

Anyone having access to the full question set for this certification. I'm also in the learning phase for the same certification. Could you pls help me with the dumps available by sending them in mail hfas.accemp@gmail.com. It wuld be very very helpful. Many thanks in advance.

upvoted 3 times

  **prachisri** 10 months, 2 weeks ago

Selected Answer: C

Select a public cloud provider that guarantees data location in the required geo area

upvoted 1 times

  **evazaf** 1 year ago

Selected Answer: C

C is correct

upvoted 1 times

Your organization needs a large amount of extra computing power within the next two weeks.
After those two weeks, the need for the additional resources will end.
Which is the most cost-effective approach?

- A. Use a committed use discount to reserve a very powerful virtual machine
- B. Purchase one very powerful physical computer
- C. Start a very powerful virtual machine without using a committed use discount
- D. Purchase multiple physical computers and scale workload across them

Correct Answer: C

Community vote distribution

C (100%)

- Syd

Highly Voted

2 years, 5 months ago

C is correct because we don't know the compute use for 2 weeks and discount is available on 1-3 years not for 2 weeks.

<https://cloud.google.com/compute/docs/instances/signing-up-committed-use-discounts>

upvoted 30 times
- ucsdmiami2020

1 year, 6 months ago

Agreed C. "When you purchase a committed use contract, you purchase Compute Engine resources—such as vCPUs, memory, GPUs, local SSDs and sole-tenant nodes—at a discounted price in return for committing to paying for those resources for 1 year or 3 years."

upvoted 4 times
- bifef

Highly Voted

5 months, 1 week ago

I took the Cloud-Digital-Leader exam and scored 88% for which I am thankful to this platform as it had all the information needed.

<https://shorturl.at/qPQR3>

upvoted 6 times
- levepapa

Most Recent

4 months, 3 weeks ago

C is Right Answer!

Get Up-to-date: <https://www.pinterest.com/pin/937522847419094013>

upvoted 1 times
- Rajk1999

6 months ago

C is absolutely correct option.

upvoted 1 times
- Tunny

7 months, 1 week ago

Selected Answer: C

C is correct

upvoted 1 times
- Tunny

7 months, 1 week ago

Selected Answer: C

C is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: C

Committed Discount is only applicable for 1-3 years of commitment.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: C

C is correct because we don't know the compute use for 2 weeks and discount is available on 1-3 years not for 2 weeks.

upvoted 1 times

🗨️ 👤 **mdsarfraz69** 9 months ago

C is correct because we don't know the compute use for 2 weeks and discount is available on 1-3 years not for 2 weeks.
upvoted 2 times

🗨️ 👤 **dienm** 9 months, 2 weeks ago

Selected Answer: C

a large amount of extra computing power: start a very powerful virtual machine
After those two weeks, the need for the additional resources will end: without using a committed use discount
upvoted 1 times

🗨️ 👤 **firdaus27** 9 months, 4 weeks ago

Anyone having access to the full question set for this certification. I'm also in the learning phase for the same certification. Could you pls help me with the dumps available by sending them in mail hfas.accemp@gmail.com. It wuld be very very helpful. Many thanks in advance.
upvoted 1 times

🗨️ 👤 **prachisri** 10 months, 2 weeks ago

Selected Answer: C

C is correct
upvoted 1 times

🗨️ 👤 **cookieMr** 1 year ago

Selected Answer: C

C is correct
upvoted 1 times

🗨️ 👤 **Ritesh_Sahu** 1 year, 2 months ago

D. Purchase multiple physical computers and scale workload across them
upvoted 2 times

🗨️ 👤 **Sourav_Gmail** 1 year, 3 months ago

Selected Answer: C

Committed use discount is for 1 or 3 years , purchasing additional resource is not required for 2 weeks.
upvoted 3 times

🗨️ 👤 **star2anand** 1 year, 3 months ago

C. Start a very powerful virtual machine without using a committed use discount
upvoted 1 times

Your organization needs to plan its cloud infrastructure expenditures.
Which should your organization do?

- A. Review cloud resource costs frequently, because costs change often based on use
- B. Review cloud resource costs annually as part of planning your organization's overall budget
- C. If your organization uses only cloud resources, infrastructure costs are no longer part of your overall budget
- D. Involve fewer people in cloud resource planning than your organization did for on-premises resource planning

Correct Answer: A

Community vote distribution

A (79%)

C (19%)

-   **desertlotus1211**

Highly Voted 

 2 years, 4 months ago

Answer is A...

Don't be fooled by 'Cloud Infrastructure' as being only physical. the resources like Cloud VPN, Internconnects, GCE, etc are all 'cloud infrastructure'... resources and cloud infrastructure are used interchangeable...

upvoted 28 times
-   **kummimhk** 2 years ago

Totally agree

upvoted 2 times
-   **ahsangh**

Highly Voted 

 2 years, 5 months ago

A. Review cloud resource costs frequently, because costs change often based on use

upvoted 24 times
-   **Vedjha** 1 year, 10 months ago

correct

upvoted 2 times
-   **ucsdmiami2020** 1 year, 6 months ago

Agreed A. See the below article for an understanding of Google cloud platform infrastructure services, aka IaaS
<https://www.globalknowledge.com/us-en/resources/resource-library/articles/understanding-gcp-infrastructure-services/>

upvoted 5 times
-   **Nigel_123**

Most Recent 

 1 month ago

Selected Answer: A

Staying on top of cloud costs, which can scale up and down easily, is important to happen frequently

upvoted 1 times
-   **76ffe88** 1 month, 3 weeks ago

Selected Answer: A

As cloud expenditure depends on usage.

upvoted 1 times
-   **ES_Capgemini_landD** 2 months, 1 week ago

Selected Answer: A

'A' - Correcto.

upvoted 1 times
-   **PRADTOMB** 2 months, 3 weeks ago

Answer is A : its not about maintenance its about expenditure

upvoted 2 times
-   **chuckc66** 2 months, 2 weeks ago

Agreed, the question is about cloud infrastructure expenditures

upvoted 1 times
-   **azrap** 3 months, 2 weeks ago

Agree cost change frequently. Question is "needs to plan" so answer should be B as part of Budget

upvoted 1 times
-   **levepapa** 4 months, 3 weeks ago

C is Right Answer!

Get Up-to-date: <https://www.pinterest.com/pin/937522847419094013>
upvoted 1 times

  **PassForSure007** 5 months ago

Selected Answer: A

A is the correct answer. Cloud Infrastructure means cloud resources in general.
upvoted 1 times

  **bifef** 5 months, 1 week ago

I took the Cloud-Digital-Leader exam and scored 88% for which I am thankful to this platform as it had all the information needed.
<https://shorturl.at/qPQR3>
upvoted 4 times

  **Surek** 5 months, 3 weeks ago

Answer is A
upvoted 1 times

  **Lakshvenkat** 6 months ago

correct answer is A
upvoted 1 times

  **Tunny** 7 months, 1 week ago

Selected Answer: A

A is correct
upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

  **__rajan__** 8 months, 2 weeks ago

Selected Answer: A

Review cloud resource costs frequently, because costs change often based on use
upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: A

A. Review cloud resource costs frequently, because costs change often based on use
upvoted 1 times

  **firdaus27** 9 months, 4 weeks ago

Anyone having access to the full question set for this certification. I'm also in the learning phase for the same certification. Could you pls help me with the dumps available by sending them in mail hfas.accemp@gmail.com. It wuld be very very helpful. Many thanks in advance.
upvoted 1 times

  **vpad** 7 months, 2 weeks ago

Good Question.Did you get the dump? If yes, could you pls kindly send it to venupadarthi@gmail.com. I will very thankful to you for your help.
upvoted 1 times

  **Kgwedi** 7 months ago

Hi did you get the dump is so could you kindly send me emeraldgreencreatives@gmail.colm
upvoted 2 times

The operating systems of some of your organization's virtual machines may have a security vulnerability.
How can your organization most effectively identify all virtual machines that do not have the latest security update?

- A. View the Security Command Center to identify virtual machines running vulnerable disk images
- B. View the Compliance Reports Manager to identify and download a recent PCI audit
- C. View the Security Command Center to identify virtual machines started more than 2 weeks ago
- D. View the Compliance Reports Manager to identify and download a recent SOC 1 audit

Correct Answer: A

Community vote distribution

A (100%)

-  **levepapa** 4 months, 3 weeks ago

Selected Answer: A

C is Right Answer!

Get Up-to-date: <https://www.pinterest.com/pin/937522847419094013>
upvoted 2 times
-  **bluetata** 3 months, 1 week ago

you chose A but you typed C is the right answer?
upvoted 3 times
-  **Surek** 5 months, 3 weeks ago

Answer is A
upvoted 1 times
-  **Tunny** 7 months, 1 week ago

Selected Answer: A

A is correct
upvoted 1 times
-  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
-  **__rajan__** 8 months, 2 weeks ago

Selected Answer: A

The most effective way for your organization to identify virtual machines that do not have the latest security update is to view the Security Command Center. By utilizing the Security Command Center, you can identify virtual machines that are running vulnerable disk images. This allows you to pinpoint the specific virtual machines that require the latest security update, enabling you to take appropriate action to address the security vulnerability.
upvoted 3 times
-  **mdsarfraz69** 9 months ago

Selected Answer: A

i think its A
upvoted 1 times
-  **Walveee** 9 months, 1 week ago

i think its A
upvoted 1 times
-  **firdaus27** 9 months, 4 weeks ago

Anyone having access to the full question set for this certification. I'm also in the learning phase for the same certification. Could you pls help me with the dumps available by sending them in mail hfas.accemp@gmail.com. It wuld be very very helpful. Many thanks in advance.
upvoted 4 times
-  **cookieMr** 1 year ago

Selected Answer: A

Security Command Center is a security management and threat detection service provided by Google Cloud Platform (GCP). It offers centralized visibility into the security posture of your GCP resources, helping you identify, understand, and mitigate potential security risks.

upvoted 4 times

  **Ritesh_Sahu** 1 year, 2 months ago

D. View the Compliance Reports Manager to identify and download a recent SOC 1 audit
upvoted 1 times

  **star2anand** 1 year, 3 months ago

A. View the Security Command Center to identify virtual machines running vulnerable disk images
upvoted 1 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

The COMPUTE_IMAGE_SCANNER detector identifies vulnerabilities related to Google Cloud image configurations.
upvoted 2 times

  **gviolant** 1 year, 5 months ago

The correct answer is A
upvoted 2 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Per Google docs, the Security Command Center offers the ability to uncover Compute instance vulnerability findings
<https://cloud.google.com/security-command-center/docs/concepts-vulnerabilities-findings#compute-instance-findings>
upvoted 3 times

  **ronieto** 1 year, 7 months ago

Selected Answer: A

Security Command Center Vulnerabilities tab in the Google Cloud console
upvoted 2 times

  **grzeev** 1 year, 8 months ago

A.
You can filter findings by detector name and finding type using the Security Command Center Vulnerabilities tab in the Google Cloud console
upvoted 1 times

  **GCP72** 1 year, 9 months ago

Selected Answer: A

The correct answer is A
upvoted 1 times

You are currently managing workloads running on Windows Server for which your company owns the licenses. Your workloads are only needed during working hours, which allows you to shut down the instances during the weekend. Your Windows Server licenses are up for renewal in a month, and you want to optimize your license cost.

What should you do?

- A. Renew your licenses for an additional period of 3 years. Renew your licenses for an additional period of 3 years. Negotiate a cost reduction with your current hosting provider wherein infrastructure cost is reduced when workloads are not in use
- B. Renew your licenses for an additional period of 2 years. Negotiate a cost reduction by committing to an automatic renewal of the licenses at the end of the 2 year period
- C. Migrate the workloads to Compute Engine with a bring-your-own-license (BYOL) model
- D. Migrate the workloads to Compute Engine with a pay-as-you-go (PAYG) model

Correct Answer: C

Community vote distribution

D (80%)

C (20%)

- fpreli

Highly Voted

2 years, 5 months ago

D for me, it says "optimize your license cost" and that license will expire in one month, so I would go with a Pay-As-You-Go model

upvoted 28 times
- kitubha

Highly Voted

2 years, 5 months ago

D is correct

upvoted 14 times
- penelop

Most Recent

2 weeks, 5 days ago

You are currently managing workloads running on Windows Server for which your company owns the licenses. Your workloads are only needed during working hours, which allows you to shut down the instances during the weekend. Your Windows Server licenses are up for renewal in a month, and you want to optimize your license cost.

D is the correct answer.

The question clearly states that we are running Windows Server with our own licenses (BYOL). These instances are already on GCP, we want to optimize licensing cost. Keeping BYOL (C) won't affect our infrastructure, D will.

With PAYG, we are paying an extra fee while the VM is running, meaning if you turn off the Vm, you are not paying the license. This will optimize our costs, as we are not expending a big portion of money to pay the license, but only a fraction for usage.

upvoted 2 times
- 76ffe88

1 month, 3 weeks ago

Selected Answer: D

As the license is already expiring so PAYG model will be more efficient here.

upvoted 2 times
- ikaro460

1 month, 3 weeks ago

C is correct. Migrating your workloads to Compute Engine with a BYOL model would allow you to utilize your existing Windows Server licenses while also benefiting from the flexibility and potentially lower costs associated with cloud computing. This way, you can optimize your license cost: by using your existing licenses efficiently and take advantage of Compute Engine's pricing structure. Additionally, you can explore further cost-saving measures such as scheduling instances to shut down during weekends when they're not needed.

upvoted 2 times
- smedrano

3 months, 1 week ago

Selected Answer: D

I think it's D, because BYOL is only available with Bring Your Own Image for Windows Server (<https://cloud.google.com/blog/products/compute/compute-engine-licensing-explained>), which is only available to run on dedicated hardware, which I think would end resulting much more expensive.

Pay-As-You-Go allows you to optimize your "License" costs, because you only pay for the license as long as the instance is in use. With BYOL you're not optimizing anything as you still have to pay the full license for each instance.

upvoted 1 times
- BhooChaal

3 months, 2 weeks ago

Selected Answer: C

See the problem with D is that, any VM cost does not cover the cost of Licensing. If you want to have a VM running RHEL, then you need to get RHEL licenses on your own. Google can help you with that, but License of S/W usage is extra, above and beyond N/W, Disk, Compute and Memory

upvoted 1 times

🗨️ 👤 **PassForSure007** 5 months ago

Selected Answer: D

D is the correct answer the license will expire in 1 month.
upvoted 1 times

🗨️ 👤 **Thesalesguy** 5 months, 2 weeks ago

C, BYOL model is correct:
The cons of D PAYG are:

Higher cost: The PAYG model typically includes a premium on top of the per-second resource usage cost to cover the included license fees. This could be more expensive than BYOL, where you leverage your existing licenses.
Recurring cost: You continuously pay for the license fees as part of the PAYG model, even when your workloads are not running. This contrasts with BYOL, where you only pay for the new license once, and only pay for resource usage during active hours.
upvoted 1 times

🗨️ 👤 **Surek** 5 months, 3 weeks ago

Answer is C
upvoted 1 times

🗨️ 👤 **houdini0** 5 months, 3 weeks ago

It's D, because with D you do not need to bring your licenses (Correct me if I'm wrong), and the scenario says there is no need to have it up and running during the weekends.
upvoted 1 times

🗨️ 👤 **Lakshvenkat** 6 months ago

Option D for me
upvoted 1 times

🗨️ 👤 **Surek** 6 months ago

D is the correct answer. But GPT say answer C ?
upvoted 1 times

🗨️ 👤 **sivakarthick16** 6 months, 1 week ago

Selected Answer: D

By distributing the virtual machines across different zones within the same region, your organization can ensure redundancy in case of any zone-specific failures or disruptions. This setup provides high availability and fault tolerance.

Additionally, locating the virtual machines within the same region allows for fast communication between them, as the network latency within a region is typically lower compared to communication across regions. This ensures that the parts of the application can communicate with each other in less than 10 milliseconds, meeting your organization's requirement for fast communication.

Choosing multiple regions (Option C) or multiple zones per region (Option D) may provide additional redundancy but could potentially increase network latency, which may not meet the requirement for extremely fast communication.
upvoted 1 times

🗨️ 👤 **su321** 6 months, 1 week ago

Selected Answer: D

Pay as you go
upvoted 1 times

🗨️ 👤 **ay_india** 6 months, 3 weeks ago

Selected Answer: D

In the given scenario, the goal is to optimize license costs for workloads running on Windows Server, which are only needed during working hours. Considering this, the most cost-effective option would be to choose a pay-as-you-go (PAYG) model in Google Cloud's Compute Engine.
upvoted 1 times

🗨️ 👤 **preetham141** 6 months, 3 weeks ago

Selected Answer: C

C is Correct
upvoted 1 times

Your organization runs a distributed application in the Compute Engine virtual machines. Your organization needs redundancy, but it also needs extremely fast communication (less than 10 milliseconds) between the parts of the application in different virtual machines.

Where should your organization locate this virtual machines?

- A. In a single zone within a single region
- B. In different zones within a single region
- C. In multiple regions, using one zone per region
- D. In multiple regions, using multiple zones per region

Correct Answer: D

Community vote distribution

B (76%)

D (17%)

4%

- SemiDeus

Highly Voted

2 years, 5 months ago

I chose B since it guarantees a latency of ~10 milliseconds is maintained, well also giving the redundancy of multi-zone. If it was D, we could not guarantee 10 milliseconds between the different regions.

upvoted 46 times
- prekair0

Highly Voted

2 years, 5 months ago

Selected Answer: B

Multi zone is also redundant within the region and it provides the lowest latency.

upvoted 12 times
- Surek

Most Recent

5 months, 3 weeks ago

Answer is B

upvoted 1 times
- Surek

6 months ago

Answer B

upvoted 1 times
- xDSKick

6 months, 2 weeks ago

Selected Answer: B

Options A, C, and D involve spreading the virtual machines across different configurations, such as single zone, multiple regions with one zone per region, or multiple zones per region. While these options may offer redundancy to some extent, they might not guarantee the desired extremely fast communication due to potential increased latency between different regions or zones.

Therefore B

upvoted 3 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- mahiravan

7 months, 3 weeks ago

Selected Answer: D

<https://cloud.google.com/compute/docs/regions-zones>

Depending on how widely you want to distribute your resources, create instances across multiple zones in multiple regions for redundancy.

upvoted 1 times
- custodio

8 months, 1 week ago

Answer correct is B accordind to 'Decreased network latency

To decrease network latency, you might want to choose a region or zone that is close to your point of service. For example, if you mostly have customers on the East Coast of the US, then you might want to choose a primary region and zone that is close to that area and a backup region and zone that is also close by.' Section at <https://cloud.google.com/compute/docs/regions-zones/>

upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: B

In different zones within a single region
upvoted 1 times

  **louiselz** 8 months, 3 weeks ago

AnswerD accordind to 'Identifying a region or zone' Section at <https://cloud.google.com/compute/docs/regions-zones/>
upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct one
upvoted 1 times

  **adyanrado** 9 months, 1 week ago

Selected Answer: D

D is correct because guarantee slow latency between regions and zones in case of any zone failed, yet have slow latency between regions
upvoted 2 times

  **jaya7** 10 months ago

I chose B as the latency is less between zones in a single region
upvoted 1 times

  **tperelle** 10 months, 2 weeks ago

Selected Answer: B

I chose B. Looking to the Google Cloud performance dashboard, latency between multiple zone of a same region is ~1ms, communication between multiple regions is ~100ms. So multi-zones provides redundancy and latency < 10ms
upvoted 1 times

  **AKSArch** 11 months ago

Selected Answer: B

10ms latency is the keyword here. The multi-region configuration would definitely be affected by latency. I would go for option B
upvoted 1 times

  **MinaGohari** 11 months ago

Selected Answer: B

B is correct
upvoted 1 times

An organization decides to migrate their on-premises environment to the cloud. They need to determine which resource components still need to be assigned ownership.

Which two functions does a public cloud provider own? (Choose two.)

- A. Hardware maintenance
- B. Infrastructure architecture
- C. Infrastructure deployment automation
- D. Hardware capacity management
- E. Fixing application security issues

Correct Answer: AD

Community vote distribution

AD (95%)

5%

- azrap

3 months, 2 weeks ago

Hardware Maintenance and Infrastructure Deployment A and C. Hardware capacity is cover under Infrastructure Deployment

upvoted 1 times
- Surek

5 months, 3 weeks ago

Answer is AD

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: AD

AD is correct

upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: AD

AD is correct

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: AD

AD is correct

upvoted 1 times
- prachisri

10 months, 2 weeks ago

CD answer

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: AD

The Cloud Provider is responsible for the security "of the cloud" infrastructure. This includes the underlying hardware, data centers, networking, an the core services. It takes care of managing and protecting these foundational components to ensure the availability, reliability, and security of the cloud infrastructure.

The customer is responsible for the security "in the cloud" or the security of their applications and data running on cloud services.

upvoted 4 times
- zambonini

1 year ago

Selected Answer: AD

Solution A and D

upvoted 1 times
- Rafa_M

1 year, 2 months ago

Selected Answer: AD

Solution A and D

upvoted 1 times
- Sp4nner

1 year, 2 months ago

Selected Answer: AD

Keyword "HARDWARE". no client will touch hardware hence answers are A&D
upvoted 1 times

  **Ritesh_Sahu** 1 year, 2 months ago
A & D is the right answer
upvoted 2 times

  **star2anand** 1 year, 3 months ago
A. Hardware Maintenance and D. Hardware capacity management
upvoted 1 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: AD

Hardware maintenance & capacity mgmt is managed by cloud provider.
upvoted 1 times

  **Jackey0117** 1 year, 4 months ago

Selected Answer: AB

A. Hardware maintenance
B. Infrastructure architecture
upvoted 1 times

  **KC_go_reply** 1 year, 3 months ago
Infrastructure would mean the cloud provider would manage your VMs etc, so I disagree
upvoted 3 times

  **AnshMan** 1 year, 5 months ago
This is there in Exam
upvoted 2 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: AD

Per Google docs, at a minimum when considering IaaS, the cloud provider is responsible for Hardware, Boot, Hardened Kernel and IPS, Storage and
etc. <https://cloud.google.com/architecture/framework/security/shared-responsibility-shared-fate>
upvoted 3 times

  **Pou1ze** 1 year, 6 months ago

Selected Answer: AD

Hardware is closely related to Cloud provider (Share Responsibility Model)
upvoted 1 times

You are a program manager within a Software as a Service (SaaS) company that offers rendering software for animation studios. Your team needs the ability to allow scenes to be scheduled at will and to be interrupted at any time to restart later. Any individual scene rendering takes less than 12 hours to complete, and there is no service-level agreement (SLA) for the completion time for all scenes. Results will be stored in a global Cloud Storage bucket. The compute resources are not bound to any single geographical location. This software needs to run on Google Cloud in a cost-optimized way.

What should you do?

- A. Deploy the application on Compute Engine using preemptible instances
- B. Develop the application so it can run in an unmanaged instance group
- C. Create a reservation for the minimum number of Compute Engine instances you will use
- D. Start more instances with fewer virtual centralized processing units (vCPUs) instead of fewer instances with more vCPUs

Correct Answer: A

Community vote distribution

A (100%)

- shamg**

Highly Voted

2 years, 5 months ago

A for me -
Can be interrupted at any time to restart later, software needs to run on Google Cloud in a cost-optimized way
upvoted 17 times
- cookieMr**

Highly Voted

1 year ago

Selected Answer: A

Preemptible VMs are cost-effective instances that can be used for short-lived and non-critical workloads. Since the rendering of individual scenes takes less than 12 hours and can be interrupted and restarted, you can leverage preemptible VMs to perform the rendering tasks. Preemptible VM are available at a significantly lower cost compared to regular instances but can be preempted by GCP with a 30-second notice. By designing your rendering workflow to handle interruptions gracefully, such as checkpointing and resuming, you can take advantage of the cost savings provided by preemptible VMs.
upvoted 10 times

Surek

Most Recent

5 months, 3 weeks ago

Answer is A
upvoted 1 times

Lakshvenkat

6 months ago

Selected Answer: A

A is the correct answer
upvoted 1 times

GCPHER01234

6 months, 1 week ago

A is correct
upvoted 1 times

chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

mdsarfraz69

9 months ago

Selected Answer: A

A is correct
upvoted 1 times

zambonini

1 year ago

Selected Answer: A

Can be interrupted at any time to restart later.
upvoted 1 times

Rafa_M

1 year, 2 months ago

Selected Answer: A

Answer A:
"Preemptible VM instances are available at much lower price—a 60-91% discount—compared to the price of standard VMs. However, Compute

Engine might stop (preempt) these instances if it needs to reclaim the compute capacity for allocation to other VMs"

upvoted 1 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: A

can be interrupted = preemptible

upvoted 2 times

  **MHx1** 1 year, 3 months ago

Selected Answer: A

A - Instances with deep savings but at the cost of being interrupted

upvoted 1 times

  **star2anand** 1 year, 3 months ago

A. Deploy the application on Compute Engine using preemptible instances

upvoted 1 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

Preemptible VMs are highly affordable, short-lived compute instances suitable for batch jobs and fault-tolerant workloads. Preemptible VMs offer the same machine types and options as regular compute instances and last for up to 24 hours.

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Per Google docs, "Preemptible VM instances are available at much lower price—a 60-91% discount—compared to the price of standard VMs. However, Compute Engine might stop (preempt) these instances if it needs to reclaim the compute capacity for allocation to other VMs"

https://cloud.google.com/compute/docs/instances/preemptible#what_is_a_preemptible_instance

upvoted 1 times

  **Pou1ze** 1 year, 6 months ago

Selected Answer: A

12h ok, is compatible with the restart time palnned for preemptible VMs (24H), and also unpredictable stop of VMs

upvoted 1 times

  **allannvictor** 1 year, 9 months ago

Selected Answer: A

A in my opiniion

upvoted 2 times

  **seths** 1 year, 10 months ago

Selected Answer: A

scheduled at will and to be interrupted at any time to restart later and no SLA makes Answer A

upvoted 1 times

Your manager wants to restrict communication of all virtual machines with internet access; with resources in another network; or with a resource outside Compute Engine. It is expected that different teams will create new folders and projects in the near future. How would you restrict all virtual machines from having an external IP address?

- A. Define an organization policy at the root organization node to restrict virtual machine instances from having an external IP address
- B. Define an organization policy on all existing folders to define a constraint to restrict virtual machine instances from having an external IP address
- C. Define an organization policy on all existing projects to restrict virtual machine instances from having an external IP address
- D. Communicate with the different teams and agree that each time a virtual machine is created, it must be configured without an external IP address

Correct Answer: A

Reference:

<https://cloud.google.com/resource-manager/docs/organization-policy/overview>

All quickstarts

Quickstart using organizations

Quickstart enforcing constraints

How-to guides

All how-to guides

▶ Resource Manager

▶ Organization policy service

▶ Early Access Center

▶ Resource settings service

Concepts

All concepts

▶ Resource Manager

▶ Organization policy service

Overview

Community vote distribution

A (100%)

Introduction to the Organization Policy Service

Send feedback

The Organization Policy Service gives you centralized and programmatic control over your organization's cloud resources. As the [organization policy administrator](#), you will be able to configure constraints across your entire [resource hierarchy](#).

Benefits

- Centralize control to configure restrictions on how your organization's resources can be used.
- Define and establish guardrails for your development teams to stay within compliance boundaries.

- Surek

5 months, 3 weeks ago

Answer is A

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

Answer is A.

upvoted 1 times
- tsk9921

11 months, 2 weeks ago

A Root level will address the issue

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: A

To ensure consistency and compliance across your organization, enforce Organization Policies that restrict VMs from having external IP addresses. Organization Policies provide centralized control and can prevent users from creating VM instances with external IP addresses across different

projects and folders within your organization.

upvoted 3 times

  **Rafa_M** 1 year, 2 months ago

Selected Answer: A

Answer is A.

upvoted 1 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: A

Root is above all

upvoted 1 times

  **star2anand** 1 year, 3 months ago

A. Define an organization policy at the root organization node to restrict virtual machine instances from having an external IP address

upvoted 1 times

  **jres** 1 year, 3 months ago

A very true. this solve the problem

upvoted 2 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Refer to Google Docs explanation of Resource hierarchy, <https://cloud.google.com/resource-manager/docs/cloud-platform-resource-hierarchy>

upvoted 3 times

  **Pou1ze** 1 year, 6 months ago

Selected Answer: A

the sentence that say "in the future other team create new project...." address the answer at the top level of hierarchy

upvoted 2 times

  **grzeev** 1 year, 8 months ago

A. projects and resources inherits from root!

upvoted 2 times

  **erwin_t82** 1 year, 10 months ago

Selected Answer: A

The Organization Policy Service gives you centralized and programmatic control over your organization's cloud resources. As the organization policy administrator, you will be able to configure constraints across your entire resource hierarchy.

upvoted 2 times

  **Shashank_R** 2 years ago

A is the Answer ; root organization node is the key here to restrict virtual machine instances from having an external IP address.

upvoted 3 times

  **haroldbenites** 2 years ago

Go for A

upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: A

Because is expected that different teams will create new folders and projects in the near future. Then the restriction must be ver folders.

upvoted 3 times

Your multinational organization has servers running mission-critical workloads on its premises around the world. You want to be able to manage these workloads consistently and centrally, and you want to stop managing infrastructure.

What should your organization do?

- A. Migrate the workloads to a public cloud
- B. Migrate the workloads to a central office building
- C. Migrate the workloads to multiple local co-location facilities
- D. Migrate the workloads to multiple local private clouds

Correct Answer: A

Community vote distribution

A (100%)

- Syd**

Highly Voted

2 years, 5 months ago

A is correct. Key is manage centrally.
manage these workloads consistently and centrally, and you want to stop managing infrastructure.
upvoted 14 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: A

A iscorrect.
upvoted 1 times
- __rajan__**

8 months, 2 weeks ago

Selected Answer: A

A is correct.
upvoted 1 times
- pratsymb**

8 months, 3 weeks ago

Selected Answer: A

A is correct because workload can be move to public cloud otherwise next service which have PII which must be move to private cloud for security purposes..
upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: A

Answer is A.
upvoted 1 times
- Yuk897**

9 months, 1 week ago

Why is Option C not convenient as an answer as the company can have many local Private clouds for mission- critical workload...please suggest
upvoted 1 times
- RK_MANOJ**

8 months, 2 weeks ago

it involves maintaining and managing infrastructure across multiple locations, which does not align with the goal of avoiding infrastructure management.
upvoted 1 times
- GamudaDavid**

5 months, 2 weeks ago

You need to know that this test is Google Cloud Digital Leader. I will prefer choose all the answer related to Google Cloud instead of choosing infrastructure that need to manage by user.
upvoted 1 times
- tsk9921**

11 months, 2 weeks ago

A - Cloud is the only option for central mgmt
upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: A

To centrally manage mission-critical workloads running on servers located around the world and eliminate the need for infrastructure management, your organization should consider adopting a cloud-based solution such as Google Cloud Platform (GCP).

upvoted 1 times

  **Rafa_M** 1 year, 2 months ago

Selected Answer: A

Answer is A. Migrate to a public cloud

upvoted 1 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: A

That's why we study THE CLOUD

upvoted 1 times

  **star2anand** 1 year, 3 months ago

A. Migrate the workloads to a public cloud

upvoted 1 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

correct A

upvoted 1 times

  **Yesmin** 1 year, 6 months ago

Selected Answer: A

Es correcta

upvoted 1 times

  **SurenderGupta** 1 year, 9 months ago

Selected Answer: A

Only cloud offer the solution being asked

upvoted 1 times

  **Sunkhat** 1 year, 9 months ago

I thought B as correct answer. Since this refers to mission-critical workloads, I was thinking of centralized ops. However this doesn't say confidential workloads, so in that case A may be the correct answer

upvoted 1 times

  **DataTrainee** 1 year, 10 months ago

Selected Answer: A

A is correct

upvoted 2 times

Your organization stores highly sensitive data on-premises that cannot be sent over the public internet. The data must be processed both on-premises and in the cloud.
What should your organization do?

- A. Configure Identity-Aware Proxy (IAP) in your Google Cloud VPC network
- B. Create a Cloud VPN tunnel between Google Cloud and your data center
- C. Order a Partner Interconnect connection with your network provider
- D. Enable Private Google Access in your Google Cloud VPC network

Correct Answer: C

After the service provider provisions your connection, you can start passing traffic between your networks by using the service provider's network.

Reference:

<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/partner-overview>

Partner Interconnect overview

Send feedback

Partner Interconnect provides connectivity between your on-premises network and your Virtual Private Cloud (VPC) network through a supported service provider. A Partner Interconnect connection is useful if your data center is in a physical location that can't reach a Dedicated Interconnect colocation facility, or your data needs don't warrant an entire 10-Gbps connection.

Before you use Partner Interconnect

Ensure that you meet the following requirements:

- Be familiar with [Cloud Interconnect terminology](#).
- Work with a [supported service provider](#) to establish connectivity between their network and your on-premises network. Supported service providers offer Layer 2 connectivity, Layer 3

Community vote distribution

C (71%)	B (25%)	4%
---------	---------	----

- egvidal** Highly Voted 2 years, 4 months ago

Selected Answer: C

Not over public internet, therefore it's C

upvoted 20 times
- haroldbenites** Highly Voted 2 years ago

Go for B.

Always we have to review the statement carefully.

It says:

"highly sensitive data".

Due to Cloud VPN, which encrypt data is the better option.

Partner interconnect:

<https://cloud.google.com/network-connectivity/docs/how-to/choose-product>

The connection between your network and Google's network is not encrypted. If you require additional data security, use application-level encryption or your own VPN. Currently, you can't use Cloud VPN with Partner Interconnect, but you can use your own VPN solution.

upvoted 16 times
- SreekarM** 1 year, 7 months ago

But it is sent over public internet

upvoted 11 times
- Karpagam_S** Most Recent 1 month ago

C is the right answer

upvoted 1 times
- cloudinvader** 4 months, 1 week ago

Selected Answer: B

To obtain secure communication between on-premises and the cloud environment.

upvoted 1 times

🗨️ **sivakarthick16** 6 months, 1 week ago

Selected Answer: B

Option B allows for secure communication between your on-premises network and the Google Cloud environment, ensuring that sensitive data remains protected during transit.

upvoted 1 times

🗨️ **Guru4Cloud** 6 months, 2 weeks ago

Selected Answer: D

The correct answer is D. Enable Private Google Access in your Google Cloud VPC network.

C - This option provides a dedicated connection with high bandwidth and low latency, suitable for large data transfers. However, it may be overkill if the primary concern is securely transferring sensitive data.

upvoted 2 times

🗨️ **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

🗨️ **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

🗨️ **__rajan__** 8 months, 2 weeks ago

Selected Answer: B

B. Create a Cloud VPN tunnel between Google Cloud and your data center

To securely process highly sensitive data both on-premises and in the cloud without sending it over the public internet, your organization should create a Cloud VPN tunnel between Google Cloud and your data center. This establishes a secure, encrypted connection between your on-premise network and your Google Cloud Virtual Private Cloud (VPC). The VPN tunnel ensures that data remains protected during transit and allows for secure communication between your on-premises environment and the cloud.

upvoted 1 times

🗨️ **mdsarfraz69** 9 months ago

Selected Answer: C

C is the Answer

upvoted 1 times

🗨️ **hireshgupt** 11 months ago

Selected Answer: B

The answer to your question is B. Create a Cloud VPN tunnel between Google Cloud and your data center.

A Cloud VPN tunnel is a secure, private connection between your on-premises network and your Google Cloud Platform (GCP) network. It uses the internet to connect your networks, but it encrypts all traffic over the connection, ensuring that your data remains private and secure.

upvoted 1 times

🗨️ **Mwafrika** 11 months, 1 week ago

Selected Answer: C

Option B (Create a Cloud VPN tunnel) might seem like a plausible choice for connecting on-premises to Google Cloud, but it is essential to note that VPNs usually use the public internet to create encrypted connections, which contradicts the requirement that the data cannot be sent over the public internet. Partner Interconnect, on the other hand, establishes a private connection, making it a more suitable option in this case.

upvoted 5 times

🗨️ **cookieMr** 1 year ago

Selected Answer: C

Set up a dedicated private network connection between your on-premises infrastructure and Google Cloud Platform (GCP). This can be achieved using Dedicated Interconnect or Partner Interconnect, which provide direct, private, and secure connections between your on-premises data center and GCP's network.

upvoted 4 times

🗨️ **sjain93** 1 year, 1 month ago

C - A Partner Interconnect connection allows you to establish a dedicated, private, and secure network connection between your on-premises data center and Google Cloud. This connection operates outside of the public internet, providing a direct link with high bandwidth and low latency.

upvoted 5 times

🗨️ **deCalle** 1 year, 2 months ago

Selected Answer: C

read the question carfully. it says it can't be sent over the public internet. did the same mistake and initially voted for b

upvoted 3 times

🗨️ **Sp4nner** 1 year, 2 months ago

Selected Answer: C

VPN is private in nature however requirement is data should not be going out to the public/internet hence answer is C even though there is an involve third party Interconnect is still going directly between your equipment and google cloud server. Think of it simply as a LAN cable difference is it is situated on a data center

upvoted 3 times

  **Atmatap** 1 year, 2 months ago

Selected Answer: B

Vote for B. Because C is dedicated connection also uses public internet. It should use VPN.

upvoted 1 times

Your company's development team is building an application that will be deployed on Cloud Run. You are designing a CI/CD pipeline so that any new version of the application can be deployed in the fewest number of steps possible using the CI/CD pipeline you are designing. You need to select a storage location for the images of the application after the CI part of your pipeline has built them. What should you do?

- A. Create a Compute Engine image containing the application
- B. Store the images in Container Registry
- C. Store the images in Cloud Storage
- D. Create a Compute Engine disk containing the application

Correct Answer: B

Reference:

<https://cloud.google.com/container-registry/docs/pushing-and-pulling>

Container Registry > Container Registry documentation > Guides

Was this helpful?

Pushing and pulling images

Send feedback

Pushing (uploading) and pulling (downloading) images are two of the most common Container Registry tasks. This document focuses on pushing and pulling images with Docker.

Note: Starting with GKE node version 1.19, the default node image for Linux nodes is the Container-Optimized OS with Containerd (**cos_containerd**) variant instead of the Container-Optimized OS with Docker (**cos**) variant.

For troubleshooting and logging using GKE with Containerd, use [crictl](#). For open source Kubernetes, refer to the [containerd documentation](#).

Community vote distribution

B (87%)

13%

- cloudinvader** 4 months, 1 week ago

Selected Answer: B

Images stored in the Container Registry

upvoted 3 times
- sivakarthick16** 5 months, 1 week ago

Selected Answer: B

Container Registry is a managed Docker container image storage service provided by Google Cloud. It is specifically designed for storing and managing container images, making it the ideal choice for storing the application images in your CI/CD pipeline. By using Container Registry, you can easily manage and deploy the application images to Cloud Run in the most efficient and streamlined manner.

upvoted 1 times
- Surek** 5 months, 3 weeks ago

Answer is B

upvoted 1 times
- Surek** 6 months ago

B is the answer

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months, 2 weeks ago

Selected Answer: B

B is correct but currently Google is promoting Artifact Registry over Container Registry. You may get that in the main exam MCQ.

upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct one

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: B

Select Google Container Registry (GCR) as the storage location for your application images. GCR allows you to securely store and manage container images in a private registry. It is tightly integrated with other Google Cloud services, including Cloud Run, making it a convenient choice.

upvoted 4 times

  **Rafa_M** 1 year, 1 month ago

Selected Answer: B

B. Store the images in Container Registry

upvoted 1 times

  **star2anand** 1 year, 3 months ago

B. Store the images in Container Registry

upvoted 2 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: B

Container registry for storing images.

upvoted 1 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: B

You need to select a storage location for the images of the application after the CI part of your pipeline has built them. I believe this is referring to the image like docker image. Hence container registry.

upvoted 1 times

  **IvyGirl** 1 year, 5 months ago

It's B <https://cloud.google.com/run/docs/deploying>

upvoted 2 times

  **MultiCloudIronMan** 1 year, 6 months ago

Selected Answer: C

The question is about application images like jpg and png not profile images of an environment like docker or vm.

upvoted 4 times

  **KC_go_reply** 1 year, 2 months ago

You got it exactly wrong. 'Application image' in this case is literally a Docker image because it's a Cloud Run pipeline (serverless container service), and 'profile images' would be in formats like jpg and png

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: B

Per Google Docs, "Container Registry is a single place for your team to manage Docker images, perform vulnerability analysis, and decide who can access what with fine-grained access control. Existing CI/CD integrations let you set up fully automated Docker pipelines to get fast feedback." <https://cloud.google.com/container-registry>

upvoted 1 times

  **abunasar786** 1 year, 7 months ago

It does not say docket images, it says images. Should be C

upvoted 1 times

  **haroldbenites** 2 years ago

Go for B

upvoted 2 times

Each of the three cloud service models - infrastructure as a service (IaaS), platform as a service (PaaS), and software as a service (SaaS) - offers benefits between flexibility and levels of management by the cloud provider and the customer.

Why would SaaS be the right choice of service model?

- A. You want a balance between flexibility for the customer and the level of management by the cloud provider
- B. You want to minimize the level of management by the customer
- C. You want to maximize flexibility for the customer.
- D. You want to be able to shift your emphasis between flexibility and management by the cloud provider as business needs change

Correct Answer: B

Benefits of SaaS -

The main benefit of SaaS is that it offloads all infrastructure and application management to the SaaS vendor.

Reference:

<https://www.ibm.com/cloud/learn/iaas-paas-saas>

What are IaaS, PaaS and SaaS?

IaaS, PaaS and SaaS are the three most popular types of cloud service offerings. (They are sometimes referred to as cloud service models or cloud computing service models.)

- IaaS, or infrastructure as a service, is on-demand access to cloud-hosted physical and virtual servers, storage and networking - the backend IT infrastructure for running applications and workloads in the cloud.
- PaaS, or platform as a service, is on-demand access to a complete, ready-to-use, cloud-hosted platform for developing, running,

Community vote distribution

B (93%)

7%

- cookieMr

Highly Voted

1 year ago

Selected Answer: B

With SaaS, the cloud provider takes care of the underlying infrastructure, including hardware, networking, and software updates. The customer does not need to worry about managing or maintaining the infrastructure, allowing them to focus on using the software and meeting their business needs. This reduces the burden of IT management for the customer.

upvoted 6 times
- cloudinvader

Most Recent

4 months, 1 week ago

Selected Answer: B

SaaS model provides the capability to use service without worries about management.

upvoted 1 times
- Surek

5 months, 3 weeks ago

Answer is B

upvoted 1 times
- Surek

6 months ago

Answer B

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times

  **__rajan__** 8 months, 2 weeks ago

Selected Answer: B

SaaS (Software as a Service) would be the right choice of service model if you want to minimize the level of management required by the customer. In a SaaS model, the cloud provider takes care of managing the infrastructure, platform, and software, allowing the customer to focus solely on using the software application. The provider handles tasks such as maintenance, updates, security, and scalability, relieving the customer of those responsibilities. This allows the customer to minimize their management efforts and instead concentrate on utilizing the software to meet their business needs.
upvoted 1 times

  **[Removed]** 9 months ago

Selected Answer: B

My answer is B bec SaaS
upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

My answer is B bec SaaS will help me in minimizing management by customer. All management will be done by Service provider
upvoted 1 times

  **prachisri** 10 months, 2 weeks ago

Selected Answer: B

My answer is B bec SaaS will help me in minimizing management by customer. All management will be done by Service provider
upvoted 1 times

  **hireshgupt** 11 months ago

Selected Answer: A

SaaS is a good choice for businesses that want to minimize the amount of IT management they need to do. With SaaS, the cloud provider takes care of all of the software maintenance, updates, and security. This frees up businesses to focus on their core competencies. However, SaaS is not as flexible as IaaS or PaaS. With SaaS, businesses are limited to the features and functionality that are provided by the cloud provider. If a business needs a specific feature that is not offered by the SaaS provider, they will need to find a different solution. Overall, SaaS is a good choice for businesses that want a balance between flexibility and management
upvoted 2 times

  **star2anand** 1 year, 3 months ago

C. You want to maximize flexibility for the customer.
upvoted 1 times

  **AnshMan** 1 year, 5 months ago

This was there in Exam but asked differently
upvoted 4 times

  **EXAM_FEAR** 1 year, 6 months ago

Selected Answer: B

The main benefit of SaaS is that it offloads all infrastructure and application management to the SaaS vendor.
upvoted 1 times

  **Pavanipravallika** 1 year, 6 months ago

B is correct
upvoted 1 times

  **gp7066** 1 year, 8 months ago

B is correct as it involves no management and only user configuration
upvoted 1 times

  **Rakesh1976** 1 year, 8 months ago

B is correct.
upvoted 1 times

As your organization increases its release velocity, the VM-based application upgrades take a long time to perform rolling updates due to OS boot times. You need to make the application deployments faster.
What should your organization do?

- A. Migrate your VMs to the cloud, and add more resources to them
- B. Convert your applications into containers
- C. Increase the resources of your VMs
- D. Automate your upgrade rollouts

Correct Answer: B

Community vote distribution

B (93%)

7%

- rikinetysix

Highly Voted

1 year, 8 months ago

Selected Answer: B

Since the problem with the current infrastructure base on VM is with OS boot time in particular, so upgrading the workload to containers would be a right choice as then OS boot time problem is eliminated which means a faster workload execution. So, correct answer should be 'B'.
upvoted 15 times
- cloudinvader

Most Recent

4 months, 1 week ago

Selected Answer: B

we need to get rid of "OS boot times" that can be only achieved by migrating to container
upvoted 1 times
- Surek

5 months, 3 weeks ago

Answer is B
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: B

To make application deployments faster and improve release velocity, your organization should consider converting the applications into containers. Containers provide a lightweight and isolated runtime environment for applications, allowing for faster deployment and scaling.
upvoted 2 times
- mdsarfraz69

9 months ago

Selected Answer: B

My answer is B
upvoted 1 times
- prachisri

10 months, 2 weeks ago

Selected Answer: B

Converting applications to containers will reduce time taken for OS reboots and it will make it faster
upvoted 1 times
- EulaGaf

1 year ago

B. Convert your applications into containers. Containers are light weight and does not need a dedicated OS to boot every time
upvoted 1 times
- cookieMr

1 year ago

Selected Answer: B

Convert your VM-based application into containerized application components using containerization technologies like Docker. Containerization allows you to package the application and its dependencies into lightweight, portable containers that can be deployed consistently across different environments.
upvoted 2 times
- Rafa_M

1 year, 1 month ago

Selected Answer: B

B. Convert your applications into containers
upvoted 1 times

  **Aravind48** 1 year, 2 months ago

Selected Answer: B

Also the purpose of Cloud Native
upvoted 1 times

  **star2anand** 1 year, 3 months ago

B. Convert your applications into containers
upvoted 1 times

  **alfantastic** 1 year, 5 months ago

Selected Answer: B

Containers do not require OS boot time overhead
upvoted 1 times

  **zellick** 1 year, 6 months ago

Selected Answer: B

B is the answer.

<https://cloud.google.com/migrate/containers/docs/anthos-migrate-benefits>
upvoted 2 times

  **muraligreetsu** 1 year, 6 months ago

Selected Answer: B

Containers are light weight and does not need a dedicated OS to boot every time
upvoted 1 times

  **Huntishi** 1 year, 6 months ago

Selected Answer: D

When you set up an automated update, the MIG rolls out a new version of an instance template automatically to all or to a subset of the group's VMs. If you have stateful configuration, the MIG also applies any unapplied per-instance configurations to the corresponding VMs.
upvoted 1 times

  **ronieto** 1 year, 7 months ago

Selected Answer: B

Containers do virtualize OS, so don't need to boot it with every new change
upvoted 3 times

Your organization uses Active Directory to authenticate users. Users' Google account access must be removed when their Active Directory account is terminated.
How should your organization meet this requirement?

- A. Configure two-factor authentication in the Google domain
- B. Remove the Google account from all IAM policies
- C. Configure BeyondCorp and Identity-Aware Proxy in the Google domain
- D. Configure single sign-on in the Google domain

Correct Answer: D

Community vote distribution

D (93%)

7%

- JCE

Highly Voted

2 years, 5 months ago

D seems correct

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-configuring-single-sign-on>

upvoted 7 times
- ucsdmiami2020

Highly Voted

1 year, 6 months ago

Selected Answer: D

Per Google Docs article, Federating Google Cloud with Active Directory. "This article describes how you can configure Cloud Identity or Google Workspace to use Active Directory as IdP and authoritative source.

The article compares the logical structure of Active Directory with the structure used by Cloud Identity and Google Workspace and describes how you can map Active Directory forests, domains, users, and groups. The article also provides a flowchart that helps you determine the best mapping approach for your scenario."

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction>

upvoted 6 times
- Surek

Most Recent

6 months ago

Answer is D

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: D

SSO is correct as deletion of AD account will remove access from GCP as well.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: D

D is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: D

SSO allows for centralized user management, where user accounts and access permissions are managed in a single identity provider (such as Active Directory). When a user's Active Directory account is terminated, SSO provides a centralized point to revoke access across multiple applications and services, including Google accounts.

upvoted 5 times
- star2anand

1 year, 3 months ago

D. Configure single sign-on in the Google domain

upvoted 1 times
- KanikaA

1 year, 4 months ago

Selected Answer: D

Using SSO would help in removing access once the account is no longer active.

upvoted 1 times

🗨️ 👤 **Pou1ze** 1 year, 6 months ago

Selected Answer: D

D is correct
upvoted 1 times

🗨️ 👤 **ronieto** 1 year, 7 months ago

Selected Answer: D

SSO means federation between AD and Cloud ID, so is the correct answer
upvoted 3 times

🗨️ 👤 **SimonIt73** 1 year, 7 months ago

The correct answer should be "Setting up federation between Active Directory and Cloud Identity or Google Workspace". To do that, you have to enable automatic users provisioning and SSO.
upvoted 3 times

🗨️ 👤 **rikininetysix** 1 year, 8 months ago

Selected Answer: C

The question asked to provide a solution to remove users' Google account access when their Active Directory account is terminated. So, option 'C' should be correct as BeyondCorp and Identity Aware Proxy are focused solutions to manage Identity and implement a Zero trust model.
upvoted 2 times

🗨️ 👤 **hogtrough** 1 year, 8 months ago

The correct answer is D. If you have SSO configured, once a user's AD account is terminated, their access is removed from all services using AD.
upvoted 3 times

🗨️ 👤 **haroldbenites** 2 years ago

Go for D
upvoted 2 times

🗨️ 👤 **Monicaarg** 2 years, 1 month ago

Selected Answer: D

Your organization uses Active Directory to authenticate users.
Then you need to use Single sign-on (SSO) is an authentication scheme that allows a user to log in with a single ID and password to different systems and software.
SSO allows IT departments to administrator a single identity that can access many machines and cloud services.
upvoted 4 times

🗨️ 👤 **NU_1234** 2 years, 2 months ago

Selected Answer: D

SSO is the answer
upvoted 1 times

🗨️ 👤 **RonRon512** 2 years, 3 months ago

Selected Answer: D

D is correct
upvoted 1 times

Your company has recently acquired three growing startups in three different countries. You want to reduce overhead in infrastructure management and keep your costs low without sacrificing security and quality of service to your customers. How should you meet these requirements?

- A. Host all your subsidiaries' services on-premises together with your existing services.
- B. Host all your subsidiaries' services together with your existing services on the public cloud.
- C. Build a homogenous infrastructure at each subsidiary, and invest in training their engineers.
- D. Build a homogenous infrastructure at each subsidiary, and invest in hiring more engineers.

Correct Answer: B

Community vote distribution

B (100%)

-  **Surek** 5 months, 3 weeks ago

Answer is B

upvoted 1 times
-  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 2 times
-  **__rajan__** 8 months, 2 weeks ago

Selected Answer: B

To reduce overhead in infrastructure management, keep costs low, and maintain security and quality of service, hosting all your subsidiaries' services together with your existing services on the public cloud

upvoted 1 times
-  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times
-  **afeff** 9 months ago

Selected Answer: B

optimize infra on the existing cloud provider

upvoted 1 times
-  **prachisri** 10 months, 2 weeks ago

Selected Answer: B

Requirement is to reduce overhead in infrastructure management and keep costs low without sacrificing security and quality of service. This can be achieved by migrating to cloud where cloud provider will take care of all infrastructure expenses and will be secure at the same time

upvoted 1 times
-  **cookieMr** 1 year ago

Selected Answer: B

: Migrate your infrastructure to a cloud computing platform, such as Google Cloud Platform (GCP). Cloud providers handle the underlying infrastructure management, allowing you to focus on your applications and services. This reduces the overhead of maintaining physical infrastructure, hardware upgrades, and capacity planning.

upvoted 2 times
-  **star2anand** 1 year, 3 months ago

B. Host all your subsidiaries' services together with your existing services on the public cloud.

upvoted 1 times
-  **alfantastic** 1 year, 5 months ago

Selected Answer: B

makes efficiency

upvoted 1 times
-  **Noboss** 1 year, 7 months ago

Only public cloud can reduce infrastructure cost.

upvoted 3 times

  **rikininetysix** 1 year, 8 months ago

Selected Answer: B

Both maintenance cost becomes cheaper and management becomes easier when subsidiaries workloads along with parent organisation's workloads are deployed on the public cloud.

upvoted 1 times

  **iamPhD** 1 year, 9 months ago

Selected Answer: B

Cloud is cheaper the on prem infrastructure

upvoted 1 times

  **haroldbenites** 2 years ago

Go for B

upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: B

hostee todos los servicios en la nube.

upvoted 1 times

  **NU_1234** 2 years, 2 months ago

Selected Answer: B

B is the answer

upvoted 2 times

  **RonRon512** 2 years, 3 months ago

Selected Answer: B

Public cloud is the solution for consolidation

upvoted 2 times

  **kirri** 2 years, 3 months ago

B is the best answer

upvoted 1 times

What is the difference between Standard and Coldline storage?

- A. Coldline storage is for data for which a slow transfer rate is acceptable.
- B. Standard and Coldline storage have different durability guarantees.
- C. Standard and Coldline storage use different APIs.
- D. Coldline storage is for infrequently accessed data.

Correct Answer: D

Reference:

<https://www.msp360.com/resources/blog/google-cloud-nearline-storage-vs-coldline-vs-standard/>

The main characteristics of Coldline are as follows:

- SLA guarantees 99% data availability (of ten thousand hours, data can be offline for up to 100 hours).
- Monthly fee per GB stored is between \$0.004-\$0.014, depending on a region.
- The minimum storing period is 90 days. If you delete data earlier, you have to pay for the remaining time. For example, if you had uploaded 100GB of data and then deleted it after 30 days, you need to pay extra \$1.4 as an early deletion fee ($\$0.004$ (or $\$0.014$) * 100 GB * 2 months).
- Data retrieval fee is mandatory and costs \$0.05 per GB.

Its low price and a long minimum storing period make Nearline the best solution for data that is unlikely to be accessed more than once a year, if ever:

- Data archive.
- Disaster recovery storage.
- Outdated backups storage.

Community vote distribution

D (95%)

5%

- Surek** 5 months, 3 weeks ago

Answer is D

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__** 8 months, 2 weeks ago

Selected Answer: D

D is correct.

upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: D

D. Coldline storage is for infrequently accessed data.

upvoted 1 times
- afeff** 9 months ago

Selected Answer: D

it's not about latency of transfer. It's about the access pattern for the coldline storage

upvoted 1 times
- prachisri** 10 months, 2 weeks ago

Selected Answer: D

Coldline storage is for infrequently accessed data whereas Standard storage is for frequently accessed data

upvoted 1 times

🗨️ 👤 **cookieMr** 1 year ago

Selected Answer: D

The primary difference between Standard and Coldline storage in Google Cloud Platform's Cloud Storage service lies in their intended use cases, access patterns, and associated costs.

Standard storage is designed for frequently accessed data that requires low-latency access. It offers faster read and write performance compared to Coldline storage. On the other hand, Coldline storage is optimized for infrequently accessed data that can tolerate higher retrieval times. Accessing data from Coldline storage may have higher latency due to the retrieval process.

upvoted 3 times

🗨️ 👤 **Rafa_M** 1 year, 1 month ago

Selected Answer: D

D. Coldline storage is for infrequently accessed data.

upvoted 1 times

🗨️ 👤 **star2anand** 1 year, 3 months ago

D. Coldline storage is for infrequently accessed data.

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 3 months ago

Selected Answer: B

Now I think about this a bit more, the question is "what's the difference between..." and not "what's the definition of..."
The only answers that provide a "difference between" are B and C...and the answer is definitely not C! I might go with B here.

upvoted 1 times

🗨️ 👤 **KC_go_reply** 1 year, 3 months ago

Only D is correct.

The different storage tiers have (officially) the same durability: "eleven 9s", see <https://cloud.google.com/blog/products/storage-data-transfer/understanding-cloud-storage-11-9s-durability-target?hl=en>.

In other providers like AWS for example, there is so-called 'One-Zone IA', which means the bucket is really just in one zone, without regional replication. This then results in lower durability guarantees / SLA

upvoted 1 times

🗨️ 👤 **jres** 1 year, 3 months ago

D for sure, coldline talks to long time access

upvoted 1 times

🗨️ 👤 **arindamIndia** 1 year, 6 months ago

question is slightly focused on the difference and not particularly on what standard or coldline refers individually. So, option B seems right to me. Please let me know if any one of you think otherwise.

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 3 months ago

I agree with you but I think this is just a poorly phrased question with confusing answers. D is my vote but it doesn't explain the "difference" between the two but is rather the definition of Coldline. I don't like this one...very bad.

upvoted 1 times

🗨️ 👤 **sfg** 1 year, 6 months ago

Selected Answer: D

D is correct

upvoted 1 times

🗨️ 👤 **Pou1ze** 1 year, 6 months ago

Selected Answer: D

D is correct

upvoted 1 times

🗨️ 👤 **haroldbenites** 2 years ago

Go for D

upvoted 1 times

🗨️ 👤 **Monicaarg** 2 years, 1 month ago

Selected Answer: D

Standard – when you are frequently using files.

upvoted 1 times

🗨️ 👤 **NU_1234** 2 years, 2 months ago

Selected Answer: D

D is the answer

upvoted 1 times

What would provide near-unlimited availability of computing resources without requiring your organization to procure and provision new equipment?

- A. Public cloud
- B. Containers
- C. Private cloud
- D. Microservices

Correct Answer: A

Reference:

<https://cloud.google.com/docs/overview>

Community vote distribution

A (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 2 weeks ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

public cloud is unlimited resources

upvoted 1 times
- afeff

9 months ago

Selected Answer: A

public cloud is unlimited resources

upvoted 1 times
- tsk9921

11 months, 2 weeks ago

A public cloud addresses the need

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: A

To achieve near-unlimited availability of computing resources without the need to procure and provision new equipment, your organization can leverage the cloud computing model, specifically Infrastructure as a Service (IaaS) offered by cloud service providers like Google Cloud Platform (GCP).

upvoted 3 times
- Rafa_M

1 year, 1 month ago

Selected Answer: A

A. Public clou

upvoted 1 times
- star2anand

1 year, 3 months ago

D. Microservices

upvoted 1 times
- saifelmadani

1 year, 3 months ago

A. Public Cloud

upvoted 1 times
- AnshMan

1 year, 5 months ago

This was there in Exam

upvoted 1 times

  **sfg** 1 year, 6 months ago

Selected Answer: A

Public cloud, so A
upvoted 2 times

  **haroldbenites** 2 years ago

Go for A
upvoted 2 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: A

Public cloud
upvoted 1 times

  **ravisinha** 1 year, 11 months ago
In today exams, this question was asked
upvoted 2 times

  **NU_1234** 2 years, 2 months ago

Selected Answer: A

That is the genesis for public cloud
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: A

A. Public cloud
upvoted 2 times

  **beginnercloud** 2 years, 4 months ago

Selected Answer: A

Public cloud
upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: A

A. Public cloud
upvoted 1 times

You are a program manager for a team of developers who are building an event-driven application to allow users to follow one another's activities in the app. Each time a user adds himself as a follower of another user, a write occurs in the real-time database. The developers will develop a lightweight piece of code that can respond to database writes and generate a notification to let the appropriate users know that they have gained new followers. The code should integrate with other cloud services such as Pub/Sub, Firebase, and Cloud APIs to streamline the orchestration process. The application requires a platform that automatically manages underlying infrastructure and scales to zero when there is no activity. Which primary compute resource should your developers select, given these requirements?

- A. Google Kubernetes Engine
- B. Cloud Functions
- C. App Engine flexible environment
- D. Compute Engine

Correct Answer: B

Reference:
<https://firebase.google.com/docs/functions/use-cases>

Cloud Functions gives developers access to Firebase and Google Cloud events, along with scalable computing power to run code in response to those events. While it's expected that Firebase apps will use Cloud Functions in unique ways to meet their unique requirements, typical use cases might fall into these areas:

- [Notify users when something interesting happens.](#)
- [Perform database sanitization and maintenance.](#)
- [Execute intensive tasks in the cloud instead of in your app.](#)
- [Integrate with third-party services and APIs.](#)

Review the use cases and examples for each category that interests you, and then proceed to our [Get Started](#) tutorial or to specific how-to guides for [authentication events](#), [analytics events](#), and more. See the [eventType API reference](#) for the complete list of supported event types.

Community vote distribution

B (100%)

- mickeyadlg

Highly Voted

2 years, 5 months ago

B is correct, app engine flexible doesn't scale to zero

upvoted 13 times
- cookieMr

Highly Voted

1 year ago

Selected Answer: B

Given the requirements of automatically managing underlying infrastructure, scaling to zero when there is no activity, and integrating with other cloud services like Pub/Sub, Firebase, and Cloud APIs, the appropriate compute resource for your developers to select would be Cloud Functions in Google Cloud Platform (GCP).

Cloud Functions is a serverless compute service that allows you to run event-driven code in a fully managed environment. It is designed to execute short-lived functions in response to various events, such as database writes, file uploads, HTTP requests, and Pub/Sub messages.

upvoted 5 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: B

Bis correct

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct, app engine flexible doesn't scale to zero

upvoted 1 times
- afeff

9 months ago

Selected Answer: B

Cloud Function cant App Engine flexible cannot run at zero scale (standard configuraiton does)
upvoted 1 times

  **Rafa_M** 1 year, 1 month ago

Selected Answer: B

B. "Cloud Functions for Firebase is a serverless framework that lets you automatically run backend code in response to events triggered by Firebase features and HTTPS requests"
upvoted 2 times

  **star2anand** 1 year, 3 months ago

B. Cloud Functions
upvoted 1 times

  **adityanarayan** 1 year, 4 months ago

But companies are using AKS/GKE/EKS for that also.
upvoted 1 times

  **Saptarsi_Mondal** 1 year, 7 months ago

Cloud Function scale from 0 to global level. Is Serverless. Compute for event driven apps. App engine flexible minimum have 1 instance not 0, So answer is B.
upvoted 4 times

  **ravikiran37** 1 year, 8 months ago

infrastructure and scales to zero when there is no activity. is key point to select B
upvoted 1 times

  **haroldbenites** 2 years ago

Go for B
upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: B

Cloud Function scale from 0 to global level. Is Serverless. Compute for event driven apps.
upvoted 2 times

  **albert_kuo** 2 years, 2 months ago

Selected Answer: B

Because of event-driven application, it should be B.
upvoted 1 times

  **NU_1234** 2 years, 2 months ago

Selected Answer: B

Cloud function - Event driven, light weight..
upvoted 1 times

  **RonRon512** 2 years, 3 months ago

Selected Answer: B

Lightweight code- Cloud functions
upvoted 1 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: B

B. Cloud Functions
upvoted 1 times

  **beginnercloud** 2 years, 4 months ago

Selected Answer: B

B is correct
upvoted 1 times

Your organization is developing an application that will capture a large amount of data from millions of different sensor devices spread all around the world. Your organization needs a database that is suitable for worldwide, high-speed data storage of a large amount of unstructured data. Which Google Cloud product should your organization choose?

- A. Firestore
- B. Cloud Data Fusion
- C. Cloud SQL
- D. Cloud Bigtable

Correct Answer: D

Reference:

<https://cloud.google.com/bigtable>

Cloud Bigtable

A fully managed, scalable NoSQL database service for large analytical and operational workloads with up to 99.999% availability.

Try Google Cloud free Contact sales

- ✓ Consistent sub-10ms latency—handle millions of requests per second
- ✓ Ideal for use cases such as personalization, ad tech, fintech, digital media, and IoT

Community vote distribution

D (95%) 5%

- Gracika** 5 months, 3 weeks ago

D is the correct

upvoted 1 times
- Surek** 6 months ago

D is the answer

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: D

BigTable is correct.

upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: D

Go for D

upvoted 1 times
- RobertPL** 9 months ago

Selected Answer: A

If you check <https://cloud.google.com/products/databases>, then IoT is indicated as a use case for Firestore.
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

For worldwide, high-speed data storage of a large amount of unstructured data captured from millions of sensor devices, the suitable Google Cloud product for your organization would be Cloud Bigtable.

Cloud Bigtable is a fully managed, highly scalable, and high-performance NoSQL database service designed for large-scale applications that require low-latency and high-throughput data storage.

upvoted 4 times

  **star2anand** 1 year, 3 months ago

D. Cloud Bigtable

upvoted 1 times

  **AkshayRamnani** 1 year, 5 months ago

I think the answer is A (check the document page)

upvoted 1 times

  **MKPandian** 1 year, 6 months ago

Selected Answer: D

unstructure Data so D

upvoted 1 times

  **Serchi** 1 year, 8 months ago

Go for D

upvoted 1 times

  **Shashank_R** 2 years ago

Huge Data, Low Access time and Analytics then it should be Big Table D is Answer

upvoted 4 times

  **haroldbenites** 2 years ago

Go for D

upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: D

high-speed data storage of a large amount of unstructured data...

Bigtable is special for stream data IOT.

upvoted 4 times

  **albert_kuo** 2 years, 2 months ago

Selected Answer: D

"capture a large amount of data from millions of different sensor devices spread all around the world" -> IoT solution

It should be D (Cloud Bigtable)

upvoted 1 times

  **NU_1234** 2 years, 2 months ago

Selected Answer: D

Large volume, unstructured, IOT

upvoted 2 times

Your organization needs to build streaming data pipelines. You don't want to manage the individual servers that do the data processing in the pipelines. Instead, you want a managed service that will automatically scale with the amount of data to be processed.

Which Google Cloud product or feature should your organization choose?

- A. Pub/Sub
- B. Dataflow
- C. Data Catalog
- D. Dataprep by Trifacta

Correct Answer: B

Reference:

<https://cloud.google.com/dataflow/docs/guides/deploying-a-pipeline>

Deploying a pipeline

Send feedback

★ This document explains in detail how Dataflow deploys and runs a pipeline, and covers advanced topics like optimization and load balancing. If you are looking for a step-by-step guide on how to create and deploy your first pipeline, use Dataflow's quickstarts for [Java](#), [Python](#) or [templates](#).

After you construct and test your Apache Beam pipeline, you can use the Dataflow managed service to deploy and execute it. Once on the Dataflow service, your pipeline code becomes a Dataflow job.

The Dataflow service fully manages Google Cloud services such as [Compute Engine](#) and [Cloud Storage](#) to run your Dataflow job, automatically spinning up and tearing down the necessary resources. The Dataflow service provides visibility into your job through tools like the [Dataflow monitoring interface](#) and the [Dataflow command-line interface](#).

Community vote distribution

B (97%)

- chai_gpt 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- anilchandane 7 months, 4 weeks ago

B is correct

upvoted 1 times
- __rajan__ 8 months, 1 week ago

Selected Answer: B

Dataflow is correct.

upvoted 1 times
- mdsarfraz69 9 months ago

Selected Answer: B

Dataflow is a fully managed service that allows you to build and run streaming data pipelines without having to manage the underlying infrastructure. Dataflow will automatically scale your pipelines to meet the demands of your data.

upvoted 3 times
- hireshgupt 11 months ago

Selected Answer: B

Dataflow is a fully managed service that allows you to build and run streaming data pipelines without having to manage the underlying infrastructure. Dataflow will automatically scale your pipelines to meet the demands of your data.

upvoted 2 times

🗄️ 👤 **cookieMr** 1 year ago

Selected Answer: B

For building streaming data pipelines without the need to manage individual servers and with automatic scalability, Google Cloud offers a managed service called Cloud Dataflow.

Cloud Dataflow is a fully managed, serverless data processing service that allows you to build and execute data pipelines for both batch and streaming data processing. It provides a unified programming model for developing data processing logic and automatically handles the underlying infrastructure and resource management.

upvoted 3 times

🗄️ 👤 **star2anand** 1 year, 3 months ago

B. Dataflow

upvoted 1 times

🗄️ 👤 **sfg** 1 year, 6 months ago

Selected Answer: B

Processing is done by dataflow, pub/sub is for data ingestion.

upvoted 4 times

🗄️ 👤 **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: B

Per Google Docs one of the key features of GCP dataflow is, " Minimize pipeline latency, maximize resource utilization, and reduce processing cost per data record with data-aware resource autoscaling. Data inputs are partitioned automatically and constantly rebalanced to even out worker resource utilization and reduce the effect of "hot keys" on pipeline performance"

<https://cloud.google.com/dataflow>

upvoted 1 times

🗄️ 👤 **Alfred_Magno** 1 year, 8 months ago

Selected Answer: B

Answer is B, Dataflow stream data processing pipelines

upvoted 3 times

🗄️ 👤 **AMMARABO** 1 year, 8 months ago

Selected Answer: A

Instead, you want a managed service that will automatically scale with the amount of data to be processed.

upvoted 1 times

🗄️ 👤 **haroldbenites** 2 years ago

Go for B

upvoted 1 times

🗄️ 👤 **Monicaarg** 2 years, 1 month ago

Selected Answer: B

Dataflow Develop real-time batch and stream data processing pipelines. (Apache Beam)

upvoted 3 times

🗄️ 👤 **NU_1234** 2 years, 2 months ago

Selected Answer: B

Batch and Streaming pipelines

upvoted 4 times

🗄️ 👤 **beginnercloud** 2 years, 4 months ago

Selected Answer: B

Dataflow

upvoted 2 times

🗄️ 👤 **victory108** 2 years, 4 months ago

Selected Answer: B

B. Dataflow

upvoted 2 times

🗄️ 👤 **leyunjohn** 2 years, 4 months ago

Selected Answer: B

B. Dataflow

upvoted 1 times

Your organization is building an application running in Google Cloud. Currently, software builds, tests, and regular deployments are done manually, but you want to reduce work for the team. Your organization wants to use Google Cloud managed solutions to automate your build, testing, and deployment process.

Which Google Cloud product or feature should your organization use?

- A. Cloud Scheduler
- B. Cloud Code
- C. Cloud Build
- D. Cloud Deployment Manager

Correct Answer: C

Deploy your application to App Engine using the gcloud app deploy command. This command automatically builds a container image by using the Cloud Build service and then deploys that image to the App Engine flexible environment.

Reference:

<https://cloud.google.com/appengine/docs/flexible/nodejs/testing-and-deploying-your-app>

- chai_gpt7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- anilchandane7 months, 4 weeks ago

C is correct Answer : Cloud Build

upvoted 1 times
- __rajan__8 months, 1 week ago

Selected Answer: C

CloudBuild is Correct.

upvoted 1 times
- gcp_prep8 months, 2 weeks ago

c - Cloud Build

upvoted 1 times
- mdsarfraz699 months ago

Selected Answer: C

Cloud Build is a fully managed CI/CD platform that automates the build, test, and deployment of your applications. It provides a reliable and scalable infrastructure to compile source code, run unit tests, and package artifacts for deployment. Cloud Build integrates with popular source code repositories like GitHub and Bitbucket, allowing you to trigger builds automatically on code changes or schedules.

upvoted 2 times
- cookieMr1 year ago

Selected Answer: C

Cloud Build is a fully managed CI/CD platform that automates the build, test, and deployment of your applications. It provides a reliable and scalable infrastructure to compile source code, run unit tests, and package artifacts for deployment. Cloud Build integrates with popular source code repositories like GitHub and Bitbucket, allowing you to trigger builds automatically on code changes or schedules.

upvoted 4 times

  **star2anand** 1 year, 3 months ago

A. Cloud Scheduler

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Per Google docs, Cloud Build allows for "building, testing, and deploying on a serverless CI/CD platform."

<https://cloud.google.com/build>

upvoted 2 times

  **md3danish** 1 year, 7 months ago

Selected Answer: C

But the description is wrong as it for appengine not for cloud build

upvoted 2 times

  **francescofusco** 1 year, 8 months ago

Cloud build is correct. the last one is used for automating the creation/deploy of infrastructures/vm's

upvoted 2 times

  **yellow28** 1 year, 11 months ago

Selected Answer: C

Cloud Build is CI/CD tool for GCP

upvoted 3 times

  **albert_kuo** 2 years, 2 months ago

Selected Answer: C

Cloud Build helps build, test and deploy on serverless CI/CD platforms

upvoted 4 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. Cloud Build

upvoted 2 times

  **beginnercloud** 2 years, 4 months ago

Selected Answer: C

C- Cloud Build

upvoted 1 times

  **Parv7221** 2 years, 4 months ago

Cloud Build

upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: C

C. Cloud Build

upvoted 2 times

  **VijaySVK** 2 years, 5 months ago

C. Cloud build

upvoted 1 times

Which Google Cloud product can report on and maintain compliance on your entire Google Cloud organization to cover multiple projects?

- A. Cloud Logging
- B. Identity and Access Management
- C. Google Cloud Armor
- D. Security Command Center

Correct Answer: B

Community vote distribution

D (82%)B (18%)

-   **SURJK**

Highly Voted 

 2 years, 5 months ago

Answer is D: SCC - Gain centralized visibility and control
Discover misconfigurations and vulnerabilities
Report on and maintain compliance
Detect threats targeting your Google Cloud assets
upvoted 37 times
-   **fz98** 11 months, 1 week ago

is there Link to refer?
upvoted 4 times
-   **Monicaarg**

Highly Voted 

 2 years, 1 month ago

Selected Answer: D

Security Command Center is a centralized security and risk management platform for your Google Cloud resources. It is a single tool that offers a variety of security features including:
1. Gain centralized visibility and control
2. Discover misconfigurations and vulnerabilities
3. Report on and maintain compliance
4. Detect threats targeting your Google Cloud assets

<https://cloud.google.com/security-command-center>
upvoted 22 times

  **Nigel_123**

Most Recent 

 1 month ago

Selected Answer: D

Answer should be D
IAM if for access management
upvoted 1 times

  **4a98421** 3 months, 2 weeks ago

Answer is D.
IAP is to manage who access and grant required privilege.
upvoted 1 times

  **aaabhi** 3 months, 3 weeks ago

Selected Answer: B

Streamline compliance with a built-in audit trail
A full audit trail history of permissions authorization, removal, and delegation gets surfaced automatically for your admins. IAM lets you focus on business policies around your resources and makes compliance easy.
upvoted 1 times

  **Barzoom** 4 months, 2 weeks ago

Selected Answer: D

The correct answer is D
upvoted 1 times

  **Gabriella_Pacheco** 5 months ago

Selected Answer: D

Answer is D for me
upvoted 1 times

  **sivakarthick16** 5 months, 1 week ago

Selected Answer: D

Security Command Center provides centralized visibility and control over security and compliance across your Google Cloud environment. It offers comprehensive security and compliance monitoring, threat detection, and vulnerability management capabilities to help you maintain a secure and compliant infrastructure.

upvoted 1 times

  **Surek** 6 months ago

D IS THE ANSWER

upvoted 1 times

  **f73aba5** 6 months, 3 weeks ago

Selected Answer: D

Correct answer should be D. If B is correct some explanation should be provided.

upvoted 2 times

  **f73aba5** 6 months, 3 weeks ago

Correct answer should be D

upvoted 1 times

  **Aleksan93** 6 months, 3 weeks ago

Selected Answer: D

Security Command Center allows to take care of compliance topic

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

It should be Security Command Center.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times

  **afeff** 9 months ago

Selected Answer: D

D OR B ?

I think the SCC is the correct answer but the provided answer is IAM. can somebody help?

upvoted 4 times

  **afeff** 9 months ago

Selected Answer: D

Security Command Center

upvoted 1 times

Your organization needs to establish private network connectivity between its on-premises network and its workloads running in Google Cloud. You need to be able to set up the connection as soon as possible. Which Google Cloud product or feature should you use?

- A. Cloud Interconnect
- B. Direct Peering
- C. Cloud VPN
- D. Cloud CDN

Correct Answer: C

Community vote distribution

C (91%)

9%

- anirban7172**

Highly Voted

2 years, 5 months ago

Selected Answer: C

Cloud VPN setup can be done instantly and though it travels over internet, it is private in nature. Cloud interconnect needs physically dedicated connection.

upvoted 38 times
- SURJK**

Highly Voted

2 years, 5 months ago

Answer is C: Cloud VPN

upvoted 13 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: C

The fastest way to connect is to use Cloud VPN.

upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: C

Answer is C: Cloud VPN

upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: C

To establish private network connectivity between your on-premises network and workloads running in Google Cloud, you should use Google Cloud VPN (Virtual Private Network).

Google Cloud VPN allows you to securely connect your on-premises network to your virtual private cloud (VPC) network in Google Cloud. It enables encrypted communication over the public internet, providing a secure and private connection between your on-premises environment and your cloud resources.

upvoted 4 times
- Rafa_M**

1 year, 1 month ago

Selected Answer: C

My answer is C, because You need to be able to set up the connection as soon as possible.

upvoted 1 times
- star2anand**

1 year, 3 months ago

A. Cloud Interconnect

upvoted 1 times
- [Removed]**

1 year, 3 months ago

Selected Answer: C

"Your organization needs to establish private network..."

VPN = Virtual Private Network.

upvoted 3 times

☒  **KanikaA** 1 year, 4 months ago

Selected Answer: C

Cloud VPN set up can be done instantly & connection is private, even though data flows through public internet.
upvoted 2 times

☒  **PHD_CHENG** 1 year, 5 months ago

Selected Answer: C

C is correct. As connection need to setup "as soon as possible", Cloud VPN can create any time. Cloud Interconnect need to order from 3rd party
upvoted 3 times

☒  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: C

Per Google's doc titled, Choosing a Network Connectivity product, it states the following, "If you need to encrypt traffic to Google Cloud, or you need a lower throughput solution, or you are experimenting with migrating your workloads to Google Cloud, you can choose Cloud VPN."
<https://cloud.google.com/network-connectivity/docs/how-to/choose-product>
upvoted 1 times

☒  **Pou1ze** 1 year, 6 months ago

Selected Answer: C

A and C are solution for private network interconnection, but C is more quick to implement, so the answer is C
upvoted 1 times

☒  **grzeev** 1 year, 8 months ago

Selected Answer: C

B. soon as possible only on internet
upvoted 1 times

☒  **Rakesh1976** 1 year, 8 months ago

C is correct Answer, Cloud VPN
upvoted 1 times

☒  **pspandher** 1 year, 8 months ago

I will go with Cloud VPN, as per Google the Cloud VPN can be established using a private and public IP both , Now here the main point or thing to focus is Bandwidth requirement if it is Private and your connection is in Google POP then we use Dedicated Interconnect or Partner Interconnect. Else for Public IP we use Direct peering Or carrier peering.

<https://cloud.google.com/blog/products/networking/google-cloud-network-connectivity-options-explained>

I hope this answer everybody questions (refer to diagram)
upvoted 2 times

☒  **pspandher** 1 year, 8 months ago

Cloud VPN is not Private In nature and it is over the Internet, This is as per Google Docs Then ... ?? How can this be the answer ?

Can some one clarify please ??
upvoted 2 times

☒  **geeroylenkins** 1 year, 2 months ago

VPN = Virtual PRIVATE Network ...
upvoted 1 times

Your organization is developing a mobile app and wants to select a fully featured cloud-based compute platform for it. Which Google Cloud product or feature should your organization use?

- A. Google Kubernetes Engine
- B. Firebase
- C. Cloud Functions
- D. App Engine

Correct Answer: B

Community vote distribution

B (75%)

D (25%)

- cyberzone213197**

Highly Voted

2 years, 5 months ago

Selected Answer: B

Firestore is the mobile platform

upvoted 30 times
- ahsangh**

Highly Voted

2 years, 5 months ago

Selected Answer: D

D. App Engine

upvoted 9 times
- mabkhan**

2 years, 5 months ago

why Answer D is correct?

upvoted 2 times
- ahsangh**

2 years, 5 months ago

<https://cloud.google.com/appengine>

upvoted 1 times
- A_A_AB**

2 years, 5 months ago

<https://cloud.google.com/firebase>

upvoted 10 times
- anverdor**

2 years, 5 months ago

I've passed my exam todasy and I chose APP Engine also

upvoted 5 times
- mdmahanti**

1 year, 10 months ago

But that doesnt mean that App Engine is correct. Unless you scores 100%

upvoted 12 times
- csrazdan**

1 year, 7 months ago

You got it wrong. Maybe you are confusing it with Firestore.

upvoted 1 times
- [Removed]**

2 years, 2 months ago

this doesnt mean you choose correct answer, because exam wont show you where you made misstakes after passing and you dont even know how many misstakes you made, iam still leaning to B for Firebase because only diffrence between app engine and firebase is that Firebase is recommended for mobile app development

upvoted 22 times
- mtbmtb**

1 year, 10 months ago

TY for saying that because Im seeing too many folks saying they passed so this or that answer HAS to be correct! But how do you know this particular answer you chose was the RIGHT one??? hmmmmm

upvoted 5 times
- GenAI**

3 months ago

Ha ha ha

upvoted 1 times
- szefco**

1 year, 11 months ago

Although it is possible to build an app in App Engine, Firebase is dedicated suite for mobile apps. B is correct

upvoted 6 times

  **waihtut** 1 year, 2 months ago

App Engine is develop and deploy application.
Firebase is the mobile platform.

upvoted 3 times

  **nghadyali** Most Recent 4 months ago

Selected Answer: B

Mobile app -> Firebase

upvoted 2 times

  **findingNemo31** 5 months ago

D: App Engine

Firebase is a mobile and web application development platform, and while it provides features for backend services, authentication, and real-time database, it may not be a fully featured cloud-based compute platform for all types of applications. App Engine, on the other hand, is a fully managed serverless platform that allows you to build and deploy applications without managing the underlying infrastructure. If your organization is developing a mobile app and wants a cloud-based compute platform, App Engine would be a more suitable choice. Firebase could still be used for specific features within the app, but it may not cover all the compute needs for the application.

upvoted 2 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: D

App Engine is a platform-as-a-service (PaaS) offering that allows developers to build and deploy applications easily, without worrying about infrastructure management. It provides automatic scaling, load balancing, and supports multiple programming languages, making it suitable for mobile app development.

upvoted 1 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: D

App Engine is a platform-as-a-service (PaaS) offering that allows developers to build and deploy applications easily, without worrying about infrastructure management. It provides automatic scaling, load balancing, and supports multiple programming languages, making it suitable for mobile app development.

upvoted 1 times

  **GenAI** 3 months ago

Why do you have to repeat?

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

  **lemonfan123** 8 months ago

Selected Answer: B

I think "Fully featured" is the keyword here, which would indicate towards Firebase than App Engine

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

"App Engine." Google App Engine is a platform-as-a-service (PaaS) offering that simplifies application deployment and scaling, making it well-suited for web and mobile app backends, web services, and APIs. It allows you to focus on writing code while Google manages the infrastructure, making it a good choice for many app development scenarios.

upvoted 1 times

  **RobertPL** 9 months ago

Selected Answer: B

It's a tricky question, because it refers to compute engine, and according to Google cheatsheet, Compute services do not include Firebase, but App Engine is included. However, one must notice, that Firebase is part of a completely separate ecosystem of services, dedicated to mobile apps, which also includes specialized compute services. So, the correct answer is B. See <https://googlecloudcheatsheet.withgoogle.com/>

upvoted 4 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D. App Engine

upvoted 1 times

  **afeff** 9 months ago

Selected Answer: B

firebase is the fully managed platform for rapidly developing and deploying web and mobile applications (Paas, serverless technology)

upvoted 1 times

  **afeff** 9 months ago

Selected Answer: D

app engine for mobile and web backend applications dev
upvoted 1 times

  **ayush_1995** 10 months, 2 weeks ago

Selected Answer: B

Firebase is a fully managed platform for building iOS, Android and web apps that provides automatic data synchronization and authentication services.
upvoted 1 times

  **dccb2023** 10 months, 3 weeks ago

Even when Firebase is good for mobile developments, It is not described as a "compute solution" by google. I would trust in the questions that specifies "compute solution" and will select App Engine.
upvoted 1 times

  **aaron773** 11 months, 3 weeks ago

The keyword here is mobile app. If it wasn't mentioned that app engine would be the answer.
upvoted 1 times

  **cookieMr** 1 year ago

Firebase is a comprehensive platform provided by Google that offers a wide range of features and services specifically designed for building mobile and web applications. Firebase focuses on providing a suite of services specifically tailored for mobile and web applications.
upvoted 2 times

Your company has been using a shared facility for data storage and will be migrating to Google Cloud. One of the internal applications uses Linux custom images that need to be migrated.

Which Google Cloud product should you use to maintain the custom images?

- A. App Engine flexible environment
- B. Compute Engine
- C. App Engine standard environment
- D. Google Kubernetes Engine

Correct Answer: B

Reference:
<https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images>

Community vote distribution

B (100%)

- x6207**

Highly Voted

2 years, 3 months ago

Selected Answer: B

B Comput Engine
You can use one of the following image types: Public and Custom images.
<https://cloud.google.com/compute/docs/images>
upvoted 10 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: B

B is correct.
upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: B

B Comput Engine
upvoted 1 times
- afeff**

9 months ago

Selected Answer: B

Compute Engine to maintain the OS cumtomization
upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: B

You can use Google Compute Engine (GCE) to create and manage virtual machines (VMs) based on custom images. GCE enables you to create new VM instances and specify your custom image as the boot disk for those instances.
upvoted 2 times
- star2anand**

1 year, 3 months ago

A. App Engine flexible environment
upvoted 1 times
- cloudmind**

1 year, 6 months ago

Compute Engine
upvoted 1 times
- Animesh165**

1 year, 10 months ago

Compute Engine
upvoted 1 times

  **ergauravtaneja** 1 year, 10 months ago
Selected Answer: B
B is correct
upvoted 1 times

  **beginnercloud** 2 years, 4 months ago
Selected Answer: B
B. Compute Engine
upvoted 1 times

  **Parv7221** 2 years, 4 months ago
B Comput Engine
upvoted 1 times

  **desertlotus1211** 2 years, 4 months ago
The true correct answer is Cloud Storage....
The reason it is A is because the application is being migrated.. so you'll being using M4CE to migrate which will put it in compute engine...
upvoted 2 times

  **AbdulioUnited** 2 years, 4 months ago
App Engine is an entirely managed, serverless platform for creating and hosting web apps at scale.

Firebase is a platform created by Google for developing mobile and web apps.
upvoted 2 times

  **spvarapu** 1 year, 10 months ago
hmm..this comment incorrectly added here
upvoted 6 times

  **victory108** 2 years, 4 months ago
Selected Answer: B
B. Compute Engine
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago
Selected Answer: B
Compute Engine
upvoted 1 times

Your organization wants to migrate its data management solutions to Google Cloud because it needs to dynamically scale up or down and to run transactional SQL queries against historical data at scale. Which Google Cloud product or service should your organization use?

- A. BigQuery
- B. Cloud Bigtable
- C. Pub/Sub
- D. Cloud Spanner

Correct Answer: D

Reference:
<https://cloud.google.com/terms/services>

Cloud Spanner: Cloud Spanner is a fully-managed, mission-critical relational database service. It is designed to provide a scalable online transaction processing (OLTP) database with high availability and strong consistency at global scale.

Community vote distribution

D (88%)

12%

-   **SURJK**

Highly Voted 

 2 years, 5 months ago

Answer: D --> Q? states transactional SQL, so its Spanner
upvoted 9 times
-   **oncleled**

Highly Voted 

 1 year, 6 months ago

Selected Answer: D

Google BigQuery does not support transactions and does not allow updating of existing records. On the other hand, Google Cloud Spanner supports OLTP along with scalability and high availability.
upvoted 7 times

  **trustud**

Most Recent 

 5 days, 18 hours ago

A, analyze HISTORICAL data, not real time DB processing data. Its BQ. They want to analyze data mgt. It's not Spanner.
upvoted 1 times

  **Surek** 6 months ago

Go for D
upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

Cloud Spanner is a globally distributed, strongly consistent, and horizontally scalable database service on Google Cloud. It's designed for high availability, scalability, and transactional consistency, making it a great choice for applications that require both transactional and analytical capabilities.
upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct
upvoted 1 times

  **Chinaar** 10 months, 1 week ago

Correct Answer is A Big query
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

Cloud Spanner is a globally distributed, horizontally scalable, and strongly consistent relational database service provided by Google Cloud. Cloud Spanner is well-suited for transactional workloads with strong consistency requirements across multiple regions.

If your focus is more on performing transactional SQL queries against historical data, ensuring strong consistency, and managing relational database workloads, Cloud Spanner would be a suitable choice.

upvoted 3 times

  **star2anand** 1 year, 3 months ago

A. BigQuery

upvoted 1 times

  **omgitsele** 1 year, 2 months ago

Wrong. BigQuery is OLAP while the question asks you for OLTP. So the answer is D.

upvoted 7 times

  **PacoR** 1 year, 5 months ago

Selected Answer: D

CLOUD Spanner

upvoted 1 times

  **grzeev** 1 year, 8 months ago

Selected Answer: D

D. transactional SQL queries

upvoted 2 times

  **ttossy** 1 year, 9 months ago

Selected Answer: A

BigQuery - as the main objective is to query against historical data

upvoted 3 times

  **csrazdan** 1 year, 7 months ago

True, but not for transactional queries rather analytic queries. BigQuery is a columnar database and is not fit for transactional queries,

upvoted 3 times

  **KC_go_reply** 1 year, 3 months ago

Exactly, BigQuery is for OLAP (online analytical processing), while Spanner is meant for OLTP (online transactional processing)

upvoted 4 times

  **Jay1312** 1 year, 10 months ago

Hello, Can A -> Big Query be an option also, since you can fire SQL queries at scale?

upvoted 3 times

  **csrazdan** 1 year, 7 months ago

Yes, but for analytic queries and not transactional queries

upvoted 1 times

  **satya_136** 1 year, 10 months ago

Selected Answer: D

Spanner

upvoted 1 times

  **Anu1120** 1 year, 11 months ago

Cloud Spanner is the right answer

upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: D

Spanner is transaccional & scalable

upvoted 1 times

Your organization needs to categorize objects in a large group of static images using machine learning. Which Google Cloud product or service should your organization use?

- A. BigQuery ML
- B. AutoML Video Intelligence
- C. Cloud Vision API
- D. AutoML Tables

Correct Answer: C

Reference:
<https://cloud.google.com/vision/automl/docs/create-datasets>

Community vote distribution
C (97%)

-   **mvdn1111**

Highly Voted 

 2 years, 5 months ago

this should be C Vision API

upvoted 14 times
-   **ergauravtaneja**

Highly Voted 

 1 year, 10 months ago

Selected Answer: C

Vision API can assign labels to images and quickly classify them into millions of predefined categories. Detect objects and faces, read printed and handwritten text, and build valuable metadata into your image catalog.

upvoted 7 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: C

C is the answer

upvoted 1 times
-   **cookieMr** 1 year ago

Selected Answer: C

For categorizing objects in static images using machine learning, your organization can utilize Google Cloud Vision API. Google Cloud Vision API offers pre-trained machine learning models for image analysis and recognition tasks.

By leveraging the Google Cloud Vision API, your organization can automate the process of categorizing objects in large sets of static images, reducing the manual effort required for image analysis. The API provides a simple and scalable solution for image recognition tasks, enabling you to extract valuable insights and metadata from your image data.

upvoted 4 times
-   **star2anand** 1 year, 3 months ago

B. AutoML Video Intelligence

upvoted 2 times
-   **ajhak** 1 year, 5 months ago

Selected Answer: C

Key word: static images

upvoted 3 times
-   **Animesh165** 1 year, 10 months ago

AutoML Tables

upvoted 1 times
-   **Jay1312** 1 year, 10 months ago

I agree with those who say that AutoML Tables is the incorrect answer as it is used for structured data. Static images is unstructured data.
<https://cloud.google.com/automl-tables/docs>
AutoML Tables enables your entire team to automatically build and deploy state-of-the-art machine learning models on structured data at massively increased speed and scale.
upvoted 4 times

  **yellow28** 1 year, 11 months ago

Selected Answer: C

C is the Correct Answer
D is incorrect -> AutoML Table is used for building ML for structured data (<https://cloud.google.com/automl-tables/docs>)
upvoted 1 times

  **Wondrer** 1 year, 11 months ago

Selected Answer: C

For Static images Cloud Vision API is more appropriate..
upvoted 1 times

  **Wondrer** 1 year, 11 months ago

Answer C.
Vision API offers powerful pre-trained machine learning models through REST and RPC APIs. Assign labels to images and quickly classify them into millions of predefined categories. Detect objects and faces, read printed and handwritten text, and build valuable metadata into your image catalog.
upvoted 2 times

  **EricNgaiKahSeng** 2 years ago

Selected Answer: C

Should be C Vision API...
upvoted 1 times

  **EricNgaiKahSeng** 2 years ago

Should be C Vision API...
upvoted 1 times

  **db1290** 2 years ago

Selected Answer: C

AutoML Table is for structured data.
"AutoML Tables enables your entire team to automatically build and deploy state-of-the-art machine learning models on structured data at massively increased speed and scale."
upvoted 2 times

  **sachinxplore** 2 years ago

Selected Answer: C

Vision API
upvoted 1 times

Your organization runs all its workloads on Compute Engine virtual machine instances. Your organization has a security requirement: the virtual machines are not allowed to access the public internet. The workloads running on those virtual machines need to access BigQuery and Cloud Storage, using their publicly accessible interfaces, without violating the security requirement.

Which Google Cloud product or feature should your organization use?

- A. Identity-Aware Proxy
- B. Cloud NAT (network address translation)
- C. VPC internal load balancers
- D. Private Google Access

Correct Answer: B

Reference:

<https://cloud.google.com/nat>



High performance NAT

As a software-defined solution with no managed middle proxy, Cloud NAT's [chokepoint-free design](#) delivers high reliability, performance, and scalability.



Works with all your workloads

At Google Cloud, we are committed to being the cloud of choice for all workloads, so we designed the Cloud NAT service to work with both Compute Engine and [Google Kubernetes Engine](#) (GKE).



Scalable by design

Designed specifically for highly scalable application deployments, a single Cloud NAT gateway can be configured to handle multiple NAT IP addresses and can scale based on the size of your network, without the need for multiple NAT gateways.



Regional high availability

We understand the importance of keeping your applications up and running at all times. That's why we designed Cloud NAT with reliability in mind: even if a zone goes down, Cloud NAT stays available across the region.



Flexible IP address allocation

Choose your [NAT IP allocation](#) based on your specific requirements. Manual mode gives you full control when specifying IPs, while auto mode enables the NAT IPs to be allocated and scaled automatically, based on the number of instances.

Community vote distribution

D (88%)

12%

A_A_AB Highly Voted 2 years, 5 months ago

Selected Answer: D

Agree with fpreli, the answer is D.

According to the Google Documents >>> By default, when a Compute Engine VM lacks an external IP address assigned to its network interface, it can only send packets to other internal IP address destinations. You can allow these VMs to connect to the set of external IP addresses used by Google APIs and services by enabling Private Google Access on the subnet used by the VM's network interface.

Besides, the security rules say no Internet access while NAT provide internet access.

Google Private Access is like AWS VPC endpoint where you access GCP Public services without using Public Internet.

upvoted 28 times

harsh5kalsait Most Recent 2 months, 3 weeks ago

Selected Answer: D

Agree with the answer is D. Virtual Private Google Access.

Ref - <https://cloud.google.com/vpc/docs/private-google-access>

upvoted 1 times

- Pearl81

5 months, 1 week ago

If it is a group of VMs with a specific IP range, then NAT works better. If there are VMs which doesn't fall with IP ranges, then PGA is best choice. Remember PGA can be enabled via subnet change through a button click. So it is easy simple and free of cost too.

upvoted 2 times
- Surek

6 months ago

D is the answer

upvoted 1 times
- madcloud32

6 months ago

Selected Answer: D

D is correct.
NAT is option for securing Internet connection which is not ask. VM and other components are in cloud.

upvoted 2 times
- sivakarthick16

6 months, 1 week ago

Selected Answer: D

Private Google Access allows virtual machine instances in a VPC network to access Google Cloud services like BigQuery and Cloud Storage using internal IP addresses, without requiring public IP addresses or access to the public internet. By enabling Private Google Access, your organization can ensure that the workloads can securely access these services while adhering to the security requirement of not allowing access to the public internet.

upvoted 3 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- shares1998

8 months, 1 week ago

Selected Answer: D

D is right

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: D

Private Google Access (PGA) allows you to access Google APIs and services from your on-premises network without exposing your workloads to the public internet. This is achieved by creating a private connection between your on-premises network and Google Cloud Platform (GCP).

upvoted 2 times
- krischait

8 months, 3 weeks ago

Cloud NAT is a good option for organizations that need to allow their instances to access the internet, but want to protect them from unauthorized access. PGA is a good option for organizations that need to access Google APIs and services from their on-premises network without exposing their workloads to the public internet. Answer should be "D"

upvoted 3 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- Lufly

10 months ago

Selected Answer: D

why its not cloud nat? -> Cloud NAT: Cloud NAT is a service that allows you to access the public internet from your VMs. However, it does not allow you to control which services your VMs can access.

upvoted 3 times
- prachisri

10 months, 2 weeks ago

Selected Answer: D

The document shared previously in this discussion is very useful to understand difference between PGA and NAT. This document clearly mentions "if we know the VMs in the subnet would never connect to internet except Google's services, GPA is a better choice due to quick setup and better security." Hence, my answer is also D

upvoted 3 times
- dnagasree

10 months, 3 weeks ago

Selected Answer: D

https://medium.com/@larry_nguyen/comparing-google-private-access-and-cloud-nat-cc43ddc9ce61
As per this article PGA is sufficient to access the Google services and NAT internally used PGA

upvoted 1 times
- kmeena

10 months, 4 weeks ago

Answer D. Google documentation - <https://cloud.google.com/vpc/docs/private-google-access>
"VM instances that only have internal IP addresses (no external IP addresses) can use Private Google Access. They can reach the external IP addresses of Google APIs and services" -> since they need to access BigQuery and Cloud Storage services

upvoted 3 times

  **MinaGohari** 11 months ago

Selected Answer: D

D is correct

upvoted 1 times

  **yadusaxena** 11 months ago

Correct answer is B. Dont ask me why but

upvoted 2 times



Which Google Cloud product is designed to reduce the risks of handling personally identifiable information (PII)?

- A. Cloud Storage
- B. Google Cloud Armor
- C. Cloud Data Loss Prevention
- D. Secret Manager

Correct Answer: C

Reference:
<https://cloud.google.com/blog/products/gcp/take-charge-of-your-sensitive-data-with-the-cloud-dlp-api>

Community vote distribution

C (100%)

- AbdulioUnited**

Highly Voted

2 years, 4 months ago

Must be C.
Cloud Data Loss Prevention:
Fully managed service designed to help you discover, classify, and protect your most sensitive data.
upvoted 8 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: C

Cloud Data Loss Prevention (DLP) is a cloud-based service that helps you discover, classify, and protect sensitive data across your Google Cloud Platform (GCP) environment. DLP can help you identify and protect PII, financial information, intellectual property, and other sensitive data.
upvoted 3 times

mdsarfraz69

9 months ago

Selected Answer: C

Must be C.
upvoted 1 times

tsk9921

11 months, 2 weeks ago

C. CDLP Cloud Data Loss Prevention
upvoted 1 times

cookieMr

1 year ago

Selected Answer: C

Cloud Data Loss Prevention (DLP) is a fully managed service that helps you discover, classify, and protect sensitive data, including PII, within your Google Cloud environment. It offers a set of powerful tools and techniques to identify and protect sensitive information, helping you meet data privacy and compliance requirements.

By leveraging Cloud Data Loss Prevention (DLP), your organization can proactively identify, classify, and protect personally identifiable information (PII) within your Google Cloud environment. It helps reduce the risks associated with handling sensitive data, ensuring compliance and data privacy.
upvoted 3 times

star2anand

1 year, 3 months ago

C. Cloud Data Loss Prevention
upvoted 1 times

Monicaarg

2 years, 1 month ago

Selected Answer: C

Cloud Data Loss Prevention (DLP) detect and protect sensitive information within GCP storage repositories
upvoted 3 times

leyunjohn

2 years, 4 months ago

Selected Answer: C

C. Cloud Data Loss Prevention
upvoted 2 times

  **victory108** 2 years, 4 months ago

Selected Answer: C

C. Cloud Data Loss Prevention
upvoted 3 times

  **VijaySVK** 2 years, 5 months ago

C. Cloud Data Loss Prevention
upvoted 2 times



Your organization is migrating to Google Cloud. As part of that effort, it needs to move terabytes of data from on-premises file servers to Cloud Storage. Your organization wants the migration process to be automated and to be managed by Google. Your organization has an existing Dedicated Interconnect connection that it wants to use. Which Google Cloud product or feature should your organization use?

- A. Storage Transfer Service
- B. Migrate for Anthos
- C. BigQuery Data Transfer Service
- D. Transfer Appliance

Correct Answer: A

Reference:

<https://cloud.google.com/architecture/migration-to-google-cloud-transferring-your-large-datasets>

Where you're moving data from	Scenario	Suggested products
Another cloud provider (for example, Amazon Web Services or Microsoft Azure) to Google Cloud	—	Storage Transfer Service
Cloud Storage to Cloud Storage (two different buckets)	—	Storage Transfer Service
Your private data center to Google Cloud	Enough bandwidth to meet your project deadline for less than 1 TB of data	gsutil
Your private data center to Google Cloud	Enough bandwidth to meet your project deadline for more than 1 TB of data	Storage Transfer Service for on-premises data
Your private data center to Google Cloud	Not enough bandwidth to meet your project deadline	Transfer Appliance

Community vote distribution

A (100%)

- harsh5kalsait** 2 months, 3 weeks ago

Selected Answer: A

A is correct. Storage Transfer Service is a fully managed service that helps you transfer data from on-premises file servers to Cloud Storage. It can transfer terabytes of data quickly and securely, using your existing Dedicated Interconnect connection.

upvoted 1 times
- Azkeban** 6 months ago

Selected Answer: A

A is correct

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: A

Storage Transfer Service is a fully managed service that helps you transfer data from on-premises file servers to Cloud Storage. It can transfer terabytes of data quickly and securely, using your existing Dedicated Interconnect connection.

upvoted 2 times
- mdsarfraz69** 9 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **cookieMr** 1 year ago

The Storage Transfer Service is a managed service provided by Google Cloud that simplifies the process of transferring data to and from Cloud Storage.

The Storage Transfer Service simplifies the data migration process by automating many of the steps involved. It provides a reliable and scalable solution for migrating large amounts of data from on-premises file servers to Cloud Storage. However, it's important to note that the Storage Transfer Service relies on internet connectivity for data transfer, so consider your network bandwidth and transfer speed requirements when using this service.

upvoted 3 times

  **star2anand** 1 year, 3 months ago

A. Storage Transfer Service

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Per Google docs, Storage Transfer Service allows for "moving data cloud-to-cloud, on-prem to cloud, and between cloud storage buckets."
<https://cloud.google.com/storage-transfer-service>

upvoted 1 times

  **MadCode** 1 year, 10 months ago

Selected Answer: A

A, se usa para mover mucho tamaño de fichero

upvoted 1 times

  **Vedjha** 1 year, 10 months ago

Answer is A, Storage Transfer Services is used upto 10 TBs and depends on n/w speed. Storage Appliances used if file size is > 10 TB and transfer time takes > 1 week

upvoted 2 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: A

Storage Transfer services. require on line conection.

upvoted 4 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: A

A. Storage Transfer Service

upvoted 3 times

  **Parv7221** 2 years, 4 months ago

A Storage transfer services

upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: A

A. Storage Transfer Service

upvoted 2 times

  **anirban7172** 2 years, 5 months ago

Selected Answer: A

Answer A: It is storage transfer Service as the customer wants to use its existing dedicated Internet.

upvoted 4 times

Your organization needs to analyze data in order to gather insights into its daily operations. You only want to pay for the data you store and the queries you perform. Which Google Cloud product should your organization choose for its data analytics warehouse?

- A. Cloud SQL
- B. Dataproc
- C. Cloud Spanner
- D. BigQuery

Correct Answer: D

Community vote distribution

D (100%)

- harsh5kalsait

2 months, 3 weeks ago

Selected Answer: D

Option D. BigQuery is the correct.
1. needs to analyze data
2.in order to gather insights into its daily operations.
3. You only want to pay for the data you store and the queries you perform.
4.its data analytics warehouse.
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: D

D is correct.
upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: D

D is correct
upvoted 1 times
- tsk9921

11 months, 2 weeks ago

D. BigQuery - OLAP - Analytical processing
upvoted 1 times
- cookieMr

1 year ago

Selected Answer: D

Google BigQuery is a fully managed, serverless data warehouse that offers high-performance analytics on large datasets. It is designed to handle massive amounts of data and provides a scalable solution for your data analytics needs.

By leveraging Google BigQuery, your organization can efficiently analyze large volumes of data, gain actionable insights, and only pay for the resources and queries consumed, making it a cost-effective and scalable solution for your data analytics warehouse needs.
upvoted 4 times
- star2anand

1 year, 3 months ago

D. BigQuery
upvoted 1 times
- Sujeeshtk

1 year, 10 months ago

i have read somewhere that Bigquery is not recommended for day to day data analysis. So the alternative option available is Cloud SQL
upvoted 3 times
- KC_go_reply

1 year, 3 months ago

Cloud SQL is *not* recommended for analytical workloads. Imagine having a high volume of transactions on your database, and then you want to run a large analytical query that has to read through TBs of data. Sike!
upvoted 1 times
- Kathyhuang

1 year, 10 months ago

D,BigQuery

upvoted 1 times

  **MadCode** 1 year, 10 months ago

Selected Answer: D

D, BigQuery sirve para analizar datos ademas es serverless. Data warehouse == BigQuery

upvoted 2 times

  **gdntcs** 1 year, 10 months ago

Selected Answer: D

Bigquery is analytical and serverless

upvoted 3 times

  **toni90ns** 1 year, 10 months ago

Selected Answer: D

Answer is

upvoted 1 times

  **FahmiPahlevi** 1 year, 11 months ago

Cloud SQL stored structured data base, you'll need BigQuery to analyze multiple data

upvoted 2 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: D

insights is the key.

upvoted 3 times

  **RamNBamal** 2 years, 1 month ago

Cloud SQL > transactional

BigQuery > analytical

upvoted 2 times

  **albert_kuo** 2 years, 2 months ago

Selected Answer: D

BigQuery is an enterprise data warehouse for large amounts of relational structured data

upvoted 1 times

  **yogi1508** 2 years, 2 months ago

Pls change suggested solution to D. it cannot be A

upvoted 3 times

Your organization wants to run a container-based application on Google Cloud. This application is expected to increase in complexity. You have a security need for fine-grained control of traffic between the containers. You also have an operational need to exercise fine-grained control over the application's scaling policies.

What Google Cloud product or feature should your organization use?

- A. Google Kubernetes Engine cluster
- B. App Engine
- C. Cloud Run
- D. Compute Engine virtual machines

Correct Answer: A

Reference:

<https://cloud.google.com/blog/products/serverless/cloud-run-gets-enterprise-grade-network-security-with-vpc-sc>

Community vote distribution

A (97%)

- fpeli

Highly Voted

2 years, 5 months ago

I would say A. GKE seems a better fit since the requirement is for "security need for fine-grained control of traffic between the containers" and "fine-grained control over scaling policies". Such level of control is easier on GKE that Cloud Run.

upvoted 21 times
- kimkimkimkimkim

Highly Voted

2 years, 5 months ago

Selected Answer: A

I think the answer should be A because it says that there is a need for "fine-grained control".
Check out this brief comparison:
<https://cloud.google.com/blog/products/containers-kubernetes/when-to-use-google-kubernetes-engine-vs-cloud-run-for-containers>

upvoted 11 times
- manishmakhija

2 years, 1 month ago

Agreed .. GKE is for much much customization and control.. Cloud run is for Easy containerization

upvoted 3 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct

upvoted 1 times
- hireshgupta

11 months, 1 week ago

It should be GKE

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: A

Google Kubernetes Engine is a managed Kubernetes service that provides a secure, reliable, and scalable environment for deploying and managing containerized applications.

By using Google Kubernetes Engine, your organization can deploy and manage container-based applications with fine-grained control over traffic

and scaling policies. GKE offers a secure, scalable, and reliable platform for running your applications, while leveraging the power and flexibility of Kubernetes for container orchestration.

upvoted 2 times

  **Rishabh_0108** 1 year, 3 months ago

Selected Answer: A

A. GKE is right
upvoted 1 times

  **star2anand** 1 year, 3 months ago

Google Kubernetes Engine cluster
upvoted 1 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: A

Container based, and expected to increase in complexity, hence GKE.
upvoted 1 times

  **Micheal_Scott** 1 year, 8 months ago

Why not App Engine ?
upvoted 1 times

  **alfantastic** 1 year, 5 months ago

It is not built to handle containers at first
upvoted 1 times

  **grzeev** 1 year, 8 months ago

Selected Answer: A

A. <https://cloud.google.com/blog/products/containers-kubernetes/when-to-use-google-kubernetes-engine-vs-cloud-run-for-containers>
upvoted 1 times

  **Rakesh1976** 1 year, 8 months ago

GKE is right answer
upvoted 1 times

  **ergauravtaneja** 1 year, 10 months ago

Selected Answer: A

A

GKE is direct answer
upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: A

"fine-grained control of traffic between containers"
upvoted 2 times

  **Hikmatyar** 2 years, 2 months ago

GKE, as stated required for "fine-grained control of traffic between containers"
upvoted 1 times

Which Google Cloud product or feature makes specific recommendations based on security risks and compliance violations?

- A. Google Cloud firewalls
- B. Security Command Center
- C. Cloud Deployment Manager
- D. Google Cloud Armor

Correct Answer: B

Reference:
<https://cloud.google.com/security-command-center>

Asset discovery and inventory

Discover and view your assets in near-real time across App Engine, BigQuery, Cloud SQL, Cloud Storage, Compute Engine, Cloud Identity and Access Management, Google Kubernetes Engine, and more. Review historical discovery scans to identify new, modified, or deleted assets.

Threat prevention

Understand the security state of your Google Cloud assets. Uncover common web application vulnerabilities such as cross-site scripting or outdated libraries in your web applications running on App Engine, GKE, and Compute Engine. Quickly resolve misconfigurations by clicking directly on the impacted resource and following the prescribed steps on how to fix it.

Threat detection

Detect threats using logs running in Google Cloud at scale. Detect some of the most common container attacks, including suspicious binary, suspicious library, and reverse shell.

Community vote distribution

B (100%)

- harsh5kalsait** 2 months, 3 weeks ago

Selected Answer: B

B. Security Command Center
Requirement: security risks and compliance violations that's why
upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: B

Security Command Center (SCC) is a security information and event management (SIEM) solution that provides insights into your security posture across Google Cloud and hybrid environments. SCC collects security signals from a variety of sources, including Google Cloud audit logs, system logs, and network traffic. It then analyzes these signals to identify security threats and compliance violations.
upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: B

Google Cloud Security Command Center (Cloud SCC) is a comprehensive security management and data risk assessment platform.

By using Google Cloud Security Command Center, your organization can gain valuable insights into the security risks and compliance violations in your Google Cloud environment. It provides actionable recommendations to improve your security posture, mitigate vulnerabilities, and ensure compliance with industry standards.

upvoted 3 times

  **star2anand** 1 year, 3 months ago

B. Security Command Center

upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: B

<https://cloud.google.com/security-command-center>

upvoted 1 times

  **albert_kuo** 2 years, 2 months ago

Selected Answer: B

<https://cloud.google.com/security-command-center>

upvoted 1 times

  **RonRon512** 2 years, 3 months ago

Selected Answer: B

Security command center

upvoted 1 times

  **beginnercloud** 2 years, 3 months ago

Selected Answer: B

Security Command Center

upvoted 1 times

  **rajvivan** 2 years, 4 months ago

Option B

upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: B

B. Security Command Center

upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: B

B. Security Command Center

upvoted 1 times

  **SURJK** 2 years, 5 months ago

Answer: B

upvoted 2 times

Which Google Cloud product provides a consistent platform for multi-cloud application deployments and extends other Google Cloud services to your organization's environment?

- A. Google Kubernetes Engine
- B. Virtual Public Cloud
- C. Compute Engine
- D. Anthos

Correct Answer: D

Community vote distribution

D (95%)

5%

-   **cookieMr** Highly Voted  1 year ago

Selected Answer: D

Anthos is an open hybrid and multi-cloud platform offered by Google Cloud. It enables you to build, deploy, and manage applications consistently across multiple clouds, including Google Cloud, on-premises environments, and other public cloud providers. Anthos brings Google Cloud's services and capabilities to your organization's infrastructure, providing a unified and consistent experience across different environments.

By using Anthos, your organization can leverage Google Cloud's services and capabilities across multiple cloud environments. It provides a consistent and unified platform for application development, deployment, and management, regardless of the underlying infrastructure.

upvoted 5 times
-   **Lakshvenkat** Most Recent  6 months ago

D: Anthos is a hybrid and multi-cloud platform provided by Google Cloud. It enables you to manage and deploy applications consistently across on-premises, multi-cloud, and Google Cloud environments.

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: D

Anthos is a hybrid and multi-cloud platform provided by Google Cloud. It enables you to manage and deploy applications consistently across on-premises, multi-cloud, and Google Cloud environments.

upvoted 2 times
-   **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct

upvoted 1 times
-   **hireshgupt** 11 months ago

Selected Answer: D

Anthos is a Google Cloud product that provides a consistent platform for multi-cloud application deployments and extends other Google Cloud services to your organization's environment.

upvoted 3 times
-   **star2anand** 1 year, 3 months ago

B. Virtual Public Cloud

upvoted 1 times
-   **manashbaruah** 1 year, 6 months ago

Selected Answer: D

Keywords: multi-cloud application deployments

upvoted 2 times
-   **jdeal** 1 year, 6 months ago

multi-cloud is key word. Answer is D..

upvoted 1 times
-   **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Per Google docs, "Anthos is the leading cloud-centric container platform to run modern apps anywhere consistently at scale."

<https://cloud.google.com/anthos>

Below video is from Google Cloud Tech, titled, How to run multicloud apps consistently at scale with Anthos

https://www.youtube.com/watch?v=UZmSM_3LJBE

upvoted 1 times

  **Nakon** 1 year, 7 months ago

Selected Answer: D

Multi cloud -> Anthos

upvoted 3 times

  **ravikiran37** 1 year, 8 months ago

Selected Answer: D

MULTI --Anthod

upvoted 1 times

  **TienHoang** 1 year, 9 months ago

Selected Answer: D

Multi cloud migration -> Anthos

upvoted 1 times

  **kinoko1330** 1 year, 9 months ago

Selected Answer: D

anthos

upvoted 1 times

Your organization is developing an application that will manage payments and online bank accounts located around the world. The most critical requirement for your database is that each transaction is handled consistently. Your organization anticipates almost unlimited growth in the amount of data stored.

Which Google Cloud product should your organization choose?

- A. Cloud SQL
- B. Cloud Storage
- C. Firestore
- D. Cloud Spanner

Correct Answer: C

Features of Cloud Spanner -

Reference:

<https://k21academy.com/google-cloud/cloud-sql-vs-cloud-spanner/>

Google Cloud SQL is a fully managed service offered by Google Cloud Platform. **Google Cloud SQL is a MySQL database inside Google Cloud.** There is no need to install, maintain and create admin accounts because it is fully managed by Google Cloud. It helps you create, modify, configure and utilize a relational database, same as MySQL.

Google sends constant updates and adds new features to its services to fulfill the business requirements of its users.

Let's have a look at the improvements done in Google Cloud SQL.

- Google deliberately increased the storage space to **100 Gigabytes** which was 10 Gigabytes earlier.
- The modified version is loaded with the capacity of **16GB RAM** to run instances hassle-free.
- Increased RAM has helped users to keep **four times more cache** than earlier.
- Now Google provides **both replicated and non-replicated databases.**

Community vote distribution

D (90%)

10%

- vinibisineli

Highly Voted

2 years, 5 months ago

Selected Answer: D

CloudSpanner is the correct answer. Firestore is for mobile applications developed from Firebase, GCS is not a Database and Cloud SQL is to simpl for this amout of data.

upvoted 26 times
- omgitsele

1 year, 2 months ago

Good answer but Firestore is not only or mobile apps. The reason it is not the correct answer is the limitations analyzed here. Essentially it does not cover the unlimited growth. https://firebase.google.com/docs/firestore/quotas#collections_documents_and_fields

upvoted 4 times
- Kessel

Most Recent

3 weeks, 3 days ago

Selected Answer: D

"...an application that will manage payments and online bank accounts located around the world." This gives us several hints: it's a transactional workload, because it's for bank and managing payments, and it's almost certainly a SQL database we are talking about. Finally it needs to scale horizontally, because it says "around the world", global expansion. Only Spanner satisfies all conditions. D is Correct. (Also Firestore is dor noSQL document databases, mainly for web apps, not a good solution at all)

upvoted 1 times
- PanosPeris

4 weeks, 1 day ago

Selected Answer: D

Firestore is a noSQL database, not a good choice when transactional consistency is a requirement

upvoted 2 times
- harsh5kalsait

2 months, 3 weeks ago

Selected Answer: D

- D. Cloud Spanner is the correct answer.
1. Globally distributed,
 2. horizontally scalable,
 3. strongly consistent relational database service,
 4. ensuring consistent transaction handling across multiple regions,
 5. Transactional, mission-critical, and unlimited growth.

upvoted 2 times

  **Radhouanev12** 4 months ago

Selected Answer: D

Cloud Spanner

upvoted 1 times

  **sivakarthick16** 5 months, 1 week ago

Selected Answer: D

Cloud Spanner is a globally distributed, horizontally scalable, and strongly consistent relational database service offered by Google Cloud. It provides the ability to scale both storage and compute resources as needed, ensuring consistent transaction handling across multiple regions. This makes it an ideal choice for applications that require strong consistency and high scalability.

While options A (Cloud SQL), B (Cloud Storage), and C (Firestore) are also valuable Google Cloud products, they may not meet the specific requirements mentioned in the context as effectively as Cloud Spanner. Cloud SQL is a managed relational database service, but it may not provide the same level of scalability and global consistency as Cloud Spanner. Cloud Storage is a scalable object storage solution, and Firestore is a NoSQL document database, but they may not offer the same level of transaction consistency required for managing payments and online bank accounts.

upvoted 4 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

  **shares1998** 8 months, 1 week ago

Selected Answer: D

cloud spanner is accurate

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

Cloud Spanner is a fully managed, mission-critical, relational database service that offers transactional consistency at global scale, schemas, SQL (ANSI 2011 with extensions), and automatic, synchronous replication for high availability.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times

  **afeff** 9 months ago

CloudSpanner is the correct answer. Cloud Spanner is well-suited for transactional workloads with strong consistency requirements across multiple regions. If your focus is more on performing transactional SQL queries against historical data, ensuring strong consistency, and managing relational database workloads, Cloud Spanner would be a suitable choice.

upvoted 2 times

  **Lufly** 10 months ago

Selected Answer: D

The other options are not as suitable for this purpose.

Cloud SQL: Cloud SQL is a managed database service that offers a variety of relational database engines, including MySQL, PostgreSQL, and SQL Server. However, Cloud SQL does not offer transactional consistency at global scale.

Cloud Storage: Cloud Storage is a object storage service that is designed for storing large amounts of unstructured data. It is not a relational database and does not offer transactional consistency.

Firestore: Firestore is a NoSQL document database that is designed for mobile and web applications. It is not a relational database and does not offer transactional consistency.

upvoted 4 times

  **Sam420** 10 months, 3 weeks ago

Cloud Spanner as it is database recommended for any banking or gaming related industry where the volumes of data is high

upvoted 3 times

  **Luizcjr** 1 year ago

Cloud Spanner

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: C

Cloud Spanner is a fully managed, globally distributed relational database service provided by Google Cloud. It combines the benefits of traditional relational databases, such as ACID transactions and SQL queries, with the scalability and high availability of a distributed system. It is specifically

designed to meet the requirements of handling globally consistent transactions and scaling to handle large amounts of data.

With its strong consistency guarantees, scalability, and global distribution capabilities, Cloud Spanner is well-suited for applications that require consistent transaction handling and anticipate significant data growth. It provides a robust and reliable database solution for managing payments and online bank accounts located worldwide.

upvoted 4 times

  **slrjgd** 1 year, 2 months ago

Thought the correct answer is Cloud spanner. But when clicked on reveal solution it says Answer is C which is Firestore ... And again the in the explanation boxed opened , Correct Answer is mentioned as C , but explains cloud spanner.

upvoted 1 times

  **MHx1** 1 year, 3 months ago

Selected Answer: D

CloudSpanner is fully managed relationship database. A proprietary relational database designed for scale.

(Firestore is No-SQL document database)

upvoted 1 times

Your organization wants an economical solution to store data such as files, graphical images, and videos and to access and share them securely. Which Google Cloud product or service should your organization use?

- A. Cloud Storage
- B. Cloud SQL
- C. Cloud Spanner
- D. BigQuery

Correct Answer: A

Reference:

<https://www.myamberlife.com/learn/top-ways-to-store-your-digital-files/>

Desktop Storage

Despite many external solutions for digital files, some people still store their photos, videos, and content files on their desktop or laptop. The only problem with this method is that your computer can quickly become cluttered with thousands of files. It slows your prized piece of hardware (computer) down.

When you want to find a digital file you probably *expect* that file to come flying up on your screen in an instant. Yet — anyone who keeps a lot of photos on a computer knows it can take minutes, sometimes hours, to find one – even if you keep it on your desktop. It’s just not all that convenient to store things this way. Most importantly, just storing these digital files on a desktop leaves them vulnerable to viruses, damage, or theft. Folks who rely on this also generally don’t have a back-up plan.

Community vote distribution

A (100%)

- A_A_AB

Highly Voted

2 years, 5 months ago

Selected Answer: A

If you're familiar with AWS, Google Storage is GCP's version of AWS Simple Storage Service (S3) and an S3 bucket would be equivalent to a Google Storage bucket across the two clouds.

upvoted 8 times
- preetham141

Most Recent

6 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: A

Google Cloud Storage is a scalable and durable object storage service that allows you to store and retrieve any amount of data from anywhere on the web. It provides a cost-effective solution for storing large volumes of unstructured data while ensuring high durability, availability, and security

By using Google Cloud Storage, your organization can store and manage large amounts of data economically, while ensuring secure access and sharing of files, images, and videos. It provides the scalability, durability, and cost-efficiency needed for such data storage requirements.

upvoted 4 times

  **saifelmadani** 1 year, 2 months ago

Cloud Storage is the correct answer

upvoted 1 times

  **star2anand** 1 year, 3 months ago

A. Cloud Storage

upvoted 1 times

  **manashbaruah** 1 year, 6 months ago

Selected Answer: A

Keywords: store data, economically

upvoted 3 times

  **Vazhma** 2 years ago

Selected Answer: A

Cloud storage

upvoted 2 times

  **beardedlensman** 2 years, 2 months ago

Selected Answer: A

Cloud Storage is the answer.

upvoted 2 times

  **Yashg8** 2 years, 2 months ago

Selected Answer: A

Cloud Storage

upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: A

A. Cloud Storage

upvoted 3 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: A

A. Cloud Storage

upvoted 1 times

  **astood** 2 years, 5 months ago

Yes, A is the answer

upvoted 2 times

Your organization wants to predict the behavior of visitors to its public website. To do that, you have decided to build a machine learning model. Your team has database-related skills but only basic machine learning skills, and would like to use those database skills. Which Google Cloud product or feature should your organization choose?

- A. BigQuery ML
- B. LookML
- C. TensorFlow
- D. Cloud SQL

Correct Answer: A

Reference:
<https://cloud.google.com/architecture/predicting-customer-propensity-to-buy>

Predicting customer propensity to buy by using BigQuery ML and AI Platform

Send feedback

Learn how to build a system to predict customer propensity to purchase by using BigQuery ML and AI Platform.

You can use a propensity to purchase system to predict customers who are most likely to make a purchase, so that you can personalize communications with them. Use online predictions to take real-time action based on user behavior on your website, or batch predictions to inform less time-sensitive communications like email.

This tutorial shows you how to create a [logistic regression](#) model to determine whether a customer will make a purchase. This type of model is used because it is good for evaluating the probability of an outcome. The model evaluates metrics that reflect customer behavior on a website, and assigns the customer a probability to purchase value between 0 and 1 based on this data. The model then sets a label indicating "likely to purchase" for any customer with a probability

Community vote distribution

A (100%)

- albert_kuo

Highly Voted

2 years, 2 months ago

Selected Answer: A

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries. BigQuery ML democratizes machine learning by letting SQL practitioners build models using existing SQL tools and skills. BigQuery ML increases development speed by eliminating the need to move data.

<https://cloud.google.com/bigquery-ml/docs/introduction>

upvoted 10 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: A

BigQuery ML is a machine learning service that makes it easy for anyone to create and train machine learning models, even if they don't have a lot of machine learning experience. BigQuery ML is built on top of BigQuery, so you can use your existing database skills to train and deploy machine learning models.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: A

BigQuery ML is a feature of Google BigQuery that allows you to build and deploy machine learning models using SQL queries directly within the BigQuery environment. It enables data analysts and database developers to utilize their existing SQL skills for machine learning tasks without requiring advanced machine learning expertise.

By using BigQuery ML, your team can leverage their database skills and work within the familiar SQL environment to build and deploy machine learning models for predicting visitor behavior. This eliminates the need for advanced machine learning knowledge and enables a smooth transition from data analysis to machine learning tasks.

upvoted 4 times

  **star2anand** 1 year, 3 months ago

A. BigQuery ML
upvoted 1 times

  **manashbaruah** 1 year, 6 months ago

Selected Answer: A

Keywords: predict, ML model, database skills
upvoted 3 times

  **Pou1ze** 1 year, 6 months ago

Selected Answer: A

BigQuery have some ML functions
upvoted 1 times

  **Vazhma** 2 years ago

Selected Answer: A

Build model
upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: A

build models using queries.
upvoted 2 times

  **RJchamp** 2 years, 1 month ago

BigQuery ML
upvoted 1 times

  **beginnercloud** 2 years, 3 months ago

Selected Answer: A

BigQuery ML
upvoted 2 times

  **victory108** 2 years, 4 months ago

Selected Answer: A

A. BigQuery ML
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: A

BigQuery ML
upvoted 1 times

  **astood** 2 years, 5 months ago

A is the answer
upvoted 3 times

Your organization needs to restrict access to a Cloud Storage bucket. Only employees who are based in Canada should be allowed to view the contents.

What is the most effective and efficient way to satisfy this requirement?

- A. Deploy the Cloud Storage bucket to a Google Cloud region in Canada
- B. Configure Google Cloud Armor to allow access to the bucket only from IP addresses based in Canada
- C. Give each employee who is based in Canada access to the bucket
- D. Create a group consisting of all Canada-based employees, and give the group access to the bucket

Correct Answer: B

Community vote distribution

D (87%)

13%

- Halimb

Highly Voted

2 years, 4 months ago

Selected Answer: D

Correct answer is D. Question is tricky, but it says "based" in Canada. That is not the same as restricting access to "from Canada". An employee can for instance be based in Canada, but access the services while on business trip to Singapore.

upvoted 35 times
- PanosPeris

Most Recent

4 weeks, 1 day ago

Selected Answer: D

IP is a poor way to restrict access. Employees based in Canada could be working from anywhere.

upvoted 1 times
- Wr5050

2 months, 3 weeks ago

Selected Answer: D

I found an excellent explanation on this site, the questions seem to be verified there
<https://techcertificationhelp.com/cloud-digital-leader/only-employees-who-are-based-in-canada-should-be-allowed-to-view-the-c>

upvoted 3 times
- MMeena

3 months, 3 weeks ago

Selected Answer: B

I think 'B' may be the option. question says "Only employees who are based in Canada" and considering the Google's security policy of 'Least Privilege Access' , option D, will give access to all Canada Employees, where they need to have access or not, which may be a security threat

upvoted 1 times
- Pearl81

5 months, 1 week ago

If read carefully, question is granting access for "employees based in Canada" and not "employees in Canada". This makes a lot of different. Correc answer is D.

upvoted 1 times
- Guru4Cloud

6 months, 1 week ago

Selected Answer: D

Imagine a lock on your bucket. You want only Canadian employees to have keys. Here's the easiest way:

Make a key club: Create a group called "Canada Keys".
Add all Canadian employees: Give everyone in that group a key.
Keep outsiders out: No key, no entry to the bucket!
This way, you manage one key club instead of many individual keys, making it easier to add/remove people and keeping your bucket secure. Clear as day?

upvoted 3 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: D

D is correct.

upvoted 1 times
- Giuliano72

9 months ago

Selected Answer: D

D is correct
upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct
upvoted 1 times

  **oezgan** 9 months ago

Although a bit old I found this on Serverfault:
"But, IP deny list/allow list for HTTP(S) Load Balancing is not supported for Cloud Storage backends. See Security Policy Concepts - Restrictions for details."

Thus, the answer must be D. (I Hope).
Reference:
<https://serverfault.com/questions/992666/using-google-cloud-armor-to-block-requests-to-google-cloud-storage>
upvoted 2 times

  **hireshgupt** 11 months ago

Selected Answer: D

Option D is the most effective and efficient way to restrict access to the bucket. Creating a group consisting of all Canada-based employees and giving the group access to the bucket will allow you to easily manage access to the bucket. You can add or remove employees from the group as needed, and you can give the group different levels of access to the bucket.
upvoted 3 times

  **ihavenonickname** 11 months, 4 weeks ago

Selected Answer: D

ip restrictions can be bypassed
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

To restrict access to a Cloud Storage bucket and ensure that only employees based in Canada can view its contents, you can use Cloud Identity and Access Management (Cloud IAM) in combination with Identity-Aware Proxy (IAP).

By combining Cloud IAM and IAP, you can enforce fine-grained access control to the Cloud Storage bucket. Only employees based in Canada, as defined in the Cloud IAM roles and IAP access policy, will be able to view the bucket's contents. This provides an effective and efficient way to satisfy the access restriction requirement while leveraging Google Cloud's built-in identity and access management capabilities.
upvoted 3 times

  **Laura93** 1 year ago

The answer should be B. In this case, you can create a rule that allows access to the Cloud Storage bucket only from IP addresses based in Canada. This will ensure that only employees who are based in Canada will be able to access the bucket. D is not the most effective way to restrict access to the bucket. If an employee is added to the group, they would be able to access the bucket, even if they are not based in Canada.
upvoted 2 times

  **MBNelo** 1 year, 1 month ago

Selected Answer: D

"based", not "in"
upvoted 1 times

  **omgitsele** 1 year, 2 months ago

Selected Answer: B

Correct answer is B. Even if a user based in Canada travels abroad, you do NOT want them to be able to have access. The question is badly phrased but essentially it says "restrict any access outside of Canada" thus correct answer is B.
upvoted 1 times

Your organization is moving an application to Google Cloud. As part of that effort, it needs to migrate the application's working database from another cloud provider to Cloud SQL. The database runs on the MySQL engine. The migration must cause minimal disruption to users. Data must be secured while in transit.

Which should your organization use?

- A. BigQuery Data Transfer Service
- B. MySQL batch insert
- C. Database Migration Service
- D. Cloud Composer

Correct Answer: C

Reference:

<https://aws.amazon.com/dms/>

Community vote distribution

C (100%)

- albert_kuo

Highly Voted

2 years, 2 months ago

Selected Answer: C

Database Migration Service Can be used only to perform the migration of PostgreSQL and MySQL databases. SQL Server database migration is currently not supported and is coming soon.

upvoted 8 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: C

DMS can migrate your database with minimal disruption to users. It uses a continuous replication process to migrate your data to Cloud SQL. This means that your application can continue to read and write data to the database during the migration.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: C

To securely migrate your application's working database from another cloud provider to Cloud SQL on Google Cloud, while causing minimal disruption to users and ensuring data security in transit, you can use Cloud SQL for MySQL in combination with Database Migration Service (DMS)

Use the Database Migration Service to perform the actual migration of data from the source database to Cloud SQL. DMS provides a managed service for database migration with minimal downtime and provides continuous replication during the migration process.

upvoted 2 times
- star2anand

1 year, 3 months ago

C. Database Migration Service

upvoted 2 times
- manashbaruah

1 year, 6 months ago

Selected Answer: C

Keywords: minimal disruption to users, data must be secured

upvoted 2 times
- 0010111

1 year, 8 months ago

C is correct. Reference in solution is linking AWS content lol
upvoted 4 times

  **Govindaraj** 1 year, 11 months ago

Selected Answer: C

<https://cloud.google.com/database-migration>
Database Migration Service supports multiple secure, private connectivity methods to protect your data in transit. Once migrated, all data is encrypted by default, and Cloud SQL provides multiple layers of security to meet even the most stringent requirements.
upvoted 4 times

  **KouFa** 2 years, 1 month ago

Selected Answer: C

Database Migration Service
upvoted 1 times

  **Monicaarg** 2 years, 1 month ago

Selected Answer: C

C. Database Migration Service
upvoted 1 times

  **RJchamp** 2 years, 1 month ago

C is correct
upvoted 1 times

  **victory108** 2 years, 4 months ago

Selected Answer: C

C. Database Migration Service
upvoted 3 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. Database Migration Service
upvoted 1 times

  **Kcaher** 2 years, 5 months ago

C: Database migration service
<https://cloud.google.com/database-migration>
upvoted 3 times

  **astood** 2 years, 5 months ago

C: Database migration service
upvoted 2 times

Your organization is developing and deploying an application on Google Cloud. Tracking your Google Cloud spending needs to stay as simple as possible.

What should you do to ensure that workloads in the development environment are fully isolated from production workloads?

- A. Apply a unique tag to development resources
- B. Associate the development resources with their own network
- C. Associate the development resources with their own billing account
- D. Put the development resources in their own project

Correct Answer: D

Community vote distribution

D (95%)

5%

- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: D

D is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: D

D is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: D

Create separate projects for your development and production environments. Projects provide a logical isolation boundary and allow you to manage resources independently.

upvoted 3 times
- star2anand

1 year, 3 months ago

A. Apply a unique tag to development resources

upvoted 3 times
- manashbaruah

1 year, 6 months ago

Selected Answer: D

Keywords: fully isolated, as simple as possible

upvoted 2 times
- Huntishi

1 year, 6 months ago

Selected Answer: D

Development resource should be in different project!!

upvoted 2 times
- Vish_Iyer

1 year, 7 months ago

How does D address the "Tracking your Google Cloud spending needs to stay as simple as possible" aspect of the question?

upvoted 4 times
- alfantastic

1 year, 5 months ago

you can add alerts and quotas and project level

upvoted 1 times
- theKingslayer

1 year, 6 months ago

maybe cause you can monitor spending by project!
upvoted 2 times

  **jaxclain** 1 year, 9 months ago

Selected Answer: D

D is the correct answer, A, B and C makes no sense at all.
upvoted 1 times

  **bnikunj** 1 year, 9 months ago

Selected Answer: B

It should be B, development resources should be create din their own network for isolation
upvoted 1 times

  **omgitsele** 1 year, 2 months ago

"as simple as possible" so not B.
upvoted 1 times

  **sedado77** 2 years, 1 month ago

Selected Answer: D

D is right
upvoted 2 times

  **AJohnny** 2 years, 1 month ago

Selected Answer: D

Development resource should be in different project!!
upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: D

D. Put the development resources in their own project
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: D

D. Put the development resources in their own project
upvoted 1 times

  **mintupintu** 2 years, 5 months ago

D is the answer.
upvoted 2 times

  **astood** 2 years, 5 months ago

D is the answer. Put it in their own project
upvoted 2 times

Your company is running the majority of its workloads in a co-located data center. The workloads are running on virtual machines (VMs) on top of a hypervisor and use either Linux or Windows server editions. As part of your company's transformation strategy, you need to modernize workloads as much as possible by adopting cloud-native technologies. You need to migrate the workloads into Google Cloud. What should you do?

- A. Export the VMs into VMDK format, and import them into Compute Engine
- B. Export the VMs into VMDK format, and import them into Google Cloud VMware Engine
- C. Migrate the workloads using Migrate for Compute Engine
- D. Migrate the workloads using Migrate for Anthos

Correct Answer: D

Community vote distribution

D (60%)

C (40%)

- fpreli

Highly Voted

2 years, 5 months ago

Agree with C, Compute Engine. From GCP documentation "For the workloads that will benefit from containers, Migrate for Anthos and GKE deliver a fast, smooth path to modernization. For other workloads that are better suited as a VM, simply move them as is with Migrate for Compute Engine"

upvoted 22 times
- Jack456

8 months, 3 weeks ago

but as per i saw a official video google virtual machine migration need vSphere.
https://cloud.google.com/migrate/virtual-machines?utm_source=bing&utm_medium=cpc&utm_campaign=japac-IN-all-en-dr-bkws-all-pkws-trial-b-dr-1009882&utm_content=text-ad-none-none-DEV_c-CRE_-ADGP_Hybrid%20%7C%20BKWS%20-%20PHR%20%7C%20Txt%20~%20Compute%20~%20Migrate%20for%20Compute%20Engine_migrate%20for%20compute%20engine-general%20-%20Products-KWID_43700071936703088-kwd-74698340556763%3Aloc-90&userloc_149209-network_o&utm_term=KW_google%20migrate%20for%20compute%20engine&gclid=3967b163c3031bcea7a9d442d7d20982&gclsrc=3p.ds

upvoted 2 times
- fpreli

2 years, 5 months ago

But, right, it says "cloud-native" and "modernize", so probably Migrate for Anthos is correct.

upvoted 7 times
- Raj_sai

1 year, 9 months ago

You can directly migrate VM loads to Anthos. It is a 2 step process

upvoted 6 times
- x6207

Highly Voted

2 years, 3 months ago

Selected Answer: D

I would choose D. Migrate the workloads using Migrate for Anthos, because it is similar to Question #43 and this is C. Migrate the workloads to Compute Engine.
I agree with Tom666

upvoted 11 times
- shanwford

Most Recent

3 weeks ago

Selected Answer: C

This Use Case fits very well with Compute engine - why using Anthos here (its core technology developed by Google empowering enterprise container platforms)? If anthos is solution, then you would have to be very creative in interpreting this question

upvoted 1 times
- rathm

2 months, 3 weeks ago

Selected answer C

upvoted 1 times
- Thesalesguy

5 months, 2 weeks ago

Selected Answer: C

Migrate for Anthos is designed for migrating containerized applications, not VMs. It wouldn't be suitable for your current VM-based workloads. Therefore, Migrate for Compute Engine is the most suitable choice for your scenario. It enables you to migrate your VMs to Google Cloud while also modernizing them with cloud-native technologies, minimizing disruption and automating the process for a smooth and efficient transition.

upvoted 2 times
- soumya763

6 months ago

Selected Answer: C

Migrate by Compute Engine

upvoted 1 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: C

By using Migrate for Compute Engine, you can seamlessly migrate your Linux or Windows server edition VMs from your co-located data center to Google Cloud. This service provides automated migration capabilities, including replication, testing, and cutover, to ensure a smooth transition of your workloads. It also offers features like continuous data replication and minimal downtime during the migration process.

Option D (migrating workloads using Migrate for Anthos) is more focused on modernizing applications by containerizing them and leveraging Kubernetes. While it is a valid option for certain scenarios, it may require additional effort to containerize your workloads and might not be the most straightforward approach for VM migration.

upvoted 5 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: C

Migrate for Compute Engine is a service provided by Google Cloud that allows you to migrate your existing virtual machines (VMs) from on-premises or other cloud environments directly into Compute Engine. It provides a seamless migration experience, allowing you to retain your VM configurations and data while taking advantage of the scalability and flexibility offered by Google Cloud. This approach will help you adopt cloud-native technologies and modernize your workloads effectively.

upvoted 1 times

  **Erasable_Mak** 7 months ago

The correct option is C.

upvoted 1 times

  **Erasable_Mak** 7 months ago

C is correct because Migrate for Compute Engine is a service specifically designed to migrate VMs from on-premises environments or other clouds into Compute Engine. It provides automated migration capabilities and allows for the modernization of workloads during the migration process.

D is incorrect because Migrate for Anthos is focused on migrating workloads into Google Kubernetes Engine (GKE) or Anthos clusters, which are container-based environments. It may not be suitable for migrating traditional VM-based workloads.

upvoted 2 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

Migrate for Anthos uses a container-based approach to migration. It packages your workloads into containers and then deploys the containers to Anthos. This allows you to migrate your workloads to Google Cloud without having to make any changes to your workloads.

upvoted 1 times

  **Kavis** 8 months, 2 weeks ago

Migrate for Anthos - migrate to containers

Migrate for Compute Engine - migrate to VM

these are the two options. Here it says VM, so C is right answer

upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct

upvoted 1 times

  **himanshu3004** 10 months, 1 week ago

Selected Answer: D

Difference between option C & D based on Purpose and Focus:

Migrate for Anthos: This solution is focused on modernizing and migrating applications to containers and Kubernetes. It helps you containerize and manage your workloads using Google Kubernetes Engine (GKE) and provides a path to adopting cloud-native technologies.

Migrate for Compute Engine: This solution is designed for straightforward lift-and-shift migrations of virtual machines (VMs) from on-premises or other clouds to Google Compute Engine (GCE) instances.

In this question it's specifically mentioned that you need to modernize workloads as much as possible by adopting cloud-native technologies. Hence, I feel Anthos is a better choice here.

upvoted 9 times

  **Preetii** 10 months, 2 weeks ago

i think since it has not mentioned anything related to Containerization or any upcoming plans for containerization, I think we can stick to C.also here it is also mentioned about the separate OS which we usually do on Compute engine.

upvoted 1 times



Nirumal 10 months, 3 weeks ago

Selected Answer: C

When migrating workloads to take advantage of containers, Migrate for Anthos offers a fast and seamless path to modernization. As for the workloads for which virtual machines are best suited, simply move them with Migrate for Compute Engine and use VPC network integration with GKE.

<https://cloud.google.com/migrate/anthos?hl=es>

upvoted 1 times

Your organization is running all its workloads in a private cloud on top of a hypervisor. Your organization has decided it wants to move to Google Cloud as quickly as possible. Your organization wants minimal changes to the current environment, while using the maximum amount of managed services Google offers.

What should your organization do?

- A. Migrate the workloads to Google Cloud VMware Engine
- B. Migrate the workloads to Compute Engine
- C. Migrate the workloads to Bare Metal Solution
- D. Migrate the workloads to Google Kubernetes Engine

Correct Answer: B

Migrate for Compute Engine enables you to lift and shift workloads at scale to Google Cloud Compute Engine with minimal changes and risk.

Reference:

<https://dataintegration.info/simplify-vm-migrations-with-migrate-for-compute-engine-as-a-service>

Customers across industries are **migrating applications** at scale to the cloud. Migrate for Compute Engine enables you to lift and shift workloads at scale to Google Cloud Compute Engine with minimal changes and risk. Recently, we delivered our latest release of **Migrate for Compute Engine** – Version 5 or V5. This version delivers a Google-managed cloud service that enables simple, frictionless, and large-scale enterprise migrations of virtual machines to Google Compute Engine with minimal downtime and risk. API-driven and integrated into your Google Cloud console for ease of use, this service uses agent-less replication to copy data without manual intervention and without VPN requirements. It also enables you to launch non-disruptive validations of your VMs prior to cutover. Rapidly migrate a single application or execute a sprint with hundred systems using migration groups with confidence.

Community vote distribution

B (59%)

A (41%)

-   **invalient**

Highly Voted 

 1 year, 3 months ago

Selected Answer: B

If you use Google Cloud VMware Engine, youre restricted to VMware vSphere as the source of workloads, so B is correct.

upvoted 6 times
-   **harsh5kalsait**

Most Recent 

 2 months, 3 weeks ago

Selected Answer: A

A. VMware Engine.

Due to -

1. workloads in a private cloud on top of a hypervisor.

2. Your organization has decided it wants to move to Google Cloud as quickly as possible.

3. Your organization wants minimal changes to the current environment.

upvoted 1 times
-   **o_l_o** 2 months, 3 weeks ago

Selected Answer: A

VMware Engine provide ability migrate AS IS, so it should be faster and without additional overheads

upvoted 1 times
-   **Surek** 5 months, 3 weeks ago

Answer is A

upvoted 1 times
-   **sivakarthick16** 6 months, 1 week ago

Selected Answer: A

This service allows you to run your existing VMware workloads natively on Google Cloud, providing a seamless transition with minimal modifications. By migrating to Google Cloud VMware Engine, you can take advantage of Google's managed services while maintaining

compatibility with your current hypervisor-based private cloud setup. This approach enables a quick migration to Google Cloud while maximizing the utilization of managed services.

upvoted 1 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: A

VMware Engine allows you to run your existing VMware workloads natively on Google Cloud, providing a seamless transition with minimal changes. It offers managed services such as automated patching, monitoring, and scaling, while leveraging the benefits of Google Cloud's infrastructure.

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

  **lemonfan123** 8 months ago

Selected Answer: A

"minimal changes to the current env" Thus A would be correct

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: A

Google Cloud VMware Engine provides a managed VMware service that allows you to run your VMware workloads in Google Cloud without significant modifications. It provides a highly compatible and familiar environment for your existing VMware-based workloads while offering the benefits of Google Cloud's infrastructure.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times

  **Lufly** 10 months ago

Selected Answer: A

its A

why not the other choices?

Option B: Migrate the workloads to Compute Engine

Compute Engine is a Infrastructure as a Service (IaaS) offering that provides virtual machines. This means that your organization would need to manage the underlying infrastructure, such as the virtual machines, networking, and storage.

Option C: Migrate the workloads to Bare Metal Solution

Bare Metal Solution is a service that provides dedicated physical servers. This is the most hands-on option, and your organization would need to manage the underlying infrastructure.

Option D: Migrate the workloads to Google Kubernetes Engine

Google Kubernetes Engine is a managed Kubernetes service. Kubernetes is a container orchestration platform that can be used to deploy and manage containerized applications. This option would require your organization to containerize its applications.

upvoted 2 times

  **Mwafrika** 11 months, 1 week ago

Selected Answer: A

The mention of a hypervisor + the only answer with a hypervisor is A

upvoted 1 times

  **aaron773** 1 year ago

Selected Answer: A

The keywords here are: "minimal changes to the current environment"

upvoted 1 times

  **Kapahnke** 1 year ago

Selected Answer: A

A. Migrate the workloads to Google Cloud VMware Engine: This option enables you to migrate your workloads to a fully managed VMware environment on Google Cloud. It allows you to quickly and easily move your existing VMs to the cloud without significant modifications or re-architecting. You can leverage your existing VMware tools and processes while taking advantage of Google Cloud's infrastructure and managed services.

upvoted 2 times

  **ChiefArch** 12 months ago

This assumes the current VMs are VMware, there is no indication of this.

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: B

You can use Compute Engine to migrate your workloads from a private cloud to Google Cloud. Compute Engine provides virtual machine instance that can host your existing workloads, allowing you to lift and shift your applications with minimal changes.

It's important to note that while Compute Engine provides flexibility and control over your infrastructure, it also requires you to manage and maintain the virtual machine instances, including tasks such as patching, scaling, and monitoring.

upvoted 3 times

  **sjain93** 1 year, 1 month ago

To minimize changes to the current environment while utilizing the maximum amount of managed services offered by Google, your organization should migrate the workloads to Google Cloud VMware Engine (option A).

upvoted 1 times

  **einchkrein** 1 year, 2 months ago

Selected Answer: A

The key is the following sentence : "maximum amount of managed services"

Option B, Compute Engine is not a fully managed service and would require users to manage their own virtual machines and infrastructure.

Then it's option A

upvoted 2 times

Your organization is releasing its first publicly available application in Google Cloud. The application is critical to your business and customers and requires a 2- hour SLA.

How should your organization set up support to minimize costs?

- A. Enroll in Premium Support
- B. Enroll in Enhanced Support
- C. Enroll in Standard Support
- D. Enroll in Basic Support

Correct Answer: B

Reference:

<https://www.secureauth.com/enhanced-support-offering/>

SecureAuth is dedicated to providing the **industry-leading** enhanced **support** ensuring the long term success of your SecureAuth SaaS IAM deployment.

While our **basic** support offers industry leading coverage and response times for some customers, SecureAuth protects critical applications meaning delays and extended downtime is simply not an option. For these customers our **Enhanced Support** offerings provide **24*7 coverage** and the most responsive and **complete SLA's** available. That's why we offer three different levels of support, so you can choose the level of support that best works for your needs.

Community vote distribution

B (100%)

- Alexanderswor

Highly Voted

2 years, 3 months ago

Selected Answer: B

B - Enhanced:
https://cloud.google.com/support/docs?hl=en#support_offerings_overview

standard=4hr
enhanced=1hr
premium=15min+tam
upvoted 39 times
- grzeev

1 year, 8 months ago

Right!
upvoted 3 times
- GamudaDavid

5 months, 2 weeks ago

cheers
upvoted 1 times
- HKeshavarz

Most Recent

7 months, 2 weeks ago

It seems B is the right answer.
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B
B is correct
upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: B

Enhanced Support is designed for organizations that require higher levels of support for critical applications and services. It offers faster response times, 24/7 coverage, and can help you meet stringent SLAs while providing access to technical experts. It strikes a balance between cost and support quality, making it a suitable choice for a mission-critical application with a 2-hour SLA requirement.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: B

The Enhanced Support Plan can be a suitable choice for ensuring the required level of support for your critical application with a 2-hour SLA. The Enhanced Support Plan offers faster response times, extended coverage hours, and technical account management services, which can be valuable for organizations running mission-critical workloads.

upvoted 3 times

  **star2anand** 1 year, 3 months ago

C. Enroll in Standard Support

upvoted 2 times

  **C8H10N4O2** 1 year, 7 months ago

Selected Answer: B

you have to think about both respond time and minimize costs. So B is the answer

upvoted 1 times

  **Govindaraj** 1 year, 11 months ago

Selected Answer: B

B-Enhanced Support

upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: B

B. Enroll in Enhanced Support

upvoted 1 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: B

B. Enroll in Enhanced Support

upvoted 1 times

  **desertlotus1211** 2 years, 4 months ago

Answer is B:

<https://cloud.google.com/support>

upvoted 1 times

  **VenRam** 2 years, 4 months ago

Enhanced Support offers P1 - 1 hour Initial response time. (while Basic offers P2 - 4hr & Premium offers P1 - 15mins Initial Response time)

link: <https://cloud.google.com/support>

upvoted 4 times

  **Pravin269** 2 years, 5 months ago

<https://cloud.google.com/support>

upvoted 2 times

Your organization offers public mobile apps and websites. You want to migrate to a Google Cloud-based solution for checking and maintaining your users' usernames and passwords and controlling their access to different resources based on their identity. Which should your organization choose?

- A. VPN tunnels
- B. Identity Platform
- C. Compute Engine firewall rules
- D. Private Google Access

Correct Answer: B

An identity platform is a modern solution for managing the identities of users and devices in a centralized fashion.

Reference:

<https://www.okta.com/blog/2021/07/what-is-an-identity-platform/#:~:text=An%20identity%20platform%20is%20a,%2C%20integrations%2C%20and%20platform%20services>

An identity platform is a modern solution for managing the identities of users and devices in a centralized fashion. It enables organizations to securely authorize workforce and customer users to access their ecosystem using access management tools, programmable components, integrations, and platform services.

Today's organizations have a wide range of identity requirements. As they expand, embrace new innovations, and meet new customer demands, they need an identity solution that can grow with them. And as they grow, these organizations are also looking for offerings that centralize and consolidate identity, reducing the need for multiple access management, governance, and authentication products that don't necessarily talk to each other.

In short, although the concept of identity platforms is still new, the growth of disruptive technologies and highly personalized products and services has made a platform approach the key for effectively unifying identity management. But how did we get here?

Community vote distribution

B (100%)

-   **desertlotus1211**

Highly Voted 

 2 years, 4 months ago

Answer: B

<https://cloud.google.com/identity-platform>

Identity Platform is a customer identity and access management (CIAM) platform that helps organizations add identity and access management functionality to their applications, protect user accounts, and scale with confidence on Google Cloud.

upvoted 7 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times
-   **cookieMr** 1 year ago

Selected Answer: B

Google Cloud Identity Platform provides a comprehensive set of features and capabilities to handle user authentication and authorization across different applications and platforms.

upvoted 3 times

  **star2anand** 1 year, 3 months ago

B. Identity Platform

upvoted 2 times

  **victory108** 2 years, 3 months ago

Selected Answer: B

B. Identity Platform

upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: B

B. Identity Platform

upvoted 1 times

  **mintupintu** 2 years, 5 months ago

B is the correct answer

upvoted 1 times

Which Google Cloud service or feature lets you build machine learning models using Standard SQL and data in a data warehouse?

- A. BigQuery ML
- B. TensorFlow
- C. AutoML Tables
- D. Cloud Bigtable ML

Correct Answer: A

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries.

Reference:

<https://cloud.google.com/bigquery-ml/docs/introduction#:~:text=BigQuery%20ML%20lets%20you%20create,the%20need%20to%20move%20data>

What is BigQuery ML?

Send feedback

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries. BigQuery ML democratizes machine learning by letting SQL practitioners build models using existing SQL tools and skills. BigQuery ML increases development speed by eliminating the need to move data.

BigQuery ML functionality is available by using:

- The Google Cloud Console
- The `bq` command-line tool
- The BigQuery REST API
- An external tool such as a Jupyter notebook or business intelligence platform

Community vote distribution

A (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: A

A is correct
upvoted 1 times
- Kavis** 8 months, 2 weeks ago

BigQueryML uses ML + Data Science Library and is a datawarehouse that can be used for unstrctred data and for making predictions. Correct answer - A

upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: A

A is correct
upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: A

BigQuery ML is a powerful tool within Google BigQuery that enables you to create and execute machine learning models directly using SQL queries. It allows you to build and train models using familiar SQL syntax and leverage the data stored in BigQuery for training and prediction.

By using BigQuery ML, you can leverage the power of SQL and the scalability of Google BigQuery to build and deploy machine learning models directly on your data warehouse, without the need for complex data transformations or separate ML platforms.

upvoted 4 times

  **manashbaruah** 1 year, 6 months ago

Selected Answer: A

Keywords: Standard SQL, data warehouse

upvoted 1 times

  **Daffy16** 2 years ago

Selected Answer: A

A. BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries.

upvoted 2 times

  **victory108** 2 years, 3 months ago

Selected Answer: A

A. BigQuery ML

upvoted 1 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: A

A. BigQuery ML

upvoted 1 times

  **desertlotus1211** 2 years, 4 months ago

Answer A:

<https://cloud.google.com/bigquery-ml/docs/introduction#:~:text=BigQuery%20ML%20lets%20you%20create,the%20need%20to%20move%20data>.

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries. BigQuery ML democratizes machine learning by letting SQL practitioners build models using existing SQL tools and skills. BigQuery ML increases development speed by eliminating the need to move data.

upvoted 3 times

  **mintupintu** 2 years, 5 months ago

A is the correct answer.

upvoted 1 times

Your organization runs an application on virtual machines in Google Cloud. This application processes incoming images. This activity takes hours to create a result for each image. The workload for this application normally stays at a certain baseline level, but at regular intervals it spikes to a much greater workload.

Your organization needs to control the cost to run this application.

What should your organization do?

- A. Purchase committed use discounts for the baseline load
- B. Purchase committed use discounts for the expected spike load
- C. Leverage sustained use discounts for your virtual machines
- D. Run the workload on preemptible VM instances

Correct Answer: C

The idea of the Sustained Use discount is that the longer you run a VM instance in any given month, the bigger discount you will get from the list price.

Reference:

<https://www.parkmycloud.com/blog/google-sustained-use-discounts/>

price. The following shows the incremental discount, and its c
impact on a hypothetical \$100/month VM instance, where the
are against the baseline 730-hour month.

Usage level		Hours usage before price break	discount	Base cost/mo
From	To			\$ 100.00
0%	25%	-	0%	\$ 25.00
25%	50%	182.5	20%	\$ 20.00
50%	75%	365.0	40%	\$ 15.00
75%	100%	547.5	60%	\$ 10.00
			Net cost	\$ 70.00

Community vote distribution

C (52%)

A (33%)

14%

- fpreli

Highly Voted

2 years, 5 months ago

A for me, with CUD you get a greater discount and it makes totally sense to buy a CUD for the baseline traffic since it will always be consumed

upvoted 25 times
- AAD1

Highly Voted

1 year, 9 months ago

Selected Answer: C

Will gop with C. as CUD is applicable for 1 or 3 year period. In the ques, there is no mention of that info. So SUD is appropriate as we cannot assume that the VMs will be used for 1 or 3 years..

upvoted 15 times
- PanosPeris

Most Recent

4 weeks ago

Selected Answer: A

The best option to control cost for your organization's image processing application is:

A. Purchase committed use discounts for the baseline load

Here's why the other options are less suitable:

- B. Purchase committed use discounts for the expected spike load: Committed use discounts are most beneficial for predictable, ongoing resource usage. Since the spike load is temporary, committing to those resources wouldn't be cost-effective.
- C. Leverage sustained use discounts for your VMs: Sustained use discounts reward consistent resource usage at a discount compared to on-demand pricing. However, they might not be ideal for workloads with significant spikes that could push you outside the sustained use window.
- D. Run the workload on preemptible VM instances: Preemptible VMs are a good option for highly flexible workloads that can be interrupted. However, for image processing tasks that take hours, interruptions during processing could lead to wasted work and require restarts.

upvoted 1 times

  **alilikpo** 4 months, 2 weeks ago

Selected Answer: C

your workload has a predictable baseline level. Google Cloud offers sustained use discounts for Compute Engine commitments that reserve VM instances for a period of 1 year or 3 years. This can significantly reduce your costs for the baseline usage.

upvoted 1 times

  **sivakarthick16** 6 months, 1 week ago

Selected Answer: C

Sustained use discounts provide cost savings for running virtual machines continuously over a month. Since the workload normally stays at a certain baseline level, leveraging sustained use discounts can help reduce costs for the consistent usage. This approach is suitable when the workload is predictable and doesn't involve significant spikes. However, if the expected spike load is known in advance, option B: Purchasing committed use discounts for the expected spike load could also be considered to optimize costs during those periods. It is important to analyze the workload patterns and cost implications to make an informed decision.

upvoted 3 times

  **ankiiiiiii** 7 months ago

Selected Answer: C

C for me

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

  **chai_gpt** 7 months, 2 weeks ago

Sorry. A is correct.

upvoted 1 times

  **lemonfan123** 8 months ago

Selected Answer: A

A would be correct, since "controlling the cost" is a requirement

upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: A

I think it should be A.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times

  **Lufly** 10 months ago

Selected Answer: C

why not to purchase committed use discounts for the baseline load. Here are the reasons why this is not a good option:
It would be more expensive than sustained use discounts because you would be paying for the baseline load even when you are not using it.
It would not take into account the spike load. If the spike load is significant, then you would still be paying a lot for the resources that you are using during the spike load.
It would be inflexible. If your workload changes, then you would have to change your committed use commitments, which can be difficult and time-consuming.

upvoted 2 times

  **himanshu3004** 10 months, 1 week ago

Selected Answer: A

I feel CUD is a better choice here considering that SUD is automatically applied by Google incase if the VM don't qualify for CUD. so we don't need to explicitly select SUD. However selecting CUD makes more sense because if the VM qualifies for CUD, you won't be losing on the additional cloud savings that CUD offers compared to SUD.

It's more of like with CUD, we have better chances of higher cloud savings at no additional effort/cost which makes it best choice out of all.

upvoted 1 times

  **chinmay_100** 12 months ago

C will be the right choice , as for CUD we require an commitment of 1-3 years whereas for SUD you don't and still get discount.
upvoted 2 times

  **cookieMr** 1 year ago

Selected Answer: A

If your workload has a consistent baseline level and you can accurately estimate the resources needed, purchasing committed use discounts for those resources can provide substantial cost savings. By committing to using the resources for a longer duration, you can secure discounted pricing.
upvoted 2 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: A

I went for A.
upvoted 1 times

  **stoker35** 1 year, 3 months ago

There is no information about the length this service will be used, hence Committed discount is not the best answer here.
upvoted 3 times

  **tracyonair** 1 year, 4 months ago

Selected Answer: A

The answer is A: CUD.
With committed use discounts, you have to commit to a minimum amount of usage, which is possible in this case because you know your baseline load.
Sustained use discounts is a discount automatically applied by Google when compute engine instances do not qualify for Committed-Use discounts.And they are for SPECIFIC COMPUTE ENGINE RESOURCES.
SUD is 30% discount. CUD is 57-70% discount
upvoted 2 times

Your organization is developing a plan for migrating to Google Cloud.
What is a best practice when initially configuring your Google Cloud environment?

- A. Create a project via Google Cloud Console per department in your company
- B. Define your resource hierarchy with an organization node on top
- C. Create projects based on team members' requests
- D. Make every member of your company the project owner

Correct Answer: B

Community vote distribution

B (100%)

- cookieMr**

Highly Voted

1 year ago

Selected Answer: B

Establish a well-structured organization hierarchy using folders and projects to effectively manage and organize your resources. This helps with access control, resource management, and billing.

upvoted 5 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: B

B is correct.

upvoted 1 times
- Kavis**

8 months, 2 weeks ago

Org > folders > projects
Correct answer - B

upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- trullera**

1 year, 7 months ago

Selected Answer: B

B. with no doubts

upvoted 1 times
- victory108**

2 years, 3 months ago

Selected Answer: B

B. Define your resource hierarchy with an organization node on top

upvoted 1 times
- DatTTran**

2 years, 3 months ago

B. Define your resource hierarchy with an organization node on top

upvoted 1 times
- leyunjohn**

2 years, 4 months ago

Selected Answer: B

B. Define your resource hierarchy with an organization node on top

upvoted 1 times
- AbdulioUnited**

2 years, 4 months ago

is the correct answer "B" because we are still talking here about to develop a plan?? Thank you..

upvoted 2 times
- desertlotus1211**

2 years, 4 months ago

Answer B

upvoted 1 times



Your organization runs many workloads in different Google Cloud projects, each linked to the same billing account. Each project's workload costs can vary from month to month, but the overall combined cost of all projects is relatively stable. Your organization needs to optimize its cost. What should your organization do?

- A. Purchase a commitment per project for each project's usual minimum
- B. Create a billing account per project, and link each project to a different billing account
- C. Turn on committed use discount sharing, and create a commitment for the combined usage
- D. Move all workloads from all different projects into one single consolidated project

Correct Answer: C

Reference:
<https://cloud.google.com/compute/docs/instances/signing-up-committed-use-discounts>

Community vote distribution

C (100%)

- deCalle**

Highly Voted

1 year, 2 months ago

wtf are those questions? HAVE ONE BILLING ACCOUNT FOR EVERYTHING?! really guys? I d say B because we're human. then C bc costs. i really really hate how google phrases its questions. after directing with its sentences into a ceratin direction it just makes a uturn and asks about costs, wtf google? asperger/adhd much? "2 miles down the road is an old workshop, each hour tinkering costs \$20. Which way is africa?"

upvoted 6 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: C

Committed use discount sharing allows you to apply a single commitment to multiple projects linked to the same billing account, which is a useful way to aggregate your usage and get the most benefit from your commitments. It's especially effective if your overall combined cost is relatively stable, as you mentioned.

upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: C

If the overall combined cost of all projects is relatively stable, consolidating billing under a single billing account can provide better visibility and management of costs. This allows you to analyze and optimize costs across all projects collectively.

If certain projects have similar resource requirements or workloads, you can consider consolidating them into a single project. By doing so, you can optimize the usage and commitment of resources, allowing you to purchase CUDs more effectively.

upvoted 4 times
- gcp_prep**

1 year, 3 months ago

Answer would be C, D doesnt make sense

upvoted 1 times
- A84-64**

1 year, 3 months ago

https://cloud.google.com/compute/docs/instances/signing-up-committed-use-discounts#sharing_committed_use_discounts_across_projects explains how to share CUD across multiple projects.

upvoted 1 times
- manashbaruah**

1 year, 6 months ago

Selected Answer: C

Keypoints: already same billing account is being used and optimize cost

upvoted 1 times
- trullera**

1 year, 7 months ago

Selected Answer: C

C. all other options does not make sense
upvoted 1 times

  **claudio_acosta** 1 year, 7 months ago

Selected Answer: C

better management
upvoted 1 times

  **raja3l49** 1 year, 8 months ago

Committed Use Discount is applicable only for 1 or 3 years term agreement. But the question didn't talk about the duartion. So the answer will be D
upvoted 1 times

  **raja3l49** 1 year, 8 months ago

I will go with Option D
upvoted 1 times

  **Loved** 1 year, 8 months ago

If you move everything into a single project you lose the isolation between environment
upvoted 3 times

  **OctavioBatera** 1 year, 6 months ago

And also this doesn't help with the costs control, which is the subject of the question.
upvoted 1 times

  **jaxclain** 1 year, 9 months ago

Selected Answer: C

I would also go for C
upvoted 1 times

  **WanessaZisa** 1 year, 9 months ago

Selected Answer: C

C. Turn on committed use discount sharing, and create a commitment for the combined usage
upvoted 1 times

How should a multinational organization that is migrating to Google Cloud consider security and privacy regulations to ensure that it is in compliance with global standards?

- A. Comply with data security and privacy regulations in each geographical region
- B. Comply with regional standards for data security and privacy, because they supersede all international regulations
- C. Comply with international standards for data security and privacy, because they supersede all regional regulations
- D. Comply with regional data security regulations, because they're more complex than privacy standards

Correct Answer: A

Community vote distribution

A (93%)

4%

- ahsangh

Highly Voted

2 years, 5 months ago

international standards do not always supersede regional regulations

A. Comply with data security and privacy regulations in each geographical region

upvoted 23 times
- astood

2 years, 5 months ago

Agreed

upvoted 1 times
- cookieMr

Highly Voted

1 year ago

Selected Answer: A

For a multinational organization migrating to Google Cloud is to comply with data security and privacy regulations in each geographical region where it operates. This ensures compliance with local laws and regulations specific to each jurisdiction. Different regions may have their own set of data protection and privacy standards that organizations need to adhere to.

upvoted 6 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct

upvoted 1 times
- oezgan

9 months ago

standards do not supersede regulations. Never. Regulations: you must adhere to, standards: are like best practices.

upvoted 2 times
- JudgeLI

1 year, 5 months ago

Selected Answer: A

A is correct

upvoted 1 times
- alfantastic

1 year, 5 months ago

Selected Answer: A

International standard (or agreements) are usually weaker, regional are more likely to be enforced

upvoted 1 times
- someshseghal

1 year, 6 months ago

Selected Answer: C

I checked other links too. They are suggesting C.

upvoted 1 times

  **someshseghal** 1 year, 6 months ago

Selected Answer: B

Regional always supersede international
upvoted 1 times

  **manashbaruah** 1 year, 6 months ago

Selected Answer: A

Keywords: each geographical region
upvoted 1 times

  **Pou1ze** 1 year, 6 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **AnithaPal** 1 year, 7 months ago

A
Comply with data security and privacy regulations in each geographical regio
upvoted 1 times

  **raja3l49** 1 year, 8 months ago

I will go with Option C, since International Standards will superceed with Regional Standards
upvoted 1 times

  **Undertaker5566** 1 year, 9 months ago

Selected Answer: A

A for sure
upvoted 1 times

  **ravisinha** 1 year, 12 months ago

Ans: A make sense. C cover international Std but regionals should have been take care too.
upvoted 1 times

  **juangt** 2 years, 1 month ago

Selected Answer: A

Answer is A, Comply with data security and privacy regulations in each geographical region
upvoted 2 times

Your organization wants to optimize its use of Google Cloud's discounts on virtual machine-based workloads. You plan to use 200 CPUs constantly for the next 3 years, and you forecast that spikes of up to 300 CPUs will occur approximately 30% of the time. What should you choose?

- A. 1-year committed use discount for 200 CPUs
- B. 3-year committed use discount for 300 CPUs
- C. 3-year committed use discount for 200 CPUs
- D. Regular pay-as-you-go pricing

Correct Answer: C

Reference:
<https://cloud.google.com/compute/all-pricing>

Community vote distribution

C (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: C

C is correct.
upvoted 1 times
- Kavis** 8 months, 2 weeks ago

If i use committed use discount for 200 CPUs that will be used for 60% of the year, that should take care of the spikes that occur 30% of time.
Answer - C
upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: C

A is correct
upvoted 1 times
- rahul_nimbalkar** 9 months, 2 weeks ago

D
Pay as you Go
upvoted 1 times
- Erinias74** 11 months, 1 week ago

What happen with "the spikes of up to 300 CPUs will occur approximately 30%"?
The option C) could cover it?
upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: C

Based on the provided information, you plan to use 200 CPUs constantly for the next 3 years, which makes it suitable for purchasing Committed Use Discounts. CUDs offer significant cost savings for committing to use specific machine types and sizes for a specified term. By purchasing CUDs for the 200 CPUs that are consistently utilized, you can take advantage of discounted pricing compared to On-Demand rates.
upvoted 4 times

manashbaruah 1 year, 6 months ago

Selected Answer: C

CUD gives more discount
upvoted 2 times

MurugeshBabu 1 year, 9 months ago

C is right
upvoted 1 times

Undertaker5566 1 year, 9 months ago

C for sure

upvoted 1 times

  **mikelabs** 1 year, 10 months ago

Selected Answer: C

The answer is C

upvoted 1 times

  **kummimhk** 2 years ago

Selected Answer: C

When committed discount option is available, no question of pay as you go model.

upvoted 4 times

  **pauldoit** 2 years, 1 month ago

How is D the answer..

upvoted 1 times

  **sedado77** 2 years, 1 month ago

Selected Answer: C

Answer is clearly C, not D

upvoted 1 times

  **EastSideCloudGuy** 2 years, 2 months ago

Selected Answer: C

Please update your site Examtopics.

Answer is C and not D.

D: PayAsYouGo gets no discount, least efficient option!

C: You get a deeper discount for a 3 year commitment (Answer C) vs 1 year commitment (Answer A). See: https://cloud.google.com/compute/vm-instance-pricing#general-purpose_machine_type_family for an example of that.

upvoted 3 times

  **kyat** 2 years, 2 months ago

it should be C. 3 yrs and predicted compute. 30% spike is only expected .may or may not come.

upvoted 2 times

  **victory108** 2 years, 3 months ago

Selected Answer: C

C. 3-year committed use discount for 200 CPUs

upvoted 2 times

Your organization needs to minimize how much it pays for data traffic from the Google network to the internet. What should your organization do?

- A. Choose the Standard network service tier.
- B. Choose the Premium network service tier.
- C. Deploy Cloud VPN.
- D. Deploy Cloud NAT.

Correct Answer: A

Reference:
<https://cloud.google.com/blog/products/networking/networking-cost-optimization-best-practices>

Community vote distribution

A (80%)

D (17%)

-   **Begum** Highly Voted 2 years, 4 months ago

C& D are mechanisms , Use premium tier for performance and standard tier for cost optimizations.
upvoted 42 times
-   **Alexanderswor** 2 years, 3 months ago

Major appreciation for straight forward guidance. LOVE this direction, Begum.

Validated here: <https://cloud.google.com/network-tiers?hl=en>
upvoted 4 times
-   **mnsait** 3 days, 23 hours ago

Absolutely, well explained by Begum.

Another resource to validate this:
<https://cloud.google.com/network-tiers/docs/overview#:~:text=Premium%20Tier%20delivers%20traffic%20on,Tier%20to%20optimize%20for%20cost.&text=Outbound%20data%20transfer%20pricing%20for,Network%20Service%20Tiers%20is%20different.>

"Premium Tier delivers traffic on Google's premium backbone, while Standard Tier uses regular ISP networks.

Use Premium Tier to optimize for performance, and use Standard Tier to optimize for cost."
upvoted 1 times
-   **Tom666** Highly Voted 2 years, 4 months ago

Selected Answer: A

While Premium tier is the default for all egress traffic and offers the highest performance, it might not be the best option when cost is a consideration. Standard tier is the more economical , So option is A
upvoted 12 times
-   **chai_gpt** Most Recent 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: A

The Standard network service tier provides lower egress (outbound) data transfer costs when compared to the Premium network service tier.
upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: A

A is correct
upvoted 1 times
-   **Lufly** 10 months ago

Selected Answer: A

The answer is A. Choose the Standard network service tier.

The Standard network service tier offers the best price-performance ratio for most workloads. It provides a predictable and reliable network with low latency and high throughput.

The Premium network service tier offers a higher level of performance and reliability, but it also comes with a higher price tag. It is a good option for workloads that require the lowest possible latency or the highest possible throughput.

Cloud VPN and Cloud NAT are both networking services that can be used to connect your Google Cloud resources to the internet. However, they do not directly affect the cost of data traffic from the Google network to the internet.

upvoted 2 times

  **cookieMr** 1 year ago

Selected Answer: D

To minimize the cost of data traffic from the Google network to the internet, your organization should choose option D. Deploy Cloud NAT.

Cloud NAT (Network Address Translation) is a service provided by Google Cloud that allows your virtual machine instances to send outbound traffic to the internet using a single shared external IP address. By using Cloud NAT, you can reduce the number of external IP addresses needed for your instances and consolidate the outbound traffic through a single IP, which can help lower costs.

Cloud VPN (Virtual Private Network) is used for secure connectivity between your on-premises network and your Google Cloud resources, but it does not directly address the cost of data traffic to the internet.

The choice between the Standard and Premium network service tiers does not directly impact the cost of outbound data traffic to the internet. It primarily affects network performance, availability, and features.

Therefore, deploying Cloud NAT is the most appropriate option to minimize the cost of data traffic from the Google network to the internet.

upvoted 3 times

  **Jeremy11** 10 months ago

This says Standard vs Premium doesn't impact the cost, that seems wrong. Here it says Standard is cheaper: <https://cloud.google.com/network-tiers>

upvoted 1 times

  **Jeremy11** 10 months ago

And here it's even more clear: <https://cloud.google.com/network-tiers/pricing>

upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: D

Choosing the Standard or Premium network service tier will not directly minimize the cost of data traffic from the Google network to the internet, as these tiers mainly differ in terms of network performance and reliability, not cost.

Deploying Cloud VPN may allow the organization to securely connect its on-premises network to its Google Cloud environment, but it will not directly minimize the cost of data traffic from the Google network to the internet.

upvoted 2 times

  **Matro71** 1 year, 3 months ago

Standard only applies if less than 500TB per month.

upvoted 1 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: A

Standard tier for cost optimization

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

Per Google docs, "Premium Use Case - Performance, reliability, global footprint, and user experience are your main considerations. Standard Use Case - Cost is your main consideration, and you're willing to trade off some network performance."

<https://cloud.google.com/network-tiers#tab2>

upvoted 1 times

  **sfsdeniso** 1 year, 6 months ago

when using Premium Tier traffic stays within google network - it does not go into internet

https://cloud.google.com/network-tiers/docs/overview#choosing_a_tier

upvoted 1 times

  **grzeev** 1 year, 8 months ago

Selected Answer: A

Network Service Tiers lets you optimize connectivity between systems on the internet and your Google Cloud instances. Premium Tier delivers traffic on Google's premium backbone, while Standard Tier uses regular ISP networks. Use Premium Tier to optimize for performance, and use Standard Tier to optimize for cost

upvoted 1 times

  **S132** 1 year, 9 months ago

So A is answer

upvoted 1 times

  **MarcosBarreto** 2 years, 1 month ago

Selected Answer: B

Standard Network Service is less costly so it reduces costs
upvoted 1 times

  **EastSideCloudGuy** 2 years, 2 months ago

Selected Answer: A

A has a direct impact on minimizing cost. See: <https://cloud.google.com/blog/products/networking/networking-cost-optimization-best-practices>

Per that article answer C does not fit, see this excerpt:
If you push large amounts of data (think TBs/PBs) on a consistent basis, it can be cheaper to establish a dedicated connection vs. accruing costs associated with your traffic traversing the public internet or using a VPN.

Therefore the VPN does not minimize traffic, for such a case, you would want to use a dedicated partner connection.

Examtopics, please provide justification for answer C or update to a more appropriate answer.
upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: A

A. Choose the Standard network service tier.
upvoted 1 times

Your organization wants to migrate your on-premises environment to Google Cloud. The on-premises environment consists of containers and virtual machine instances. Which Google Cloud products can help to migrate the container images and the virtual machine disks?

- A. Compute Engine and Filestore
- B. Artifact Registry and Cloud Storage
- C. Dataflow and BigQuery
- D. Pub/Sub and Cloud Storage

Correct Answer: B

Reference:

<https://cloud.google.com/compute/docs/import/importing-virtual-disks>

If you have virtual disks in your on-premises environment with software and configurations that you need (sometimes referred to as *golden disks* or *golden images*), you can save time by importing those virtual disks into Compute Engine and using the resulting image to [create virtual machines](#). The import tool supports most virtual disk file formats, including VMDK and VHD.

If you exported your disk from Compute Engine, you can [create](#) images from the disk.

For information about how to create an automated system for migrating several virtual machines (VMs), see [Migrating VMs to Compute Engine](#).

Community vote distribution

B (88%)

13%

- Daffy16

Highly Voted

2 years ago

Selected Answer: B

B: Artifact Registry can manage container image migration and Cloud Storage can store virtual disk image
upvoted 18 times
- S00999

Highly Voted

1 year, 10 months ago

Selected Answer: B

I'm voting for B, not A because Filestore is more user files oriented when Cloud Storage is the best choice for disk images. Artifact Registry (aka Container Registry) is THE storage service for container images.
upvoted 8 times
- sivakarthick16

Most Recent

5 months ago

Selected Answer: A

Compute Engine allows you to create and manage virtual machine instances, making it ideal for migrating virtual machine disks. Filestore, on the other hand, provides a managed file storage service that can be used to migrate container images.

Therefore, option A (Compute Engine and Filestore) is the correct choice for migrating both container images and virtual machine disks to Google Cloud.
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: B

Artifact Registry is a fully managed registry service for storing container images. Cloud Storage is a highly scalable object storage service. Together they can be used to migrate container images from an on-premises environment to Google Cloud.
upvoted 2 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct
upvoted 1 times
- Lufly

10 months ago

Selected Answer: B

Here are the reasons why the other options are not as good as Artifact Registry and Cloud Storage for migrating container images and virtual machine disks to Google Cloud:

Compute Engine and Filestore are both compute services that can be used to create and run virtual machines. However, they do not provide a good way to store and manage container images.
Dataflow and BigQuery are both data processing services. They can be used to process and analyze data, but they are not designed for storing and managing container images.
Pub/Sub and Cloud Storage are both messaging and storage services. They can be used to store and transmit data, but they are not designed for storing and managing container images.

upvoted 3 times

  **cookieMr** 1 year ago

Selected Answer: B

When migrating your containers and virtual machines to Google Cloud, you can utilize Artifact Registry and Cloud Storage to streamline the process and ensure a smooth transition.

Artifact Registry can be used to store your container images as "golden images." These are pre-configured and optimized container images that serve as a starting point for deploying your containers in Google Cloud. You can push your golden images to Artifact Registry, where they will be securely stored and versioned. This allows you to easily manage and track different versions of your container images.

On the other hand, Cloud Storage can be used to store your virtual machine disk images, which can also be considered as "golden images." These disk images contain the operating system and pre-installed software configurations that you want to use for your virtual machine instances. By uploading these disk images to Cloud Storage, you can ensure their availability and accessibility during the migration process.

upvoted 4 times

  **skarths** 1 year, 7 months ago

Selected Answer: B

B is the answer.
Migration of your container images includes the following steps:

Set up prerequisites.
Identify images to migrate.
Search your Dockerfile files and deployment manifests for references to third-party registries
Determine pull frequency of images from third-party registries using Cloud Logging and BigQuery.
Copy identified images to Artifact Registry.
Verify that permissions to the registry are correctly configured, particularly if Artifact Registry and your Google Cloud deployment environment are in different projects.
Update manifests for your deployments.
Re-deploy your workloads.

upvoted 2 times

  **Nitheeshsk** 1 year, 8 months ago

Selected Answer: B

<https://cloud.google.com/compute/docs/import/importing-virtual-disks#bootable>,
It says to import the virtual disks to cloudstorage. Then only it is migrated to compute engine

upvoted 3 times

  **Dan_W** 1 year, 8 months ago

Answer is A. Migrate for Compute Engine is how to migrate. B has to do with managing and storage, not migration.

upvoted 1 times

  **omgitsele** 1 year, 2 months ago

Firestore has nothing to do with migration.

upvoted 1 times

  **KNACK** 1 year, 8 months ago

No mention of Migrate for Compute Engine in the options. And for container images, Artifact registry is used. So, correct answer is B.

upvoted 5 times

  **ergauravtaneja** 1 year, 10 months ago

Selected Answer: B

refer to link <https://cloud.google.com/compute/docs/import/import-existing-image> and check the overview section point no 4 which says "Upload the image file to Cloud Storage and import the image to Compute Engine as a new custom image."

upvoted 2 times

  **therealpumato** 1 year, 11 months ago

Selected Answer: A

The point is that you want to run it in GCP after migrating, not only manage the resources...?

upvoted 1 times

  **sedado77** 2 years, 1 month ago

Selected Answer: B

It's B: Artifact Registry can manage container image migration and Cloud Storage can store virtual disk image

upvoted 3 times

  **kyat** 2 years, 2 months ago
B.migrate container image: container registry, migrate vm disk: cloud storage.
upvoted 2 times

  **ckintheocean** 2 years, 2 months ago
Selected Answer: B
Its B because it's about the container images and the virtual machine disks
upvoted 1 times

  **victory108** 2 years, 3 months ago
Selected Answer: B
B. Artifact Registry and Cloud Storage
upvoted 2 times

Your company security team manages access control to production systems using an LDAP directory group.
How is this access control managed in the Google Cloud production project?

- A. Assign the proper role to the Service Account in the project's IAM Policy
- B. Grant each user the roles/iam.serviceAccountUser role on a service account that exists in the Google Group.
- C. Assign the proper role to the Google Group in the project's IAM Policy.
- D. Create the project in a folder with the same name as the LDAP directory group.

Correct Answer: C

Reference:

<https://cloud.google.com/blog/products/identity-security/achieving-identity-and-access-governance-on-google-cloud>

Community vote distribution

C (100%)

- Kessel** 3 weeks, 2 days ago

Selected Answer: C

A is about Service accounts, which is for Apps or VMs, so it's out. B is the typical wrong reply in these kinds of exams where they suggest doing something manually one by one and not efficiently. D sounds wrong to me, just putting the same name and calling it a day? C is the only option, because Google Groups work with LDAP directories.

upvoted 1 times
- ceecreeve** 6 months, 3 weeks ago

Selected Answer: C

Assigning roles to groups is the recommended best practice for controlling access

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times
- Kavis** 8 months, 2 weeks ago

why not C, All production access is given to service accounts with right access using IAM

upvoted 1 times
- Kavis** 8 months, 2 weeks ago

Typo - *why not A

upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- samirvarude** 10 months, 4 weeks ago

Ans is A

upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: C

Identity and Access Management (IAM) allows you to manage access to Google Cloud resources. Within your production project, you can use IAM to grant specific permissions to the LDAP directory group. IAM roles can be assigned to the group, defining the level of access they have within the project. This way, you can control who can perform certain actions on the project's resources.

upvoted 4 times
- AMMARABO** 1 year, 8 months ago

Selected Answer: C

the answer is c for more information <https://www.youtube.com/watch?v=ZTpOAmJWhAk>
upvoted 2 times

  **victory108** 2 years, 3 months ago

Selected Answer: C

C. Assign the proper role to the Google Group in the project's IAM Policy
upvoted 3 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. Assign the proper role to the Google Group in the projects IAM Policy.
upvoted 2 times

  **desertlotus1211** 2 years, 4 months ago

Could be A or C...
upvoted 2 times

  **JoeBK** 2 years, 4 months ago

Selected Answer: C

Assigning permissions to groups is considered a best practice.
upvoted 4 times

Your organization wants to be sure that its expenditures on cloud services are in line with the budget. Which two Google Cloud cost management features help your organization gain greater visibility into its cloud resource costs? (Choose two.)

- A. Billing dashboards
- B. Resource labels
- C. Sustained use discounts
- D. Financial governance policies
- E. Payments profile

Correct Answer: AB

Reference:

<https://cloud.google.com/cost-management#section-6>

Resource hierarchy

Structure and organize your [resource hierarchy](#) for fine-grained management and cost allocation using organizations, folders, projects, and labels.

Billing access control

Enforce organizational policies with granular [permissions](#) at different levels in the resource hierarchy to control who can spend and who has administrative and cost-viewing permissions.

Budgets and alerts

[Set budgets](#) to closely monitor your costs and alert stakeholders through email or [Pub/Sub](#) when exceeding defined budget thresholds.

Community vote distribution

AB (93%)

7%

- fpreli

Highly Voted

2 years, 5 months ago

I think it's AB. Financial governance policies are not related to cost visualization

upvoted 20 times
- ap6491

Highly Voted

2 years, 3 months ago

A and B. <https://cloud.google.com/billing/docs/how-to/visualize-data>

upvoted 7 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: AB

AB is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: AB

Billing dashboards: Billing dashboards provide a graphical overview of your cloud spending, organized by project, service, and other dimensions. You can use billing dashboards to track your spending over time, identify trends, and troubleshoot any unexpected spikes in costs.

Resource labels: Resource labels are key-value pairs that you can attach to your cloud resources. You can use resource labels to organize and group your resources, and to filter and analyze your cost data. For example, you could label all of your production resources with the label environment=production. This would allow you to easily see how much you are spending on production resources, and to compare your production costs to your development or testing costs.

upvoted 6 times
- mdsarfraz69

9 months ago

Selected Answer: AB

A and B. <https://cloud.google.com/billing/docs/how-to/visualize-data>
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: AB

The two Google Cloud cost management features that help organizations gain greater visibility into their cloud resource costs are:

A. Billing dashboards: Billing dashboards provide detailed information and insights into your organization's cloud resource costs. They display usage data, cost breakdowns, and trends, allowing you to analyze and monitor your expenditure. You can customize the dashboards to track specific metrics and visualize the cost data in a way that suits your organization's needs.

B. Resource labels: Resource labels allow you to categorize and tag your cloud resources based on attributes that are relevant to your organization's cost management. By applying labels to resources, you can track and analyze costs based on different dimensions such as project, department, environment, or application. This helps in understanding the cost allocation and identifying areas where optimization can be applied.
upvoted 6 times

  **alfantastic** 1 year, 5 months ago

Selected Answer: AB

payment profile is not only cloud related
upvoted 1 times

  **pspandher** 1 year, 8 months ago

AB as financial Governance policies are "Benefits", Not features.
upvoted 1 times

  **Mr_zach** 1 year, 10 months ago

Selected Answer: AB

A. Billing dashboards
B. Resource labels
upvoted 1 times

  **sedado77** 2 years, 1 month ago

Selected Answer: AB

A for Visualization and B for grouping
upvoted 4 times

  **SomeshSule** 2 years, 2 months ago

A and B for cost visibility
upvoted 2 times

  **victory108** 2 years, 3 months ago

Selected Answer: AB

A. Billing dashboards
B. Resource labels
upvoted 5 times

  **kohlakos** 2 years, 4 months ago

Selected Answer: AB

Haha, financial governance policy is not a google feature
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: BD

B. Resource labels
D. Financial governance policies
upvoted 2 times

Your organization needs to process large amounts of data from an online application that operates continuously. You do not want to be required to provision infrastructure or create server clusters. What should your organization choose?

- A. Compute Engine with BigQuery
- B. Dataproc
- C. Google Kubernetes Engine with Cloud Bigtable
- D. Dataflow

Correct Answer: D

Community vote distribution

D (100%)

- mvdm1111**

Highly Voted

2 years, 5 months ago

Here it is D

upvoted 12 times
- madcloud32**

Most Recent

6 months ago

Selected Answer: D

D is correct

Batch data + continuing stream is for Dataflow managed service.

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: D

Dataflow is a fully managed streaming data processing service that can handle large amounts of data from continuous sources. It is serverless, so you do not need to provision infrastructure or create server cluster

upvoted 2 times
- mdsarfraz69**

9 months ago

Selected Answer: D

D is correct

upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: D

To process large amounts of data from an online application continuously without the need to provision infrastructure or create server clusters, your organization should choose Google Cloud Dataflow.

Google Cloud Dataflow is a fully managed and serverless data processing service. It allows you to process and analyze large volumes of data in real-time or batch mode without the need to manage the underlying infrastructure. Dataflow automatically scales the processing resources based on the volume of data being processed, ensuring high performance and efficiency.

upvoted 3 times
- atoledo**

1 year, 8 months ago

Selected Answer: D

D: "You do not want to be required to provision infrastructure or create server clusters".

upvoted 4 times
- KC_go_reply**

1 year, 3 months ago

Yep. Compute Engine, Kubernetes Engine, and Dataproc are all IaaS offerings.

upvoted 3 times
- sedado77**

2 years, 1 month ago

Selected Answer: D

Yup, A is not correct in any universe. Do with D

upvoted 1 times
- EastSideCloudGuy**

2 years, 2 months ago

Selected Answer: D

ExamTopics team, why do you consider this A?

I and the rest of the commenters voted D, are we missing something?
upvoted 1 times

  **EastSideCloudGuy** 2 years, 2 months ago

I would add that the question also states: You do not want to be required to provision infrastructure or create server clusters.

Compute Engine by its nature requires provisioning, while DataFlow is mean to be serverless: <https://cloud.google.com/dataflow>
- Excerpt: Unified stream and batch data processing that's serverless, fast, and cost-effective.
upvoted 3 times

  **kyat** 2 years, 2 months ago

D dataflow for processing continuos stream and serverless too..
upvoted 3 times

  **victory108** 2 years, 3 months ago

Selected Answer: D

D. Dataflow
upvoted 1 times

  **ap6491** 2 years, 3 months ago

D : <https://cloud.google.com/dataflow#section-7>
upvoted 4 times

  **francozalamena** 2 years, 3 months ago

Selected Answer: D

D makes so much more sense
upvoted 1 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: D

D. Dataflow
upvoted 1 times

  **ahsangh** 2 years, 5 months ago

why not D. Dataflow for continuous data
upvoted 3 times

Your organization needs to ensure that the Google Cloud resources of each of your departments are segregated from one another. Each department has several environments of its own: development, testing, and production. Which strategy should your organization choose?

- A. Create a project per department, and create a folder per environment in each project.
- B. Create a folder per department, and create a project per environment in each folder.
- C. Create a Cloud Identity domain per department, and create a project per environment in each domain.
- D. Create a Cloud Identity domain per environment, and create a project per department in each domain.

Correct Answer: B

Reference:
<https://cloud.google.com/identity/docs/setup>

Community vote distribution
B (97%)

-   **Daffy16**

Highly Voted 

 2 years ago

Selected Answer: B

Its will be B. Organization > Folder > Project > Resource
upvoted 15 times
-   **ravisinha** 1 year, 11 months ago

today , this question was asked in exam
upvoted 7 times
-   **cookieMr**

Highly Voted 

 1 year ago

Selected Answer: B

The recommended strategy for ensuring segregation of Google Cloud resources for each department with multiple environments (development, testing, and production) would be to create a folder per department, and create a project per environment in each folder.

This strategy provides a clear and organized structure for managing resources. By creating a folder per department, you can logically group projects and enforce access controls and policies at the folder level. Within each folder, you can then create separate projects for each environmer (development, testing, and production). This allows for better isolation and control of resources within each environment.
upvoted 5 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: B

B is correct.
upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct
upvoted 1 times
-   **An1990** 1 year, 3 months ago

<https://cloud.google.com/resource-manager/docs/creating-managing-folders>
upvoted 1 times
-   **someshseghal** 1 year, 6 months ago

Selected Answer: B

B WITHOUT DOUBT
upvoted 1 times
-   **sedado77** 2 years, 1 month ago

Selected Answer: B

Yup, B is correct
upvoted 2 times

🗨️ 👤 **merlock8** 2 years, 1 month ago
B - <https://cloud.google.com/resource-manager/docs/creating-managing-folders>
upvoted 1 times

🗨️ 👤 **victory108** 2 years, 3 months ago
Selected Answer: B
B. Create a folder per department, and create a project per environment in each folder.
upvoted 2 times

🗨️ 👤 **leyunjohn** 2 years, 4 months ago
Selected Answer: B
B. Create a folder per department, and create a project per environment in each folder.
upvoted 2 times

🗨️ 👤 **Halimb** 2 years, 4 months ago
Selected Answer: B
I vote D. This solution given here is nonsense.
upvoted 3 times

🗨️ 👤 **Halimb** 2 years, 4 months ago
I mean B!!!
upvoted 4 times

🗨️ 👤 **Aman1001** 2 years, 5 months ago
Answer should be B, a folder for each department and projects for diff environments
upvoted 2 times

🗨️ 👤 **bulbsutr** 2 years, 5 months ago
B is correct
upvoted 1 times

🗨️ 👤 **ahsangh** 2 years, 5 months ago
Selected Answer: A
A. Create a project per department, and create a folder per environment in each project.
upvoted 1 times

🗨️ 👤 **ahsangh** 2 years, 5 months ago
correction: B. Create a folder per department, and create a project per environment in each folder.

<https://cloud.google.com/resource-manager/docs/cloud-platform-resource-hierarchy>
upvoted 9 times

Your organization is defining the resource hierarchy for its new application in Google Cloud. You need separate development and production environments. The production environment will be deployed in Compute Engine in two regions. Which structure should your organization choose?

- A. Create a single project for all environments. Use labels to segregate resources by environment.
- B. Create a single project for all environments. Use tags to segregate resources by environment.
- C. Create one project for the development environment and one project for the production environment.
- D. Create two projects for the development environment and two projects for the production environment (one for each region).

Correct Answer: C
Reference:
<https://cloud.google.com/resource-manager/docs/creating-managing-projects>
Community vote distribution
C (100%)

-   **fpreli** Highly Voted 2 years, 5 months ago

C is easier and you won't get less than D. I think creating one project per region wouldn't create benefit.
upvoted 9 times
-   **cookieMr** Highly Voted 1 year ago

Selected Answer: C

The recommended structure for your organization's resource hierarchy in Google Cloud would be C. Create one project for the development environment and one project for the production environment.

This structure provides clear separation between the development and production environments, allowing for better management and control of resources. By having separate projects, you can enforce different access controls, configurations, and deployments for each environment. This ensures that changes made in the development environment do not impact the production environment.

Option A and B suggest using a single project for all environments and relying on labels or tags to segregate resources. While labels and tags can be useful for organizing and categorizing resources, they may not provide the same level of isolation and control as separate projects.

Option D suggests creating multiple projects for each environment, with separate projects for the development and production environments in each region. While this approach could work, it may introduce unnecessary complexity, especially if the two regions have similar configurations and requirements.

upvoted 5 times
-   **chai_gpt** Most Recent 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
-   **__rajan__** 8 months, 1 week ago

Selected Answer: C

C is correct.
upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct
upvoted 1 times
-   **ajhak** 1 year, 5 months ago

Selected Answer: C

The part of compute engine being in two regions is meant to confuse you. It's not relevance.
upvoted 4 times
-   **deCalle** 1 year, 2 months ago

welcome to gcp
upvoted 3 times
-   **Pou1ze** 1 year, 6 months ago

Selected Answer: C

C is correct
upvoted 1 times

-   **ravisinha** 1 year, 12 months ago
C . you can create two VM & in two regions in one proejct
upvoted 2 times
-   **tcca** 2 years ago
D - <https://cloud.google.com/appengine/docs/standard/php/creating-separate-dev-environments>
upvoted 1 times
-   **juangt** 2 years, 1 month ago

Selected Answer: C

Answer is C, <https://cloud.google.com/resource-manager/docs/creating-managing-projects>
upvoted 1 times
-   **sedado77** 2 years, 1 month ago

Selected Answer: C

D is INCORRECT. correct answer is C
upvoted 1 times
-   **merlock8** 2 years, 1 month ago

C - <https://cloud.google.com/resource-manager/docs/creating-managing-folders>
upvoted 1 times
-   **SomeshSule** 2 years, 2 months ago

Resources can be deployed across regions irrespective of their Project ID. No need to create separate Projects Region-wise.
upvoted 2 times
-   **kyat** 2 years, 2 months ago

C it is .no concept of creating project for each region.
upvoted 3 times
-   **victory108** 2 years, 3 months ago

Selected Answer: C

C. Create one project for the development environment and one project for the production environment.
upvoted 2 times
-   **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. Create one project for the development environment and one project for the production environment.
upvoted 2 times
-   **desertlotus1211** 2 years, 4 months ago

There is no such thing as creating a project by region... Resources are either global or regional....A project is just a 'project'...
When you set up a project - where do you configure the region?

<https://cloud.google.com/resource-manager/docs/creating-managing-projects>
upvoted 4 times

Your organization meant to purchase a 3-year Committed Use Discount, but accidentally purchased a 1-year Committed Use Discount instead. What should your organization do?

- A. Contact your financial institution.
- B. Contact Trust and Safety.
- C. Contact Cloud Billing Support.
- D. Contact Technical Support.

Correct Answer: C
Reference:
<https://cloud.google.com/compute/docs/instances/signing-up-committed-use-discounts>
Community vote distribution
C (100%)

- mintupintu**

Highly Voted

2 years, 5 months ago

Totally agree that correct answer is C. Disappointing to see back to back wrong/misleading answers to the questions

upvoted 14 times
- GerryY**

Highly Voted

2 years, 1 month ago

Wonder why answers by Exam Cloud has been providing bad answers for the last 20 or so questions.

upvoted 6 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__**

8 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: C

If your organization accidentally purchased a 1-year Committed Use Discount (CUD) instead of the intended 3-year CUD you can reach out to Google Cloud Support as soon as possible to explain the situation and request assistance. They can guide you through the process and provide options for resolving the issue.

If the purchase is recent and you haven't started using the committed usage, you may have the option to cancel the 1-year CUD and repurchase the 3-year CUD.

upvoted 4 times
- SomeshSule**

1 year, 10 months ago

Such a bad answer given in the exam.

upvoted 1 times
- CDL_Learner**

1 year, 11 months ago

Selected Answer: C

C is the correct answer.

upvoted 1 times
- brainsaki**

1 year, 11 months ago

Selected Answer: C

C is the correct answer.

upvoted 1 times
- fernandinhoo15**

2 years ago

Selected Answer: C

C. Contact Cloud Billing Support

upvoted 2 times
- sedado77**

2 years, 1 month ago

Selected Answer: C

... Again correct is C, not the highlighted one
Thank you Examtopics for the voting system!
upvoted 2 times

  **kyat** 2 years, 2 months ago
C it is ..billing support for billing related resolutions.
upvoted 2 times

  **victory108** 2 years, 3 months ago
Selected Answer: C
C. Contact Cloud Billing Support
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago
Selected Answer: C
C. Contact Cloud Billing Support.
upvoted 1 times

  **kapomony** 2 years, 5 months ago
Selected Answer: C
is a billing error
upvoted 2 times

  **fpreli** 2 years, 5 months ago
I would go with C, I don't think Security and Compliance has nothing to do with CUD
upvoted 4 times

  **ahsangh** 2 years, 5 months ago
C. Contact Cloud Billing Support.
upvoted 1 times

  **kiranm2021** 2 years, 5 months ago
???? Contact Trust and Safety for billing issues ?
upvoted 3 times

Your organization needs to allow a production job to have access to a BigQuery dataset. The production job is running on a Compute Engine instance that is part of an instance group.

What should be included in the IAM Policy on the BigQuery dataset?

- A. The Compute Engine instance group
- B. The project that owns the Compute Engine instance
- C. The Compute Engine service account
- D. The Compute Engine instance

Correct Answer: C

Reference:
<https://cloud.google.com/compute/docs/instance-groups>

Instance groups

Send feedback

An instance group is a collection of virtual machine (VM) instances that you can manage as a single entity.

Compute Engine offers two kinds of VM instance groups, managed and unmanaged:

- **Managed instance groups** (MIGs) let you operate apps on multiple identical VMs. You can make your workloads scalable and highly available by taking advantage of automated MIG services, including: autoscaling, autohealing, regional (multiple zone) deployment, and automatic updating.
- **Unmanaged instance groups** let you load balance across a fleet of VMs that you manage yourself.

Community vote distribution

C (100%)

- fpreli

Highly Voted

2 years, 5 months ago

C also for me, also because MIG uses the Compute Engine default service account for its own operation

upvoted 9 times
- AbdulioUnited

Highly Voted

2 years, 4 months ago

It seems to be "C"

When an identity calls a Google Cloud API, BigQuery requires that the identity has the appropriate permissions to use the resource. You can grant permissions by granting roles to a user, a group, or a service account.

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 5 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: C

To allow a production job running on a Compute Engine instance group to have access to a BigQuery dataset, you should include the service account associated with the Compute Engine instance group to the IAM policy of the BigQuery dataset. The service account represents the Compute Engine instances and is used to authenticate the job's access to the dataset.

By including the service account associated with the Compute Engine instance group in the IAM policy and granting the appropriate BigQuery

roles or permissions, you ensure that the production job running on the Compute Engine instances has the necessary access to the BigQuery dataset.

upvoted 3 times

  **manashbaruah** 1 year, 6 months ago

Selected Answer: C

Keypoint: Production job needs access and not human

upvoted 1 times

  **Sheryb** 1 year, 6 months ago

Selected Answer: C

Service accounts are for bots and applications

upvoted 2 times

  **SomeshSule** 2 years, 2 months ago

Why are the site Admins correcting the Answers? Even though the "Question Bank updated date" shows latest, still answers are not corrected.

upvoted 2 times

  **SomeshSule** 2 years, 2 months ago

*not correcting

upvoted 1 times

  **kyat** 2 years, 2 months ago

C :service accounts are used to access resources by other resources.

upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: C

C. The Compute Engine service account

upvoted 2 times

  **Alexanderswor** 2 years, 3 months ago

Selected Answer: C

https://cloud.google.com/bigquery/docs/dataset-access-controls#granting_access_to_a_dataset

upvoted 3 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. The Compute Engine service account

upvoted 1 times

  **Aman1001** 2 years, 5 months ago

C should be the answer

upvoted 3 times

  **mintupintu** 2 years, 5 months ago

Answer is C.

upvoted 2 times

  **kapomony** 2 years, 5 months ago

answer C --> <https://cloud.google.com/bigquery/docs/dataset-access-controls>

upvoted 3 times

  **mvdm11111** 2 years, 5 months ago

It should be C as this is what the service account is about

upvoted 2 times

Your team is publishing research results and needs to make large amounts of data available to other researchers within the professional community and the public at minimum cost.
How should you host the data?

- A. Use a Cloud Storage bucket and enable "Requester Pays."
- B. Use a Cloud Storage bucket and provide Signed URLs for the data files.
- C. Use a Cloud Storage bucket and set up a Cloud Interconnect connection to allow access to the data.
- D. Host the data on-premises, and set up a Cloud Interconnect connection to allow access to the data.

Correct Answer: A

Community vote distribution

A (71%)B (29%)

- jvg637** Highly Voted 2 years, 5 months ago

A. Enabling Requester Pays is useful, for example, if you have a lot of data you want to make available to users, but you don't want to be charged for their access to that data.
<https://cloud.google.com/storage/docs/requester-pays>
upvoted 19 times
- VenRam** 2 years, 4 months ago

Agree, Answer A Enabling Requester Pays is useful.
Signed URLs are more about limiting the window of making the request and not about minimizing cost.
upvoted 5 times
- pspandher** 1 year, 8 months ago

The requester pays works on Projects Inside your GCP Env. in that the requester will pay for request not the primary project from where the request is made. This is not Public as per Google Documentation I believe.
upvoted 3 times
- cookieMr** Highly Voted 1 year ago

Selected Answer: A

Enabling "Requester Pays" on a Cloud Storage bucket allows you to shift the cost of accessing the data to the requester. This means that the individuals or organizations accessing the data will be responsible for the cost of data transfer and operations, rather than your organization incurring the expenses.

With this approach, you can make the data publicly accessible in the Cloud Storage bucket while ensuring that the cost of data access is covered by the requesters. This can help reduce your organization's expenses in hosting and serving the data.

Option B (providing Signed URLs) would require you to generate and manage unique URLs for each requester, which might not be practical for making the data available to a wide audience.
upvoted 16 times

prachisri 10 months, 2 weeks ago

Very useful comment. Explains well
upvoted 3 times

chai_gpt Most Recent 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

__rajan__ 8 months, 1 week ago

Selected Answer: A

Requester Pays is a feature of Cloud Storage that allows you to transfer the cost of data egress to your users. This means that users who download your data will be charged for the egress traffic, not you.
upvoted 1 times

RK_MANOJ 8 months, 2 weeks ago

Selected Answer: B

Cloud Storage Bucket: Host your data in a Google Cloud Storage bucket, which provides scalable and reliable storage.

Signed URLs: Create and provide signed URLs to users. Signed URLs grant temporary access to specific objects in your Cloud Storage bucket without the need for users to have Google Cloud accounts or pay for egress costs. This approach allows you to maintain control over who accesses your data while minimizing costs for your users.

Option B is a better choice for sharing data with a wider public audience while keeping costs low and ensuring accessibility without requiring users to pay for egress fees.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: A

A is correct

upvoted 1 times

  **Atmatap** 1 year, 2 months ago

Selected Answer: A

Minimum cost means using GCS instead of on premise server. And use Requester Pay also help reduce the cost

upvoted 1 times

  **MHx1** 1 year, 3 months ago

The answer is A

Enabling Requester Pays is useful, for example, if you have a lot of data you want to make available to users, but you don't want to be charged for their access to that data.

<https://cloud.google.com/storage/docs/requester-pays>

upvoted 1 times

  **skybd3hel** 1 year, 3 months ago

Selected Answer: A

<https://cloud.google.com/storage/docs/requester-pays>

upvoted 1 times

  **alfantastic** 1 year, 5 months ago

Selected Answer: B

making requester to pay may defeat the sharing purpose, answer B does not create readblocks

upvoted 1 times

  **someshseghal** 1 year, 6 months ago

Selected Answer: A

<https://cloud.google.com/storage/docs/requester-pays>

upvoted 1 times

  **vicjean** 1 year, 6 months ago

Selected Answer: A

https://www.youtube.com/watch?v=nDm11wg2Fn4&ab_channel=GoogleCloudTech

upvoted 1 times

  **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: A

In reading through Google docs on Signed URLs there is no mention of cost minimization.

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 1 times

  **jellybiscuit** 1 year, 8 months ago

Selected Answer: B

I just don't see how it's practical to ask everyone to have a google account to bill against. A might be what they're looking for, but I'm just seeing i as a practical solution.

If you want to make them pay... do so through an app, then provide them a signed url for the download. It's more complex to implement, but it actually works.

upvoted 6 times

  **Dan_W** 1 year, 8 months ago

It's B. Signed URL's are available to the public, and you control who has access. Requester pays - With Requester Pays enabled on your bucket, you can require requesters to include a billing project in their requests, thus billing the requester's project

Requestor pays is not for the public

upvoted 4 times

  **brainysaki** 1 year, 8 months ago

Selected Answer: A

A: Use a Cloud Storage bucket and enable Requester Pays.

upvoted 2 times

  **zelck** 1 year, 9 months ago

Selected Answer: A

A is my answer.

upvoted 1 times

Your company needs to segment Google Cloud resources used by each team from the others. The teams' efforts are changing frequently, and you need to reduce operational risk and maintain cost visibility. Which approach does Google recommend?

- A. One project per team.
- B. One organization per team.
- C. One project that contains all of each team's resources.
- D. One top-level folder per team.

Correct Answer: A
Reference:
<https://cloud.google.com/security/infrastructure/design>
Community vote distribution

D (75%)

A (25%)

- mamitos

Highly Voted

1 year, 8 months ago

Selected Answer: D

It should be D, a top-level folder for each team give the ability to have multiple projects(I guess here goes the "teams efforts change frequently) and also cost is closely being watched and separated by Team->Team_Project.

I hate GCP, so many poorly worded and vague questions...shame

upvoted 38 times
- geeroylenkins

1 year, 2 months ago

In a company of 80,000, you might have 8,000 teams. Maybe 16,000. You want a top level folder for EACH? You don't need a TOP level folder fo each. Maybe 1 per BU, then 1 per BU Department, then 1 per BU Department Division, then 1 per BU Department Division Team, etc. Nested folders, like Russian Dolls.

upvoted 3 times
- desertlotus1211

Highly Voted

2 years, 4 months ago

Answer is A.... The Teams need to segmented to have visibility on the resources each team consumes

upvoted 7 times
- 4a98421

Most Recent

3 months, 2 weeks ago

A is correct, there is only 1 top level for Organization, then folders as per department/team and then projects for each department/team.

upvoted 1 times
- sivakarthick16

5 months, 4 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- sanidh

6 months, 4 weeks ago

Selected Answer: A

.....

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: D

D. One top-level folder per team.

Folders in Google Cloud provide a way to organize and manage resources hierarchically, making it easier to allocate resources to specific teams, manage access control, and maintain cost visibility. Each top-level folder can represent a different team, and you can create a hierarchy of folders to further organize resources as needed.

upvoted 2 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct

upvoted 1 times

  **Lufly** 10 months ago

Selected Answer: D

Here are some reasons why using one project per team is not a good idea for resource segmentation:

Inflexibility: It would prevent teams from changing their resources frequently. If a team needs to add or remove resources, they would need to create a new project or delete the old project. This can be time-consuming and disruptive.

Cost: It would be more expensive to manage multiple projects. Each project has its own set of costs, such as storage costs, compute costs, and networking costs. This can add up quickly if you have a lot of projects.

Complexity: It would be more complex to manage multiple projects. You would need to keep track of the resources in each project, as well as the costs of each project. This can be difficult to do, especially if you have a lot of projects.

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

One project per team approach offers clear separation and isolation of resources for each team, providing enhanced security, compliance, and operational control. Each team can have its own dedicated project, allowing them to independently manage their resources, control access permissions, and maintain cost visibility. Changes made by one team will not directly impact the resources of other teams, reducing operational risks and providing a more stable environment.

upvoted 3 times

  **ptoul74** 1 year, 1 month ago

Have a look at this image https://cloud.google.com/static/resource-manager/img/cloud-hierarchy.svg?dcb_=0.14515370615705625.

Then, the correct answer is obviously D.

upvoted 6 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: A

Option D, "One top-level folder per team" is a possible approach to organizing Google Cloud resources, but it is not the recommended approach for the scenario described in the question.

Folders are a way to organize projects and resources hierarchically within a Google Cloud organization. They are useful for creating custom organizational structures, managing access controls, and delegating administrative responsibilities.

However, if the goal is to separate resources used by each team from the others and maintain cost visibility, organizing resources into one project per team or one organization per team would provide more clear separation and better cost tracking.

Therefore, option A, "One project per team," is the recommended approach for this scenario.

upvoted 3 times

  **santoshh** 1 year, 3 months ago

It should be 'A' as it says "segment Google Cloud resources". if we follow the best practices of resource hierarchy Org->folders-subfolder->Project >resources. so the idea here is to create two project for each team to segment the resources with clear visibility

upvoted 2 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

Answer A

upvoted 3 times

  **Jackey0117** 1 year, 5 months ago

Selected Answer: A

Google recommends using one project per team as the best approach to segmenting Google Cloud resources used by each team. Each project is a separate container for resources, with its own set of IAM policies and firewall rules. This allows each team to have full control over their own resources and also allows for cost visibility for each team. By creating a project for each team, you can also reduce operational risk, as changes made by one team will not affect the resources of another team.

upvoted 2 times

  **someshseghal** 1 year, 6 months ago

Selected Answer: A

GCP projects allows tracking and controlling: <https://cloud.google.com/security/infrastructure/design>

upvoted 1 times

  **sfg** 1 year, 6 months ago

Selected Answer: A

Correct answer A.

D is not correct, as having folders per team is overkill. Folders are meant for separate departments, whereas projects are for separating out teams.

upvoted 1 times

How do Migrate for Compute Engine and Migrate for Anthos differ?

- A. Unlike Migrate for Anthos, Migrate for Compute Engine assumes that the migration source is VMware vSphere.
- B. Migrate for Compute Engine charges for ingress, but Migrate for Anthos does not.
- C. Migrate for Compute Engine is closed source, and Migrate for Anthos is open source.
- D. Migrate for Anthos migrates to containers, and Migrate for Compute Engine migrates to virtual machines.

Correct Answer: D

Reference:

<https://cloud.google.com/migrate/anthos>

Migrate for [Anthos](#) and [GKE](#) makes it fast and easy to modernize traditional applications away from virtual machines and into native containers. Our unique automated approach extracts the critical application elements from the VM so you can easily insert those elements into containers in [Google Kubernetes Engine](#) or [Anthos clusters](#) without the VM layers (like Guest OS) that become unnecessary with containers.

This significantly reduces the cost and labor that would be required for a manual application modernization project. Upon completion, it also empowers your teams to deploy, operate, and maintain existing applications more efficiently and cost effectively, using the same services, policies, and methodologies of a modern platform.



The diagram illustrates a migration process. On the left, a red container with a blue Kubernetes logo represents a source environment. A dashed line connects it to a central cloud icon. From the cloud, another dashed line leads to a stack of server racks on the right, representing the destination environment. A colorful triangle logo is positioned between the cloud and the server racks.

Craft your ideal migration journey

Most digital transformations will be a mix of strategies. For the workloads that will benefit from containers, Migrate for Anthos and GKE delivers a fast, smooth path to modernization. For other workloads that are better suited as a VM, simply move them as is with [Migrate for Compute Engine](#) and leverage VPC network integration with GKE. Don't settle for being locked into existing infrastructure or one migration path. With Google, run your workloads how you want, where you want.



The illustration shows a laptop on the left and a server rack on the right, both connected by a dashed line that passes through a cloud icon in the center, symbolizing cloud-based migration.

Community vote distribution

D (100%)

-  **kohlakos**

Highly Voted 

2 years, 4 months ago
- Selected Answer: D

You might decide to break your migration journey into two distinct phases, even for workloads that are suitable for containers:

Migrate workloads to Compute Engine with Migrate for Compute Engine.
Migrate from Compute Engine to containers with Migrate for Anthos and GKE.
This method makes sense, for instance, in cases where you want to conduct a data-center migration and migrate all workloads into Compute Engine, and only at a second stage selectively modernize suitable workloads to containers.

<https://cloud.google.com/migrate/anthos/docs/architecture>
upvoted 5 times
-  **chai_gpt**

Most Recent 

7 months, 3 weeks ago
- Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months, 1 week ago

Selected Answer: D

D is correct.
upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

Migrate for Compute Engine is focused on migrating VM-based workloads to Compute Engine, while Migrate for Anthos extends the migration capabilities to containerized workloads and enables modernization through the adoption of containerization and Kubernetes-based orchestration.
upvoted 2 times

  **Rakesh1976** 1 year, 8 months ago

My Answer also is D
upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: D

D. Migrate for Anthos migrates to containers, and Migrate for Compute Engine migrates to virtual machines.
upvoted 3 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: D

D. Migrate for Anthos migrates to containers, and Migrate for Compute Engine migrates to virtual machines.
upvoted 1 times

  **mintupintu** 2 years, 5 months ago

D seems the correct answer.
upvoted 1 times

Your large and frequently changing organization's user information is stored in an on-premises LDAP database. The database includes user passwords and group and organization membership.

How should your organization provision Google accounts and groups to access Google Cloud resources?

- A. Replicate the LDAP infrastructure on Compute Engine
- B. Use the Firebase Authentication REST API to create users
- C. Use Google Cloud Directory Sync to create users
- D. Use the Identity Platform REST API to create users

Correct Answer: C

You can run a single instance of Google Cloud Directory Sync to synchronize user accounts and groups to Google Cloud.

Reference:

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction>

Federating Google Cloud with Active Directory

[Send feedback](#)

This article describes how you can configure Cloud Identity or Google Workspace to use [Active Directory as IdP and authoritative source](#).

The article compares the logical structure of Active Directory with the structure used by Cloud Identity and Google Workspace and describes how you can map Active Directory forests, domains, users, and groups. The article also provides a [flowchart](#) that helps you determine the best mapping approach for your scenario.

This article assumes that you're familiar with Active Directory.

Community vote distribution

C (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: C

To provision Google accounts and groups to access Google Cloud resources based on your on-premises LDAP database, you can set up Google Cloud Directory Sync (GCDS). GCDS is a tool provided by Google Cloud that synchronizes user and group information between your on-premises LDAP directory and Google Workspace (formerly G Suite) directory. Configure GCDS to connect to your LDAP server and sync user accounts, passwords, groups, and organization membership information to Google Cloud.

upvoted 4 times
- sfg

1 year, 6 months ago

Selected Answer: C

C is correct

upvoted 2 times

  **skarths** 1 year, 7 months ago

Selected Answer: C

C is the answer
upvoted 1 times

  **tcca** 2 years ago

C -
<https://support.google.com/a/answer/106368?hl=en>
upvoted 1 times

  **victory108** 2 years, 3 months ago

Selected Answer: C

C. Use Google Cloud Directory Sync to create users
upvoted 2 times

  **leyunjohn** 2 years, 4 months ago

Selected Answer: C

C. Use Google Cloud Directory Sync to create users
upvoted 1 times

  **mintupintu** 2 years, 5 months ago

C seems the correct answer.
upvoted 1 times

Your organization recently migrated its compute workloads to Google Cloud. You want these workloads in Google Cloud to privately and securely access your large volume of on-premises data, and you also want to minimize latency. What should your organization do?

- A. Use Storage Transfer Service to securely make your data available to Google Cloud
- B. Create a VPC between your on-premises data center and your Google resources
- C. Peer your on-premises data center to Google's Edge Network
- D. Use Transfer Appliance to securely make your data available to Google Cloud

Correct Answer: B

Community vote distribution

B (61%)

C (37%)

- rga91

Highly Voted

2 years, 1 month ago

"privately and securely" doesnt fit with peering.... Answer is B

upvoted 15 times
- CiscoTerminator

3 months, 3 weeks ago

A VPC does not establish a connection its just a container not a VPN

upvoted 1 times
- Loved

1 year, 8 months ago

Direct Peering or Carrier peering are both private and secure

upvoted 10 times
- KC_go_reply

1 year, 3 months ago

You know B says 'VPC' and not 'VPN', right? A VPC is a virtual network, not a connection itself. You can't use a VPC alone to connect between on-prem and GCP.

upvoted 7 times
- jexmtropicscheatchatya

Highly Voted

1 year, 10 months ago

Answer is C: Peer your on-premises data center to Googles Edge Network

<https://cloud.google.com/blog/products/networking/understanding-google-cloud-network-edge-points>

Scroll down to or Find "Cloud Interconnect POPs" where it says the following:

"Dedicated Interconnect provides direct physical connections between your on-premises network and Google's network. Dedicated Interconnect enables you to efficiently transfer large amounts of data between networks."

upvoted 11 times
- rudi009

6 months, 3 weeks ago

But the question doesn't want you to transfer the data rather make it available for access.

upvoted 3 times
- ceeceevee

Most Recent

6 months, 3 weeks ago

Selected Answer: B

Direct Peering documentation mostly talk about use-cases where on-prem resources need to hit Google services and in fact specifically calls out that it "Does not provide direct access to VPC network resources that have only internal IP addresses." Therefore, the correct answer is B.

upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: C

I think it should be C. As there is large amiu nt of data that require low latency for which C is correct.

upvoted 1 times
- Jack456

8 months, 2 weeks ago

Selected Answer: C

C IS RIGHT

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct
upvoted 1 times

  **kushalesh** 10 months, 2 weeks ago

Selected Answer: B

create a VPC between your on-premises data center and your Google resources. This can be achieved by using Private Google Access for on-premises hosts, which provides a way for on-premises systems to connect to Google APIs and services by routing traffic through a Cloud VPN tunnel or a VLAN attachment for Cloud Interconnect. This solution should minimize latency as traffic is routed through a Cloud VPN tunnel or a VLAN attachment for Cloud Interconnect, instead of traversing the public internet.
upvoted 1 times

  **Chemssou** 10 months, 2 weeks ago

Selected Answer: B

"privately and securely"
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: B

Creating a Virtual Private Cloud (VPC) between your on-premises data center and Google Cloud allows for a private and secure network connection. It provides a dedicated network environment for your organization, allowing your compute workloads in Google Cloud to privately and securely access your on-premises data. This option helps minimize latency and ensures a direct and controlled connection between your on-premises infrastructure and Google Cloud resources.
upvoted 4 times

  **ptoul74** 1 year, 1 month ago

The correct answer is C.
This can't be B for these two reasons. VPC alone doesn't provide a secure connection. It requires a VPN. Then, the question also states to minimize latency which is actually the main weakness of any VPN. Therefore, B is not an option
upvoted 4 times

  **Atmatap** 1 year, 2 months ago

Selected Answer: C

Like jexmtropicscheatchatya
upvoted 1 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: B

Option C is incorrect because peering your on-premises data center to Google's Edge Network is only used to improve the network performance between your on-premises network and Google's cloud services. It does not provide a secure way to access on-premises data from Google Cloud. A Virtual Private Cloud (VPC) with a VPN or Cloud Interconnect provides a secure way to privately access on-premises data from Google Cloud.
upvoted 3 times

  **AKhaled** 1 year, 2 months ago

Selected Answer: C

VPC is a virtual version of a physical network that is implemented inside of Google's production network.
upvoted 1 times

  **MHx1** 1 year, 3 months ago

B. Cloud VPN securely extends your on-prem network to your Google's network through an IPSec VPN tunnel
upvoted 1 times

  **kjkjkj_kk** 1 year, 3 months ago

Selected Answer: B

Creating a VPC between your on-premises data center and your Google Cloud resources is the best option to privately and securely access your on-premises data. This approach provides a dedicated, private connection between your on-premises data center and your Google Cloud resources, which ensures security and minimizes latency. The connection is established using a VPN tunnel or Cloud Interconnect, depending on your specific requirements. Using a VPC also allows you to control network traffic between your on-premises data center and Google Cloud, providing greater control and security.
upvoted 1 times

  **kjkjkj_kk** 1 year, 3 months ago

Answer is B, Creating a VPC between your on-premises data center and your Google Cloud resources is the best option to privately and securely access your on-premises data. This approach provides a dedicated, private connection between your on-premises data center and your Google Cloud resources, which ensures security and minimizes latency. The connection is established using a VPN tunnel or Cloud Interconnect, depending on your specific requirements. Using a VPC also allows you to control network traffic between your on-premises data center and Google Cloud, providing greater control and security.
upvoted 1 times

Your organization consists of many teams. Each team has many Google Cloud projects. Your organization wants to simplify the management of identity and access policies for these projects.
How can you group these projects to meet this goal?

- A. Group each team's projects into a separate domain
- B. Assign labels based on the virtual machines that are part of each team's projects
- C. Use folders to group each team's projects
- D. Group each team's projects into a separate organization node

Correct Answer: B

Community vote distribution

C (96%)4%

- astood

Highly Voted

2 years, 5 months ago

c is the answer

upvoted 10 times
- aUMMDqZmjLBx

Most Recent

3 months, 1 week ago

Selected Answer: C

Org > Folder > Project >Resource

upvoted 1 times
- the_var

5 months, 4 weeks ago

Selected Answer: C

Org > Folder > Project >Resource

upvoted 2 times
- 1Style

6 months ago

Selected Answer: C

c is the ans

upvoted 1 times
- ceeevee

6 months, 3 weeks ago

Selected Answer: C

B does not cover all use-cases, since it refers specifically to VMs. C is the correct answer IMHO.

upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months ago

C is correct.

upvoted 1 times
- dboy711

8 months, 2 weeks ago

Selected Answer: C

Folders in Google Cloud provide a hierarchical structure that allows you to organize and manage resources and projects effectively. You can apply identity and access policies at the folder level, making it easier to control permissions and policies for all projects within a folder.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- afeff

9 months ago

Selected Answer: C

this is confusing, the votes say C is the solution whereas the website shows B as correct.

upvoted 2 times

🗨️ 👤 **Bkt123** 9 months, 1 week ago

Correct answer is C
upvoted 1 times

🗨️ 👤 **hireshgupta** 11 months ago

The correct answer is C. Use folders to group each team's projects.

Folders in Google Cloud Platform (GCP) provide a way to group projects together. This can be helpful for simplifying the management of identity and access policies, as well as for organizing your GCP resources.
upvoted 2 times

🗨️ 👤 **reachmalhotras** 11 months, 2 weeks ago

Answer is C But.... confusing part was IAM
upvoted 1 times

🗨️ 👤 **the_var** 11 months, 3 weeks ago

Selected Answer: C

C is the answer.
upvoted 1 times

🗨️ 👤 **cookieMr** 1 year ago

Selected Answer: C

To simplify the management of identity and access policies for multiple Google Cloud projects, you can group them using folders in the Google Cloud resource hierarchy.

By organizing your projects into folders, you can apply common IAM (Identity and Access Management) policies to all projects within a folder. This helps ensure consistent access controls and permissions across teams and projects. You can also delegate administrative responsibilities at the folder level, allowing team leads or project managers to manage the projects within their respective folders.

Folders provide a way to logically group projects based on organizational structure, business units, or any other criteria that aligns with your organization's needs. This helps streamline management and reduces the need to apply individual policies to each project separately.
upvoted 4 times

🗨️ 👤 **pillais** 1 year, 2 months ago

Selected Answer: C

C Multiple Projects to One Folder.
upvoted 1 times

🗨️ 👤 **tbolick6** 1 year, 3 months ago

Selected Answer: C

Multiple projects to a single folder.
upvoted 1 times

An organization needs to categorize text-based customer reviews on their website using a pre-trained machine learning model. Which Google Cloud product or service should the organization use?

A. Cloud Natural Language API

B. Dialogflow

C. Recommendations AI

D. TensorFlow

Correct Answer: C

Community vote distribution

A (97%)

- mamitos

Highly Voted

1 year, 8 months ago

Selected Answer: A

Seriously who marks the "correct" answers. So many wrongs....
It's ok if the question is vague or weird but like here, it's obvious...
upvoted 28 times
- jellybiscuit

1 year, 8 months ago

Agree.
Imagine how bad it was before the voting feature was implemented.
upvoted 14 times
- Kozato

Most Recent

3 days, 13 hours ago

Selected Answer: A

A is correct.
upvoted 1 times
- 9b8feb6

2 months, 1 week ago

Natural Language API discovers syntax, entities and sentiment in text and classifies texts into a predefined set of categories. It could capture comments about your business, ie, complaints, praises, attempt to learn about your business and more;
upvoted 1 times
- 4a98421

3 months, 2 weeks ago

Selected Answer: A

Cloud Natural Language API categorize text based input, it can do sentiment analytics on top of it. While Recommendation AI is to recommend customer another product
upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
- dboy711

8 months, 2 weeks ago

Selected Answer: A

Cloud Natural Language API is a powerful and easy-to-use service that provides pre-trained machine learning models for various natural language processing tasks, including sentiment analysis, entity recognition, and text classification. It's specifically designed for tasks like categorizing text or analyzing sentiment in reviews.
upvoted 2 times
- katyk

8 months, 2 weeks ago

please change correct answer to A. The Cloud Natural Language API is a pre-trained machine learning model that can analyze and categorize text-based data. It can identify entities, sentiment, and syntax in text, making it an ideal choice for categorizing customer reviews on a website. The other options provided (Dialogflow, Recommendations AI, and TensorFlow) are not specifically designed for text analysis and categorization, so they would not be the best choice for this task.
upvoted 4 times
- chandangupta8dec

9 months ago

Selected Answer: A

A is the right answer.
upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct
upvoted 1 times

  **Andromida23** 9 months ago

Selected Answer: A
upvoted 1 times

  **the_var** 11 months, 3 weeks ago

Selected Answer: A

Moderator, please change the Selected Answer. It is wrongly showing as C!!
upvoted 1 times

  **Learner721** 1 year ago

Selected Answer: A

A should be the correct answer as NLP can detect sentiment and categorize according to it
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: A

Google Cloud Natural Language API provides pre-trained machine learning models for various natural language processing tasks, including sentiment analysis, entity recognition, and text classification. In this case, the organization can leverage the text classification feature of the Natural Language API to categorize customer reviews into different categories or topics.
upvoted 2 times

  **pillais** 1 year, 2 months ago

Selected Answer: A

NL API is pretrained, and can be applied to conversation etc.
upvoted 1 times

  **[Removed]** 1 year, 3 months ago

Selected Answer: A

Natural Language API "pre-trained to recognize sentiment and categorize accordingly". Wish they were all this easy!
upvoted 1 times

  **kk2531** 1 year, 4 months ago

Selected Answer: B

Dialogflow is an NLP (Natural Language Processing) platform, which is used to develop an application related to the conversations and experience for the customers of the company in different languages on numerous platforms. Dialogflow is mainly used to build actions for most of the Google Assistant devices.
upvoted 1 times

  **BTAB** 1 year, 4 months ago

Wrong, dude. Dialogflow is chatbot agent you can use in an app or website, not for sentiment analysis or customer review.
upvoted 4 times

An organization is planning its cloud expenditure.
What should the organization do to control costs?

- A. Consider cloud resource costs as capital expenditure in annual planning.
- B. Use only cloud resources; they have no cloud infrastructure costs.
- C. Review cloud resource costs frequently because costs depend on usage.
- D. Assess cloud resources costs only when SLO is not met by their cloud provider.

Correct Answer: C

Community vote distribution

C (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: C

C is correct

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- the_var

11 months, 3 weeks ago

Selected Answer: C

C is the answer!

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: C

Continuously review your cost management strategy and make adjustments as needed. Stay informed about new cost optimization features and best practices provided by your cloud provider to ensure you are utilizing the latest cost-saving techniques.

upvoted 2 times
- AAD1

1 year, 9 months ago

Selected Answer: C

C is the answer

upvoted 3 times
- Govindaraj

1 year, 9 months ago

Selected Answer: C

C is the answer

upvoted 1 times
- Sav94

1 year, 9 months ago

I think C.

upvoted 1 times

An organization is searching for an open-source machine learning platform to build and deploy their own custom machine learning applications using TPUs.

Which Google Cloud product or service should the organization use?

- A. TensorFlow
- B. BigQuery ML
- C. Vision API
- D. AutoML Vision

Correct Answer: D

Community vote distribution

A (66%)

D (34%)

- Lufly

Highly Voted

10 months ago

Selected Answer: D

TensorFlow is an open-source machine learning framework that can be used to build and deploy custom machine learning applications. It is also compatible with TPUs, which are Google's custom-designed machine learning accelerators.

The other options are not as suitable for this scenario. Option B, BigQuery ML, is a machine learning service that can be used to build and deploy machine learning models, but it is not compatible with TPUs. Option C, Vision API, is a machine learning API that can be used to extract information from images, but it is not a platform for building and deploying custom machine learning applications. Option D, AutoML Vision, is a machine learning service that can be used to build and deploy machine learning models for image classification, but it is not compatible with TPUs.

upvoted 13 times
- the_var

Highly Voted

11 months, 3 weeks ago

TPU = TensorFlow = TPU = TensorFlow! Admin, please correct the selected answer!

upvoted 9 times
- vic.bonto.umlaut

Most Recent

3 weeks, 3 days ago

even Gemini says TensorFlow

upvoted 1 times
- Gabriella_Pacheco

4 months, 3 weeks ago

Selected Answer: A

TensorFlow is the correct answer for me

upvoted 1 times
- GamudaDavid

5 months, 2 weeks ago

Selected Answer: A

Bro, TPU = TensorFlow, this answer is weird.

upvoted 1 times
- Riise

5 months, 2 weeks ago

Selected Answer: A

AutoML Vision is limited to pictures only. It does not say anywhere in the question that ML training is needed for pictures. Because it does not mention the type of data needed to be processed I consider A the right answer, it is general and covers all scenarios. At a certain point, all 3 google products listed below (B,C and D) are using TensorFlow.

upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct. (Tensorflow)

upvoted 2 times
- __rajan__

8 months ago

Selected Answer: D

D is correct. TensorFlow is an open-source machine learning framework.

upvoted 2 times
- mdsarfraz69

9 months ago

Selected Answer: D

D is correct

upvoted 2 times

🗨️ 👤 **Gopi382** 11 months, 3 weeks ago

Tensorflow
upvoted 3 times

🗨️ 👤 **cookieMr** 1 year ago

Selected Answer: A

For building and deploying custom machine learning applications using TPUs, the organization can consider using Google Cloud's TensorFlow platform. TensorFlow is an open-source machine learning framework that provides comprehensive support for developing and deploying machine learning models, including integration with TPUs (Tensor Processing Units).
upvoted 5 times

🗨️ 👤 **telosd** 1 year, 1 month ago

Selected Answer: A

TPU = tensorflow. So many wrong answers on this module :(
upvoted 2 times

🗨️ 👤 **omgitsele** 1 year, 2 months ago

Selected Answer: A

whoever manages the answers, needs to sit for the exam, they probably will fail..When we talk about OPEN-Source using TPUs it's Tensorflow.
upvoted 1 times

🗨️ 👤 **tbolick6** 1 year, 4 months ago

Selected Answer: A

TensorFlow is a free and open-source software library for machine learning and artificial intelligence. TPUs used for TensorFlow
upvoted 2 times

🗨️ 👤 **Jackey0117** 1 year, 5 months ago

Selected Answer: A

TensorFlow is an open-source machine learning platform that can be used to build and deploy custom machine learning applications.
upvoted 1 times

🗨️ 👤 **thewalker** 1 year, 5 months ago

If you read the question completely, It says which Google Cloud product. Though TPUs goes with TensorFlow; in GCP, AutoML Vision is the product to achieve the mentioned requirement in the question.
upvoted 2 times

🗨️ 👤 **KC_go_reply** 1 year, 3 months ago

It says open-source, however. AutoML is not open-source, Tensorflow is
upvoted 2 times

🗨️ 👤 **edoo** 5 months, 4 weeks ago

But also it says machine learning platform, however, Tensorflow is not a platform nor is in GCP.
upvoted 1 times

🗨️ 👤 **CoffeeCloud** 7 months, 1 week ago

TPUs are a google product. You can only access TPUs through google as they developed them
upvoted 1 times

🗨️ 👤 **dhp88** 1 year, 5 months ago

Selected Answer: A

TPU = TensorFlow
upvoted 1 times

What is an example of unstructured data that organizations can capture from social media?

- A. Post comments
- B. Tagging
- C. Profile picture
- D. Location

Correct Answer: C

Community vote distribution

A (58%)	C (37%)	5%
---------	---------	----

rrgonzalez1992

Highly Voted

1 year, 4 months ago

It's A and C,

Maybe in the actual exam this question is a multiple choice.

Greetings

upvoted 12 times

cookieMr

Highly Voted

1 year ago

Selected Answer: A

Post comments are an example of unstructured data that organizations can capture from social media because they typically consist of free-text responses written by users. Unlike structured data that follows a predefined format or schema, unstructured data does not have a fixed structure or organization. In the case of post comments, users can express their opinions, provide feedback, ask questions, or engage in conversations without any specific format or constraint.

Tagging, profile pictures, and location, on the other hand, can be considered as structured or semi-structured data. Tagging refers to the process of labeling or categorizing content, which typically follows a predefined set of tags or labels. Profile pictures are static images associated with user profiles, and while they may contain visual information, they are not typically considered as unstructured data. Location data, although it can have variations, is often structured or semi-structured and can be represented using specific formats such as latitude and longitude coordinates.

upvoted 10 times

Cloudhungery

Most Recent

1 week ago

A is correct answer because in post comment we can use any thing like text, image and emoji etc which is not structure data

upvoted 1 times

Kessel

3 weeks, 2 days ago

Selected Answer: C

Unsure between A and C, but since it's not a multiple-choice question I will say C because photos are definitely unstructured data, while written posts can sometimes be semi-structured. Still, I hope they don't throw this curveball at us during the exam.

upvoted 1 times

9b8feb6

2 months, 1 week ago

Unstructured data can also come in the form of a photo, video, engineering CAD drawing, social media text stream, HTML document, or any form of data that is not captured as a fixed record, field-defined data format.

upvoted 1 times

4a98421

3 months, 2 weeks ago

Selected Answer: A

A and C are unstructured datasets. It should be mulitple choice question as per me.

upvoted 2 times

DenD

5 months ago

Selected Answer: C

C. Profile pictures are really unstructured data. The other options - structured or semistructured data.

upvoted 2 times

sivakarthick16

5 months, 4 weeks ago

Selected Answer: A

A correct answer

upvoted 1 times

AZ_Master

6 months, 3 weeks ago

Selected Answer: C

Whether you can categorized or not is besides the point. Images are considered unstructured data.
<https://cloud.google.com/blog/products/data-analytics/build-an-image-data-classification-model-with-bigquery-ml>
upvoted 5 times

  **Goodwin777** 7 months ago

Selected Answer: A

A. Correct, because: Post comments are considered unstructured data as they come in various forms (like text, emoticons, etc.) and do not adhere to a predefined data model or structure.

B. Incorrect, because: Tagging is usually a categorical variable and can be easily structured, making it not an example of unstructured data.

C. Incorrect, because: A profile picture, although it might seem unconventional compared to traditional structured data (like numerical or categorical data), it can also be organized and structured (as an object or a file link/path), hence it's not counted as unstructured data.

D. Incorrect, because: Locations, which could be stored as geo-coordinates or addresses, can be structured in a database and thus are not considered unstructured data.

upvoted 1 times

  **GamudaDavid** 5 months, 2 weeks ago

Nope Picture considered as Unstructured data, Comment is structured, it can giving meaningful data, like positive or negative comment.

upvoted 2 times

  **Goodwin777** 7 months ago

Selected Answer: A

A. Correct, because: Post comments are considered unstructured data as they come in various forms (like text, emoticons, etc.) and do not adhere to a predefined data model or structure.

B. Incorrect, because: Tagging is usually a categorical variable and can be easily structured, making it not an example of unstructured data.

C. Incorrect, because: A profile picture, although it might seem unconventional compared to traditional structured data (like numerical or categorical data), it can also be organized and structured (as an object or a file link/path), hence it's not counted as unstructured data.

D. Incorrect, because: Locations, which could be stored as geo-coordinates or addresses, can be structured in a database and thus are not considered unstructured data.

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct. (C is also correct. But generally comments of social media are more useful for performing analysis compared to profile pic. That's why comments is preferrable answer here)

upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: A

A. Post comments

Post comments are an example of unstructured data from social media. Unlike structured data (e.g., databases), unstructured data doesn't fit neatly into predefined categories, and it often includes text, images, or other forms of content that may not have a fixed format. In the context of social media, comments on posts can vary widely in content, making them unstructured data that organizations may want to capture and analyze for insights.

upvoted 1 times

  **subhadeep_ghosh** 8 months, 3 weeks ago

Selected Answer: A

Profile Image is structured - why? Because it can have attributes such as file type extension, file size, pixel dimensions, etc. + image file name. But 'Post Comments' are unstructured when it comes to 'social media'.

Here the question specifically mentions "social media" - where post comments are raw text, from which social media engineers need to figure out the tokens, keywords, meta tags, likes, dislikes, suggestions, opinions, political ideology, social strata, etc. This requires natural language processing + pre-trained transformations.

If satellite images are mentioned - then its unstructured - because from the image color, texture, engineers have to figure out topography, geography, climate pattern, temperature, etc.

But profile images of social media platform - always structured.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times

  **himanshu3004** 10 months, 1 week ago

Selected Answer: C

I have a whole new perspective for this answer. I feel the correct answer is C because we all know that BigQuery can't be used for unstructured data. However we can store comments in text fields whereas storing profile pictures is not possible in BigQuery which makes it qualify for unstructured data. However if this question comes out to be multiple choice than Post comments will be right options as well along with profile picture, considering the fact that people do use emotes, gifs etc as well in comments which clearly don't qualify for structured data.

upvoted 3 times

  **harithacML** 11 months ago

Selected Answer: C

Image, videos are unstructured
upvoted 1 times

  **harithacML** 11 months ago

Images and videos are unstructured but c is sensitive info. a comment is something you post in public hence that can be acquired and text is also an unstructured data. Hence A it is.
upvoted 1 times

  **RemoMx** 7 months, 3 weeks ago

profile pic is a public information, if you lookup the apis of any social media platform like FB
upvoted 1 times

An organization relies on online seasonal sales for the majority of their annual revenue.
Why should the organization use App Engine for their customer app?

- A. Automatically adjusts physical inventory in real time
- B. Autoscales during peaks in demand
- C. Runs maintenance during seasonal sales
- D. Recommends the right products to customers

Correct Answer: C

Community vote distribution

B (97%)

- caprice**

Highly Voted

1 year, 7 months ago

All of these brain dumps are copied from the one flawed original version...it's identical to others sites I have paid for.

upvoted 14 times
- omgitsele**

Highly Voted

1 year, 2 months ago

After seeing all these horrible answers I don't get why I would pay 40\$ a month for this fiasco..As for the answer it's obviously B

upvoted 10 times
- CoffeeCloud**

7 months, 1 week ago

I just paid for the year access to GCP Digital Leader and can't believe that the 'expert' assigning the answers is getting more wrong than me who's just started learning about cloud platforms

upvoted 2 times
- 9b8feb6**

Most Recent

2 months, 1 week ago

<https://cloud.google.com/appengine/docs/legacy/standard/python/how-instances-are-managed>

upvoted 1 times
- Goodwin777**

7 months ago

Selected Answer: B

A. Incorrect, because: Although an app could be built on App Engine to adjust physical inventory in real time, this is not a built-in feature of App Engine itself. It's rather a feature of the application built on top of it.

B. Correct, because: App Engine is designed to scale automatically in response to increases or decreases in traffic. This feature is particularly useful in the case of seasonal sales, where traffic can have significant spikes.

C. Incorrect, because: While it's possible to schedule maintenance tasks, it's not a key feature that helps in handling the high load during seasonal sales.

D. Incorrect, because: Recommending the right products to customers is a feature of personalized recommendation services, which would need to be built into the application. App Engine doesn't inherently provide such features.

upvoted 5 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 2 times
- __rajan__**

8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- afeff**

9 months ago

B is the right answer. C is wrong

another wrong answer !

upvoted 1 times

🗄️ 👤 **Lanundso** 10 months, 1 week ago
Always deploy during peak hours guys
upvoted 2 times

🗄️ 👤 **cookieMr** 1 year ago
Selected Answer: B
App Engine is a fully managed platform that automatically scales up or down based on the demand. During peak seasons when online sales are high, App Engine can handle increased traffic and ensure that the app remains responsive and available to customers.
upvoted 2 times

🗄️ 👤 **telosd** 1 year, 1 month ago
Selected Answer: B
another wrong answer on this dump.... the only real option is B
upvoted 1 times

🗄️ 👤 **JJ221** 1 year, 1 month ago
Selected Answer: B
B is correct
upvoted 1 times

🗄️ 👤 **gcp_prep** 1 year, 3 months ago
I don't even know why anyone would choose C. B is what I'd choose
upvoted 2 times

🗄️ 👤 **tbolick6** 1 year, 4 months ago
Selected Answer: B
Please change the correct answer to B.
upvoted 2 times

🗄️ 👤 **Jackey0117** 1 year, 5 months ago
Selected Answer: B
App Engine is a platform-as-a-service (PaaS) that allows an organization to deploy and scale their web applications automatically. It can automatically adjust the number of resources (such as servers) allocated to an application based on the amount of traffic it receives.
upvoted 2 times

🗄️ 👤 **PHD_CHENG** 1 year, 5 months ago
Selected Answer: B
B is correct. Autoscales during peaks in demand
upvoted 1 times

🗄️ 👤 **Sukanya_R** 1 year, 5 months ago
App Engine is a fully managed, serverless platform for developing and hosting web applications at scale. You can choose from several popular languages, libraries, and frameworks to develop your apps, and then let App Engine take care of provisioning servers and scaling your app instances based on demand. AUTOSCALE DURING SEASONAL SALES is CRITICAL where most of revenue is earned and and this feature is supported by App Engine. While since fully managed App Engine does support maintenance but this will be during entire life cycle and there is no correlation as such to seasonal sales while maintenance during seasonal sales also supported. B fits best as the only correct answer of choice and i we can choose more than 1 then B and C could apply
upvoted 2 times

An organization is using machine learning to make predictions. One of their datasets mistakenly includes mislabeled data. How will the prediction be impacted?

- A. Increased risk of privacy leaks
- B. Increased risk of inaccuracy
- C. Decreased model compatibility
- D. Decreased model training time

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B is correct

upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- tsk9921** 11 months, 2 weeks ago

B - Mislabeled or wrong data will lead to inaccurate result

upvoted 1 times
- Chai007** 1 year ago

Selected Answer: B

B is the answer

upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: B

Mislabeled data can introduce errors and inconsistencies into the training process, which can lead to inaccurate predictions. When the model learns from mislabeled examples, it may develop incorrect patterns or associations, compromising its ability to make accurate predictions on new, unseen data.

upvoted 2 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

It is B for sure

upvoted 2 times
- kinoko1330** 1 year, 9 months ago

Selected Answer: B

B. Increased risk of inaccuracy

upvoted 1 times
- Vin1975** 1 year, 9 months ago

B. Increased risk of inaccuracy

upvoted 2 times

A global organization is developing an application to manage payments and online bank accounts in multiple regions. Each transaction must be handled consistently in their database, and they anticipate almost unlimited growth in the amount of data stored.

Which Google Cloud product should the organization choose?

- A. Cloud SQL
- B. Cloud Spanner
- C. Cloud Storage
- D. BigQuery

Correct Answer: B

Community vote distribution

B (100%)

- cookieMr**

Highly Voted

1 year ago

Selected Answer: B

For an application that requires consistent handling of transactions and anticipates unlimited data growth, the organization should consider using Google Cloud Spanner.

Google Cloud Spanner is a horizontally scalable, globally distributed relational database service. It offers strong consistency and high availability across multiple regions, ensuring that transactions are handled consistently regardless of the geographic location of the data. It is designed to handle large amounts of data and scales automatically as the data grows, providing unlimited storage capacity.

upvoted 5 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 2 times
- __rajan__**

8 months ago

Selected Answer: B

For highly scalable application Cloud Spanner.

upvoted 1 times
- mdsarfraz69**

9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- omgitsele**

1 year, 2 months ago

Selected Answer: B

Unlimited growth -> Spanner

upvoted 2 times
- ucsdmiami2020**

1 year, 6 months ago

Selected Answer: B

Per Google docs, Cloud Spanner is a "Full managed relational database with unlimited scale, strong consistency, and up to 99.999% availability. Additionally, it delivers high-performance ACID transactions with strong consistency over regions and continents."

<https://cloud.google.com/spanner/>

upvoted 3 times
- Pank_G**

1 year, 6 months ago

B - cloud spanner - for multi regional and transactional database

upvoted 1 times
- Sheryb**

1 year, 6 months ago

Selected Answer: B

Spanner whenever unlimited growth is mentioned as it is unlimitedly scalable according to google.

upvoted 1 times
- pushpendras21**

1 year, 6 months ago

Cloud Spanner global RDB

upvoted 1 times

  **grzeev** 1 year, 8 months ago

B.

The main difference between Cloud Spanner and Cloud SQL is the horizontal scalability + global availability of data over 10TB. Spanner isn't for generic SQL needs, Spanner is best used for massive-scale opportunities. 1000s of writes per second, globally. 10,000s - 100,000s of reads per second, globally

upvoted 3 times

  **GanaL** 1 year, 9 months ago

Selected Answer: B

B is correct answer

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

B. Cloud Spanner Global relational db with unlimited scale and very high external consistency

upvoted 3 times

An organization has servers running mission-critical workloads on-premises around the world. They want to modernize their infrastructure with a multi-cloud architecture.

What benefit could the organization experience?

- A. Ability to disable regional network connectivity during cyber attacks
- B. Ability to keep backups of their data on-premises in case of failure
- C. Full management access to their regional infrastructure
- D. Reduced likelihood of system failure during high demand events

Correct Answer: C

Community vote distribution

D (85%)

C (15%)

- hireshgupta**

Highly Voted

11 months ago

The answer is D: Reduced likelihood of system failure during high demand events.

A multi-cloud architecture can help organizations to reduce the likelihood of system failure during high demand events by spreading their workloads across multiple clouds. This can help to ensure that even if one cloud experiences an outage, the organization's applications will still be available to users.

upvoted 8 times
- 76ffe88**

Most Recent

2 months, 1 week ago

Selected Answer: D

C don't seem correct to me

upvoted 2 times
- aUMMDqZmjLBx**

3 months, 1 week ago

Selected Answer: D

D is correct

upvoted 2 times
- Riise**

5 months, 2 weeks ago

Selected Answer: D

I do not see C as the right answer because this is a benefit that the organization already has. The question ask what are the benefits of partnering with a public cloud provider so it should be D.

upvoted 1 times
- azezdev**

6 months, 2 weeks ago

Selected Answer: D

As the are servers running in mission-critical workloads, so most advantageous organization benefit would be reduced system failure during high demand and peak.

upvoted 1 times
- JJSENRIQUE**

7 months, 2 weeks ago

Selected Answer: D

The ability to reduce the likelihood of system failure during high demand events is a benefit of a multi-cloud architecture for mission-critical workloads.

With a multi-cloud architecture, the organization can distribute their workload across multiple cloud providers. This means that if one cloud provider is experiencing high demand or an outage, the other cloud providers can continue to serve the workload. This reduces the likelihood of system failure and ensures that the organization's mission-critical workloads are always available.

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__**

8 months ago

D is correct .

upvoted 1 times
- katyk**

8 months, 2 weeks ago

Answer should be D - By adopting a multi-cloud architecture, the organization can distribute its workload across multiple cloud providers and regions, which can help reduce the risk of system failure during high demand events. Can the admin team please change the correct answer? thanks
upvoted 1 times

  **Gagan12345** 8 months, 3 weeks ago

Selected Answer: D

Option C, "Full management access to their regional infrastructure," might be important in a multi-cloud architecture, but in the context of the given scenario, the organization is specifically looking to modernize their infrastructure with a multi-cloud approach.

Option D, "Reduced likelihood of system failure during high demand events," is a significant benefit of a multi-cloud architecture. By distributing workloads across multiple cloud providers, the organization can better handle spikes in demand and reduce the risk of system failures.

While option C is relevant for management access, it doesn't directly address the benefits of a multi-cloud approach in terms of modernizing and improving infrastructure resilience. So, in this scenario, option D seems to be the more directly applicable benefit.

upvoted 3 times

  **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times

  **afeff** 9 months ago

Selected Answer: D

D is correct answer

upvoted 1 times

  **Anuj_Agrawal** 11 months, 3 weeks ago

Selected Answer: C

correct answer

upvoted 1 times

  **Anuj_Agrawal** 11 months, 3 weeks ago

Selected Answer: D

D is obvious answer

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

By distributing workloads across multiple cloud providers, the organization can enhance performance and availability. They can leverage different cloud providers' data centers in various regions to bring the workloads closer to end-users, reducing latency and improving responsiveness. Additionally, if one cloud provider experiences an outage or performance issues, the organization can redirect traffic and maintain operations using alternative cloud providers.

A multi-cloud architecture can improve disaster recovery and business continuity capabilities. By spreading workloads across different cloud providers and geographic regions, the organization can implement geo-redundancy and ensure data replication and failover mechanisms in case of disruptions or disasters. This helps mitigate the risk of a single point of failure and provides resiliency for critical workloads.

upvoted 3 times

  **[Removed]** 1 year, 3 months ago

Selected Answer: D

Key phrase is "mission-critical workloads". Answer is D.

upvoted 2 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: D

D looks like an obvious choice

upvoted 1 times

An organization needs to run frequent updates for their business app.
Why should the organization use Google Kubernetes Engine (GKE)?

- A. Customer expectations can be adjusted without using marketing tools
- B. Seamless changes can be made without causing any application downtime.
- C. GKE handles version control seamlessly and out of the box
- D. GKE is well suited for all monolithic applications

Correct Answer: D

Community vote distribution

B (98%)

- bohraa**

Highly Voted

1 year, 9 months ago

Selected Answer: B

Who is choosing correct answers?
upvoted 27 times
- AZ_Master**

6 months, 3 weeks ago

Some idiot who has never logged into Google Cloud Console, apparently
upvoted 12 times
- Iacopo22**

Highly Voted

1 year, 9 months ago

Selected Answer: B

<https://cloud.google.com/architecture/migrating-a-monolithic-app-to-microservices-gke>
upvoted 9 times
- 9b8feb6**

Most Recent

2 months, 1 week ago

<https://www.examtopycs.com/exams/google/cloud-digital-leader/view/8/>
answer: C
Apigee API Management has 3 key Pillars of Managing APIs:
1. Developer Portal - Supportive portal for other developers discover your APIs, rapidly understand them with intention that they could build applications with them
2. Monitoring & Analytics – available so your API program managers can see which APIs are used most and how. Also, could receive alerts when appropriate.
3. Policies = Runtime – runtime executes the polices you configure for security, traffic management and efficiency of your APIs. You can make changes without disrupting users
upvoted 1 times
- DenD**

5 months ago

Selected Answer: B

Oh, who is made this "correct" answer? :(
upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: B

B is correct.
upvoted 1 times
- dboy711**

8 months, 2 weeks ago

Selected Answer: B

GKE, which is a managed Kubernetes service, allows for rolling updates and canary deployments, making it possible to update the application without causing downtime. GKE provides features like auto-scaling, load balancing, and self-healing, which ensure that updates can be deployed with minimal disruption to the application's availability. This helps meet customer expectations by minimizing downtime during updates.

GKE is primarily designed for containerized applications and microservices. It is not specifically tailored for monolithic applications. While you can run monolithic applications in containers on GKE, it may not be the best choice for all monolithic applications, especially if the application doesn't benefit from container orchestration features.
upvoted 2 times

🗨️ 👤 **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct
upvoted 1 times

🗨️ 👤 **afeff** 9 months ago

GKE allows to perform frequent updates
B is the correct answer
upvoted 1 times

🗨️ 👤 **Chinaar** 10 months, 1 week ago

bsdk answer sai krde kam se kam fail krwaega ky
upvoted 6 times

🗨️ 👤 **Nikki1897** 10 months, 2 weeks ago

Fail krwayga admin bhai tu
upvoted 2 times

🗨️ 👤 **hireshgupta** 11 months ago

he answer to the question is B. Seamless changes can be made without causing any application downtime.

GKE provides a number of features that make it well-suited for organizations that need to run frequent updates for their business apps. These features include:

Automatic upgrades: GKE can be configured to automatically upgrade clusters to the latest Kubernetes version. This means that you can be confident that your apps are always running on the latest and greatest software.

Rolling updates: GKE supports rolling updates, which allow you to update your clusters without causing any application downtime. This is done by updating one node at a time, and then rolling the update out to the rest of the cluster once the first node has been successfully updated.

upvoted 4 times

🗨️ 👤 **cookieMr** 1 year ago

Selected Answer: B

GKE supports rolling updates, which means new versions of the application can be deployed gradually, one by one, without interrupting the overa
availability of the app. This approach ensures zero or minimal downtime during updates, keeping the app accessible to users while updates are
being applied.

upvoted 3 times

🗨️ 👤 **omgitsele** 1 year, 2 months ago

Selected Answer: B

One of GKE's best characteristics is zero downtime. Whoever chose D instead of B, please stop giving answers

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 3 months ago

Selected Answer: B

@Exam topics - I'm volunteering to screen these questions for you. Answer is B, not D. I am not a robot.

upvoted 4 times

🗨️ 👤 **tbolick6** 1 year, 4 months ago

Selected Answer: B

I think it's B, as well.

upvoted 2 times

🗨️ 👤 **Jackey0117** 1 year, 5 months ago

Selected Answer: B

GKE allows them to deploy new versions of the app without affecting the existing one

upvoted 1 times

An organization wants to use Apigee to manage all their application programming interfaces (APIs).
What will Apigee enable the organization to do?

- A. Increase application privacy
- B. Measure and track API performance
- C. Analyze application development speed
- D. Market and sell APIs

Correct Answer: C

Community vote distribution

B (95%)

2%

- hireshgupta**

Highly Voted

11 months ago

Selected Answer: B

The answer is B. Measure and track API performance.

Apigee is an API management platform that helps organizations to develop, secure, and manage APIs. It provides a number of features that can help organizations to measure and track API performance, including:

upvoted 6 times
- anuj9882**

Most Recent

3 months ago

b IS CORRECT answer. why most of answers are wrong here

upvoted 1 times
- Tsan_Nake**

6 months ago

Selected Answer: B

<https://www.coursehero.com/file/p70ieg8f/Question-72-An-organization-is-using-machine-learning-to-make-predictions-One-of/>

QUESTION 76

upvoted 1 times
- azezdev**

6 months, 2 weeks ago

Selected Answer: B

answers would (B)

upvoted 1 times
- ba1e4a3**

6 months, 3 weeks ago

Selected Answer: B

B is the right solution

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

B is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: B

B. Measure and track API performance

Apigee is an API management platform that helps organizations design, secure, deploy, monitor, and scale APIs. One of the key functionalities is the ability to measure and track API performance. Apigee provides analytics and monitoring tools that allow organizations to gain insights into how their APIs are being used, track performance metrics, and ensure that APIs are operating efficiently. This helps in optimizing API usage and identifying areas for improvement.

upvoted 4 times
- SauravAmuze**

8 months, 1 week ago

Selected Answer: B

B is the correct answer here. Admins, pls correct the selections.

upvoted 3 times
- mdsarfraz69**

9 months ago

Selected Answer: C

C is correct

upvoted 1 times

  **afeff** 9 months ago

Selected Answer: B

track API perofmance

upvoted 2 times

  **cookieMr** 1 year ago

Selected Answer: B

Apigee provides comprehensive analytics and monitoring capabilities for the organization's APIs. It tracks API usage, performance, and error rates, allowing the organization to gain insights into how their APIs are being utilized. This data helps in making informed decisions, optimizing API performance, and identifying potential issues or bottlenecks.

upvoted 2 times

  **telosd** 1 year, 1 month ago

Selected Answer: B

exam topics crew please fix those answers.... 20 out of 71 wrong answers is too much to be called a reliable brain dump.

upvoted 3 times

  **omgitsele** 1 year, 2 months ago

Selected Answer: B

So many wrong answers...Apigee is a comprehensive API management platform. Deployment options: Cloud, on-premises and hybrid. It can manage COMPLETE API lifecycle (Design, Secure, Monitor and Monetize APIs)

upvoted 4 times

  **Matro71** 1 year, 3 months ago

Selected Answer: B

There are many obivous wrong answers. Please let someone read the comments.

upvoted 4 times

  **Jackey0117** 1 year, 5 months ago

Selected Answer: B

Apigee enables organizations to measure and track the performance of their APIs, including traffic, usage, and response time

upvoted 2 times

  **Sukanya_R** 1 year, 5 months ago

Monitor and analyze your traffic to gain reliable performance and actionable insights is ta feature quoted for Apigee in GC documentation. B is the right answer . Since API Management it would also track measure and track API performance. Application development speed would be influenced if right APIs are built quickly and managed through Apigee

upvoted 1 times

  **glk48** 1 year, 6 months ago

Selected Answer: B

Where are these incorrect "correct" answers being captured and then fed to the masses? Answer is B, D would be a close second.

upvoted 2 times

An e-commerce organization is reviewing their cloud data storage.
What type of raw data can they store in a relational database without any processing?

- A. Product inventory
- B. Product photographs
- C. Instructional videos
- D. Customer chat history

Correct Answer: A

Community vote distribution

A (100%)

- aTTager

3 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: A

A is correct

upvoted 1 times
- Preetii

10 months, 2 weeks ago

Selected Answer: A

Ans is A

upvoted 1 times
- hireshgupta

11 months ago

Selected Answer: A

Obvious

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: A

Product inventory data, such as product names, descriptions, pricing, and inventory levels, can be stored in a relational database without any processing. This type of structured data is commonly stored in a database to track and manage the organization's product inventory.

Product photographs are typically large files in image format (JPEG, PNG, etc.) and are considered unstructured data. Storing these files directly in relational database is not efficient or recommended.

Similar to product photographs, instructional videos are typically large files in video format (MP4, AVI, etc.) and are considered unstructured data. Storing videos directly in a relational database is not efficient.

Customer chat history often consists of unstructured text data, including messages, timestamps, and metadata. Relational databases are designed for structured data, and storing unstructured text data directly in a database is not optimal.

upvoted 4 times
- KC_go_reply

1 year, 3 months ago

B + C: images and videos could only be saved in BLOB columns, and those have size limits of course. In any case, it would be far from optimal to store them in a relational DB.

D: The chat history itself is semi-structured at best. Also, chat texts probably are longer strings, which a relational DB is also not optimized for.

A: Product inventory probably references to data points such as product ID, stock etc., and is a typical use case for relational DB.

upvoted 2 times

  **atoledo** 1 year, 8 months ago

Selected Answer: A

A. Product Inventory

upvoted 2 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: A

A Product Inventory

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A product inventory

upvoted 1 times

A hotel wants to modernize their legacy systems so that customers can make reservations through a mobile app. What's the benefit of using an application programming interface (API) to do this?

- A. They do not have to develop the end-user application
- B. They can deprecate their legacy systems
- C. They can transform their systems to be cloud-native
- D. They do not have to rewrite the legacy system

Correct Answer: C

Community vote distribution

D (81%)

Other

- Ej24356

Highly Voted

1 year, 9 months ago

Selected Answer: D

Organizations can package existing legacy services in a way that makes them more consumable by the business, by either wrapping services with RESTful APIs, or exposing access to constituent web service operations through APIs.

upvoted 14 times
- Krishna_7176

Most Recent

5 months, 2 weeks ago

Selected Answer: D

d is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: D

D is correct.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times
- afeff

9 months ago

D is the correct answer

upvoted 1 times
- adyanrado

9 months, 1 week ago

obviously is D, API allows communication between legacy system and the new APP

upvoted 1 times
- Preetii

10 months, 2 weeks ago

Selected Answer: C

I think it is C, as it can become cloud native

upvoted 1 times
- hireshgupta

11 months ago

Selected Answer: D

The answer is D. They do not have to rewrite the legacy system.

An API is a software intermediary that allows two applications to talk to each other. In this case, the API would allow the hotel's legacy systems to talk to the mobile app. This means that the hotel would not have to rewrite their legacy systems to make them compatible with the mobile app

upvoted 3 times
- cookieMr

1 year ago

Selected Answer: D

By providing an API, the hotel can offer a standardized and well-documented interface for developers to integrate their reservation system with th mobile app. This allows for seamless communication and data exchange between the app and the hotel's backend systems (legacy).

upvoted 3 times

  **Sp4nner** 1 year, 2 months ago

Selected Answer: A

Using an API allows the hotel to provide a standard interface for the mobile app to communicate with their legacy systems. This means they can create a new application (in this case, the mobile app) without having to rewrite their entire legacy system. By providing an API, the hotel can also control the data and functionality that the mobile app can access, ensuring security and privacy. Additionally, using an API allows the hotel to provide a better user experience by streamlining the reservation process and making it easier for customers to book rooms.

upvoted 1 times

  **MHx1** 1 year, 3 months ago

C

Extract from Google page "...you can modernize legacy services via RESTFul interfaces, unlocking them to be consumed by new cloud services and apps"

upvoted 3 times

  **Tazingo** 1 year, 3 months ago

Selected Answer: C

(Sorry for reposting, forgot to vote) Not sure why everyone thinks it's possible to stick together and API with a legacy system and have them magically play together without a massive rewrite. It's like placing a window frame next to an untouched solid wall and expecting to see inside the house. I'd say the closest answer is C.

upvoted 3 times

  **Tazingo** 1 year, 3 months ago

Not sure why everyone thinks it's possible to stick together and API with a legacy system and have them magically play together without a massive rewrite. It's like placing a window frame next to an untouched solid wall and expecting to see inside the house. I'd say the closest answer is C.

upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: A

All of the options mentioned could be potential benefits of using an API to modernize the hotel's legacy systems for mobile app reservations, but the most relevant one is option A - "They do not have to develop the end-user application. By using an API, the hotel can expose certain functions of their legacy system to the mobile app, allowing users to make reservations without having to develop a brand new application from scratch. This saves time and resources for the hotel, as they do not have to invest in the development of a new application and can instead focus on improving the existing system to work with the API.

upvoted 3 times

  **CrashoveR** 1 year, 4 months ago

Selected Answer: C

How you are imagine to add API to existing legacy system without rewriting it?

upvoted 1 times

  **Jackey0117** 1 year, 4 months ago

The best answer for this question is D. They do not have to rewrite the legacy system. By using an API, the hotel can expose the functionality of their legacy systems through a secure and manageable interface, making it possible to access the system's functionality from the mobile app without having to make major changes to the underlying system. This can save time and resources compared to rewriting the legacy system from scratch.

upvoted 1 times

  **Jackey0117** 1 year, 5 months ago

Selected Answer: D

Both A and D are correct.

Using an API allows the hotel to expose their reservation system to external developers, allowing them to develop the end-user mobile app without the hotel having to develop it themselves. And also, the hotel does not have to rewrite their legacy system to make it accessible to the mobile app instead, they can use the API as a bridge between their legacy system and the mobile app, allowing them to modernize their infrastructure without disrupting their service to customers.

upvoted 1 times

  **Jackey0117** 1 year, 4 months ago

The best answer for this question is D.

upvoted 1 times

An organization wants to digitize and share large volumes of historical text and images.
Why is a public cloud a better option than an on-premises solution?

- A. In-house hardware management
- B. Provides physical encryption key
- C. Cost-effective at scale
- D. Optimizes capital expenditure

Correct Answer: C

Community vote distribution

C (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C is correct.
upvoted 1 times
- mdsarfraz69** 9 months ago

Selected Answer: C

C is correct
upvoted 1 times
- Preetii** 10 months, 2 weeks ago

Selected Answer: C

Ans is C
upvoted 1 times
- dnagasree** 10 months, 3 weeks ago

Selected Answer: C

Cost effective
upvoted 1 times
- tsk9921** 11 months, 1 week ago

C
Cost-effective at scale is the only benefit listed
upvoted 2 times
- cookieMr** 1 year ago

Selected Answer: C

Public cloud services operate on a pay-as-you-go model, meaning you only pay for the resources you use. This eliminates the need for upfront infrastructure investments, hardware maintenance, and ongoing operational costs associated with on-premises solutions.

Additionally, you can take advantage of the cloud provider's economies of scale, reducing the overall cost of storage and processing.
upvoted 2 times

Ninad_Raj 1 year, 5 months ago

C is the answer
upvoted 1 times

najmdin 1 year, 8 months ago

C is the answer
upvoted 1 times

Govindaraj 1 year, 9 months ago

Selected Answer: C

C is the answer
upvoted 1 times

  **lax797** 1 year, 9 months ago
@Vin1975 please comment your answers for qstn no. 80 - 100.
upvoted 3 times

  **Vin1975** 1 year, 9 months ago
C. Cost-effective at scale
upvoted 4 times

An organization wants to develop an application that can be personalized to user preferences throughout the year. Why should they build a cloud-native application instead of modernizing their existing on-premises application?

- A. Developers can rely on the cloud provider for all source code
- B. Developers can launch new features in an agile way
- C. IT managers can migrate existing application architecture without needing updates
- D. IT managers can accelerate capital expenditure planning

Correct Answer: B

Community vote distribution

B (100%)

- Londonkiwi

Highly Voted

1 year, 6 months ago

Agile development is not only used in the cloud, so do not fully agree with B. However, it can be done quickly but only if the functionality is preserved in the cloud.

upvoted 7 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: B

B. Developers can launch new features in an agile way

Building a cloud-native application provides the flexibility and agility needed to adapt to changing user preferences throughout the year. Cloud-native applications are designed to take advantage of cloud services and architectures, enabling developers to deploy updates and new features more quickly and efficiently. This agility is essential for meeting evolving user needs and preferences in a timely manner.

Options A, C, and D do not directly address the agility and flexibility required for rapidly changing user preferences in this context.

upvoted 1 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct

upvoted 1 times
- cookieMr

1 year ago

Selected Answer: B

Cloud-native applications are designed with agility and flexibility in mind. They are built using modular and loosely coupled components, making it easier to introduce new features, make updates, and iterate on the application based on evolving user preferences. The cloud-native approach allows for faster development cycles, enabling the organization to respond quickly to changing user needs and market dynamics.

upvoted 2 times
- BryRob

1 year, 7 months ago

Selected Answer: B

B for me

upvoted 2 times
- atoledo

1 year, 8 months ago

Selected Answer: B

B is correct.

upvoted 2 times
- Govindaraj

1 year, 9 months ago

Selected Answer: B

B seems to be correct

upvoted 1 times



Vin1975 1 year, 9 months ago

B. Developers can launch new features in an agile way
upvoted 2 times



Which technology allows organizations to run multiple computer operating systems on a single piece of physical hardware?

- A. Hypervisor
- B. Containers
- C. Serverless computing
- D. Open source

Correct Answer: A

Community vote distribution

A (100%)

-   **edoo** 5 months, 4 weeks ago

Selected Answer: A

Correct
upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: A

A is correct
upvoted 1 times
-   **SelinaZh** 1 year ago

Selected Answer: A

Hypervisor in GCP refers to the software that creates and manages virtual machines on Google Cloud Platform. So both VM or hypervisor should be the correct answer.
upvoted 2 times
-   **cookieMr** 1 year ago

Selected Answer: A

A hypervisor is a specific type of software or firmware that enables virtualization by creating and managing virtual machines (VMs). It is the underlying technology that allows multiple operating systems to run on a single physical server.

The hypervisor sits between the hardware and the virtual machines, providing an abstraction layer that allows the operating systems running in the virtual machines to interact with the underlying hardware resources. It manages the allocation and sharing of physical resources such as CPU, memory, storage, and networking among the virtual machines.
upvoted 3 times
-   **OctavioBatera** 1 year, 6 months ago

Selected Answer: A

I agree with A, but IMHO I think the answer is malformed, because the name of the technology itself, for me is "virtualization", hypervisor is more like one of the layers that operates on this technology to make viable the hardware sharing for the different virtual machines.
upvoted 2 times
-   **tsk9921** 11 months, 2 weeks ago

Totally agree with Virtualizaton reference
upvoted 1 times
-   **Londonkiwi** 1 year, 6 months ago

I think it is A, as Hypervisors refer to physical machines whereas Containerisation applies more to Docker.
upvoted 1 times
-   **BryRob** 1 year, 7 months ago

letter A

upvoted 1 times

  **atoledo** 1 year, 8 months ago

Selected Answer: A

A is the answer

upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: A

A is the answer

upvoted 1 times

  **Kamlakar** 1 year, 9 months ago

hypervisor abstracts away hardware for the virtual machines so they can run an operating system

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A. Hypervisor

upvoted 1 times

An organization is making a strategic change to customer support in response to feedback. They plan to extend their helpline availability hours. Why is the organization making this change?

- A. Users expect professional expertise
- B. Users require personalization
- C. Users expect always-on services
- D. Users require regional access

Correct Answer: C

Community vote distribution

C (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
- chandangupta8dec** 9 months ago

Selected Answer: C

I guess, who gave the answer is taking high quality stuff.

upvoted 2 times
- mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: C

Extending helpline availability hours ensures that the organization's customer support services are available for a longer duration, providing users with access to assistance and support whenever they need it. By offering an "always-on" service, the organization can meet the expectations of users who require support outside of regular business hours.

upvoted 3 times
- BryRob** 1 year, 7 months ago

letter C

upvoted 1 times
- shuvs** 1 year, 7 months ago

Selected Answer: C

Always on

upvoted 2 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: C

C is the correct one

upvoted 2 times
- Vin1975** 1 year, 9 months ago

C. Users expect always-on services

upvoted 1 times

An organization is migrating their business applications from on-premises to the cloud.
How could this impact their operations and personnel costs?

A. Reduced on-premises infrastructure management costs

B. Increased on-premises hardware maintenance costs

C. Reduced cloud software licensing costs

D. Increased cloud hardware management costs

Correct Answer: B

Community vote distribution

A (92%)8%

- WesThompson

Highly Voted

1 year, 5 months ago

This worries me that examtopics has the most basic question wrong. Cloud 101 - "A". I am basing all other answers on popular vote assuming 75%+

upvoted 15 times
- BimJean

1 year, 4 months ago

The answers are wrong on purpose, so they don't get in trouble with exam question creators.

upvoted 1 times
- omgitsele

1 year, 2 months ago

Then they shouldn't expect 40\$/month ffs

upvoted 5 times
- 00f7048

Most Recent

1 month ago

Apparently this is a new studying technique where you realize that their answers are wrong you have truly mastered the material and should Ace the test! And here I thought all my studying was a waste at first blush! Now I can write the test and the answers.....If you need help Exam topics I would be available to resolve the issues! ;)

upvoted 1 times
- tikounass

5 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- madcloud32

6 months ago

Selected Answer: A

A is obvious and correct.

upvoted 1 times
- Sourabhsoni

6 months, 1 week ago

A is correct

upvoted 1 times
- Groodle10

7 months ago

i was wondering it is intentional to generate discussion?

upvoted 2 times
- Groodle10

7 months, 1 week ago

At the very least provide reference to your answer

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
- Boca1

8 months, 3 weeks ago

Let’s move everything to the cloud so our on-prem hardware maintenance costs go up....dah.. B is wrong, A is correct
upvoted 2 times

  **chandangupta8dec** 9 months ago

Selected Answer: A

Could you please tell how you're getting high quality stuff while answering?
upvoted 2 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct
upvoted 1 times

  **adyanrado** 9 months, 1 week ago

Selected Answer: A

A is correct answer, obviously
upvoted 1 times

  **the_var** 11 months, 3 weeks ago

This "correct answer" is blasphemy!!! Examtopics, please correct this basic answer ASAP
upvoted 3 times

  **cookieMr** 1 year ago

Selected Answer: A

By moving to the cloud, the organization can offload the responsibility of managing and maintaining physical infrastructure. This can lead to cost savings in terms of hardware maintenance, data center expenses, and personnel required to manage the infrastructure.
upvoted 4 times

  **telosd** 1 year, 1 month ago

Selected Answer: A

another wrong answer from ET... sadly this is too costly for having so many wrong answer.
upvoted 1 times

  **pillais** 1 year, 2 months ago

Selected Answer: A

Come on Its A. Please change the answer !
upvoted 1 times

A retail company stores their product inventory in a legacy system. Often, customers find products on the company's website and want to purchase them in-store. However, when they arrive, they discover that the products are out of stock. How could the company benefit from using an application programming interface (API)?

- A. To create personalized product recommendations for customers
- B. To optimize their on-premises legacy system stability
- C. By manually linking each inventory system to the website on a case-by-case basis
- D. By programmatically connecting the inventory system to their website

Correct Answer: B

Community vote distribution

D (98%)

-   **jexmtropicscheatchatya**

Highly Voted 

 1 year, 9 months ago

Selected Answer: D

D. By programmatically connecting the inventory system to their website
The issue is the website shows an item is available at the store, but when the customer gets to the store, they find out that item is out of stock.
upvoted 17 times
-   **Kritika28**

Most Recent 

 5 months ago

Selected Answer: D

By programmatically connecting the inventory system to their website
The issue is the website shows an item is available at the store, but when the customer gets to the store, they find out that item is out of stock.
upvoted 1 times
-   **cloudmind** 5 months ago

D is correct
upvoted 1 times
-   **madcloud32** 6 months ago

Selected Answer: D

D is correct.
API will help keep their in store stock synced with online orders.
upvoted 3 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D. By programmatically connecting the inventory system to their website

Using an Application Programming Interface (API) to connect the website to the inventory system allows for real-time or near-real-time updates on product availability. This means that when a customer checks the product availability online and decides to purchase it in-store, the website can provide accurate and up-to-date information about the product's stock status in the physical store. This helps prevent situations where customers arrive to find out-of-stock items, improving customer satisfaction and experience.
upvoted 2 times
-   **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct
upvoted 1 times
-   **michael435** 1 year ago

Answer is D
upvoted 1 times
-   **cookieMr** 1 year ago

Selected Answer: D

By integrating the legacy inventory system with an API, the company can provide real-time inventory updates on its website or mobile app. When customer finds a product online and wants to purchase it in-store, they can check the availability status through the API. This ensures that the inventory information is accurate and reduces the likelihood of customers encountering out-of-stock products.

upvoted 3 times

  **telosd** 1 year, 1 month ago

Selected Answer: D

geez... 5 wrong answers on the last 10 questions.... I thought this was updated frequently.

upvoted 1 times

  **adityanarayan** 1 year, 4 months ago

when team members are not able to understand the answer they choose them as answer.

upvoted 1 times

  **raj_gets_google** 1 year, 5 months ago

The right answer is D.

upvoted 3 times

  **LouCie** 1 year, 5 months ago

These are terrible

upvoted 1 times

  **arna5** 1 year, 5 months ago

Selected Answer: D

D. By programmatically connecting the inventory system to their website M

upvoted 2 times

  **ronieto** 1 year, 7 months ago

Selected Answer: D

D describes exactly what an API is for.

upvoted 2 times

  **BryRob** 1 year, 7 months ago

Selected Answer: D

its letter D

upvoted 1 times

  **SoftSami** 1 year, 8 months ago

Selected Answer: D

Definitely D

upvoted 1 times

An organization is training a machine learning model to predict extreme weather events in their country. How should they collect data to maximize prediction accuracy?

- A. Collect all weather data evenly across all cities
- B. Collect all weather data primarily from at-risk cities
- C. Collect extreme weather data evenly across all cities
- D. Collect extreme weather data primarily from at-risk cities

Correct Answer: B

Community vote distribution

A (78%)

D (20%)

-   **jexmtropicscheatchatya**

Highly Voted 

 1 year, 9 months ago

Selected Answer: A

My vote is for A. Collect all weather data evenly across all cities. Mainly because it seems that the emphasis for data collection for ML is to make sure there are no holes in your data collection.

upvoted 21 times
-   **Pennycheng**

Most Recent 

 3 months ago

A is correct

upvoted 1 times
-   **zebroo** 4 months ago

Selected Answer: D

D is the correct answer as find in Google Gemini AI bot

The best option for collecting data to maximize prediction accuracy for extreme weather events

The goal is to train the model on data that closely resembles the real-world situations it will encounter when predicting extreme events.

upvoted 1 times
-   **himel2024** 4 months, 4 weeks ago

Selected Answer: A

Collect all weather data evenly across all citie

upvoted 1 times
-   **himel2024** 4 months, 4 weeks ago

Collect all weather data evenly across all citie

upvoted 1 times
-   **cloudmind** 5 months ago

D is correct

upvoted 1 times
-   **madcloud32** 6 months ago

Selected Answer: A

A is right.

More relevant input training data for ML will help better predictions.

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D. Collect extreme weather data primarily from at-risk cities.

To maximize the prediction accuracy of a machine learning model for extreme weather events, it is important to collect data from the locations and times where these events are most likely to occur. This can be done by identifying the cities and regions that are most at risk for extreme weather events, and then collecting data from those areas.

upvoted 3 times
-   **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct
upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: D

By collecting extreme weather data primarily from at-risk cities, the organization can target areas where the occurrence of extreme weather events is higher. This approach allows the machine learning model to learn from data that is most relevant to the specific problem of predicting extreme weather events. It enables the model to capture the unique characteristics and patterns associated with extreme weather in those areas, leading to improved prediction accuracy.
upvoted 2 times

  **YNES** 1 year, 2 months ago
3 / 3

The best approach to maximize prediction accuracy is to collect extreme weather data primarily from at-risk cities. This approach will allow the machine learning model to learn from data that is more relevant to the problem being solved. By focusing on at-risk cities, the model will have a better chance of accurately predicting extreme weather events, which is the goal of the project. Collecting extreme weather data evenly across all cities or collecting all weather data primarily from at-risk cities may not provide enough data for the model to accurately predict extreme weather events.
upvoted 4 times

  **Flexip** 1 year, 2 months ago

Selected Answer: D

To maximize prediction accuracy, an organization training a machine learning model to predict extreme weather events in their country should collect extreme weather data primarily from at-risk cities (Option D). This is because collecting data evenly across all cities (Option A) or collecting all weather data primarily from at-risk cities (Option B) may not provide enough data on extreme weather events to train the model effectively. Collecting all extreme weather data evenly across all cities (Option C) may provide too much data on extreme weather events in areas that are not at risk, which could lead to overfitting of the model.
upvoted 2 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: A

Definitely A or you skew the model
upvoted 2 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: A

A is the correct answer for data completeness and coverage
upvoted 2 times

  **Jackey0117** 1 year, 5 months ago

Selected Answer: D

To maximize prediction accuracy for extreme weather events, the organization should prioritize collecting data from at-risk cities. By the way, A is 'wrong', because collecting extreme weather data evenly across all cities doesn't make sense as this data is not evenly distributed and collecting all weather data primarily from at-risk cities will not provide enough data on extreme weather events.
upvoted 1 times

  **JudgeLI** 1 year, 5 months ago

Selected Answer: D

D. Collect extreme weather data primarily from at-risk cities.

It is important to have a good representation of the events that the model is trying to predict, so collecting data on extreme weather events primarily from at-risk cities will likely lead to the most accurate predictions. Additionally, collecting data from all cities evenly may not be necessary or practical and may lead to a lot of irrelevant data that the model has to process.
upvoted 1 times

  **KC_go_reply** 1 year, 3 months ago

But this introduces two kinds of bias for the data set. If your model never sees 'normal' weather or other cities, how would it differentiate? The 'knowledge domain' must be covered as much as possible. Knowledge domain in this case means different weather and cities
upvoted 5 times

An organization needs to search an application's source code to identify a potential issue. The application is distributed across multiple containers.

Which Google Cloud product should the organization use?

- A. Google Cloud Console
- B. Cloud Trace
- C. Cloud Monitoring
- D. Cloud Logging

Correct Answer: C

Community vote distribution

D (44%)

B (32%)

A (19%)

5%

- VijayendraBaahubali**

Highly Voted

1 year, 9 months ago

Selected Answer: B

B. Cloud Trace is supposed to be the correct answer. It's an application performance management tool. It's a Google solution for monitoring application performance. It is a distributed tracing system that helps developers debug or fix and optimize their code. Since the question talks about distributed tracing, this is the correct answer.

Cloud Monitoring on the other hand is not correct because it's an operations-focused tool

upvoted 17 times
- aaron773**

1 year ago

That doesn't make sense, like you said, it's an application performance management tool. It even says it primary use is to collect latency data and to inspect performance bottlenecks. <https://cloud.google.com/trace/>

Whereas cloud logging integrates with cloud tracing, error reporting and cloud monitoring and with the help of these 3 you can find potential issues and not just performance bottlenecks.

upvoted 3 times
- Jackey0117**

Highly Voted

1 year, 4 months ago

Selected Answer: A

A. Google Cloud Console, specifically the Code Search feature in Google Cloud Source Repositories, can be used to search an application's source code in Google Cloud. The Code Search feature allows you to search across all your source code repositories, including code stored in containers, to quickly find and identify potential issues.

upvoted 7 times
- 4a98421**

Most Recent

3 months, 2 weeks ago

Selected Answer: C

GCP Cloud Monitoring helps indentify root cause analysis

upvoted 1 times
- sivakarthick16**

5 months ago

Selected Answer: D

Cloud Logging is a Google Cloud product that allows organizations to store, search, analyze, and monitor logs from various sources, including containers. By configuring the appropriate logging options, the organization can capture logs from the containers and then search the logs to identify potential issues within the application's source code.

Therefore, option D (Cloud Logging) is the suitable Google Cloud product for searching an application's source code distributed across multiple containers.

upvoted 5 times
- madcloud32**

6 months ago

Selected Answer: D

D is more appropriate here.

B. Cloud Trace : used for analysis of overall performance trace, latency of service or app.

C. Cloud Monitoring: used for overall cloud monitoring of cloud.

D. Cloud Logging: is correct to debug an app module and find about cause of issue.

upvoted 4 times
- madcloud32**

6 months ago

D is correct here. Logging is potentially used to find the issues. Trace is for latency measure and monitoring is used for cloud full monitor.

upvoted 2 times

🗨️ **mohammeddigital** 6 months, 1 week ago

Selected Answer: D

Cloud Logging is a centralized logging service that collects logs from various sources across your Google Cloud Platform (GCP) infrastructure. It provides a searchable and indexed store for logs, enabling you to analyze and troubleshoot issues.

In this scenario, the organization needs to search an application's source code to identify a potential issue. Since the application is distributed across multiple containers, Cloud Logging is the most suitable choice. It can collect logs from containers and other sources, making it easier to identify the source of the issue.

upvoted 2 times

🗨️ **befeyac438** 6 months, 1 week ago

Selected Answer: D

Use Cloud Logging for monitoring individual events and understanding application behavior at specific points in time.
Use Cloud Trace when you need to understand the flow of requests across your distributed application and diagnose performance issues.

upvoted 2 times

🗨️ **Dardeery** 7 months, 2 weeks ago

D is the correct answer ,
Cloud Logging is the Google Cloud product that allows organizations to search and analyze log data, including logs from distributed applications running across multiple containers. It provides a centralized location to view, filter, and search logs, making it suitable for identifying potential issues in an application's source code. You can use it to track and troubleshoot issues efficiently across your distributed environment. It's like having a detective tool that helps you find clues about what's happening inside your containers.

upvoted 5 times

🗨️ **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times

🗨️ **Ab1975** 8 months ago

Selected Answer: D

Cloud Trace does not collect or store source code. Cloud Trace only collects latency data and metadata about the application's traces.

Cloud Logging can be used to search the source code of a distributed application to identify a potential issue. Therefore D.
upvoted 4 times

🗨️ **__rajan__** 8 months ago

Selected Answer: D

D. Cloud Logging

Cloud Logging is a Google Cloud product that collects and stores logs from applications and infrastructure. It provides a unified view of logs from all sources, making it easy to search and analyze logs to identify potential issues.

upvoted 1 times

🗨️ **RK_MANOJ** 8 months, 2 weeks ago

- A. Google Cloud Console: This is the management interface for Google Cloud services, but it is not designed for searching application source code.
- B. Cloud Trace: Cloud Trace is a service for tracing and profiling application performance but does not help in searching source code.
- C. Cloud Monitoring: Cloud Monitoring allows you to monitor the performance and availability of your applications, but it is not designed for searching source code.
- D. Cloud Logging (formerly known as Stackdriver Logging): Cloud Logging is designed for collecting and analyzing log data, including logs from your containers, making it the most suitable option for searching application source code issues.

So, the correct choice is D. Cloud Logging.
upvoted 4 times

🗨️ **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct
upvoted 1 times

🗨️ **himanshu3004** 10 months, 1 week ago

Selected Answer: D

Going with Cloud Logging because the question specifically asks about identifying a potential issue with the source code where the logs can prove helpful to figure out the cause of the problem.
However if the question was asking about talking about optimizing or enhancing performance, in that case Cloud Trace would have made more sense.

upvoted 4 times

🗨️ **ayush_1995** 10 months, 2 weeks ago

Selected Answer: A

Question is about source code, not during run time. Cloud logging is for Log analytics after the code has performed some action.

upvoted 1 times

  **ayush_1995** 10 months, 2 weeks ago
Srry B...by mistake selected A option.
upvoted 1 times

  **Preetii** 10 months, 2 weeks ago

Selected Answer: D

Ans is D

Cloud Logging is the Google Cloud product that allows organizations to search and analyze the logs generated by various services and application running in their environment, including those distributed across multiple containers. In this case, the organization needs to identify a potential issue in an application's source code, and logs are often a valuable source of information to troubleshoot and diagnose issues.

Cloud Logging allows you to collect, store, and analyze logs from different sources, including containers, virtual machines, and other Google Cloud services. By searching through logs using filters, you can pinpoint specific events, errors, or patterns that might indicate the source of the issue you're investigating. This is particularly useful when dealing with distributed systems like those running in containers, where logs from different containers can be aggregated and analyzed in one place.

upvoted 3 times

An organization's web developers and operations personnel use different systems.
How will increasing communication between the teams reduce issues caused by silos?

- A. By assigning blame for failures and establishing consequences
- B. By combining job role responsibilities to ensure that everyone has shared access
- C. By increasing data encryption to strengthen workflows
- D. By emphasizing shared ownership of business outcomes

Correct Answer: D

Community vote distribution

D (76%)

B (24%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: D

D. By emphasizing shared ownership of business outcomes

When web developers and operations personnel emphasize shared ownership of business outcomes, they are more likely to communicate and collaborate effectively. This is because they are all working towards the same goal, and they recognize that they need to work together in order to achieve that goal.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: D

D is correct
upvoted 1 times

  **Lufly** 10 months ago

Selected Answer: D

By emphasizing shared ownership of business outcomes, organizations can reduce the issues caused by silos and improve their overall performance.

The other options are not as relevant to this scenario. Option A, assigning blame for failures and establishing consequences, is not a good way to reduce silos. This approach can create a blame culture and discourage communication and collaboration. Option B, combining job role responsibilities to ensure that everyone has shared access, is not necessary. Silos can occur even when teams have shared access to information. Option C, increasing data encryption to strengthen workflows, is not a way to reduce silos. This approach is more about security than communication and collaboration.

upvoted 2 times

  **cookieMr** 1 year ago

Selected Answer: D

Increasing communication between web developers and operations personnel can help reduce issues caused by silos. By emphasizing shared ownership of business outcomes, both teams are encouraged to work collaboratively towards common goals and take collective responsibility for the success of the application and its operations.

This approach fosters a sense of shared accountability, where individuals from both teams recognize that their actions and decisions impact the overall performance, reliability, and user experience of the application. It promotes a culture of collaboration and teamwork, where the focus is on achieving business outcomes rather than assigning blame or working in isolation.

upvoted 2 times

  **tbolick6** 1 year, 4 months ago

Selected Answer: D

Shared ownership
upvoted 1 times

  **KanikaA** 1 year, 4 months ago

Selected Answer: D

Basic devops principle
upvoted 2 times

  **gviolant** 1 year, 5 months ago

Selected Answer: B

B. , reduce silos, together they defined service level objectives or SLOS, calculate error budgets and determine order reliability and work priorities. Culturally, shared ownership promotes shared vision and knowledge and the need for better collaboration and communication.

upvoted 2 times

  **USalo** 1 year, 4 months ago

shared ownership isn't equal to shared access

D is the right answer

upvoted 1 times

  **Vijaydv** 1 year, 7 months ago

Selected Answer: D

Shared access and combining job responsibilities are two different things. Answer is D, By emphasizing shared ownership of business outcomes

upvoted 3 times

  **C8H10N4O2** 1 year, 7 months ago

Selected Answer: B

B.

The content of the question is the cause of DevOps. To reduce silos, of course, information must be shared. B's 'everyone has shared access' is suitable. In D, simply 'emphasize', so the effect will not be greater than mixing the responsibilities of the job role.

upvoted 1 times

  **minmin2020** 1 year, 8 months ago

Selected Answer: D

D. By emphasizing shared ownership of business outcomes

upvoted 2 times

  **Klajdi_98** 1 year, 8 months ago

Selected Answer: B

One of DevOps principles

upvoted 2 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: D

D seems to be correct as it is required unified vision

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

D. By emphasizing shared ownership of business outcomes

upvoted 1 times

How does a large hotel chain benefit from storing their customer reservation data in the cloud?

- A. On-premises hardware access to transaction data
- B. Real-time data transformation at scale within an on-premises database
- C. Real-time business transaction accuracy at scale
- D. Physical hardware access during peak demand

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **mdsarfraz69** 9 months ago

Selected Answer: C

C is correct

upvoted 1 times
-   **cookieMr** 1 year ago

Selected Answer: C

A large hotel chain can benefit from storing their customer reservation data in the cloud by achieving real-time business transaction accuracy at scale.

Overall, by leveraging the cloud for storing customer reservation data, the hotel chain can achieve real-time accuracy, scalability, and data insights, leading to improved customer service, streamlined operations, and better business decision-making.

upvoted 3 times
-   **jaxclain** 1 year, 9 months ago

Selected Answer: C

Easy answer here, just discard the other 3 A, B and D which are no benefits, the only benefit between the 4 if you move to the Cloud is Real time business transaction accuracy at scale, Correct answer C

upvoted 2 times
-   **kinchit9k** 1 year, 9 months ago

Selected Answer: C

C. Real-time business transaction accuracy at scale

upvoted 2 times
-   **Vin1975** 1 year, 9 months ago

C. Real-time business transaction accuracy at scale

upvoted 1 times

An organization wants to migrate legacy applications currently hosted in their data center to the cloud. The current architecture dictates that each application needs its own operating system (OS) instead of sharing an OS. Which infrastructure solution should they choose?

- A. Virtual machines
- B. Open source
- C. Serverless computing
- D. Containers

Correct Answer: B

Community vote distribution

A (88%)

8%

- jellybiscuit

Highly Voted

1 year, 8 months ago

Selected Answer: A

Sometimes I wonder why this site even bothers to provide answers.
upvoted 24 times
- shuvs

1 year, 7 months ago

Lol. And who provide the answers on their behalf
upvoted 5 times
- Kasica96

1 year, 3 months ago

It's mentioned on the top of the each page :) - "Expert Verified" HAHAHAHA
upvoted 4 times
- maariacasasola

7 months, 1 week ago

I read in another answer discussion that many of the answers are wrong because they don't want to get in trouble with GC. So I always trust in 'Most voted' option xD
upvoted 3 times
- Ej24356

Highly Voted

1 year, 9 months ago

Selected Answer: A

Virtualization enables you to run multiple operating systems on the hardware of a single physical server, while containerization enables you to deploy multiple applications using the same operating system on a single virtual machine or server. Serverless computing would be no OS require and the open source operating system allows the use of code that is freely distributed and available to anyone and for commercial purposes such as Linux and Free BSD.
upvoted 9 times

himel2024

Most Recent

4 months, 4 weeks ago

Selected Answer: A

Virtual machines
upvoted 1 times

VenVPS

5 months, 2 weeks ago

Selected Answer: A

Virtual Machines
upvoted 1 times

pandeypc

6 months ago

A. Virtual machines

If the current architecture dictates that each application needs its own operating system (OS) instead of sharing an OS, the organization should choose Virtual Machines (VMs). Virtual machines provide a level of isolation that allows each application to run on its own virtualized instance with a dedicated operating system.

Options B (Open source), C (Serverless computing), and D (Containers) are not directly related to the requirement of each application having its own operating system. Open source refers to the development model, serverless computing abstracts infrastructure details, and containers share the OS kernel. Virtual machines align with the specified architectural constraint of separate operating systems for each application.
upvoted 1 times

chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times

  **mdsarfraz69** 9 months ago

Selected Answer: B

B is correct

upvoted 2 times

  **zeezeez** 9 months, 2 weeks ago

Selected Answer: D

having containers with different operating systems for each application is much efficient and lightweight solution that provisioning VMs for each application just to have different OS.

upvoted 1 times

  **Lufly** 10 months ago

Selected Answer: D

Containers are a lightweight virtualization technology that allows you to package an application and its dependencies into a single unit that can be run on any infrastructure. This makes them a good choice for migrating legacy applications that need their own OS.

The other options are not as relevant to this scenario. Option A, virtual machines, are a more heavyweight virtualization technology that can be used to run multiple operating systems on a single physical server. However, they are not as portable as containers and can be more difficult to manage. Option B, open source, is a term for software that is freely available for anyone to use and modify. It is not an infrastructure solution. Option C, serverless computing, is a cloud computing model that allows you to run code without having to manage servers. It is not a good choice for legacy applications that need their own OS.

upvoted 1 times

  **Brooklyn_** 10 months, 1 week ago

Selected Answer: A

Given that the organization's current architecture dictates that each application needs its own operating system (OS) rather than sharing an OS, using virtual machines (VMs) would be a more suitable choice for migrating the legacy applications to the cloud. Here's why:

Virtual Machines (VMs):

VMs provide a way to run multiple instances of operating systems on a single physical server. Each VM operates as a separate and isolated environment with its own dedicated OS, applications, libraries, and resources. In this case, where each application requires its own OS, VMs align well with the existing architecture.

Containers:

Containers are designed to share a single OS kernel while isolating applications from one another. In the case where each application requires its own OS, containers might not be the best fit, as they are designed to share a common OS.

upvoted 1 times

  **Erinias74** 11 months, 1 week ago

Selected Answer: D

If each application requires its own operating system, using virtual machines would be less efficient and more resource-intensive, as each VM would run its own full OS.

The correct option is D. Containers

upvoted 2 times

  **Azwscp2023** 11 months, 3 weeks ago

Selected Answer: A

surely A

upvoted 1 times

  **peterm42** 1 year ago

THIS WEBSITE IS USELESS

upvoted 1 times

  **cookieMr** 1 year ago

Selected Answer: A

Google Compute Engine provides a flexible and scalable virtual machine (VM) infrastructure that allows organizations to run their applications in the cloud. It offers a wide range of VM configurations, including various operating system options, to meet the specific requirements of different applications.

In the case of legacy applications that require separate operating systems for each application, Google Compute Engine allows the organization to create individual VM instances with dedicated operating systems. Each application can have its own isolated environment, ensuring compatibility and independence.

upvoted 1 times

  **pillais** 1 year, 2 months ago

Selected Answer: A

Virtual machines - One Os
Container - OS shared

upvoted 1 times

  **Tan1002_** 1 year, 2 months ago

Why are the answers incorrect provided by the site?

upvoted 1 times

An organization operates their entire IT infrastructure from Google Cloud.
What should they do to prepare for data breaches?

- A. Reduce reliance on multi-factor authentication
- B. Data security is Google's responsibility, so preparation is minimal
- C. Create an incident plan to mitigate impacts
- D. Strengthen their data center perimeter security

Correct Answer: B

Community vote distribution

C (97%)

- Govindaraj

Highly Voted

1 year, 9 months ago

Selected Answer: C

C seems to be correct as data security is responsible for both cloud vendor and customer
upvoted 7 times
- himel2024

Most Recent

4 months, 4 weeks ago

Selected Answer: C

Create an incident plan to mitigate impacts
upvoted 1 times
- Millervilla

5 months, 2 weeks ago

C IS CORRECT
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 2 times
- __rajan__

8 months ago

Selected Answer: C

C. Create an incident plan to mitigate impacts

Even though Google Cloud provides a secure platform, it is still important for organizations to have an incident plan in place in case of a data breach.
upvoted 3 times
- mdsarfraz69

9 months ago

Selected Answer: B

B is correct
upvoted 1 times
- Lufly

10 months ago

Selected Answer: C

By creating an incident plan, you can help to protect your organization from the impacts of a data breach.

The other options are not as relevant to this scenario. Option A, reducing reliance on multi-factor authentication, is not a good idea. Multi-factor authentication is an important security measure that can help to prevent unauthorized access to your data. Option B, data security is Google's responsibility, so preparation is minimal, is not correct. Google Cloud provides a secure platform, but it is still important for organizations to take steps to protect their data. Option D, strengthening their data center perimeter security, is a good step, but it is not the only thing that organizations need to do to prepare for data breaches.
upvoted 2 times
- cookieMr

1 year ago

Selected Answer: C

It's important to note that data breaches can never be completely eliminated, but by implementing strong security measures and following best practices, an organization can significantly reduce the risk and impact of a data breach in their Google Cloud infrastructure.

Create an incident response plan that outlines the steps to be taken in the event of a data breach. This should include procedures for containing the breach, assessing the impact, notifying affected parties, and initiating recovery processes.
upvoted 3 times

🗨️ 👤 **omgitsele** 1 year, 2 months ago

Selected Answer: C

Shared responsibility
upvoted 1 times

🗨️ 👤 **adityanarayan** 1 year, 4 months ago

bhai tum kar kya rahe ho
upvoted 1 times

🗨️ 👤 **Mukesh21** 1 year, 4 months ago

B says, minimal preparation that means, so least preparation is definitely needed on the customer end but majority of it will be handled by cloud provider.
upvoted 2 times

🗨️ 👤 **Overcast8509** 1 year, 5 months ago

I think the answer will be Option C. If google completely take care of security then why google will provide custom encryption key option instead of google's default one. If the custom key which we had provided may be easy to hack then we can't blame GCP regarding this. So it must be C as per my opinion
upvoted 1 times

🗨️ 👤 **Jackey0117** 1 year, 5 months ago

Selected Answer: C

C. Create an incident plan to mitigate impacts. Additionally, they should also consider implementing security best practices such as multi-factor authentication, regular security audits and assessments, and monitoring for unusual activity on their systems.
upvoted 2 times

🗨️ 👤 **Londonkiwi** 1 year, 6 months ago

Google is fully responsible for hardware maintenance, so it must be C.
upvoted 1 times

🗨️ 👤 **manashbaruah** 1 year, 6 months ago

Selected Answer: C

Shared responsibility
upvoted 4 times

🗨️ 👤 **minmin2020** 1 year, 8 months ago

Selected Answer: C

C. Create an incident plan to mitigate impacts
upvoted 2 times

🗨️ 👤 **Ashish_01** 1 year, 9 months ago

Selected Answer: C

option C is the correct answer, since it is a shared responsibilty
upvoted 2 times

An organization wants to transform multiple types of structured and unstructured data in the cloud from various sources. The data must be readily accessible for analysis and insights.

Which cloud data storage system should the organization use?

- A. Relational database
- B. Private data center
- C. Data field
- D. Data warehouse

Correct Answer: D

Community vote distribution

D (100%)

- Naree**

Highly Voted

1 year, 5 months ago

Can Unstructured data be stored in DWH..? Or, the question should be Structured and Semi-Structured..?

upvoted 11 times
- Millervilla**

Most Recent

5 months, 2 weeks ago

D is correct

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: D

D is correct.

upvoted 1 times
- cookieMr**

1 year ago

Selected Answer: D

A data warehouse is specifically designed to store and manage large volumes of structured and unstructured data from different sources. It provides a centralized repository for data storage, processing, and analysis. Data warehouses are optimized for querying and analyzing data, allowing organizations to perform complex analytics, generate insights, and make data-driven decisions.

A data warehouse can integrate with various data sources, including relational databases, file systems, streaming data, and more. It supports data transformation and aggregation, making it easier to prepare the data for analysis. Additionally, data warehouses often provide advanced features such as columnar storage, compression, indexing, and parallel processing to enhance performance and scalability.

By using a cloud data warehouse, the organization can leverage the flexibility, scalability, and agility of cloud computing to store and analyze their data effectively. It enables them to perform comprehensive analytics across different data types and sources, uncover valuable insights, and make informed business decisions.

upvoted 2 times
- Chayan_101**

1 year, 1 month ago

D. I think it will be Data warehouse

upvoted 2 times
- ArjunKennedy**

1 year, 2 months ago

I have the same query as Naree. But chose DWH as it seems only correct answer from options.

upvoted 3 times
- GSK25**

1 year, 3 months ago

D. Data warehouse

upvoted 1 times
- someshsehgal**

1 year, 6 months ago

Selected Answer: D

Datawarehouse

upvoted 2 times
- shuvs**

1 year, 7 months ago

Selected Answer: D

All types of data and ready for analysis
upvoted 1 times

  **juangt** 1 year, 8 months ago

Selected Answer: D

D. data warehouse is used for the analysis and reporting of structured and semi-structured data from multiple sources.
upvoted 3 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: D

D. It supports real-time insights. A data warehouse is an enterprise system used for the analysis and reporting of structured and semi-structured data from multiple sources,
<https://cloud.google.com/learn/what-is-a-data-warehouse>
upvoted 2 times

  **Vin1975** 1 year, 9 months ago

D. Data warehouse
upvoted 1 times

An organization wants to use all available data to offer predictive suggestions on their website that improve over time. Which method should the organization use?

- A. Data automation
- B. Trends analysis
- C. Machine learning
- D. Multiple regression

Correct Answer: C

Community vote distribution

C (100%)

- Pennycheng** 3 months ago

C is correct

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 2 times
- __rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
- Abhijeet_Patel** 8 months, 2 weeks ago

C is the best suitable answer.

upvoted 1 times
- cookieMr** 1 year ago

Selected Answer: C

Machine learning algorithms can process large volumes of data and identify complex patterns that may not be apparent to humans. As the organization continues to collect more data and gather feedback on the effectiveness of the suggestions, they can refine and improve the predictive models over time through a process known as "learning."

By employing machine learning techniques, the organization can leverage the power of their data to provide personalized and relevant suggestion to their website users. These suggestions can be based on user behavior, historical data, preferences, and other relevant factors. The organization can continuously analyze and update their machine learning models to adapt to changing user preferences and improve the accuracy of the suggestions provided.

upvoted 4 times
- Chayan_101** 1 year, 1 month ago

C seems to be correct answer

upvoted 1 times
- juangt** 1 year, 8 months ago

Selected Answer: C

C. ML provides modern machine learning services, with pre-trained models and a service to generate your own tailored models.

upvoted 2 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: C

C seems to be correct answer

upvoted 1 times
- Vin1975** 1 year, 9 months ago

C. Machine learning

upvoted 1 times

After rolling out a new update, an organization found a minor bug in its online video game.
How should the organization approach this bug while following SRE principles?

- A. Accept and learn from the bug because failure is normal
- B. Accept and ignore the bug because it is only minor
- C. Hold a postmortem to reprimand the employee responsible for the bug
- D. Document bug correction to eliminate all future bugs

Correct Answer: C

Community vote distribution

A (68%)

D (32%)

- VijayendraBaahubali

Highly Voted

1 year, 9 months ago

Selected Answer: A

Option A is correct. Accepting failure as normal is one of the SRE principles. SREs believe that accepting failure as normal helps to build an iterative collaborative culture. One way this is done is by holding a blameless "lessons learned" discussion after an incident occurs.

upvoted 13 times
- WellsExam

Highly Voted

1 year, 8 months ago

Selected Answer: D

The primary goals of writing a postmortem are to ensure that the incident is documented, that all contributing root cause(s) are well understood, and, especially, that effective preventive actions are put in place to reduce the likelihood and/or impact of recurrence

Blameless postmortems are a tenet of SRE culture. For a postmortem to be truly blameless, it must focus on identifying the contributing causes of the incident without indicting any individual or team for bad or inappropriate behavior. A blamelessly written postmortem assumes that everyone involved in an incident had good intentions and did the right thing with the information they had. If a culture of finger pointing and shaming individuals or teams for doing the "wrong" thing prevails, people will not bring issues to light for fear of punishment.

<https://sre.google/sre-book/postmortem-culture/>

upvoted 12 times
- 4a98421

Most Recent

3 months, 2 weeks ago

Selected Answer: A

Certainly not C, I wonder why such laughable answer as been tagged as correct answer.

upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A. Accept and learn from the bug because failure is normal

One of the core principles of Site Reliability Engineering (SRE) is that failure is normal. SRE teams understand that things will go wrong, and they focus on building systems that are resilient to failure.

upvoted 2 times
- GargiKekre

8 months, 2 weeks ago

Selected Answer: D

Option A is as per DevOps and Option D is as per SRE.

upvoted 1 times
- tsk9921

11 months, 2 weeks ago

A is right for me

upvoted 1 times
- Iz0001

1 year, 2 months ago

Option A - Accept failure as normal

upvoted 1 times
- tbolick6

1 year, 4 months ago

Selected Answer: A

Not D because you can't eliminate ALL future bugs.
upvoted 3 times

  **pushpendras21** 1 year, 6 months ago
Option A is more correct
upvoted 1 times

  **ronieto** 1 year, 7 months ago

Selected Answer: A

Accepting failure as normal is one of the SRE principles
upvoted 1 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: A

OPTION A
upvoted 1 times

  **SomeshSule** 1 year, 9 months ago
Why not D, document bug correction for future?
upvoted 1 times

  **TOLETIHARIPRIYASRUJANA** 1 year, 9 months ago
OPTION A
upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: A

A si the right answer. <https://www.blameless.com/sre/sre-principles>
upvoted 2 times

  **IvanKhru** 1 year, 9 months ago
100% A
upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: A

Option A is the correct answer. Option C is the funniest answer.
upvoted 3 times

Why do organizations often struggle to scale their on-premises application infrastructure?

- A. Scaling compute instances could breach compliance and/or regulation
- B. Increasing compute capacity is time-consuming and costly
- C. Their serverless compute functions struggle to meet the demand
- D. Their multi-cloud architecture is complex and expensive

Correct Answer: A

Community vote distribution

B (93%)

7%

- Bhavi0909

5 months, 4 weeks ago

B is correct

upvoted 2 times
- gyenesge

7 months, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- EeBs

1 year, 6 months ago

Selected Answer: A

In large organizations, there are people involved across multiple functions. Technology and line-of-business teams are often using cloud resources, but they don't necessarily factor cost into their decision-making. Finance teams control cloud costs, but may struggle to understand or to keep up with cloud spend on a daily, weekly, or monthly basis.

upvoted 1 times
- cristianBBBB

1 year, 6 months ago

Selected Answer: B

obviously B

upvoted 1 times
- shuvs

1 year, 7 months ago

Selected Answer: B

Easy one

upvoted 1 times
- Cert_IT

1 year, 8 months ago

Selected Answer: B

Seriously, How can scaling compute instance be a breach. The correct answer is B

upvoted 3 times
- minmin2020

1 year, 8 months ago

Selected Answer: B

B. Increasing compute capacity is time-consuming and costly

upvoted 1 times
- TOLETIHARIPRIYASRUJANA

1 year, 9 months ago

B OPTION

upvoted 1 times
- MurugeshBabu

1 year, 9 months ago

Option B is the right answer
upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: B

B. For Sure
upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: B

B is the correct answer
upvoted 1 times

  **POOJA3808** 1 year, 9 months ago

B is the correct answer
upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: B

B is the correct answer
upvoted 1 times

  **Pichikik** 1 year, 9 months ago

B - it is about on-premise... so will be time consuming and costly
upvoted 2 times

  **Vin1975** 1 year, 9 months ago

B. Increasing compute capacity is time-consuming and costly
upvoted 1 times

An organization wants to use BigQuery data analytics to understand their website performance, but wants to move only some data into the cloud. Which environment should the organization use?

- A. Private cloud
- B. On-premises
- C. Multi-cloud
- D. Hybrid cloud

Correct Answer: C

Community vote distribution

D (95%)

5%

- Loved

Highly Voted

1 year, 8 months ago

D. Hybrid Cloud allows to establish a "connection" between cloud and on-premises.

upvoted 6 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: D

In a hybrid cloud environment, the organization can keep their website data on-premises and then use BigQuery to analyze the data in the cloud. This allows the organization to take advantage of the scalability and performance of BigQuery without having to move all of their data to the cloud.

upvoted 2 times
- GargiKekre

8 months, 2 weeks ago

Selected Answer: C

Since there is no option for using on-premises data in Bigquery.

upvoted 1 times
- MCERT_TEST

1 year, 3 months ago

D Hybrid Cloud

upvoted 1 times
- Jackey0117

1 year, 5 months ago

Selected Answer: D

A hybrid cloud environment combines on-premises infrastructure with cloud resources, allowing organizations to keep sensitive data on-premises while still taking advantage of the scalability and flexibility of the cloud

upvoted 3 times
- grzeev

1 year, 8 months ago

Selected Answer: D

Multi-clouds always include more than one public cloud service

upvoted 1 times
- minmin2020

1 year, 8 months ago

Selected Answer: D

D. Hybrid cloud

upvoted 1 times
- grzeev

1 year, 8 months ago

Selected Answer: D

D. <https://www.techtarget.com/searchcloudcomputing/feature/Multi-cloud-vs-hybrid-cloud-and-how-to-know-the-difference>

upvoted 2 times
- TOLETIHARIPRIYASRUJANA

1 year, 9 months ago

HYBRID CLOUD --> OPTION D

upvoted 1 times
- MurugeshBabu

1 year, 9 months ago

is the question to move only some data into Google Cloud and other data in other cloud? If so, Multi-Cloud will be correct. Question is still not clear.

upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: D

D. Hybrid Cloud

upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: D

Option D (Hybrid cloud) is correct

upvoted 1 times

  **POOJA3808** 1 year, 9 months ago

D Hybrid Cloud

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: D

D Hybrid Cloud

upvoted 1 times

  **Ej24356** 1 year, 9 months ago

Selected Answer: D

The assumption should be made that there is still a private network involved. Hybrid clouds always include a private cloud and are typically managed as one entity. Multi-clouds always include more than one public cloud service, which often perform different functions.

upvoted 3 times

  **Vin1975** 1 year, 9 months ago

D. Hybrid cloud

upvoted 1 times

What is logging within the context of cloud technology?

- A. Writing application and operating system events as text
- B. Monitoring network and resource limitations
- C. Tracking source code across an organization
- D. Recording infrastructure and hardware expenditure

Correct Answer: A

Community vote distribution

A (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
-   **tsk9921** 11 months, 2 weeks ago

A seems correct here

upvoted 1 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: A

A. Logs are textual recording of events

upvoted 3 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: A

A. Cloud Logging is a fully managed service that allows you to store, search, analyze, monitor, and alert on logging data and events from Google Cloud and Amazon Web Services

upvoted 4 times
-   **POOJA3808** 1 year, 9 months ago

A is the correct answer

upvoted 1 times
-   **Kamlakar** 1 year, 9 months ago

You can collect logging data/events from most of the application on cloud hence A seems to be correct .

upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

A. Writing application and operating system events as text

upvoted 1 times

How does Cloud SQL help organizations create business insights?

- A. Integrates with business intelligence and analytics platforms
- B. Generates predictions using machine learning models
- C. Generates real-time charts and intelligent analytics
- D. Transforms business data from unstructured to structured

Correct Answer: A

Community vote distribution

A (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A. Integrates with business intelligence and analytics platforms

Cloud SQL is a fully managed database service that makes it easy to set up, maintain, and scale relational databases in the cloud. Cloud SQL integrates with a variety of business intelligence and analytics platforms, which allows organizations to use their data to create business insights.

upvoted 1 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: A

Can't do any of B, C, D

upvoted 1 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: A

A. <https://cloud.google.com/sql/docs/postgres/using-query-insights>

upvoted 3 times
-   **lax797** 1 year, 9 months ago

A. Integrates with business intelligence and analytics platforms

upvoted 2 times
-   **Kamlakar** 1 year, 9 months ago

ill go with A.

upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

A. Integrates with business intelligence and analytics platforms

upvoted 2 times

An organization recently launched a virtual customer support agent, generating vast amounts of text and speech data. Why should they use a cloud data warehouse to interpret this data?

- A. To natively visualize both types of data using a dashboard in real time
- B. To ingest and analyze structured and unstructured data at scale, in real time
- C. To secure data transmission between cloud and on-premises environments
- D. To transform data from structured to unstructured

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- [Removed]

1 year, 3 months ago

This one threw me. I wouldn't associate "unstructured data" with a data warehouse. The phrasing of this one is very confusing. I'm not sold on B being the answer.

upvoted 2 times
- HariSPK

1 year, 6 months ago

Selected Answer: B

Data warehouses assemble data from multiple sources including databases. An enterprise data warehouse is a system used for the analysis and reporting of structured and semi-structured data from multiple sources.

upvoted 1 times
- shuvs

1 year, 7 months ago

Selected Answer: B

A could have been a valid answer, but one can do a lot more than just visualize.

upvoted 1 times
- Govindaraj

1 year, 9 months ago

Selected Answer: B

B. Real-time data ingestion and updates. A simple and universal solution for continually ingesting your enterprise data into popular cloud-based data warehouses in real time.

https://www.qlik.com/us/cloud-data-migration/cloud-data-warehouse

upvoted 4 times
- Kamlakar

1 year, 9 months ago

B seems OK

upvoted 2 times
- Vin1975

1 year, 9 months ago

B. To ingest and analyze structured and unstructured data at scale, in real time

upvoted 1 times

An organization wants to add a new function to their application. They want to write the code and let the public cloud provider handle the infrastructure.
Which infrastructure solution should they use?

- A. Virtual machines
- B. Bare Metal Solution
- C. Serverless computing
- D. Container Registry

Correct Answer: A

Community vote distribution

C (80%)A (20%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 2 times
- __rajan__

8 months ago

Selected Answer: C

To add a new function to their application, the organization can use the serverless computing platform offered by their public cloud provider. For example, they could use AWS Lambda, Google Cloud Functions, or Azure Functions.
upvoted 2 times
- Mehrannn

11 months, 3 weeks ago

The organization should use serverless computing as the infrastructure solution to add the new function to their application.
upvoted 2 times
- manashbaruah

1 year, 6 months ago

Selected Answer: C

Keywords: Functions, serverless Computing
upvoted 4 times
- mtpro

1 year, 7 months ago

Selected Answer: A

My choice is A; Because "wants to add a new function to their application". Using VM is more appropriate
For Serverless Computing - the organization needs to fully migrate the application - not just add a new function alone
upvoted 4 times
- d11shah

3 months, 3 weeks ago

VM is a IaaS, this itself will put A as a wrong choice as infrastructure is to be handled by cloud, not user- hence no IaaS
upvoted 3 times
- gyenesge

7 months, 1 week ago

"to write the code and let the public cloud provider handle the infrastructure" I think this is the definition of serverless technology.
upvoted 3 times
- Srini_ANZ

1 year, 7 months ago

C : Serverless Computing. They have to write a function. So use Cloud Functions which is Serverless computing.
upvoted 3 times
- TonyMel

1 year, 9 months ago

C. Serverless computing
upvoted 1 times
- Govindaraj

1 year, 9 months ago

Selected Answer: C

C is the correct answer
upvoted 1 times
- jexmtropicscheatchatya

1 year, 9 months ago

Selected Answer: C

C. Serverless computing

upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: C

C is the correct answer. Virtual machines are IAAS in which user have to manage the infra.

upvoted 4 times

  **MGSS** 1 year, 6 months ago

IAAS - Infra as A Service , as a user what Infra do you manage, its managed by Cloud Provider. You just get ur required software installed for coding purpose.

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

Answer is C

upvoted 1 times

  **victorcl_111** 1 year, 9 months ago

Selected Answer: C

c es la correcta

upvoted 1 times

  **Kamlakar** 1 year, 9 months ago

Serverless computing , as public cloud prouder(eg. google) will mange the infra things . hence C .

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

C. Serverless computing

upvoted 1 times

  **soumik4u** 1 year, 9 months ago

Why the option is not C. Serverless Computing rather than A. Virtual Machines.
Admin please review and correct the answers with proper description.

upvoted 3 times

An organization wants to leverage tooling and automation as part of its new DevOps philosophy.
Which operational challenge will this resolve?

- A. Repetitive manual tasks that hinder workflows
- B. Time-consuming supervision of creative tasks
- C. Distribution and supply-chain issues
- D. Defective technical equipment that limits innovation

Correct Answer: A

Community vote distribution

A (100%)

-   **chai_gpt** 7 months, 3 weeks ago

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
-   **HariSPK** 1 year, 6 months ago

Selected Answer: A

Automation help in managing repetitive tasks

upvoted 3 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: A

Only valid answer is A

upvoted 1 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: A

A. Automation is the correct answer

upvoted 2 times
-   **Kamlakar** 1 year, 9 months ago

we use automation to avoid Repetitive manual tasks. hence A seems to be correct !

upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

A. Repetitive manual tasks that hinder workflows

upvoted 1 times

An organization needs to store structured, semi-structured, and unstructured data in its raw, native format in the same repository. Which cloud data management solution should the organization use?

- A. Data field
- B. Data lake
- C. Database
- D. Data warehouse

Correct Answer: B

Community vote distribution

B (87%)

13%

Leen8199 2 months, 1 week ago
key word is native format = Data lake
upvoted 1 times

madcloud32 5 months, 4 weeks ago
Selected Answer: B
B Data lake is correct
upvoted 1 times

chai_gpt 7 months, 3 weeks ago
Selected Answer: B
B is correct
upvoted 1 times

__rajan__ 8 months ago
Selected Answer: B
B. Data lake

A data lake is a cloud data management solution that can store structured, semi-structured, and unstructured data in its raw, native format in the same repository. Data lakes are designed to be scalable and flexible, and they can be used to store a wide variety of data

upvoted 1 times

tbolick6 1 year, 4 months ago
Selected Answer: B
Data lake for raw, native format all in one place.
upvoted 2 times

zellck 1 year, 6 months ago
Selected Answer: B
B is the answer.

<https://cloud.google.com/learn/what-is-a-data-lake>
A data lake is a centralized repository designed to store, process, and secure large amounts of structured, semistructured, and unstructured data. It can store data in its native format and process any variety of it, ignoring size limits.

upvoted 4 times

cristianBBBB 1 year, 6 months ago
Selected Answer: D
data lake can store any type of data
upvoted 1 times

ucsdmiami2020 1 year, 6 months ago
Selected Answer: B
Per Google docs, "data lake is a centralized repository designed to store, process, and secure large amounts of structured, semi structured, and unstructured data. It can store data in its native format and process any variety of it, ignoring size limits." native format aka raw
<https://cloud.google.com/learn/what-is-a-data-lake>
upvoted 3 times

shuvs 1 year, 7 months ago
Selected Answer: B
B is right. Data warehouse could also store all formats, but not in raw native form

upvoted 2 times

  **juangt** 1 year, 8 months ago

Selected Answer: B

B. A data lake is a centralized repository designed to store, process, and secure large amounts of structured, semistructured, and unstructured data.
upvoted 2 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: D

d: A data warehouse collects data from various sources, whether internal or external, and optimizes the data for retrieval for business purposes. The data is usually structured, often from relational databases, but it can be unstructured too.
upvoted 2 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: B

A data lake can store all types of data with no fixed limitation on account size or file and with no specific purpose defined yet. The data comes from disparate sources and can be structured, semi-structured, or even unstructured. Data-lake data can be queried as needed.
upvoted 2 times

  **Ej24356** 1 year, 9 months ago

Selected Answer: B

A data lake is a centralized repository designed to store, process, and secure large amounts of structured, semistructured, and unstructured data. It can store data in its native format and process any variety of it, ignoring size limits.
upvoted 2 times

  **Ej24356** 1 year, 9 months ago

<https://cloud.google.com/learn/what-is-a-data-lake>
upvoted 2 times

  **Vin1975** 1 year, 9 months ago

B. Data lake
upvoted 2 times

An organization wants its users to validate a series of new features for their app.
Why should they use App Engine?

- A. Because their app is containerized and enabled by microservices
- B. Because the updated app will only include new features
- C. To run different versions of the app for different users
- D. To run different versions of the app for the same user

Correct Answer: A

Community vote distribution

C (82%)

A (18%)

- Vin1975

Highly Voted

1 year, 9 months ago

C. To run different versions of the app for different users

AppEngine supports multiversion testing

upvoted 7 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: C

C is correct

upvoted 1 times
- reachmalhotras

11 months, 2 weeks ago

Selected Answer: C

Validate is the Key here, focus on user validation using App Engine Capability

upvoted 3 times
- gviolant

1 year, 5 months ago

Selected Answer: A

You can run your applications in App Engine by using the App Engine flexible environment or the App Engine standard environment. You can also choose to simultaneously use both environments for your application and allow your services to take advantage of each environment's individual benefits.

App Engine environments

App Engine is well suited to applications that are designed using a microservice architecture, especially if you decide to utilize both environments. Use the following sections to learn and understand which environment best meets your application's needs.

upvoted 3 times
- girishjoshi124

1 year, 5 months ago

Selected Answer: C

Option A doesn't focus on the part "wants its users to validate a series of new features for their app" .

upvoted 2 times
- ilberg

1 year, 5 months ago

Selected Answer: C

C is the correct answer. ...run different versions of apps...

upvoted 1 times
- Pankaj1978

1 year, 6 months ago

A is the correct option. For the users to test some new features of the App, that should run in a containerised environment having bifurcated into microservices so that changes in one part doesn't affect the other parts.

upvoted 3 times
- shuvs

1 year, 7 months ago

Selected Answer: C

Versioning. App version testing is done with multiple users

upvoted 1 times

  **jellybiscuit** 1 year, 8 months ago

Selected Answer: C

C - app engine versioning
upvoted 1 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: C

c : run different versions of the app for different users
upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: C

C is corret
upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

It's got to be C. A and B make assumptions.
upvoted 2 times



How would an organization benefit from using Looker?

- A. Optimal identity and access management
- B. Leading serverless warehousing technology
- C. Robust data roll-back accuracy
- D. Advanced business intelligence and analytics

Correct Answer: *D*

Community vote distribution

D (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D. Advanced business intelligence and analytics

Looker is a business intelligence (BI) and analytics platform that helps organizations to explore, analyze, and visualize their data. Looker provides a variety of features that can help organizations to improve their business insights
upvoted 1 times
-   **dectum** 9 months, 3 weeks ago

D, its a BI solution.
upvoted 1 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: D

D . Looker is a business intelligence software and big data analytics platform that helps you explore, analyze and share real-time business analytics easily.
upvoted 4 times
-   **IvanKhr** 1 year, 9 months ago

Definitely D
upvoted 2 times
-   **Kamlakar** 1 year, 9 months ago

Looker is a powerful business intelligence (BI) tool that can help a business develop insightful visualizations. hence D
upvoted 4 times
-   **Vin1975** 1 year, 9 months ago

D. Advanced business intelligence and analytic
upvoted 3 times

An organization delivers a proactive healthcare service. They want to efficiently and automatically collect patient data. What should the organization encourage the patients to do?

- A. Use at-home health screening devices and then upload their health data daily
- B. Wear Internet of Things (IoT) devices that upload their health data in real time
- C. Self-assess their health data and then document and upload it in real time
- D. Visit a nurse who will use Internet of Things (IoT) devices to collect and upload their health data

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- ArjunKennedy** 1 year, 2 months ago

it is mentioned efficient, why not visiting nurse is correct as user don't have to purchase IOT devices?

upvoted 1 times
- KirSF** 1 year, 2 months ago

Selected Answer: B

automatic is a key word

upvoted 1 times
- [Removed]** 1 year, 3 months ago

This seems to be a general tech/common sense question. I don't see how this is specific to GCP?

upvoted 1 times
- shuvs** 1 year, 7 months ago

Selected Answer: B

Most efficient and automatic

upvoted 1 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

B for Sure

upvoted 2 times
- SomeshSule** 1 year, 9 months ago

Automatic is the key word.

upvoted 1 times
- Vin1975** 1 year, 9 months ago

B. Wear Internet of Things (IoT) devices that upload their health data in real time

upvoted 2 times

How would a global organization benefit from managing their data with Cloud Spanner?

- A. Cloud Spanner is optimized for cold storage
- B. Cloud Spanner replicates data across regions in real time
- C. Cloud Spanner is optimized to ingest unstructured data
- D. Cloud Spanner visualizes and analyzes data in real time

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 2 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct

upvoted 1 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: B

Truly global service.

upvoted 1 times
-   **Cert_IT** 1 year, 8 months ago

Selected Answer: B

Cloud Spanner

upvoted 1 times
-   **juangt** 1 year, 8 months ago

Selected Answer: B

B. Cloud Spanner optimizes performance by automatically sharding the data based on request load and size of the data in real time.

upvoted 2 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: B

Spanner is Google's scalable, multi-version, globally-distributed, and synchronously-replicated database.

upvoted 2 times
-   **Vin1975** 1 year, 9 months ago

B. Cloud Spanner replicates data across regions in real time

upvoted 1 times

An employee receives an email from their internet service provider asking for their bank account number and password.
Which cybersecurity threat is this?

- A. Ransomware
- B. Distributed Denial of Service
- C. Spamming
- D. Phishing

Correct Answer: C

Community vote distribution

D (97%)

- shuvs

Highly Voted

1 year, 7 months ago

Selected Answer: D

Question 80 onwards, majority of provided answers are wrong. Admin should check. Go with most voted.
upvoted 19 times
- 4a98421

Most Recent

3 months, 2 weeks ago

Selected Answer: D

C really???
upvoted 1 times
- madcloud32

6 months ago

Marked answer is wrong. This is Phishing of data. D is correct.
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: D

D is correct.
upvoted 1 times
- Benoni

9 months, 1 week ago

Selected Answer: C

Phishing is when someone pretends to be original. In this case Internet service Provider is not pretending. He is asking for details in straight forward way. So I go with Spamming
upvoted 1 times
- reachmalhotras

11 months, 2 weeks ago

Selected Answer: D

Collecting PII Data Is Phishing
upvoted 1 times
- gviolant

1 year, 5 months ago

Selected Answer: D

Phishing
upvoted 1 times
- someshseghal

1 year, 6 months ago

Selected Answer: D

Admin PLS correct... Answer is D
upvoted 1 times
- cristianBBBB

1 year, 6 months ago

Selected Answer: D

easy D.
upvoted 1 times

berberar 1 year, 6 months ago

Selected Answer: D

Phishing is a type of social engineering where an attacker sends a fraudulent message designed to trick a person into revealing sensitive information to the attacker or to deploy malicious software on the victim's infrastructure like ransomware.

upvoted 2 times

Tecatero2001 1 year, 8 months ago

Selected Answer: D

Trying to extract confidential information. Phishing.

upvoted 2 times

jellybiscuit 1 year, 8 months ago

Selected Answer: D

d obviously

upvoted 1 times

Abhi_ma_cdl_22 1 year, 9 months ago

Selected Answer: D

D is correct.

upvoted 2 times

MurugeshBabu 1 year, 9 months ago

It should be phishing though the impacted people scammed downloading malicious attachment, giving passwords and sharing data

upvoted 2 times

erwin_t82 1 year, 9 months ago

The difference between spam and phishing is that, while they both may be inbox-clogging nuisances, only one (phishing) is actively aiming to steal login credentials and other sensitive data.

Spam is a tactic for hawking goods and services by sending unsolicited emails to bulk lists

upvoted 3 times

jexmtropicscheatchatya 1 year, 9 months ago

Selected Answer: D

D is the answer

upvoted 1 times

An organization wants to move from a tactical cloud adoption approach to a transformational approach. How should they adapt the way they lead the organization?

- A. Increase top-down visibility and foster a culture of blamelessness
- B. Shift from an operational expenditure model to capital expenditure
- C. Drive cloud adoption with an individual contributor focus
- D. Invest in on-premises infrastructure to redesign relationships between IT and employees

Correct Answer: B

Community vote distribution

A (94%)

6%

-   **stevee954** 7 months, 2 weeks ago

Oddly D could be right. to transform is a culture change, which requires true collaboration instead of just lip service to it. investing in physically co-locating stakeholders would help shift the culture.

upvoted 2 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 2 times
-   **__rajan__** 8 months ago

Selected Answer: A

A. Increase top-down visibility and foster a culture of blamelessness.

To move from a tactical cloud adoption approach to a transformational approach, an organization needs to change the way it leads the organization. This includes increasing top-down visibility and fostering a culture of blamelessness.

Top-down visibility means that leaders need to communicate the organization's cloud strategy and goals clearly and often. They need to explain why the organization is moving to the cloud and the benefits that it will bring.

upvoted 2 times
-   **Mehrannn** 11 months, 3 weeks ago

A,B and C are all correct answers. The only option which is not true is D.

upvoted 1 times
-   **WellsExam** 1 year, 8 months ago

Selected Answer: B

If your objective is to innovate with net new applications that solve new business problems (i.e., transformational) then we recommend taking the greenfield software development approach.

<https://cloud.google.com/blog/products/cloud-migration/real-world-cloud-migration-strategies>

upvoted 1 times
-   **WellsExam** 1 year, 8 months ago

Selected Answer: A

A is correct

upvoted 4 times
-   **SoftSami** 1 year, 8 months ago

Selected Answer: A

Absolutely A

upvoted 3 times
-   **minmin2020** 1 year, 8 months ago

Selected Answer: A

It has to be A as all the other answers don't make sense

upvoted 1 times
-   **grzeev** 1 year, 8 months ago

What is a transformational approach? -

Cloud transformation is the process of transitioning a business's data, apps, and software to the cloud

Answer is A!

upvoted 3 times

  **Abhi_ma_cdl_22** 1 year, 9 months ago

Selected Answer: A

A is correct.

upvoted 1 times

  **Kushman** 1 year, 9 months ago

Why are you posting wrong answers - Answer is A

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: A

Answer is A. B is wrong no matter what question is asked.

upvoted 2 times

  **Pichikik** 1 year, 9 months ago

A... seems some straight forward questions have incorrect answer...

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A. Increase top-down visibility and foster a culture of blamelessness

upvoted 2 times

Why is data stored in Google Cloud secure and private?

- A. Data is encrypted by the Security Command Center
- B. Data is encrypted by Cloud Data Loss Prevention
- C. Data is encrypted by default
- D. Data is encrypted when an appropriate tag is applied

Correct Answer: C

Community vote distribution

C (100%)

-   **cloudmind** 5 months ago

C is correct

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **gviolant** 1 year, 5 months ago

Selected Answer: C

Data is encrypted by default

upvoted 2 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: C

C. Encrypted and distributed

upvoted 1 times
-   **juangt** 1 year, 8 months ago

Selected Answer: C

C. All data that is stored by Google is encrypted at the storage layer using the Advanced Encryption Standard (AES) algorithm, AES-256

upvoted 2 times
-   **Pikapika12638** 1 year, 9 months ago

Selected Answer: C

C. Data is encrypted by default

<https://cloud.google.com/docs/security/encryption/default-encryption#:~:text=Google%20uses%20the%20Advanced%20Encryption,to%202015%20that%20use%20AES128.>

upvoted 4 times
-   **Vin1975** 1 year, 9 months ago

C. Data is encrypted by default

upvoted 2 times

An organization wants to use multiple marketing datasets to forecast user acquisition.
How should they use cloud technology to gain new insights from the data?

- A. Import the datasets into a custom data warehouse, and then archive old data
- B. Import and selectively archive the datasets in a custom data lake
- C. Separate the datasets and make predictions using machine learning
- D. Combine the datasets and make predictions using machine learning

Correct Answer: D

Community vote distribution

D (100%)

- Vin1975**

Highly Voted

1 year, 9 months ago

D. Combine the datasets and make predictions using machine learning

upvoted 6 times
- stevee954**

Most Recent

7 months, 2 weeks ago

oops- i meant C.

upvoted 1 times
- stevee954**

7 months, 2 weeks ago

Machine learning uses a test dataset to create the model and a validation dataset to test it. Combining them first makes sense if you then stipulate the validation dataset is extracted from the combined set, but in that case why combine them. i voted D.

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: D

D is correct.

upvoted 1 times
- shuvs**

1 year, 7 months ago

Selected Answer: D

ML is required for prediction. And ML models work best on large and varied data, so combining them is the best option

upvoted 4 times
- zellck**

1 year, 9 months ago

Selected Answer: D

D is my answer

upvoted 2 times

An organization wants to collect metrics and metadata from their cloud applications and put them into dashboards. Which Google Cloud tool should they use?

- A. Cloud Monitoring
- B. Cloud Trace
- C. Cloud Logging
- D. Cloud Debugger

Correct Answer: A

Community vote distribution

A (85%)

D (15%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

<https://cloud.google.com/monitoring/docs>

Google Cloud's operations suite ingests that data and generates insights via dashboards, charts, and alerts.
upvoted 3 times

  **131983** 1 year, 5 months ago

Selected Answer: D

Answer is D - Keyword here is "application"
Cloud Debugger helps monitor application performance.
Cloud Monitoring enables users to monitor the performance of their entire cloud infrastructure.
upvoted 2 times

  **131983** 1 year, 5 months ago
but not sure about the dashboard bit :/
upvoted 1 times

  **moi23** 1 year, 4 months ago
It's A
upvoted 1 times

  **ronieto** 1 year, 7 months ago

Selected Answer: A

Cloud Monitoring as it provides dashboards extracted from defined metrics
upvoted 2 times

  **C8H10N4O2** 1 year, 7 months ago

Selected Answer: A

"dashboards" >>>> Cloud Monitoring
upvoted 2 times

  **BryRob** 1 year, 7 months ago

Selected Answer: A

Cloud Monitoring
Gain visibility into the performance, availability, and health of your applications and infrastructure.

<https://cloud.google.com/monitoring>
upvoted 1 times

  **zelck** 1 year, 9 months ago

Selected Answer: A

A is the answer.

<https://cloud.google.com/monitoring>
upvoted 1 times

  **Kamlakar** 1 year, 9 months ago

Cloud Monitoring let you create custom dashboard.
upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A. Cloud Monitoring
upvoted 1 times

A manager wants to review Google Cloud data access among their employees.
Who is responsible for defining data access policies?

- A. Cloud Identity
- B. Google Cloud Customer Care team
- C. Their organization's IT team
- D. Their organization's end users

Correct Answer: A

Community vote distribution

C (87%)

13%

- Ej24356

Highly Voted

1 year, 9 months ago

Selected Answer: C

The question asks who is responsible not how it is done

upvoted 10 times
- madcloud32

Most Recent

6 months ago

C is correct.

Marked answer A is wrong. A is platform that provides the service. C is one who decides what permissions are needed.

upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: C

C. Their organization's IT team

The organization's IT team is responsible for defining data access policies in Google Cloud. This includes defining who has access to what data and how they can access it.

upvoted 1 times
- Manny_SUI

1 year, 4 months ago

Answer is C - Keyword: Who

upvoted 1 times
- ed_per

1 year, 7 months ago

Selected Answer: C

Who? The IT Team.

What Service? IAM

upvoted 2 times
- shuvs

1 year, 7 months ago

Selected Answer: C

Cloud identity is the tool (not 'who'). And end users shouldn't be defining access - security risk.

upvoted 1 times
- Alfred_Magno

1 year, 8 months ago

Easy, Who is responsible for defining data access policies?, the key is the first word they use for question, We use the question word "who" when we ask questions about a person.

upvoted 2 times
- brainysaki

1 year, 8 months ago

Selected Answer: A

IMO answer to these types of questions is always one of the google cloud services, which is why I'm going with Option A which even relates to what has been asked.

upvoted 2 times
- ed_per

1 year, 7 months ago

If they meant "Which service" they'd have used that language.

upvoted 1 times

  **jaxclain** 1 year, 9 months ago

Selected Answer: C

The only right answer here is C, those who think is Cloud Identity are wrong, remember that if having CI is optional so either way the one who is responsible for defining data access policies is always the Orgs IT team.

upvoted 4 times

  **AwesomeGCP** 1 year, 9 months ago

Its A - Cloud Identity and Access Management (IAM) helps customers to define fine-grained access policies and precisely control access to Google Cloud-hosted data.

upvoted 2 times

  **AwesomeGCP** 1 year, 9 months ago

C - Cloud Identity and Access Management (IAM) helps customers to define fine-grained access policies and precisely control access to Google Cloud-hosted data.

upvoted 2 times

  **dianprat** 1 year, 9 months ago

The question is "Who" not how or what. so answer A is wrong, (and definitely not B). So either C and D. "End users" seems low level end users (not management). so C should be correct

upvoted 3 times

  **Sbgani** 1 year, 9 months ago

Selected Answer: A

Cloud Identity and Access Management (IAM) helps customers to define fine-grained access policies and precisely control access to Google Cloud hosted data.

upvoted 1 times

  **ed_per** 1 year, 7 months ago

Yes, but who defines the policies which are then implemented in IAM?

upvoted 1 times

  **soumik4u** 1 year, 9 months ago

The term is using in the question Google Cloud Data.....we have to perform the task using Cloud Identity option A.....if the term Google Cloud Data is not mentioned in the questions then we can go for the option C.

upvoted 1 times

  **SomeshSule** 1 year, 9 months ago

It ultimately depends on end user requirements. Who wants to access what data is defined by end user needs.

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

C. Their organization's IT team

upvoted 1 times

An organization has decided to modernize their applications in the cloud to keep up with their customers' needs. What may have prompted this business decision?

- A. Their on-premises applications only autoscale to meet demand.
- B. They want to change from a pay-as-you-go model to a capital expenditure model.
- C. Their source code changes erroneously without developer interaction.
- D. Their on-premises applications take months to update and deploy.

Correct Answer: D

Community vote distribution

D (100%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: D

D. Their on-premises applications take months to update and deploy.

One of the main benefits of modernizing applications in the cloud is that it can help organizations to update and deploy applications more quickly and easily. This is because cloud platforms provide a variety of features and tools that can automate the development, deployment, and management of applications.

upvoted 1 times

  **Emmanski08** 1 year, 7 months ago

Keyword: Modernized
upvoted 3 times

  **BryRob** 1 year, 7 months ago

Selected Answer: D

D. Their on-premises applications take months to update and deploy.
upvoted 2 times

  **shuvs** 1 year, 7 months ago

Selected Answer: D

On-premise takes longer, as one has to update systems too.
upvoted 1 times

  **Abhi_ma_cdl_22** 1 year, 9 months ago

Selected Answer: D

D is correct
upvoted 1 times

  **Vin1975** 1 year, 9 months ago

D. Their on-premises applications take months to update and deploy.
upvoted 1 times

An organization needs to categorize a large group of photographs using pre-trained machine learning. Which Google Cloud product or service should the organization use?

- A. Vision API
- B. BigQuery ML
- C. AutoML Vision
- D. Looker

Correct Answer: A

Community vote distribution

A (100%)

- jellybiscuit**

Highly Voted

1 year, 8 months ago

Selected Answer: A

Correct. A and C handle image classification, but Vision API is pretrained. AutoML vision is not.
upvoted 5 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: A

A is correct
upvoted 1 times
- BryRob**

1 year, 7 months ago

Selected Answer: A

<https://cloud.google.com/vision>

Vision AI
Use our game-changing fully managed development environment Vertex AI Vision to create your own computer vision applications or derive insights from images and videos with pre-trained APIs, AutoML, or custom models.
upvoted 1 times
- juangt**

1 year, 8 months ago

Selected Answer: A

A. The Google Cloud Vision API allows developers to easily integrate vision detection features within applications, including image labeling, face and landmark detection, optical character recognition (OCR), and tagging of explicit content
upvoted 2 times
- zelck**

1 year, 9 months ago

Selected Answer: A

A is the answer.

<https://cloud.google.com/vision>
upvoted 1 times
- Kamlakar**

1 year, 9 months ago

Vision API offers powerful pre-trained machine learning models through REST and RPC APIs. hence A.
upvoted 1 times
- Vin1975**

1 year, 9 months ago

A. Vision API
upvoted 1 times

An organization provides a loyalty program for its customers. It recently partnered with other businesses so that customers can get loyalty points at a range of other stores.

Why should the organization use application programming interfaces (APIs)?

- A. To migrate all partner data for disaster recovery
- B. To analyze and publish loyalty program statistics to a dashboard
- C. To personalize recommendations for loyalty card users
- D. To connect third-party systems to ensure up-to-date information

Correct Answer: D

Community vote distribution

D (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D is correct

upvoted 1 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: D

D. To connect third-party systems to ensure up-to-date information

key words: partnered with other business

upvoted 2 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: D

Connecting multiple systems (cloud and on-prem) is what API helps with

upvoted 1 times
-   **zelck** 1 year, 9 months ago

Selected Answer: D

D is my answer.

upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

D. To connect third-party systems to ensure up-to-date information

upvoted 1 times

A travel agency wants to seamlessly share their social media content with partners.
What should the organization use?

- A. A NoSQL database
- B. Anthos Config Management
- C. The App Engine standard environment
- D. An application programming interface (API)

Correct Answer: D

Community vote distribution

D (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: D

D is correct

upvoted 1 times
- Londonkiwi

1 year, 6 months ago

D , as the aim of APIs is to create a seamless connection between different systems.

upvoted 4 times
- shuvs

1 year, 7 months ago

Selected Answer: D

API's help sharing information across systems

upvoted 2 times
- zellck

1 year, 9 months ago

Selected Answer: D

D is the answer.

upvoted 1 times
- Vin1975

1 year, 9 months ago

D. An application programming interface (API)

upvoted 2 times

How do out-of-the-box APIs make artificial intelligence and machine learning more accessible for all Google Cloud customers?

- A. They don't require technical specialists.
- B. They don't require data input and validation.
- C. They require fewer security permissions.
- D. They all require custom training models.

Correct Answer: A

Community vote distribution

A (90%)

10%

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
- Vijaydv** 1 year, 7 months ago

Selected Answer: A

Need less tech expertise to use them out of the box

upvoted 4 times
- ed_per** 1 year, 7 months ago

Selected Answer: A

Can't be D, as not *all* require custom training.

upvoted 1 times
- shuvs** 1 year, 7 months ago

Selected Answer: A

All out of the box API's like Vision API, Natural Language API are pre-trained models. They need less technical expertise and even non data scientists can train models using them. D is not the answer, as the API's are pre-trained and don't need custom training.

upvoted 2 times
- juangt** 1 year, 8 months ago

<https://cloud.google.com/ai-platform/docs/technical-overview>

upvoted 1 times
- juangt** 1 year, 8 months ago

Selected Answer: D

The AI Platform training service allows you to train models using a wide range of different customization options. You can also select different way to customize your training application. You can submit your input data for AI Platform to train using a built-in algorithm (beta)

upvoted 1 times

dabizssan 1 year, 7 months ago

It is dangerous when an answer has the word "all" in it, I have my doubts that it is correct.

upvoted 1 times

An organization needs to migrate specialized workloads to the cloud while maintaining their existing complex licensing and architecture. What Google Cloud solution should the organization use?

- A. Compute Engine
- B. Bare Metal Solution
- C. Cloud Run
- D. Cloud Functions

Correct Answer: B

Community vote distribution

B (78%)

A (22%)

- joesatriani

Highly Voted

1 year, 6 months ago

B.

There are many workloads that are easy to lift and shift to the cloud, but there are also specialized workloads (such as Oracle) that are difficult to migrate to a cloud environment due to complicated licensing, hardware, and support requirements. Bare Metal Solution provides a path to modernize these applications.

upvoted 10 times
- Ej24356

Highly Voted

1 year, 9 months ago

Selected Answer: B

"This solution provides a path to modernize your application infrastructure landscape, while maintaining your existing investments and architecture. With Bare Metal Solution, you can bring your specialized workloads to Google Cloud, allowing you access and integration with GCP services with minimal latency."

upvoted 5 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: B

The best Google Cloud solution for migrating specialized workloads to the cloud while maintaining their existing complex licensing and architecture is Bare Metal Solution.

upvoted 1 times
- [Removed]

1 year, 3 months ago

Selected Answer: B

"specialized workloads" = Bare Metal

upvoted 2 times
- BryRob

1 year, 7 months ago

Selected Answer: B

<https://cloud.google.com/blog/products/compute/bare-metal-solution-enabling-specialized-workloads-in-google-cloud>

<https://cloud.google.com/bare-metal/docs/bms-setup>

upvoted 2 times
- ed_per

1 year, 7 months ago

I guess it's based on what they don't say. They don't say that they're already all running in VMs, which means that Compute Engine may not work.

upvoted 1 times
- shuvs

1 year, 7 months ago

Selected Answer: A

My vote is A Compute Engine. Compute engine BYOL supports bringing your own licensing. Bare metal is specialized for Oracle workloads, which isn't mentioned in question.

upvoted 3 times
- AMMARABO

1 year, 8 months ago

Selected Answer: B

B as complex licensing

upvoted 2 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: A

A. Compu<https://www.examtopics.com/exams/google/cloud-digital-leader/view/117/#te> Engine
upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: B

B seems to be the correct answer
upvoted 1 times



An organization finds that the amount of cash in their vending machines doesn't match the value of items sold. They have decided to upgrade their vending machines with cloud-based mobile payment systems.
How could the organization benefit from this upgrade?

- A. They could relax data access permissions.
- B. They could reduce their error budget overspend.
- C. They could improve their perimeter security.
- D. They could view data history to see transactions.

Correct Answer: D

Community vote distribution

D (77%)14%9%

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: D

D. View data history to see transactions.

Upgrading vending machines with cloud-based mobile payment systems can help organizations track sales transactions more accurately and efficiently. This can help organizations to identify and address any discrepancies between the amount of cash in their vending machines and the value of items sold.
upvoted 1 times

  **Emma412** 10 months ago

Answer = D
explanation
A) is not the answer as relaxing data access permission is not related to fixing the mismatches sale/transaction/value
b) is not the answer as error budget has nothing to do with this the problem statement
c) is not the answer as "perimeter security" is to control the communication between Google Cloud resources such as VMs within a zone.
4) directly relates to the problem statement where transaction history can be used to verify or identify the reason for mismatch in values
upvoted 4 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: D

D to reconcile the mismatch in transaction volume and revenue.
upvoted 3 times

  **tbolick6** 1 year, 4 months ago

Selected Answer: D

D seems like an obvious benefit of cloud technology.
upvoted 2 times

  **maney381** 1 year, 7 months ago

Selected Answer: D

for me
upvoted 2 times

  **BryRob** 1 year, 7 months ago

Selected Answer: D

for me
D. They could view data history to see transactions

key word: doesn't match the value of items sold / upgrade to cloud-based and what benefit to this
not an error its not updated and tallied / data and permission (customer responsibility, you sure you can relax? or perimeter security, their physical network (since they moved to cloud, customer on the configuration and cloud provider physical responsibility) did not improved, they will never use it.
upvoted 2 times

  **shuvs** 1 year, 7 months ago

Selected Answer: D

I see many selecting B. B could be a valid answer, but IMO error budgeting is not about cash tally at the till. It is defined as "An error budget is the maximum amount of time that a technical system can fail without contractual consequences."

upvoted 1 times

  **C8H10N4O2** 1 year, 7 months ago

Selected Answer: B

B.... "The amount of cash in their vending machines doesn't match the value of items sold", the most obvious answer is B, also C is ok but the most logic answer for me is B.

Also noticed some selected D? it's obvious you can view the data history for transactions but you have to read the question: "The amount of cash in their vending machines doesn't match with the value of items sold" This only concerns security / budget overspent.

upvoted 2 times

  **AwesomeGCP** 1 year, 8 months ago

B. They could reduce their error budget overspend.

upvoted 1 times

  **zellick** 1 year, 9 months ago

Selected Answer: D

D is the answer.

upvoted 1 times

  **jaxclain** 1 year, 9 months ago

B.... "The amount of cash in their vending machines doesn't match the value of items sold", the most obvious answer is B, also C is ok but the most logic answer for me is B.

Also noticed some selected C? it's obvious you can view the data history for transactions but you have to read the question: "The amount of cash in their vending machines doesn't match with the value of items sold" This only concerns security / budget overspent.

upvoted 3 times

  **bobby8521** 1 year, 9 months ago

Selected Answer: C

Option C looks reasonable

upvoted 1 times

  **all_cloud** 1 year, 9 months ago

Selected Answer: B

after moving to cloud, they will see the data discrepancy hence it will reduce error budget

upvoted 1 times

  **MrMyself1** 1 year, 9 months ago

Selected Answer: D

Cloud mobile payment system will allow you to see transactions and history

upvoted 4 times

  **LimeCake** 1 year, 9 months ago

Selected Answer: C

I think Its C, others are not related to migrating cloud, i think.

upvoted 1 times

A food delivery service needs access to real-time menu information from all partner restaurants. They also need to share customer order information with the restaurants in real time.
What should the organization use?

- A. Site reliability engineering (SRE)
- B. An application programming interface (API)
- C. A customized machine learning model
- D. A multi-regional database

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct

upvoted 1 times
-   **ALIZHEEEEE** 1 year, 6 months ago

B:for data integration

upvoted 2 times
-   **brainysaki** 1 year, 8 months ago

Selected Answer: B

An application programming interface (API)

upvoted 2 times
-   **Pikapika12638** 1 year, 9 months ago

Selected Answer: B

B. An application programming interface (API)

upvoted 2 times

An organization is planning to create many new applications and needs a system for managing and orchestrating their containers. Why should they use Google Kubernetes Engine?

- A. It enables rapid application deployment and testing.
- B. It optimizes virtual machine performance.
- C. It connects legacy systems to an application.
- D. It provides pre-written source code for application development.

Correct Answer: A

Community vote distribution

A (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: A

A. It enables rapid application deployment and testing.
upvoted 2 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: A

A. Containerized, hence faster to deploy and test
upvoted 1 times
-   **juangt** 1 year, 8 months ago

Selected Answer: A

Enable rapid application development and iteration by making it easy to deploy, update, and manage your applications and services. Configure GKE, Cloud Source Repositories, Cloud Build, and Spinnaker for Google Cloud services to automatically build, test, and deploy an app.
<https://cloud.google.com/kubernetes-engine>
upvoted 1 times
-   **AMMARABO** 1 year, 8 months ago

Selected Answer: A

A. It enables rapid application deployment and testing.
upvoted 2 times

How does Google Cloud enable customers to leverage the power of artificial intelligence and machine learning?

- A. With scalable on-premises tensor processing units
- B. With APIs and tools for users at different levels of expertise
- C. With affordable and scalable quantum computing
- D. With machine learning models that always require custom training

Correct Answer: B

Community vote distribution

B (100%)

-   **BryRob**

Highly Voted 

 1 year, 7 months ago

Selected Answer: B

<https://cloudacademy.com/blog/what-are-the-benefits-of-machine-learning-in-the-cloud/>

What are the Benefits of Machine Learning in the Cloud?

The cloud makes intelligent capabilities accessible without requiring advanced skills in artificial intelligence or data science.

upvoted 6 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
-   **Kaycee101** 9 months, 2 weeks ago

I think the answer is C

upvoted 1 times
-   **Emma412** 10 months ago

The Answer is B

the keyword here is "Different levels of expertise" which directly resonates with google definitions for its AI products "with limited machine learning expertise" @ <https://cloud.google.com/automl>

"with minimal effort and machine learning expertise" @ <https://cloud.google.com/products/ai>

upvoted 1 times
-   **tbolick6** 1 year, 3 months ago

Selected Answer: B

I chose B, but I think you could argue C, as well.

upvoted 1 times
-   **AMMARABO** 1 year, 8 months ago

Selected Answer: B

B. With APIs and tools for users at different levels of expertise

upvoted 2 times
-   **TonyMel** 1 year, 8 months ago

My answer is B

upvoted 1 times

An online retail organization wants to optimize their service.
What is an example of unstructured data that they can use to make decisions?

A. Customer survey comments

B. Seller location coordinates

C. Product sales trends

D. Warehouse inventory records

Correct Answer: C

Community vote distribution

A (100%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **reachmalhotras** 11 months, 2 weeks ago

Selected Answer: A

A-> Customer Survey Comments makes as text based unstructured data. Numerical Trends are structured data.
upvoted 2 times

  **minmin2020** 1 year, 8 months ago

Selected Answer: A

A. Customer survey comments - unstructured data
upvoted 4 times

  **Haplo46** 1 year, 8 months ago

Structured data:
Basically, they are those that are ordered. For example in:

An Excel table.
Spreadsheets.
Databases of any other type.
Applications to carry out test-type questionnaires.
Web forms.
Standardized customer files.
Surveys of users of a service.
You will find them organized by a series of well-defined rows and columns.
A is wrong answer
upvoted 2 times

  **1209apl** 1 year, 9 months ago

Selected Answer: A

A -> Customer Survey comments.
upvoted 2 times

  **saoirse98** 1 year, 9 months ago

Selected Answer: A

Customer Survey comments
upvoted 1 times

  **zellck** 1 year, 9 months ago

Selected Answer: A

A is the answer.
upvoted 2 times

  **Bill1000** 1 year, 9 months ago

A. Customer survey comments

upvoted 1 times

  **Abhi_ma_cdl_22** 1 year, 9 months ago

A. Customer survey comments

upvoted 2 times

  **Vin1975** 1 year, 9 months ago

A. Customer survey comments

upvoted 2 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: A

A. Customer survey comments

C is structured because sales trends are structured numerical values.

upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: A

Both product sales trends and survey comments can be used. But question specifically says which unstructured data, So answer should be A - Survey Comments.

upvoted 1 times

  **lax797** 1 year, 9 months ago

Survey comments are structured. Answer should be C.

upvoted 1 times

  **_jade** 1 year, 9 months ago

Comments are text - they are unstructured

upvoted 2 times

  **gcp_prep** 8 months, 2 weeks ago

thats incorrect.

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A. Customer survey comments

upvoted 4 times

An organization's public cloud provider failed to meet their SLA of 99.99% availability.
What is the potential impact on the organization?

- A. The organization risks using up their error budget.
- B. Renegotiation of the SLA to put less emphasis on uptime could be necessary.
- C. Unexpected downtime could risk the loss of customers.
- D. All data stored in their database could be unexpectedly lost.

Correct Answer: C

Community vote distribution

C (100%)

- BryRob**

Highly Voted

1 year, 7 months ago

Selected Answer: C

C. Unexpected downtime could risk the loss of customers.

A service-level agreement (SLA) sets the expectations between the service provider and the customer and describes the products or services to be delivered, the single point of contact for end-user problems, and the metrics by which the effectiveness of the process is monitored and approved

upvoted 5 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
- AMMARABO**

1 year, 8 months ago

Selected Answer: C

c: as IT is SLA not OLA

upvoted 1 times
- zellick**

1 year, 9 months ago

Selected Answer: C

C is the answer.

upvoted 1 times
- Abhi_ma_cdl_22**

1 year, 9 months ago

C. Unexpected downtime

upvoted 2 times
- Vin1975**

1 year, 9 months ago

C. Unexpected downtime could risk the loss of customers.

upvoted 2 times

An organization needs to train a machine learning model to categorize photographs.
Which Google Cloud product or service should the organization use?

- A. Video Intelligence API
- B. AutoML Vision
- C. BigQuery ML
- D. Looker

Correct Answer: B

Community vote distribution

B (100%)

- Matiz

Highly Voted

1 year, 6 months ago

Selected Answer B.AutoML vision.
For the same task, if I have to use a PRE TRAINED model then I will go for Vision API
upvoted 7 times
- BryRob

Highly Voted

1 year, 7 months ago

Selected Answer: B

A. Video Intelligence API - obviously video
B. AutoML Vision - enables you to train machine learning models to classify your images according to your own defined labels
C. BigQuery ML - this is using database
D. Looker - easy to build insight-powered workflows and applications

<https://cloud.google.com/vision/automl/docs>
upvoted 7 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: B

B. AutoML Vision

AutoML Vision is a Google Cloud product designed for training machine learning models on image classification tasks. It allows organizations to build custom image recognition models without extensive machine learning expertise. By providing labeled training data, users can use AutoML Vision to train a model that can categorize photographs based on specific criteria.
upvoted 1 times
- juangt

1 year, 8 months ago

Selected Answer: B

AutoML Vision enables you to train machine learning models to classify your images according to your own defined labels
upvoted 1 times

A multinational organization has decided to use public cloud services to modernize their IT infrastructure.
What is one benefit of doing this?

- A. They can perform hardware maintenance outside of normal business hours.
- B. They can expect 100% service availability in all regions.
- C. Built-in security is no longer required during data migrations.
- D. Access to data and computer resources is not tied to a geographical location.

Correct Answer: D

Community vote distribution

D (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D is correct.
upvoted 1 times
-   **Mahavijay** 1 year, 6 months ago

I think no one can expect 100% availability in all regions, it can be 99.999 but not 100
upvoted 3 times
-   **TRMFudge** 1 year, 7 months ago

Selected Answer: D

I agree with D, but why not B, this also seems as a second correct answer.
upvoted 2 times
-   **moi23** 1 year, 4 months ago

There is no such thing as a 100% availability.
upvoted 3 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: D

D for me
upvoted 1 times
-   **shuvs** 1 year, 7 months ago

Selected Answer: D

D is most likely
upvoted 1 times
-   **AMMARABO** 1 year, 8 months ago

Selected Answer: D

D. Access to data and computer resources is not tied to a geographical location.
upvoted 2 times

An organization meets their service level objective (SLO) of 99.999% (`five nines`).
How much downtime do their end users experience per year?

- A. 5 minutes
- B. 500 minutes
- C. 5 hours
- D. 5 days

Correct Answer: A

Community vote distribution

A (100%)

- Vin1975

Highly Voted

1 year, 9 months ago

A ~ $((100 - \text{SLO})/100) \times 365 \times 24 \times 60$
upvoted 10 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
- MNotABot

11 months, 3 weeks ago

0.001 % of total $60 \times 24 \times 365$ minutes in a year
upvoted 2 times
- BryRob

1 year, 7 months ago

Selected Answer: A

A. 5 minutes
upvoted 2 times
- AMMARABO

1 year, 8 months ago

Selected Answer: A

a is correct
upvoted 1 times
- zellick

1 year, 9 months ago

Selected Answer: A

A is my answer.
upvoted 1 times

An organization wants to introduce a new image recognition login system.
What should the organization do to follow SRE principles?

- A. Roll out the new system to a subset of employees to test it out.
- B. Roll out the new system to all employees to collect as much data as possible.
- C. Avoid rolling out the new system because it may have security flaws.
- D. Avoid rolling out the new system because it may violate privacy policy.

Correct Answer: A

Community vote distribution

A (100%)

-   **cloudmind** 5 months ago

A is correct. When implementing gradual changes, which is the third DevOps objective' SREs aim to reduce the cost of failure by rolling out changes to a small percentage of users first. Culturally, this promotes more prototyping and launching iteratively.

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
-   **Panpn** 1 year, 7 months ago

Selected Answer: A

A - testing on a subset of users is an SRE principle

upvoted 3 times
-   **Tecatero2001** 1 year, 8 months ago

Selected Answer: A

A. Roll out the new system to a subset of employees to test it out.

upvoted 3 times
-   **grzeev** 1 year, 8 months ago

A. Roll out the new system to a subset of employees to test it out.

upvoted 4 times

What does Cloud Debugger help an organization do?

- A. Implement code updates in real time without affecting the service level objective (SLO).
- B. Inspect source code in real time without affecting user downtime.
- C. Manage code and accelerate application development.
- D. Analyze live source code during user downtime.

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

Cloud Debugger is a feature provided by some cloud platforms to allow developers to inspect the state of a running application in real-time without affecting its performance or users. It provides tools for capturing the state of an application at any code location in real time, without stopping or slowing it down.

upvoted 1 times
-   **BryRob** 1 year, 7 months ago

B. Inspect source code in real time without affecting user downtime.

<https://cloud.google.com/debugger/docs>

upvoted 2 times
-   **najmdin** 1 year, 8 months ago

Selected Answer: B

answer B

upvoted 1 times
-   **Pikapika12638** 1 year, 9 months ago

Selected Answer: B

Cloud Debugger is a feature of Google Cloud Platform that lets you inspect the state of an application, at any code location, without stopping or slowing down the running app. Cloud Debugger makes it easier to view the application state without adding logging statements.

upvoted 2 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: B

You can use Cloud Debugger with any deployment of your application, including test, development, and production. The debugger adds less than 10ms to the request latency only when the application state is captured. In most cases, this is not noticeable by users.

upvoted 2 times

How can a streaming service meet global compliance requirements using the cloud?

- A. By automatically encrypting personally identifiable information
- B. By obtaining a business license to operate in a new market
- C. By allowing users to disable two-factor authentication
- D. By adhering only to data polices of the country in which the head office is registered

Correct Answer: B

Community vote distribution

A (86%)

9%

- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A. By automatically encrypting personally identifiable information.

upvoted 1 times
- Emma412

10 months ago

i initially thought B but after research I concur that the answer is A

Explanation:

Refer - <https://cloud.google.com/blog/products/networking/benefits-of-building-media-services-on-google-cloud>.

This is not referring to default data encryption stored in Google cloud storage. This refers to "enhanced end-to-end encryption using Digital Rights Management (DRM)"

upvoted 2 times
- Iz0001

1 year, 2 months ago

Isn't all data in the cloud encrypted at rest - "All data that is stored by Google is encrypted at the storage layer using the Advanced Encryption Standard (AES) algorithm, AES-25". So it can't be A, because this already happens right?

upvoted 1 times
- serenity8468

11 months, 1 week ago

The question doesn't state the streaming service is using Google cloud.

upvoted 2 times
- Dimi_Mous

6 months ago

yup. sneaky!

upvoted 1 times
- Arimaverick

1 year, 3 months ago

A is the answer. D says about the headquarter but the importance of data should be from where its collected.

upvoted 2 times
- tbolick6

1 year, 3 months ago

Selected Answer: A

A gets my vote.

upvoted 1 times
- AK10000

1 year, 4 months ago

Selected Answer: A

Think youtube.

upvoted 1 times
- Danielki

1 year, 5 months ago

Selected Answer: C

Hey people the question it self is wrong. This should be Which of the following is NOT a way that a streaming service can meet global compliance requirements using the cloud?" And the answer should be C because least of 3options need streaming service.

upvoted 2 times
- theKingslayer

1 year, 6 months ago

Selected Answer: A

business licences are granted by the local governments and not by the cloud provider.
upvoted 1 times

  **Pankaj1978** 1 year, 6 months ago
Option B: Encrypting PII complies with security compliances, but to go global, one needs to have a licence obtained to operate in the new market for streaming dataflow.
upvoted 1 times

  **Shaik_Kousar** 1 year, 6 months ago
Option A
upvoted 1 times

  **ed_per** 1 year, 7 months ago

Selected Answer: A

D doesn't make sense. It needs to comply in all countries it's operating in, not just where the head office is.

upvoted 2 times

  **shuvs** 1 year, 7 months ago

Selected Answer: A

A. Others don't seem valid. Certainly not D as the service needs to comply with local regulations too.

upvoted 1 times

  **md3danish** 1 year, 7 months ago
D nbmbnmjbnmj
upvoted 1 times

  **minmin2020** 1 year, 8 months ago

Selected Answer: A

A - PII Encryption

upvoted 2 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: D

I think it is D

upvoted 1 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: A

A. By automatically encrypting personally identifiable information

upvoted 3 times

An organization wants full control of their virtual machine infrastructure for a custom home-grown application with a product that autoscales and automatically updates.

Which Google Cloud product or solution should the organization use?

- A. Cloud Build
- B. Cloud Run
- C. Compute Engine
- D. App Engine

Correct Answer: C

Community vote distribution

C (93%)

4%

  **BryRob** Highly Voted 1 year, 7 months ago

Selected Answer: C

A. Cloud Build - serverless
B. Cloud Run - container
C. Compute Engine - VM/instance
D. App Engine - platform as a service

upvoted 11 times

  **jaxclain** Highly Voted 1 year, 9 months ago

Selected Answer: C

Read the question properly, "Wants full control of their VM infrastructure for their custom app", this answer should be seriously easy, how can you say B or any other option, the only option here is C, Compute Engine will allow you to have full control of their VM infrastructure and you can autoscale and also apply automatic updates.

upvoted 6 times

  **kapomony** Most Recent 5 months, 4 weeks ago

Selected Answer: C

No doubt

upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: C

C is correct as the organization wants full control of their virtual machine infrastructure.

upvoted 1 times

  **NoCrapEva** 10 months ago

TBH - It could be both C or D - the question is badly worded...

C is right if the first part of the question is read "An organization wants full control of their virtual machine infrastructure"

However, D is also right for the latter part of the question "...for a custom home-grown application with a product that autoscales and automatically updates" as App engine does autoscalling and updating, whereas Compute engine (C) does NOT BY DEFAULT.

upvoted 2 times

  **Oksa** 1 year, 5 months ago

Answer is C

A. Cloud Build - containers
B. Cloud Run - serverless
C. *Compute Engine - (VM) instances *
D. App Engine - serverless

upvoted 2 times

  **shuvs** 1 year, 7 months ago

Selected Answer: C

GCE, as all other options don't allow full control

upvoted 1 times

  **grzeev** 1 year, 8 months ago

Selected Answer: D

they start with control vm infra, but end with app that needs to scale
upvoted 1 times

  **minmin2020** 1 year, 8 months ago

Selected Answer: C

C. Compute Engine is the only option that makes sense
upvoted 1 times

  **zellick** 1 year, 9 months ago

Selected Answer: C

C is the answer as compute engine will allow full control of the VM infra.
upvoted 1 times

  **all_cloud** 1 year, 9 months ago

Selected Answer: B

Answer is cloud run, which autoscales
upvoted 1 times

  **moi23** 1 year, 4 months ago

It's C, you need full control to satisfy the "homegrown" keyword
upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C. Compute Engine is the answer
upvoted 2 times

An organization wants to build an entirely new infrastructure and applications in the cloud.
Which application modernization approach should the organization use?

- A. Move the application to the cloud, and then change it.
- B. Change their application, and then move it to the cloud.
- C. Invent in greenfield.
- D. Invent in brownfield.

Correct Answer: C

Community vote distribution

C (100%)

-   **Sozan**

Highly Voted 

 1 year, 2 months ago

Selected Answer: C

In the context of business and innovation, "greenfield" and "brownfield" refer to two different approaches to innovation.

Greenfield innovation refers to creating something completely new, starting from scratch, and developing something that has never existed before. This approach is often associated with startups or brand new ventures.

Brownfield innovation, on the other hand, refers to the process of developing new products, services, or processes within an existing company or industry. This approach involves building upon existing infrastructure, technology, and resources to create something new and innovative.

Therefore, inventing in greenfield means developing a new product, service, or process in an entirely new industry or market. On the other hand, inventing in brownfield means creating something new within an existing industry or market, often by improving or modifying existing products, services, or processes.

upvoted 9 times
-   **jaxclain**

Highly Voted 

 1 year, 9 months ago

Selected Answer: C

For sure this one is C (Invent in greenfield), because there is another similar question but they added "they want to build a new application to replace the existing one" so that one was brownfield.

upvoted 7 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C. Invent in greenfield.

A "greenfield" project refers to building on a blank slate without legacy constraints. This approach is most suitable for organizations looking to build entirely new applications and infrastructure in the cloud.

upvoted 1 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: C

C

key word: build an entirely new infra / app

<https://www.capgemini.com/se-en/2020/05/confusion-in-the-cloud-greenfield-or-brownfield/>

upvoted 1 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C. A Greenfield approach is a brand-new implementation , where companies then add their needed configurations and customizations. This approach provides a clean slate to start from, does not carry over needless customizations and technical debt, and provides a solid foundation for business process re-engineering.

upvoted 2 times
-   **SomeshSule** 1 year, 9 months ago

But why "Invent"?

upvoted 1 times

  **Ej24356** 1 year, 9 months ago

Selected Answer: C

A greenfield deployment is the design, installation and configuration of computer infrastructure where none existed before, for example, in a new office. In contrast, a brownfield deployment is an upgrade or addition to existing infrastructure using legacy components.

upvoted 2 times

An organization wants to upskill their IT staff.
How can they do this in a transformational way?

- A. Prioritize training current employees instead of hiring new recruits with cloud experience.
- B. Prioritize giving privileged access to third-party partners and contractors to fill IT knowledge gaps.
- C. Create a culture of self-motivated, isolated learning with official training materials.
- D. Create a culture of continuous peer-to-peer learning with official training materials.

Correct Answer: D

Community vote distribution

D (72%)

A (28%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: D

The best option here is D.

upvoted 1 times
- tbolick6

1 year, 3 months ago

Selected Answer: D

D makes sense to me

upvoted 1 times
- dabizssan

1 year, 7 months ago

Selected Answer: D

It is part of the DevOps culture.

upvoted 3 times
- BryRob

1 year, 7 months ago

Selected Answer: D

D for me

<https://emeritus.org/blog/how-to-upskill-employees/>

upvoted 2 times
- Tecatero2001

1 year, 8 months ago

Selected Answer: D

An organization wants to upskill their IT staff = Training not new employees, so the correct answer is D. Create a culture of continuous peer-to-peer learning with official training materials.

upvoted 1 times
- Rajacpani

1 year, 8 months ago

D Peer learning and sharing is the transformational method in Learning theme

upvoted 1 times
- minmin2020

1 year, 8 months ago

Selected Answer: A

A. Prioritize training current employees instead of hiring new recruits with cloud experience. This is the only transformational option.

upvoted 1 times
- AMMARABO

1 year, 8 months ago

B Learn: The quality and scale of the learning programs you have in place to upskill your technical teams, and your ability to augment your IT staff with experienced partners. Who is engaged? How widespread is that engagement? How concerted is the effort? How effective are the results?

upvoted 1 times
- ace_abhishek

1 year, 9 months ago

Selected Answer: A

no comm

upvoted 1 times

  **jaxclain** 1 year, 9 months ago

Selected Answer: D

I also believe is D, if you did watch the videos in the guide, they mentioned the best way to achieve this is by changing its culture so the correct answer is D.

upvoted 4 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: A

Why not A . Prioritize training current employees instead of hiring new recruits with cloud experience.

upvoted 3 times

  **moi23** 1 year, 4 months ago

You always change the game and not the player to win. In this case, culture will always be the soundest answer.

upvoted 2 times

  **Gigante** 1 year, 9 months ago

In my opinion, Its because you need "a transformational way" and it is done when you change the organization's culture.

upvoted 4 times

Several departments in an organization are working together on a project. The organization wants to customize access to resources for each department.

What is the quickest and most efficient way to achieve this?

- A. By mapping IAM roles to job functions for each department
- B. By assigning IAM primitive roles to each employee
- C. By applying “least-privilege” to roles for each employee
- D. By creating a single shared service account for all departments

Correct Answer: C

Community vote distribution

A (88%)

12%

- jaxclain

Highly Voted

1 year, 9 months ago

Selected Answer: A

I would go for A, it says "What is the quickest and most efficient way to achieve this", obviously doing option C (applying least priviledge roles for each employee) would take a long time, the most efficient way is by separating the permissions by Jobs / Departments so this s done by creating Groups then apply their corresponding IAM permissions to each Group (Job / Department).

upvoted 10 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__

8 months ago

A is correct.

upvoted 1 times
- Anuska1099

12 months ago

A would be the answer

upvoted 1 times
- someshsehgal

1 year, 6 months ago

Selected Answer: A

create groups with access permissions.

upvoted 2 times
- Akshay0403

1 year, 7 months ago

Keyword is "quickest and most efficient".

Option B and C suggesting to create access for all employees which is not quickest way.

Also, Option D also suggesting to create service account for all departments which is not quickest way.

So, answer will be A. By mapping IAM roles to job functions for each department.

upvoted 3 times
- berberar

1 year, 6 months ago

Another key word is “customize” which in this case would only be achieved by mapping IAM roles to job functions.

upvoted 2 times
- jellybiscuit

1 year, 8 months ago

Selected Answer: A

A. By mapping IAM roles to job functions for each departmen

upvoted 3 times
- najmdin

1 year, 8 months ago

Selected Answer: C

C is the answer

upvoted 2 times
- minmin2020

1 year, 8 months ago

Selected Answer: A

A. By mapping IAM roles to job functions for each department

upvoted 1 times

🗨️ 👤 **zelck** 1 year, 9 months ago

Selected Answer: A

A is the answer.

C is not google best practices as you should apply IAM to group, not individual.

upvoted 3 times

🗨️ 👤 **TOLETIHARIPRIYASRUJANA** 1 year, 9 months ago

OPTION IS C --> Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

upvoted 2 times

🗨️ 👤 **MurugeshBabu** 1 year, 9 months ago

Inorder to customize, Google recommends “least privilege” model

upvoted 2 times

🗨️ 👤 **zelck** 1 year, 9 months ago

Selected Answer: A

it is always more efficient to map roles to user groups rather than to each users

upvoted 1 times

🗨️ 👤 **Govindaraj** 1 year, 9 months ago

Selected Answer: C

I think it is C. Custom roles help you enforce the principle of least privilege, because they help to ensure that the principals in your organization have only the permissions that they need.

<https://cloud.google.com/iam/docs/understanding-custom-roles>

upvoted 1 times

🗨️ 👤 **SomeshSule** 1 year, 9 months ago

A is correct

upvoted 1 times

🗨️ 👤 **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: A

A. By mapping IAM roles to job functions for each department

upvoted 1 times

🗨️ 👤 **Pichikik** 1 year, 9 months ago

my vote is A as well

upvoted 2 times

An organization notices that some of their cloud expenditures are too high.
What should the organization do to control costs?

- A. Streamline the hardware procurement process to reduce costs.
- B. Share cost views with the departments to establish more accountability.
- C. Change the cost model from operational expenditure to capital expenditure.
- D. Ensure that all cloud resources are tagged with a single tag.

Correct Answer: B

Community vote distribution

B (83%)

D (17%)

- Govindaraj**

Highly Voted

1 year, 9 months ago

Selected Answer: B

B. Share cost views with the departments to establish more accountability.
upvoted 7 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: B

B is correct.
upvoted 1 times
- tbolick6**

1 year, 3 months ago

Selected Answer: B

B - promote accountability.
upvoted 1 times
- tbolick6**

1 year, 3 months ago

Selected Answer: B

B for me
upvoted 1 times
- Akshay0403**

1 year, 7 months ago

Selected Answer: B

A foundational principle of FinOps is: "everyone takes ownership for their cloud usage." The true key to understanding total cost of ownership is built upon transparency and accuracy.
upvoted 4 times
- BryRob**

1 year, 7 months ago

Selected Answer: D

For Me this is D
By tagging the resources, its a good idea to identify / know whose responsible for that service and tagging is free

<https://cloud.google.com/resource-manager/docs/tags/tags-overview>
<https://aws.amazon.com/blogs/apn/how-better-tagging-can-help-organizations-optimize-expenses-and-improve-roi/>
upvoted 3 times
- dabizssan**

1 year, 7 months ago

Various tags are used to catalog resources for different purposes. Therefore, the single tag rules out this answer.
upvoted 4 times

What is monitoring within the context of cloud operations?

- A. Observing cloud expenditure in real time to ensure that budgets are not exceeded
- B. Collecting predefined and custom metrics from applications and infrastructure
- C. Tracking user activities to guarantee compliance with privacy regulations
- D. Tracing user location to document regional access and utilization

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- Akshay0403** 1 year, 7 months ago

Selected Answer: B

Cloud monitoring is a method of reviewing, observing, and managing the operational workflow in a cloud-based IT infrastructure. Manual or automated management techniques confirm the availability and performance of websites, servers, applications, and other cloud infrastructure.

upvoted 2 times
- minmin2020** 1 year, 8 months ago

Selected Answer: B

B. Collecting predefined and custom metrics from applications and infrastructure

upvoted 2 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

B is the correct answer

upvoted 1 times
- Vin1975** 1 year, 9 months ago

B. Collecting predefined and custom metrics from applications and infrastructure.

upvoted 1 times
- jexmtropicscheatchatya** 1 year, 9 months ago

Answer is A “within context of cloud” means monitoring costs in real time

upvoted 1 times
- Ej24356** 1 year, 9 months ago

Cloud Monitoring collects metrics, events, and metadata from Google Cloud, Amazon Web Services (AWS), hosted uptime probes, and application instrumentation. Cloud Billing would monitor budget costs in real time

upvoted 4 times
- jexmtropicscheatchatya** 1 year, 9 months ago

Ya, I think you are correct. "context of cloud OPERATIONS". If it said context of cloud spending or expenditures or something then it would be A.

upvoted 2 times

An organization wants to create a new application in the cloud to replace an existing on-premises application. Which application modernization approach should the organization use?

- A. Move the application to the cloud, and then change it.
- B. Change their application, and then move it to the cloud.
- C. Invent in greenfield.
- D. Invent in brownfield.

Correct Answer: C

Community vote distribution

D (65%)C (35%)

- Sozan**

Highly Voted

1 year, 2 months ago

Selected Answer: D

In the context of business and innovation, "greenfield" and "brownfield" refer to two different approaches to innovation.

Greenfield innovation refers to creating something completely new, starting from scratch, and developing something that has never existed before. This approach is often associated with startups or brand new ventures.

Brownfield innovation, on the other hand, refers to the process of developing new products, services, or processes within an existing company or industry. This approach involves building upon existing infrastructure, technology, and resources to create something new and innovative.

Therefore, inventing in greenfield means developing a new product, service, or process in an entirely new industry or market. On the other hand, inventing in brownfield means creating something new within an existing industry or market, often by improving or modifying existing products, services, or processes.

upvoted 9 times
- ilberg**

Highly Voted

1 year, 6 months ago

Selected Answer: C

key word is new -therefore answer is C

upvoted 5 times
- moi23**

1 year, 4 months ago

No, don't let that mislead you. Keywords are "move" and "existing".

upvoted 2 times
- moi23**

1 year, 4 months ago

I meant "replace" and "existing".

upvoted 2 times
- shanwford**

Most Recent

3 weeks, 4 days ago

Selected Answer: C

"create new application in Cloud", no hint that existing on-prem infra is used for this, so greenfield for sure

upvoted 1 times
- chai_gpt**

7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: C

C is Correct. Invent in greenfield.

A "greenfield" project refers to developing on a blank slate without legacy constraints. This is the best approach when wanting to build an entirely new application from scratch.

A "brownfield" project involves working with or around existing systems. This means there are legacy constraints, and it doesn't fit the scenario of building a completely new application to replace the old one.

upvoted 2 times
- Emma412**

10 months ago

the Answer is D

Explanation:

Contrast the way the question is framed in a different question - <https://www.examttopics.com/discussions/google/view/80318-exam-cloud-digital>

leader-topic-1-question-131-discussion/
2 keywords

- though create new application is misleading it further goes to say "in the cloud" - So it is new in the cloud but then further clarifies that with "to replace an existing application".

use this info in light of the definition for Greenfield and Brownfield development given under cloud migration strategies @
<https://successive.cloud/greenfield-vs-brownfield-development-which-strategy-is-right-for-you/>

upvoted 1 times

  **[Removed]** 1 year, 3 months ago

Selected Answer: C

This ones nasty and worded in a way to make it tricky. I'm going with C. Nothing in the phrasing suggests this will be a "parallel" run, otherwise I would have gone with D.

upvoted 2 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: D

Brownfield for sure

upvoted 1 times

  **Matro71** 1 year, 3 months ago

Selected Answer: D

Invent in brownfield - A brownfield strategy is to invent a new application in the cloud environment that will replace an existing legacy application that remains on-premises. The legacy application is only retired after the new application is built.

Invent in greenfield - Inventing in greenfield allows you to build innovative applications that will help drive the business forward

upvoted 2 times

  **Balliache520505** 1 year, 4 months ago

I go with option D. More information about the difference between greenfield and brownfield strategies is in the following link:

<https://www.capgemini.com/se-en/insights/expert-perspectives/confusion-in-the-cloud-greenfield-or-brownfield/#>

upvoted 2 times

  **gviolant** 1 year, 5 months ago

Selected Answer: D

absolutely D

upvoted 1 times

  **131983** 1 year, 5 months ago

Selected Answer: D

since they refer to an existing app on premises the answer would be D - Brownfield

upvoted 3 times

  **someshseghal** 1 year, 6 months ago

Selected Answer: C

<https://synoptek.com/insights/it-blogs/greenfield-vs-brownfield-software-development/>

Brownfield software development refers to the development and deployment of a new software system in the presence of existing or legacy software systems. Brownfield development usually happens when you want to develop or improve upon an existing application, and compels you to work with previously created code.

upvoted 1 times

  **William_the_Great** 1 year, 6 months ago

Selected Answer: C

With a greenfield strategy, a whole new infrastructure is built in the Cloud using a public, private, or hybrid model. Then data, operating systems, and applications are transferred to this new infrastructure in their entirety.

A brownfield strategy, on the other hand, only involves moving some functions from one system to the other. Often, proprietary systems that are difficult to replicate or those that hold particularly sensitive information will be kept on-prem during a brownfield migration.

upvoted 2 times

  **ShriRama** 1 year, 6 months ago

Selected Answer: D

Invent in brownfield: A brownfield strategy is to invent a new application in the cloud environment that will replace an existing legacy application that remains on-premises. The legacy application is only retired after the new application is built.

upvoted 1 times

  **khedron** 1 year, 7 months ago

Selected Answer: C

It says new application in the cloud to replace on-prem. To me its clearly a greenfield approach.

upvoted 3 times

  **ed_per** 1 year, 7 months ago

Selected Answer: C

They want to build a new application in a new set of infrastructure, (moving from on-prem to the cloud), hence it's greenfield - C.

upvoted 2 times

Why should an organization consider the total cost of ownership (TCO) when moving from on-premises to the cloud?

- A. To evaluate error budget
- B. To understand service level availability
- C. To evaluate return on investment
- D. To calculate required compute power

Correct Answer: C

Community vote distribution

C (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
- Akshay0403** 1 year, 7 months ago

Selected Answer: C

TCO demonstrates the real, complete cost of a given solution over its lifetime – usually measure in 3-year or 5-year TCO. Management consider TCO when deciding to build or buy, when deciding if a solution is worth the investment, or when deciding whether to host an application on-premises or in the cloud.

upvoted 4 times
- BryRob** 1 year, 7 months ago

Selected Answer: C

<https://www.cloudzero.com/blog/cloud-tco>
Total cost of ownership (TCO) is the sum of all costs involved in the purchase, operation, and maintenance of a given asset during its lifetime. TCO helps businesses understand the cost of a tool beyond the initial purchase price and is extremely helpful for understanding ROI.

upvoted 2 times
- shuvs** 1 year, 7 months ago

Selected Answer: C

TCO helps establish ROI in a shared resource env

upvoted 1 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: C

C. To evaluate return on investment

upvoted 2 times

How does Google Cloud ensure that customer data remains secure and private when at rest?

- A. By aggregating training data for customers within each industry
- B. By automatically locking files containing suspicious code
- C. By auditing platform privacy practices against industry standards
- D. By providing privacy reviews for critical customer applications

Correct Answer: B

Community vote distribution

C (91%)9%

-   **MurugeshBabu**

Highly Voted 

 1 year, 9 months ago

Answer C
Google Cloud commitment to keep the data secure and private:
1. Org owns the data and not Google
2. Google does not sell data to 3rd parties
3. All customer data is encrypted by default
4. Google Cloud guards insider against your data
5. No backdoor access to any govt. entity
6. Google’s privacy practices are audited against international standards
upvoted 16 times
-   **Pankaj1978** 1 year, 6 months ago

Point number 4 suggests the Option B.
upvoted 2 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

One of the ways GCP demonstrates its commitment to security and privacy is by undergoing independent third-party audits of its security, data protection, and privacy controls. This means that GCP's platform privacy practices are audited against industry standards such as ISO/IEC 27001, ISO/IEC 27017, ISO/IEC 27018, and others. These audits validate that Google Cloud has robust security and privacy practices in place, ensuring the protection of customer data at rest, in transit, and during processing.
upvoted 1 times

  **dru371993** 1 year ago

<https://cloud.google.com/docs/security/encryption/default-encryption>
Would say B is the correct answer
upvoted 1 times

  **Jackey0117** 1 year, 4 months ago

Selected Answer: D

The best answer would be D, "By providing privacy reviews for critical customer applications". While auditing platform privacy practices against industry standards (C) is also an important step in ensuring the security of customer data, providing privacy reviews for critical customer applications is a proactive approach that helps to build privacy and security into the applications themselves. 'D' can help to reduce the risk of data breaches or other security incidents and provide customers with additional peace of mind.
upvoted 1 times

  **someshseghal** 1 year, 6 months ago

Selected Answer: C

There is no such option as B.
upvoted 2 times

  **Pankaj1978** 1 year, 6 months ago

Option B suits to the nearest logic. In option C, the word 'against' spoils the scenario, a bit tricky.
upvoted 1 times

  **SoftSami** 1 year, 8 months ago

Selected Answer: C

Part of Google's commitment

upvoted 1 times

  **Cert_IT** 1 year, 8 months ago

Selected Answer: C

C. By auditing platform privacy practices against industry standards
upvoted 1 times

  **minmin2020** 1 year, 8 months ago

Selected Answer: C

C. By auditing platform privacy practices against industry standards
upvoted 1 times

  **AMMARABO** 1 year, 8 months ago

Selected Answer: C

C. By auditing platform privacy practices against industry standards
upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

C. By auditing platform privacy practices against industry standards
upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: C

I think C should be the answer
upvoted 1 times

  **Vin1975** 1 year, 9 months ago

C. By auditing platform privacy practices against industry standards
upvoted 3 times

An organization's developers are growing increasingly frustrated by the limitations of their on-premises infrastructure. How would they benefit from leveraging cloud technology?

- A. They can expect 100% service availability.
- B. They can avoid the limitations of serverless computing.
- C. They can have new tools to innovate and optimize resource usage.
- D. They can optimize maintenance for their on-premises infrastructure.

Correct Answer: C

Community vote distribution

C (92%)

8%

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 2 times
- __rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
- minmin2020** 1 year, 8 months ago

Selected Answer: C

C. They can have new tools to innovate and optimize resource usage.

upvoted 3 times
- zellick** 1 year, 9 months ago

Selected Answer: C

C is the answer.

upvoted 1 times
- jaxclain** 1 year, 9 months ago

Selected Answer: C

Correct answer is C all the way, first the option A is wrong (you cannot have 100% service availability) you need to watch the video guides from Google and it mentions that one of the benefits is exactly that devs will have new tools to innovate more so is the only answer that makes sense here, B and D are not benefits at all.

upvoted 2 times
- Pikapika12638** 1 year, 9 months ago

Selected Answer: C

A. They can expect 100% service availability. - incorrect since it's impossible to achieve true 100% availability.
C. They can have new tools to innovate and optimize resource usage. - correct. Google cloud have vast majority of products/tools that you can use to innovate. Additionally, there are products in google that scale automatically based from usage (Ex. App Engine, Cloud Run, etc.)

upvoted 2 times
- soumik4u** 1 year, 9 months ago

Option A should not be the answer rather there is a probability of option B

upvoted 1 times
- Skghost** 1 year, 9 months ago

Selected Answer: A

Service Availability is 100%

upvoted 1 times
- MurugeshBabu** 1 year, 9 months ago

Any application cannot have 100% availability

upvoted 1 times
- moi23** 1 year, 4 months ago

No! This is basic information.

upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C. They can have new tools to innovate and optimize resource usage.
upvoted 1 times

Question #140

Topic 1

How is service availability measured in the context of cloud technology?

- A. Number of available regions
- B. Percentage of uptime
- C. Speed of response time
- D. Number of downtime incidents

Correct Answer: B

Community vote distribution

B (100%)

  **Akshay0403** **Highly Voted**  1 year, 7 months ago

Selected Answer: B

Service availability is simply the measure of the service being available and accessible to the customers during the time you promised to keep the service available. It's usually calculated as a percentage.
upvoted 5 times

  **chai_gpt** **Most Recent**  7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 2 times

  **__rajan__** 8 months ago

Selected Answer: B

B is correct.
upvoted 2 times

  **Pikapika12638** 1 year, 9 months ago

Selected Answer: B

B. Percentage of uptime
upvoted 2 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: B

B is the answer
upvoted 1 times

  **Vin1975** 1 year, 9 months ago

B. Percentage of uptime
upvoted 1 times

An organization wants a cost-effective relational database.
Which Google Cloud service should the organization use?

- A. Cloud Storage
- B. BigQuery
- C. Cloud SQL
- D. Dataflow

Correct Answer: C

Community vote distribution

C (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C is Correct.

upvoted 1 times
- serenity8468** 11 months, 1 week ago

Cloud Storage - Not DB
Big Query - Dataware House
Dataflow - data processing service
Cloud SQL is the only DB!

Answer C

upvoted 2 times
- tsk9921** 11 months, 2 weeks ago

C Cloud SQL

upvoted 1 times
- Akshay0403** 1 year, 7 months ago

Selected Answer: C

Cloud SQL: Provides managed MySQL, PostgreSQL and SQL Server databases on Google Cloud. It reduces maintenance cost and automates database provisioning, storage capacity management, back ups, and out-of-the-box high availability and disaster recovery/failover.

upvoted 3 times

Pikapika12638 1 year, 9 months ago

Selected Answer: C

C. Cloud SQL

upvoted 1 times

Govindaraj 1 year, 9 months ago

Selected Answer: C

C. Cloud SQL

upvoted 1 times

An organization is altering their gaming product so that it is compatible with cloud technology. What can they expect when moving from traditional technology to cloud technology?

- A. No change to existing responsibilities
- B. A shift toward OpEx
- C. A shift toward using structured data
- D. Increased hardware maintenance

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- Akshay0403** 1 year, 7 months ago

Selected Answer: B

Computing costs are predictable and relatively fixed in traditional IT environments. An organization purchases computing capacity upfront and uses it over time. The total cost of ownership is fairly easier to calculate with this setup. By contrast, cloud computing operates on a pay-as-you-go basis, with no upfront payments. Resources and services are available on-demand, and IT spend fluctuates based on consumption. The traditional approach prioritizes capital expenditure (CapEx), whereas cloud economics favors operating expenses (OpEx).

upvoted 3 times
- BryRob** 1 year, 7 months ago

Selected Answer: B

OpEx is short for operating expenses, also known as operational expenses, and operating costs. Operating expenses refer to the money a company spends to run day-to-day operations.

<https://www.cloudzero.com/blog/capex-vs-opex#:~:text=most%20popular%20guide!-,What%20Is%20OpEx%20In%20Cloud%20Computing%3F,a%20year%20of%20purchasing%20them.>

upvoted 2 times
- Pikapika12638** 1 year, 9 months ago

Selected Answer: B

B. A shift toward OpEx

upvoted 1 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

B. A shift toward OpEx

upvoted 1 times

When an organization adopts cloud technology, how does their total cost of ownership (TCO) shift?

- A. Away from cost management toward capital expenditure
- B. Away from operational expenditure toward cost management
- C. Away from capital expenditure toward operational expenditure
- D. Away from operational expenditure toward capital expenditure

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **tsk9921** 11 months, 2 weeks ago

C - Question similar to 142, but same answer

upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: C

C. Away from capital expenditure toward operational expenditure

upvoted 1 times
-   **Akshay0403** 1 year, 7 months ago

Selected Answer: C

Computing costs are predictable and relatively fixed in traditional IT environments. An organization purchases computing capacity upfront and uses it over time. The total cost of ownership is fairly easier to calculate with this setup.
By contrast, cloud computing operates on a pay-as-you-go basis, with no upfront payments. Resources and services are available on-demand, and IT spend fluctuates based on consumption.
The traditional approach prioritizes capital expenditure (CapEx), whereas cloud economics favors operating expenses (OpEx).

upvoted 2 times
-   **Pikapika12638** 1 year, 9 months ago

Selected Answer: C

C. Away from capital expenditure toward operational expenditure

upvoted 1 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C. Away from capital expenditure toward operational expenditure

upvoted 2 times

Which policy helps Google Cloud keep customer data private?

- A. Google tests the service availability of customer applications.
- B. Google does not use customer data for advertising purposes.
- C. Google migrates customer data to an offline server when a threat is detected.
- D. Google does not allow customers to change encryption keys.

Correct Answer: B

Community vote distribution

B (100%)

-  **BryRob** Highly Voted

1 year, 7 months ago

Selected Answer: B

<https://cloud.google.com/privacy>
2. We never use your data for ads targeting
We do not process your customer data to create ads profiles or improve Google Ads products.
upvoted 6 times
-  **chai_gpt** Most Recent

7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times
-  **__rajan__**

8 months ago

Selected Answer: B

B is correct.
upvoted 1 times
-  **shuvs**

1 year, 7 months ago

Selected Answer: B

Only option that is s a policy
upvoted 1 times
-  **The_least_perfect_guy**

1 year, 7 months ago

werwerwe

upvoted 1 times
-  **Govindaraj**

1 year, 9 months ago

Selected Answer: B

B. Google does not use customer data for advertising purposes.
upvoted 1 times

A cloud-native organization is not meeting their service level objective (SLO) but has not exhausted their error budget.
What should the organization prioritize?

- A. Innovation to improve user experience
- B. Hardware reliability to improve availability
- C. Stability to avoid prolonged user downtime
- D. Speed to release new features

Correct Answer: C

Community vote distribution

C (68%)	A (27%)	5%
---------	---------	----

- chesterlp** 6 months, 2 weeks ago

Selected Answer: C

To achieve SLO, a company needs its applications and services to be resilient. Definitely C

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C. Stability to avoid prolonged user downtime

Since the organization is not meeting their SLO, they need to prioritize stability to avoid prolonged user downtime. This means focusing on fixing the underlying issues that are causing the SLO violations, rather than innovating on new features or releasing new features quickly.

upvoted 1 times
- Jkzz** 1 year, 1 month ago

C - focus on stability or lose customers.

upvoted 1 times
- Arimaverick** 1 year, 3 months ago

Answer is C. Easier than A as if its not taken quickest care the Error Budget will get exhausted and the application will fall behind SLO

upvoted 2 times
- SoftSami** 1 year, 3 months ago

Selected Answer: C

In this scenario, the organization should prioritize stability to avoid prolonged user downtime. This means focusing on ensuring that the service is stable and reliable, even if it means delaying new feature releases or reducing the rate of innovation temporarily. By prioritizing stability, the organization can prevent prolonged outages or downtime, which can negatively impact the user experience and erode customer trust. Once stability is achieved, the organization can then focus on innovation and improving the user experience.

upvoted 3 times
- Matro71** 1 year, 3 months ago

Selected Answer: A

They still have error budget left.
The error budget gives developers clarity into how many failed fixes they can attempt without affecting the end user experience.

upvoted 1 times
- tbolick6** 1 year, 3 months ago

Selected Answer: C

Minimizing downtown improves SLO.

upvoted 1 times
- mtpro** 1 year, 7 months ago

"The error budget is typically the space between the SLA and the SLO. This error budget gives developers clarity into how many failed fixes they can attempt without affecting the end user experience."

upvoted 2 times
- Akshay0403** 1 year, 7 months ago

Selected Answer: C

When we monitor SLOs, our system begins to aggregate the data, and ties that back into error budgets which determines the allowable amount of system downtime or latency over a specific timeframe.

upvoted 3 times

  **zellick** 1 year, 9 months ago

Selected Answer: C

C is my answer.

upvoted 2 times

  **Gigante** 1 year, 9 months ago

Response is correct (C). Both Devs and SRE team must ensure that the error budget does not become exhausted. To avoid it, releases have to stop for the time being until the error budget resets. The team would have to reprioritise to focus on reliability to get it back to an acceptable state.

upvoted 4 times

  **Sbgani** 1 year, 9 months ago

Selected Answer: D

Error budgets let you track how many bad individual events (like requests) are allowed to occur during the remainder of your compliance period before you violate the SLO. You can use the error budget to help you manage maintenance tasks like deployment of new versions

<https://cloud.google.com/stackdriver/docs/solutions/slo-monitoring#defn-error-budget>

upvoted 1 times

  **LimeCake** 1 year, 9 months ago

Selected Answer: C

to prevent errors like this case, should spend on stability

upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: A

“The development team can ‘spend’ this error budget in any way they like. If the product is currently running flawlessly, with few or no errors, they can launch whatever they want, whenever they want. Conversely, if they have met or exceeded the error budget and are operating at or below the defined SLA, all launches are frozen until they reduce the number of errors to a level that allows the launch to proceed.”

upvoted 3 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: A

A. Innovation to improve user experience

The organization has NOT exhausted their error budget

upvoted 2 times

  **moi23** 1 year, 4 months ago

Stability before innovation. You should’t build a house on uneven ground or rock the boat on raging seas.

upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: C

C should be the correct choice

upvoted 2 times

An organization is moving away from an on-premises infrastructure. Instead, they want to create, access, and share information virtually in the cloud.

What should the organization consider?

- A. Built-in security when moving their data to the cloud
- B. Replacing their perimeter security with data encryption keys
- C. Optimizing cost-management with a capital expenditure model
- D. Increased hardware capacity when moving their data to the cloud

Correct Answer: B

Community vote distribution

A (80%)B (20%)

- mtpro

Highly Voted

1 year, 7 months ago

B would be my choice;
Choice A is a given & comes with the cloud model
Therefore, Chioce B is something any organization should consider
upvoted 7 times
- chai_gpt

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A. Built-in security when moving their data to the cloud

When moving from an on-premises infrastructure to the cloud, security is a top priority. The organization should consider the built-in security features offered by the cloud provider, as well as their own security requirements.
upvoted 3 times
- Matro71

1 year, 3 months ago

Selected Answer: B

They mention create information in the cloud, not moving. Therefore B.
upvoted 4 times
- someshsehgal

1 year, 6 months ago

Based on my practical experience. we don't complete rely on default security. Rather uses encryptions keys to add additional security measures.
upvoted 1 times
- Akshay0403

1 year, 7 months ago

Selected Answer: A

To help protect your data, Google encrypts data at rest, ensuring that it can only be accessed by authorized roles and services, with audited access to the encryption keys.
upvoted 3 times
- Cert_IT

1 year, 8 months ago

Selected Answer: A

A is the right Answer
upvoted 1 times
- AMMARABO

1 year, 8 months ago

Selected Answer: A

A Built-in security when moving their data to the cloud
upvoted 3 times
- MurugeshBabu

1 year, 9 months ago

A is the right answer
upvoted 1 times
- Govindaraj

1 year, 9 months ago

Selected Answer: A

A seems to be correct answer
upvoted 1 times

  **Ashish_01** 1 year, 9 months ago

Selected Answer: A

A is correct
upvoted 1 times

  **Sha1** 1 year, 9 months ago

Selected Answer: A

A, Built-in security when moving their data to the cloud
upvoted 2 times

  **Prosecute** 1 year, 9 months ago

Selected Answer: A

A. is the answer
upvoted 1 times

  **Vin1975** 1 year, 9 months ago

A. Built-in security when moving their data to the cloud
upvoted 1 times

An organization has an on-premises IT infrastructure. Their customer-facing application repeatedly fails during peak usage. What could be causing this issue?

- A. A serverless compute function struggles to scale.
- B. The application contains unclean data.
- C. They don't have enough servers to meet the demand.
- D. The application is only configurable on-premises.

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **tsk9921** 11 months, 2 weeks ago

C - Since it's on-premises, the data center / server is fully managed by the company

upvoted 1 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: C

C. They don't have enough servers to meet the demand.

upvoted 2 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C is the correct answer

upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

C. They don't have enough servers to meet the demand.

upvoted 2 times

What is artificial intelligence?

- A. Any system that ingests data in real time
- B. Any system that automatically structures data
- C. Any system capable of a task that requires smart analytics to generate predictions
- D. Any system capable of a task that normally requires human cognition

Correct Answer: D

Community vote distribution

D (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: D

D is correct.
upvoted 1 times
- Emma412** 10 months ago

Selected Answer: D

Answer = D
Explanation:
"...building computers and machines that can reason, learn, and act in such a way that would normally require human intelligence..."
@ <https://cloud.google.com/learn/what-is-artificial-intelligence>
upvoted 1 times
- someshsehgal** 1 year, 6 months ago

Selected Answer: D

HUMAN INTELLIGENCE
upvoted 1 times
- BryRob** 1 year, 7 months ago

Selected Answer: D

AI
the theory and development of computer systems able to perform tasks that normally require human intelligence, such as visual perception, speech recognition, decision-making, and translation between languages
upvoted 4 times
- juangt** 1 year, 8 months ago

Selected Answer: D

Is the ability of a computer or a robot controlled by a computer to do tasks that are usually done by humans because they require human intelligence and discernment
upvoted 1 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: D

D The correct answer
upvoted 1 times
- Sha1** 1 year, 9 months ago

Selected Answer: D

D. Any system capable of a task that normally requires human cognition
upvoted 1 times

What does Cloud Logging help an organization do?

- A. Analyze live source code and log code updates.
- B. Deploy infrastructure as code.
- C. Analyze logs and accelerate application troubleshooting.
- D. Manage storage of custom VM images.

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.
upvoted 1 times
-   **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: C

Per Google Docs, Cloud Logging, provides numerous key features with one being, Logs Explorer which "enables you to search, sort, and analyze logs through flexible query statements, along with rich histogram visualization, a simple field explorer, and ability to save the queries. "
<https://cloud.google.com/logging>
upvoted 1 times
-   **BryRob** 1 year, 7 months ago

Selected Answer: C

C. Analyze logs and accelerate application troubleshooting.
upvoted 2 times
-   **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C is the answer
upvoted 1 times
-   **Vin1975** 1 year, 9 months ago

Analyze logs and accelerate application troubleshooting.
upvoted 1 times

How is privacy defined in the context of cloud technology?

- A. Restrictions on data access and sharing
- B. Procedures to authenticate user identity
- C. Susceptibility to data breaches and cyber attacks
- D. Compliance with regulatory standards

Correct Answer: A

Community vote distribution

A (82%)D (18%)

- jaxclain

Highly Voted

1 year, 9 months ago

Selected Answer: A

I believe A and D are correct but I would say A because that is how most of Google Services allow you to define the privacy by restricting data access or sharing policies.
upvoted 5 times
- 00f7048

Most Recent

1 month ago

To me this is a general question, not just Google. Privacy defined by compliance standards makes more sense. A seems to be a general response to restrictions on data not specific to cloud technology.
upvoted 1 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A. Restrictions on data access and sharing

Privacy in the context of cloud technology is the protection of personal data from unauthorized access, use, or disclosure. This includes the right to control who has access to your data and how it is used.
upvoted 1 times
- rathod1010

1 year, 3 months ago

A is right
upvoted 1 times
- someshsehgal

1 year, 6 months ago

Selected Answer: A

Privacy linked to data security
upvoted 1 times
- someshsehgal

1 year, 6 months ago

Selected Answer: A

Privacy linked to data security
upvoted 1 times
- ShriRama

1 year, 6 months ago

Selected Answer: A

What is Privacy? Privacy, in the context of cloud technology, refers to the data an organization or an individual has access to and who they can share that data with.
upvoted 1 times
- AMMARABO

1 year, 8 months ago

Selected Answer: D

Compliance with regulatory standards
upvoted 2 times
- AMMARABO

1 year, 8 months ago

A: Compliance with regulatory standards
upvoted 1 times

  **Bill1000** 1 year, 9 months ago

D is correct
upvoted 1 times

  **Chirag7** 1 year, 9 months ago

A is correct
upvoted 2 times

  **zelck** 1 year, 9 months ago

Selected Answer: A
A is my answer.
upvoted 2 times

  **all_cloud** 1 year, 9 months ago

Selected Answer: D
D is correct
upvoted 1 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: A
A seems to be correct answer.
upvoted 2 times

  **Vin1975** 1 year, 9 months ago

A seems more correct. But why not D
upvoted 1 times



An organization wants to write and run small pieces of code in a serverless way that respond to events like huge discounts.
Which Google Cloud compute solution should the organization use?

- A. Google Kubernetes Engine
- B. Cloud Functions
- C. Bare Metal Solution
- D. Compute Engine

Correct Answer: B

Community vote distribution

B (100%)

- Leen8199** 2 months, 1 week ago

key world is serverless = Cloud function

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
- yphatt** 1 year, 7 months ago

Selected Answer: B

Only B is a serverless option.

upvoted 1 times
- BryRob** 1 year, 7 months ago

Selected Answer: B

<https://cloud.google.com/functions/docs/concepts/overview>

Google Cloud Functions is a serverless execution environment for building and connecting cloud services. With Cloud Functions you write simple, single-purpose functions that are attached to events emitted from your cloud infrastructure and services.

upvoted 4 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

It is B the correct answer

upvoted 1 times
- Pichikik** 1 year, 9 months ago

Event driven functions... Cloud serverless

upvoted 1 times

An organization wants to search hundreds of scanned documents for key information like dates, names, and other specific words. Why should the organization use application programming interfaces (APIs)?

- A. To replace the scanned documents with an online survey
- B. To ingest data in real time and encrypt unmatched words
- C. To create digital versions of the documents and locate key information
- D. To transform the documents into unstructured data.

Correct Answer: C

Community vote distribution

C (60%)D (40%)

- alexandercamachop

Highly Voted

1 year, 7 months ago

D is correct.
C is wrong because it says the data is in scanned documents, which means it is already digitalized.
upvoted 8 times
- Romio2023

5 months, 2 weeks ago

I think scanned documents and digital version of the documents are not the same. Scanned are only a picture of these documents, on the other hand the digital version of the documents are converted in Text and saved in format like PDF or any other common text formats.
upvoted 2 times
- shuvs

Highly Voted

1 year, 7 months ago

Selected Answer: D

Services like Natural Language API's scan the documents and extract entities which are unstructured data.
upvoted 6 times
- IVANHOE

Most Recent

4 months ago

Very tricky
upvoted 1 times
- IVANHOE

4 months ago

Leaning towards C since extracting key information seems to be the end goal.
Why not D? Well scanned documents (images or PDFs) are already unstructured data, so you would transform unstructured data to unstructured data.
upvoted 1 times
- Adon_97

6 months, 1 week ago

Selected Answer: C

The correct answer is C.
The organization can use APIs to connect to a document processing service that takes the scanned documents, convert them into digital text, and extract the key information that the organization is looking for.

Option D is not as effective for this task because it would make the documents even more difficult to search.
upvoted 1 times

gyenesge

7 months, 1 week ago

Selected Answer: C

I think the listed key details are not unstructured data. So D can't be the right answer me. Digitalized means OCR-ed document for me.
upvoted 1 times

chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times

__rajan__

8 months ago

Selected Answer: C

C. To create digital versions of the documents and locate key information

Application programming interfaces (APIs) can be used to create digital versions of scanned documents and locate key information within those documents. This is because APIs can be used to perform optical character recognition (OCR) on scanned documents. OCR is the process of converting scanned images of text into machine-readable text.
upvoted 2 times

🗨️ **mohitms1996** 8 months, 3 weeks ago

Selected Answer: C

Seems C is correct one. Making unstructured data is not making any sense here.

upvoted 1 times

🗨️ **oezgan** 9 months, 1 week ago

Selected Answer: D

Guys, D is correct. Think of it this way: the APIs are used to search for the data required in the documents (which are already digital btw). So the APIs are queried with the documents and not build on top of the documents. Hope that makes sense.

upvoted 1 times

🗨️ **dectum** 9 months, 3 weeks ago

C. To create digital versions of the documents and locate key information: APIs can be used to interact with Optical Character Recognition (OCR) software and other document processing tools. These APIs can help convert scanned documents into digital, machine-readable formats and then extract specific information like dates, names, and keywords from them. This is a common use case for APIs when dealing with large volumes of scanned documents, as it allows for efficient data extraction and analysis.

upvoted 1 times

🗨️ **Mehrannn** 11 months, 3 weeks ago

I will go for C - This is referring to full text search, by digitalizing the scanned documents (scanned docs are already unstructured data and are not digitalized yet) and indexing them to make the search process easier.

upvoted 1 times

🗨️ **inner4444** 11 months, 4 weeks ago

Selected Answer: C

The organization should use application programming interfaces (APIs) to create digital versions of the scanned documents and locate key information. APIs can provide access to text recognition and document processing services, which can extract text from scanned documents and make it searchable. By utilizing APIs, the organization can convert the scanned documents into machine-readable formats, extract relevant information such as dates, names, and specific words, and perform searches efficiently. This enables the organization to easily locate and retrieve the desired information from a large volume of documents.

upvoted 2 times

🗨️ **Arimaverick** 1 year, 3 months ago

Question and the options are confusing or not proper. Vision API can detect text in PDF/TIFF files <https://cloud.google.com/vision/docs/pdf>. Hence I think the closest is C but it also says to convert it to digital document which already has been done.

upvoted 1 times

🗨️ **[Removed]** 1 year, 3 months ago

Terrible question and choices. The answer is def not A or B which leaves C or D. The docs are already scanned (so not C) and why would they want to convert the docs to unstructured data when searching for "specific words"?

upvoted 4 times

🗨️ **SoftSami** 1 year, 3 months ago

Selected Answer: C

APIs can be used to access and extract specific information from documents, including scanned documents. By using APIs, the organization can extract key information such as dates, names, and specific words from the scanned documents without manually reviewing each document. This can significantly reduce the time and effort required to search the documents for the desired information.

Additionally, APIs can be used to convert the scanned documents into digital versions, which can be easily searched and analyzed. Once the scanned documents are converted into digital versions, the organization can use APIs to extract the desired information and store it in a database or other system for further analysis.

upvoted 2 times

🗨️ **tbolick6** 1 year, 3 months ago

Selected Answer: C

I chose C. I agree with one comment that the documents are scanned, so they are already digitized. But the question reads to me like the data is already unstructured. I do not like this question.

upvoted 1 times

🗨️ **Jackey0117** 1 year, 4 months ago

Selected Answer: C

APIs can be used to extract information from scanned documents and convert them into digital formats that can be searched and analyzed. APIs can help automate the process of locating key information, such as dates, names, and specific words, within hundreds of scanned documents, making it easier for the organization to find what they are looking for.

upvoted 1 times

🗨️ **Jackey0117** 1 year, 4 months ago

Modify my answer to D.

upvoted 1 times

An organization needs a platform to create custom end-to-end artificial intelligence models.
Which Google Cloud product or service should the organization use?

- A. Dataproc
- B. Compute Engine
- C. Recommendations AI
- D. Vertex AI

Correct Answer: C

Community vote distribution

D (100%)

-   **Ej24356** Highly Voted 1 year, 9 months ago

Selected Answer: D

Recommendations AI enables you to build an end-to-end personalized recommendation system based on state-of-the-art deep learning ML models, without a need for expertise in ML or recommendation systems. With Vertex AI, both AutoML training and custom training are available options. Whichever option you choose for training, you can save models, deploy models, and request predictions with Vertex AI.

upvoted 8 times
-   **CoffeeCloud** Most Recent 7 months ago

Selected Answer: D

What is going on here? C isn't the best option as it reduces the level of customization it offers

upvoted 2 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D. Vertex AI

Vertex AI is a managed machine learning platform that provides a comprehensive set of tools and services for building, training, deploying, and managing machine learning models. It is a good choice for organizations that need to create custom end-to-end AI models because it offers the following features:

upvoted 1 times
-   **Azwscp2023** 11 months, 1 week ago

Selected Answer: D

Vertex AI

upvoted 1 times
-   **reachmalhotras** 11 months, 2 weeks ago

Selected Answer: D

Focus on the Qn. "Need Platform"

upvoted 1 times
-   **ucsdmiami2020** 1 year, 6 months ago

Selected Answer: D

Per Google docs, Recommendations AI "enables you to build an end-to-end personalized recommendation system based on state-of-the-art deep learning ML models, without a need for expertise in ML or recommendation systems."

<https://cloud.google.com/recommendations-ai/docs>

upvoted 2 times
-   **Akshay0403** 1 year, 7 months ago

Selected Answer: D

Vertex AI Workbench is the single environment for data scientists to complete all of their ML work, from experimentation, to deployment, to managing and monitoring models

upvoted 2 times
-   **jellybiscuit** 1 year, 8 months ago

Selected Answer: D

<https://cloud.google.com/vertex-ai>
upvoted 3 times

  **TOLETIHARIPRIYASRUJANA** 1 year, 9 months ago
vertex AI
upvoted 2 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: D

D. Vertex is the correct answer.
Key word Custom C Recommendation AI is for making Recommendations
upvoted 1 times

  **Prosecute** 1 year, 9 months ago

Selected Answer: D

CUSTOM is the key word
upvoted 1 times

  **Skghost** 1 year, 9 months ago

Selected Answer: D

<https://cloud.google.com/vertex-ai>
upvoted 1 times

  **Kamlakar** 1 year, 9 months ago

Vertex AI it is a platform provide notebooks to create ML models .
upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

D Vertex is the correct answer. Key word Custom
C Recommendation AI is for making Recommendations
upvoted 1 times

  **WorkerG** 1 year, 9 months ago

why not D?
upvoted 1 times

An organization is training a machine learning model to make predictions.
What could improve the prediction accuracy of their model?

- A. An increase in storage capacity
- B. Higher network bandwidth
- C. An increase in training data
- D. Faster CPU processors

Correct Answer: D

Community vote distribution

C (100%)

-  **CoffeeCloud** 7 months ago

Selected Answer: C

Faster CPU will not increase accuracy of the final model so can't be D

upvoted 2 times
-  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-  **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-  **bujh** 1 year ago

I select C as they have specified improve prediction accuracy.

upvoted 1 times
-  **AK10000** 1 year, 4 months ago

I vote C. More the data the more accurate the prediction is. D would just speed up the process.

upvoted 2 times
-  **BryRob** 1 year, 7 months ago

Selected Answer: C

C. An increase in training data

upvoted 3 times
-  **SoftSami** 1 year, 8 months ago

Selected Answer: C

Definitely C

upvoted 1 times
-  **KIBOU** 1 year, 8 months ago

Selected Answer: C

should be C

upvoted 1 times
-  **bobby8521** 1 year, 9 months ago

Selected Answer: C

should be C

upvoted 2 times
-  **mtbmtb** 1 year, 9 months ago

Such as life right? More data typically gives more accurate results. Answer as everyone else. C

upvoted 2 times
-  **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C is the answer

upvoted 2 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

C. An increase in training data
forgot to select C

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

C. An increase in training data
upvoted 3 times

  **Skghost** 1 year, 9 months ago

Selected Answer: C

Go for C

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

I vote for C also.

upvoted 2 times

What is an example of structured data that a healthcare facility stores in their system?

- A. X-ray images
- B. Surgery video recordings
- C. Blood pressure history
- D. Physician-written notes

Correct Answer: D

Community vote distribution

C (100%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__

8 months ago

Selected Answer: C

C is structured.

upvoted 1 times
- Riaz1094

9 months ago

Dr. Hand written notes cant be structured data so the Answer should be C

upvoted 1 times
- ArjunKennedy

1 year, 2 months ago

D definitely not true its C. Because doctor's hand writing lol.

upvoted 3 times
- CarloGiovanni

1 year, 1 month ago

lol yes

upvoted 1 times
- ilberg

1 year, 6 months ago

Selected Answer: C

C is the answer - numeric values are structure data.

upvoted 1 times
- theKingslayer

1 year, 6 months ago

D is the least structured data in the universe. Even humans cant understand it.

upvoted 4 times
- BryRob

1 year, 7 months ago

Selected Answer: C

C. Blood pressure history

upvoted 2 times
- Klajdi_98

1 year, 8 months ago

Selected Answer: C

Blood pressure history

upvoted 2 times
- KIBOU

1 year, 8 months ago

Selected Answer: C

C. Blood pressure history

upvoted 1 times
- karthik0328

1 year, 8 months ago

Selected Answer: C

Blood Pressure is structured, notes are not structured

upvoted 2 times
- Chirag7

1 year, 9 months ago

Selected Answer: C

C is the answer - Physical measures like height, weight, blood pressure, blood type, and stage of the disease can be recorded numerically and they are structured.

upvoted 3 times

  **Govindaraj** 1 year, 9 months ago

Selected Answer: C

C is the answer

upvoted 1 times

  **jexmtropicscheatchatya** 1 year, 9 months ago

Selected Answer: C

C. Blood pressure history

I forgot to choose an answer in my previous comment

upvoted 1 times

  **Vin1975** 1 year, 9 months ago

C. Blood pressure history

upvoted 3 times

  **Prosecute** 1 year, 9 months ago

Selected Answer: C

c is structured data.

upvoted 1 times

  **Ej24356** 1 year, 9 months ago

Selected Answer: C

Examples specific to healthcare, the group explains, include radiology images or text files, like a physician's notes in the electronic health record. Even anonymous data is seen as unstructured.

upvoted 2 times

  **Skghost** 1 year, 9 months ago

Selected Answer: C

Numeric values only

upvoted 1 times

An organization wants to move from a tactical cloud adoption approach to a transformational approach. How should they change their cloud security?

- A. Provide staff identities using only Google Cloud authentication.
- B. Provide multiple layers of network security using a zero-trust model.
- C. Emphasize strong perimeter security and trust in their private network.
- D. Emphasize three main Identity Access Management roles: owner, editor, and viewer.

Correct Answer: B

Community vote distribution

B (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: B

B. Provide multiple layers of network security using a zero-trust model.

A transformational cloud adoption approach requires a shift in security thinking from a perimeter-based approach to a zero-trust approach. In a zero-trust approach, all users and devices are treated as untrusted until they are verified and authorized. This means that security is applied at every layer of the cloud environment, from the network to the application.

upvoted 3 times
- vimalraj02** 1 year, 7 months ago

B is correct. Zero trust model.

upvoted 1 times
- BryRob** 1 year, 7 months ago

Selected Answer: B

<https://www.crowdstrike.com/cybersecurity-101/zero-trust-security/>

Zero Trust is a security framework requiring all users, whether in or outside the organization’s network, to be authenticated, authorized, and continuously validated for security configuration and posture before being granted or keeping access to applications and data

upvoted 4 times
- Govindaraj** 1 year, 9 months ago

Selected Answer: B

It is B

upvoted 1 times
- Kamlakar** 1 year, 9 months ago

zero-trust model . B seems to be correct .

upvoted 1 times
- WorkerG** 1 year, 9 months ago

why not Vertex AI?

upvoted 1 times
- someshsehgal** 1 year, 6 months ago

Lol. where is that option

upvoted 1 times
- grzeev** 1 year, 8 months ago

where do you see this answer?

upvoted 4 times

What can customers expect if their cloud provider doesn't meet their service level agreement (SLA)?

- A. Increase in subscription fees
- B. Cloud service shutdown
- C. Refund for service interruption
- D. Error budget expansion

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: C

An SLA is a contractual agreement between a cloud provider and a customer that specifies the level of service that the provider will deliver, such as uptime, availability, performance, and support. If the provider fails to meet the SLA, the customer may be entitled to a refund or credit for the service interruption, based on the terms of the agreement

upvoted 2 times
-   **ucsdmiami2020** 1 year, 5 months ago

Selected Answer: C

Per Google docs, specifically that of the SLA for security command center "Customer will be eligible to receive the corresponding Financial Credits <https://cloud.google.com/security-command-center/sla>

upvoted 1 times

An organization wants to ensure that their developers can easily use application programming interfaces (APIs) in future projects.

Which Google product or solution should they use?

- A. Application migration
- B. Apigee
- C. The App Engine flexible environment
- D. AppSheet

Correct Answer: B

Community vote distribution

B (100%)

-   **ucsdmiami2020**

Highly Voted 

 1 year, 5 months ago

Selected Answer: B

Per Google docs, Apigee API Management allows for a customer to "Build, manage, and secure APIs - for any use case, environment, or scale. Google Cloud's native API management to operate your APIs with enhanced scale, security, and automation."

<https://cloud.google.com/apigee>

upvoted 5 times
-   **Leen8199**

Most Recent 

 2 months, 1 week ago

key word is API =APIGEE

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B. Apigee

Apigee is a Google Cloud product that enables organizations to publish, manage, and monitor APIs. It provides a variety of features that can make it easier for developers to use APIs, such as:

upvoted 1 times

An organization needs frequent access to only a subset of their data. They want to reduce costs by depositing the rest of their data across Nearline, Coldline, and Archive repositories.

Which Google Cloud product should the organization use?

- A. Filestore
- B. Cloud Storage
- C. Data Catalog
- D. Cloud Spanner

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
-   **serenity8468** 11 months, 1 week ago

Google Cloud Storage offers different storage classes that cater to different access frequency and cost requirements

Answer: B

upvoted 4 times
-   **ucsdmiami2020** 1 year, 5 months ago

Selected Answer: B

Per Google docs, specifically for GCP Cloud Storage there exists four types of storage with one of them, standard storage, being described as "storage for data that is frequently accessed ("hot" data) and/or stored for only brief periods of time."

<https://cloud.google.com/storage>

upvoted 3 times

A customer service department wants to increase their operational efficiency while maintaining personalized dialog with their customers.

What Google Cloud product or service should the organization use?

- A. Recommendations AI
- B. Cloud Identity
- C. Contact Center AI
- D. Text-to-Speech

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C. Contact Center AI

Google Cloud Contact Center AI is a cloud-based platform that helps organizations improve their customer service operations.

upvoted 2 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: C

Contact Center AI is a cloud-based solution that provides automated customer service and support using natural language processing, speech recognition, and machine learning technologies. It enables organizations to handle customer inquiries and support requests more efficiently, by automating routine tasks and providing personalized responses based on customer data and historical interactions. With Contact Center AI, the customer service department can automate common tasks such as routing calls, responding to frequently asked questions, and processing simple requests. This can help reduce wait times and improve response times, while freeing up agents to focus on more complex and high-value interactions.

upvoted 3 times
-   **ucsdmiami2020** 1 year, 5 months ago

Selected Answer: C

Per Google docs, Contact Center AI, "delight your customers with human-like AI-powered contact center experiences, lower costs, and free up you human agents' time."

<https://cloud.google.com/solutions/contact-center>

upvoted 3 times

An organization is struggling to keep up with the growth of their application which is running on legacy infrastructure.

What might be holding them back?

- A. The time it takes their serverless compute function to scale
- B. The overreliance on platform as a service
- C. The inaccessibility of their data due to perimeter security
- D. The cost of provisioning hardware for peak usage

Correct Answer: D

Community vote distribution

D (100%)

  **Adon_97** 6 months, 1 week ago

Selected Answer: D

D is correct
upvoted 1 times

  **chesterlp** 6 months, 2 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: D

D is correct.
upvoted 1 times

  **CoGCDL** 1 year, 3 months ago

The option that is most likely holding the organization back in keeping up with the growth of their application which is running on legacy infrastructure is option D, the cost of provisioning hardware for peak usage.

Legacy infrastructure is typically based on on-premises hardware that is managed and maintained by the organization. As the application grows and the user base expands, the hardware required to support it must also grow. This can lead to significant costs associated with provisioning and maintaining hardware, particularly if the organization needs to provision for peak usage.

upvoted 3 times

  **Ronniejsantos** 1 year, 4 months ago

Selected Answer: D

D. On premises enviromnent Brings this with a main debit.
upvoted 1 times

An organization wants to build autoscaling web applications without having to manage application infrastructure

Which Google Cloud product should they use?

- A. Anthos
- B. Apigee
- C. AutoML
- D. App Engine

Correct Answer: D

Community vote distribution

D (100%)

-   **ucsdmiami2020**

Highly Voted 

 1 year, 5 months ago

Selected Answer: D

Per Google docs, App Engine, allows for "freeing up your developers with zero server management and zero configuration deployments".

<https://cloud.google.com/appengine>

upvoted 5 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D is correct.

upvoted 1 times

What is an organization exclusively responsible for when they access an application through a software as a service (SaaS) model?

- A. Maintaining overall system operability
- B. Maintaining customer-facing content
- C. Monitoring data center servers
- D. Monitoring computer networks

Correct Answer: B

Community vote distribution

B (100%)

-   **SoftSami**

Highly Voted 

 1 year, 3 months ago

Selected Answer: B

In a SaaS model, the application and its underlying infrastructure are managed and maintained by the service provider. This includes the responsibility for system operability, server maintenance, network security, and other infrastructure-related tasks. The service provider ensures that the application is always available, secure, and up-to-date, and handles any issues or incidents that arise. However, the organization is responsible for maintaining their own content that is displayed through the application. This includes managing user accounts, data, and configurations, and ensuring that the content is accurate, up-to-date, and compliant with relevant laws and regulations.

upvoted 6 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times

An organization needs to train a custom machine learning model to categorize customer responses from their website's contact form.

Which Google Cloud product or service should the organization use?

- A. Anthos
- B. Recommendations AI
- C. AutoML Natural Language
- D. Pub/Sub

Correct Answer: C

Community vote distribution

C (82%)B (18%)

- Swati511**

Highly Voted

1 year, 5 months ago

It should be C, as per the question information has to be categorised from website's contact form which will be in the form of documented text data. This has nothing to do with Recommendations AI

upvoted 7 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: C

C. AutoML Natural Language

Google Cloud AutoML Natural Language is a product that allows users to train custom machine learning models to perform natural language processing (NLP) tasks, such as text classification, sentiment analysis, and entity extraction. To train a model, users simply provide AutoML Natural Language with a dataset of labeled data.

upvoted 3 times
- tbolick6**

1 year, 3 months ago

Selected Answer: C

Natural Language

upvoted 1 times
- Kasica96**

1 year, 3 months ago

Selected Answer: C

<https://cloud.google.com/natural-language/automl/docs>

upvoted 1 times
- SoftSami**

1 year, 3 months ago

Selected Answer: C

C is the answer for sure

upvoted 1 times
- CoGCDL**

1 year, 3 months ago

The Google Cloud product or service that the organization should use to train a custom machine learning model to categorize customer response: from their website's contact form is option C, AutoML Natural Language.

AutoML Natural Language is a Google Cloud service that enables organizations to train custom machine learning models to understand and categorize text. With AutoML Natural Language, the organization can easily create a custom model that can analyze customer responses and categorize them based on their content.

upvoted 2 times
- guilhermemls**

1 year, 4 months ago

Selected Answer: C

I would go with C because in this situation, the ML model needs to infer from the text the customers write in the forms, so it needs to have a Natural Language Processing (NLP) capability, present on AutoML Natural Language.

upvoted 2 times
- ucsdmiami2020**

1 year, 5 months ago

Selected Answer: B

Per Google docs, Recommendations AI, "delivers highly personalized product recommendations at scale". Additionally, "Recommendations AI draws on that experience and expertise in machine learning to deliver personalized recommendations that suit each customer's tastes and preferences across all your touchpoints."

<https://cloud.google.com/recommendations>
upvoted 2 times

Question #165

Topic 1

Which cybersecurity threat can lead to information being stolen or damaged without a user ever being aware?

- A. SLA policy violation
- B. Malware attack
- C. Ransomware attack
- D. Incomplete data input

Correct Answer: B

Community vote distribution

B (100%)

-   **Manny_SUI**

Highly Voted 

 1 year, 4 months ago

Answer is B - Keywords: Without user being aware... Ransomware attack would mean that users would be aware.

upvoted 5 times
-   **Pepe_Dev**

Highly Voted 

 1 year, 1 month ago

Selected Answer: B

A malware attack is a type of cyber attack in which malicious software is used to gain unauthorized access to a system, steal information, or cause damage. Unlike other types of attacks, such as ransomware or incomplete data input, a user may not be aware that their system has been compromised.

upvoted 5 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times

An organization has decided to create a mobile app for their customer-facing service.

What may have prompted this business decision?

- A. Customers expect information to be easily accessible on any device.
- B. Customers expect mobile-only access to information in the cloud era
- C. Developers can apply increased data security on a mobile app.
- D. Developers can run different versions of the service to test features only through a mobile app.

Correct Answer: A

Community vote distribution

A (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
-   **Pepe_Dev** 1 year, 1 month ago

Selected Answer: A

The most likely reason that prompted the organization to create a mobile app for their customer-facing service is that customers expect information to be easily accessible on any device, as option A suggests.

upvoted 2 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: A

the decision to create a mobile app for a customer-facing service is usually prompted by the need to provide customers with easy and convenient access to information and services on any device, rather than other technical reasons such as security or testing.

upvoted 3 times

An organization suffers a major data leak only six months after upgrading its security system.

What should the organization do to ensure long-term security?

- A. Extend data retention policy lengths to at least seven years.
- B. Pay more to get the best security system available on the market.
- C. Hire cybersecurity experts to further develop their data security plan
- D. Wait for a more advanced security system to be released

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C is correct.
upvoted 1 times
-   **Pepe_Dev** 1 year, 1 month ago

Selected Answer: C

To ensure long-term security, the organization should hire cybersecurity experts to review their current security system and develop a comprehensive data security plan. It is important to continuously assess and improve security measures in response to changing threats and evolving technology. Simply waiting for a more advanced security system to be released is not a proactive approach to addressing security concerns. Additionally, extending data retention policies or paying more for a security system does not necessarily address the root cause of the data leak and may not prevent future incidents.
upvoted 2 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: C

The best approach would be to conduct a thorough investigation to identify the root cause of the breach, take immediate steps to contain the breach and minimize the impact, and then implement a comprehensive remediation plan to prevent similar breaches from occurring in the future. One critical step in this process would be to hire cybersecurity experts to help identify vulnerabilities and develop an improved data security plan. These experts can perform a thorough review of the organization's existing security measures, identify any weaknesses, and recommend improvements to strengthen the organization's security posture.
upvoted 2 times
-   **CoGCDL** 1 year, 3 months ago

Hiring cybersecurity experts to further develop their data security plan (option C) is a proactive approach to improving security. These experts can conduct a thorough review of the organization's existing security measures, identify weaknesses, and develop a comprehensive plan to address them. By continually reviewing and improving their security measures, the organization can better protect themselves against future data leaks and cyber attacks.
upvoted 2 times
-   **bda92b3** 1 year, 4 months ago

Correct answer is A
upvoted 2 times

An organization has started to develop cloud-native applications to replace their legacy applications.

What else should the organization do to adopt the cloud?

- A. Provide privileged cloud network access to third-party partners.
- B. Migrate some of their legacy infrastructure to the cloud.
- C. Conduct maintenance on their legacy hardware.
- D. Increase CapEx on existing projects.

Correct Answer: *B*

Community vote distribution

B (80%)

A (20%)

  **Leen8199** 2 months, 1 week ago
C is correct step by step to adopting
upvoted 1 times

  **cloudmind** 5 months ago
C is correct!

Since the organization has started to develop "cloud-native" applications to replace their legacy applications, I suppose the keyword is: "Cloud-Native". So I'd exclude the answer B as you do not migrate an on-prem infrastructure to the cloud. I exclude also answer D, as you want to maintain legacy infrastructure expenditures, not increase them. I exclude also A: why do you need provide privileged cloud network access to third-party partners? It does not make sense.
You want to maintain your existing legacy hardware until you finish developing new cloud-native applications
upvoted 1 times

  **PrivateHulk** 5 months, 1 week ago
B.
Migrating legacy infrastructure to the cloud is a common and feasible practice. Cloud migration offers several benefits, including increased scalability, flexibility, cost savings, and improved performance.
upvoted 1 times

  **alo4u** 6 months, 2 weeks ago
I think the Answer is A. what is the benefit of migrating LEGACY infrastructure to the cloud. You, however may need to connect to 3rd parties, say MasterCard or Visa for your application to work.
upvoted 1 times

  **gyenesge** 7 months, 1 week ago

Selected Answer: A

I don't think that B is correct as normally you do not migrate infrastructure to cloud. Even the lowest level of cloud adoption is about renting infra from cloud provider.
I don't like any of the answers, but the A is the only meaningful for me.
upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: B

B is correct.
upvoted 1 times

  **Luizcjr** 1 year ago
B ESTA CORRETO.
upvoted 1 times

  **PHD_CHENG** 1 year, 4 months ago

Selected Answer: B

B is correct. This is a logical answer
upvoted 2 times

An organization has created an application that can diagnose different medical conditions when users submit images of their affected body parts.

Which Google Cloud product or service did the organization use?

- A. Cloud Logging
- B. Cloud Profiler
- C. App Engine
- D. Machine learning

Correct Answer: D

Community vote distribution

D (75%)C (25%)

- Leen8199** 2 months, 1 week ago

D classifications
upvoted 1 times
- Adon_97** 6 months, 1 week ago

Selected Answer: D

D is correct
Option C (App Engine) can be used to host the organization's application, it has nothing to do with image recognition techonology, for that you need a Machine Learning Service
upvoted 1 times
- CoffeeCloud** 7 months ago

Selected Answer: C

Machine Learning is not a GCP service or product so it is ruled out instantly. Only thing we can say for sure if the App Engine out of the other three
For all we know the app just connects to a third part API to get the result from the photo
upvoted 2 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct.
upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: D

D is Correct.
upvoted 1 times
- inner4444** 1 year ago

Selected Answer: D

D - machine learning. It says "GCP product OR service". Hence, one of the machine learning services.
upvoted 1 times
- [Removed]** 1 year, 2 months ago

Machine learning is neither a GCP product nor a GCP service, it is rather a branch of AI, a field of science, or a technology. Out of these options, App Engine is the most suitable one for "An organization has created an application... - What GCP product or service an organization uses?".
upvoted 4 times
- SoftSami** 1 year, 3 months ago

Selected Answer: D

The organization used Machine learning on Google Cloud to create an application that can diagnose different medical conditions when users submit images of their affected body parts.
upvoted 2 times
- CarloGiovanni** 1 year, 1 month ago

Machine Learning is not a Google Cloud Product, so I'd say it's App engine.
upvoted 1 times

An organization wants to better understand the behavior of their code in production and analyze its state to identify hard-to-find programming errors.

Which Google Cloud tool should the organization use?

- A. Debugger
- B. Profiler
- C. Cloud Monitoring
- D. Trace

Correct Answer: C

Community vote distribution

A (100%)

-   **Kasica96**

Highly Voted 

 1 year, 3 months ago

Selected Answer: A

Cloud Debugger solves the problem of isolating issues that occur only in production. By letting you inspect the state of a running application in real time, without stopping or slowing it down, Debugger helps you solve problems that can be impossible to reproduce in a local environment. Debugger also saves time by eliminating the process of redeploying applications just to add logging statements.

Some issues only arise in production—and debugging locally isn't an option. At other times, redeploying resolves the issue for some period of time, but the root cause of the problem hasn't been corrected. When this happens, use Cloud Debugger to find the problem by taking snapshots of variables and the call stack and injecting debug logpoints in your running application. Your users are not affected while you capture the call stack and variables at any location in your source code.

upvoted 5 times
-   **chai_gpt**

Most Recent 

 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 2 times
-   **__rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 2 times
-   **serenity8468** 11 months, 1 week ago

Debugger - inspect states of running app after deployment

Profiler - Analyze performance

Cloud Monitoring - Analyze metrics of all apps and infrastructures

Trace - analysis to understand request flow, service topology, and latency issues

upvoted 4 times
-   **Kasica96** 1 year, 3 months ago

<https://cloud.google.com/debugger/docs/setup>

Cloud Debugger solves the problem of isolating issues that occur only in production. By letting you inspect the state of a running application in real time, without stopping or slowing it down, Debugger helps you solve problems that can be impossible to reproduce in a local environment. Debugger also saves time by eliminating the process of redeploying applications just to add logging statements.

Some issues only arise in production—and debugging locally isn't an option. At other times, redeploying resolves the issue for some period of time, but the root cause of the problem hasn't been corrected. When this happens, use Cloud Debugger to find the problem by taking snapshots of variables and the call stack and injecting debug logpoints in your running application. Your users are not affected while you capture the call stack and variables at any location in your source code.

upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: A

The organization should use Google Cloud Debugger to better understand the behavior of their code in production and analyze its state to identify hard-to-find programming errors.

upvoted 2 times
-   **CoGCDL** 1 year, 3 months ago

The tool that the organization should use to better understand the behavior of their code in production and analyze its state to identify hard-to-find programming errors is option D, Trace.

Google Cloud Trace is a distributed tracing system that helps in collecting, analyzing, and debugging distributed systems. It allows developers to trace the path of a request through a complex system, and it provides detailed information about latency and other performance metrics. With Trace, organizations can identify and diagnose hard-to-find programming errors, optimize application performance, and improve the end-user experience.

upvoted 1 times

  **PHD_CHENG** 1 year, 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

  **Bolah** 1 year, 5 months ago

A is the answer

upvoted 2 times

  **MultiCloudIronMan** 1 year, 5 months ago

Selected Answer: A

I think its A

upvoted 1 times

  **Swati511** 1 year, 5 months ago

It should be A

upvoted 1 times

  **ujuuuu** 1 year, 5 months ago

Selected Answer: A

Cloud Debugger helps monitor application performance. IT teams can inspect the state of a running application in real time, without stopping or slowing it down. This means that end users are not affected while a developer searches the source code. IT teams can use it to understand the behavior of their code in production and analyze its state to find those hard-to-find bugs.

upvoted 1 times

  **Haplo46** 1 year, 5 months ago

Profiler is the tool that an organization should use to better understand the behavior of their code in production and analyze its state to identify hard-to-find programming errors. It provides detailed performance data, including CPU and heap usage, that can be used to identify performance bottlenecks and memory leaks. Additionally, it allows you to profile your code in real time and see the source code

upvoted 2 times

  **gviolant** 1 year, 5 months ago

Selected Answer: A

the answer is Debugger.

Cloud Debugger Investigate code behavior in production instead Cloud Monitoring and Cloud Logging to monitor the health and performance

upvoted 3 times

A retail organization has moved all of their inventory data to a relational database in the cloud.

What functionality does a relational database offer?

- A. It stores large amounts of raw data in its original format
- B. It stores transactional data, which can then be accessed electronically
- C. It rapidly analyzes large and multi-dimensional datasets
- D. It analyzes unstructured data, which can then be accessed in multiple regions

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: B

A relational database offers the functionality of storing transactional data, which can then be accessed electronically. Relational databases store structured data that can be organized in tables with defined relationships between them. This makes them well-suited for transactional data, such as inventory data, that needs to be accessed and updated frequently.

upvoted 4 times
-   **MultiCloudIronMan** 1 year, 5 months ago

Selected Answer: B

Relational database is built for transactional data and processing.

upvoted 3 times

How does Google Cloud AI Hub make machine learning easy and accessible for Google Cloud customers to get started?

- A. It deploys artificial intelligence models in real time
- B. It automatically codes users’ machine learn ng models
- C. It automatically sources and prepares users’ machine learning data
- D. It includes a repository of plug-and-play artificial intelligence tools

Correct Answer: D

Community vote distribution

D (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: D

D is correct. It include Pre-trained models.
upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: D

Google Cloud AI Hub includes a repository of plug-and-play artificial intelligence tools, which makes machine learning easy and accessible for Google Cloud customers to get started. AI Hub provides pre-built machine learning models, such as image recognition and natural language processing, that customers can use as-is or customize to their needs. Additionally, it allows customers to discover, share, and reuse machine learning components, including Jupyter notebooks and TensorFlow modules, with other members of their organization or the broader community.
upvoted 3 times
-   **Haplo46** 1 year, 5 months ago

D is right, Google Cloud AI Hub is a platform designed to make machine learning easy and accessible for Google Cloud customers
upvoted 4 times

A venue with an online booking system has partnered with a catering business.

How can the venue leverage application programming interfaces (APIs) to create new business value?

- A. Use an API to redesign their booking system to appeal to targeted customers
- B. Use an API to identify opportunities for new business collaboration
- C. Use an API to migrate all their customer data to a machine learning model to predict food requests
- D. Use an API to connect their booking system with the partner's catering form

Correct Answer: D

Community vote distribution

D (100%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: D

D is correct. API's can be use to connect to the partnered catering business.
upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: D

The correct answer is D. The venue can use an API to connect their booking system with the partner's catering form, enabling customers to book both the venue and catering services seamlessly. By leveraging APIs, the venue can integrate and share data between their booking system and the catering form, ensuring that customer bookings are accurately processed and that the venue and catering business can efficiently manage their operations together. This creates new business value by offering customers a convenient and streamlined booking process for both services.
upvoted 3 times

  **bda92b3** 1 year, 4 months ago

D is correct
upvoted 1 times

A public cloud provider's service level performance has moved below the service level objective (SLO), but remains above the service level agreement (SLA).

What is an outcome of this change in service level performance?

- A. The public cloud provider is encouraged to push out new updates
- B. The public cloud provider is encouraged to prioritize service reliability
- C. End users face unexpected or prolonged downtime
- D. End users face a decrease in service latency

Correct Answer: D

Community vote distribution

B (77%)

14%

9%

MultiCloudIronMan

Highly Voted

1 year, 5 months ago

Selected Answer: B

SLO is internal and if its low it is an early warning that SLA will be breached so the provider should improve service.
upvoted 9 times

PrivateHulk

Most Recent

5 months, 1 week ago

Where service level performance has moved below the SLO but remains above the SLA:

Users may notice degraded performance or service disruptions, leading to a reduction in user satisfaction. While the service may still be within contractual obligations (SLA), the degradation signals a potential impact on the user experience.

The deviation from the SLO threshold can trigger a focus on continuous improvement. Service providers may use this information to identify areas for optimization and take proactive measures to enhance performance - (not service reliability)

Reliability
referes to the ability of a system or service to consistently perform its intended functions without failures

Performance
refers to how well a service executes its tasks, often measured in terms of speed, responsiveness, throughput, and efficiency.

So correct answer is D.
upvoted 2 times

PrivateHulk

5 months, 1 week ago

Correction - although its still pointing to D.
Where service level performance has moved below the SLO but remains above the SLA.

Users may notice degraded performance or service disruptions, leading to a reduction in user satisfaction. While the service may still be within contractual obligations (SLA), the degradation signals a potential impact on the user experience.

The deviation from the SLO threshold can trigger a focus on continuous improvement. Service providers may use this information to identify areas for optimization and take proactive measures to enhance performance - and not service reliability -

This means not B but C.
upvoted 1 times

PrivateHulk

5 months, 1 week ago

They question is asking "What is an outcome of this change in service level performance"?

When a service level performance falls below the agreed-upon Service Level Agreement (SLA), but still remains above the minimum threshold defined in the SLA, it can have several impacts on users, customers, and the overall business.

Users may experience degraded performance, slower response times, or intermittent issues with the service. This can lead to frustration and dissatisfaction among users.

Hence D is correct
upvoted 1 times

chai_gpt

7 months, 3 weeks ago

Selected Answer: D

I think D is correct answer. Because option A and B are not the outcomes of the performance (but they are actions).
Now answer can be either C or D.
But in C, prolonged downtime is mentioned. If performance is above SLA then prolonged downtime can't be there.
So, we are left with only D option.

upvoted 2 times

  **__rajan__** 8 months ago

Selected Answer: B

B. The public cloud provider is encouraged to prioritize service reliability.

A service level objective (SLO) is a target that a cloud provider sets for the performance of its services. A service level agreement (SLA) is a contract between a cloud provider and its customers that guarantees a certain level of performance.

upvoted 3 times

  **Emma412** 10 months ago

I don't think D is the correct answer.
Read the blog post from google cloud site - <https://cloud.google.com/blog/products/gcp/consequences-of-slo-violations-cre-life-lessons>
Surprisingly it has reference to both "B" and "C"
why C - "...a service in violation of its SLO is by definition making users unhappy..."
why B - "Summary: ...how responding to SLO violations moves the needle towards reliability"
read the full article and suggest which statement you agree with.

upvoted 1 times

  **Emma412** 10 months ago

Reading the document again, i am wondering there is a hidden statement in the question.
As per the article when there is a breach of SLO there is a immediate shift in priority to focus on reliability. As a result the service latency is reduced. So May be that is why Option "D" is the correct answer.

upvoted 1 times

  **andr31** 1 year, 1 month ago

Selected Answer: C

Answer should be C. "The provider is encouraged" is not an outcome of performance moving between SLO and SLA, but rather an advice.

upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: C

The outcome of the public cloud provider's service level performance moving below the service level objective (SLO) but remaining above the service level agreement (SLA) is that end users may face unexpected or prolonged downtime. This means that the cloud provider is not meeting the expected performance levels, even if they are still meeting the minimum standards set forth in the SLA. To maintain high levels of service reliability, the cloud provider should strive to meet or exceed the SLO.

upvoted 2 times

  **CoGCDL** 1 year, 3 months ago

Selected Answer: B

B is the correct choice here

upvoted 1 times

  **guilhermemls** 1 year, 4 months ago

Selected Answer: B

B because staying below SLO means the CSP needs to have attention and focus in improving the reliability and availability of its services to reach the Objective level.

upvoted 3 times

  **Swati511** 1 year, 5 months ago

B is the correct choice here

upvoted 2 times

  **ujuuuu** 1 year, 5 months ago

Selected Answer: B

B is the correct answer

upvoted 1 times

  **Haplo46** 1 year, 5 months ago

I think B is the right answer because the service level performance has moved below the service level objective (SLO), which is a higher standard than the service level agreement (SLA). SLOs are internal targets set by the service provider, while SLAs are agreements between the service provider and the customer. If the provider's performance falls below the SLO, it means that they are not meeting their own standards and should prioritize improving service reliability in order to meet or exceed their SLOs

upvoted 4 times

An organization is looking for a solution to help them build applications using their preferred programming language and tools. They also want to minimize time spent on infrastructure management.

What solution should the organization use?

- A. Cloud Run
- B. Dataproc
- C. Pub/Sub
- D. Bare Metal Solution

Correct Answer: A

Community vote distribution

A (100%)

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
- Emma412** 10 months ago

Selected Answer: A

Answer = A
Explanation:
Comparing the rest of the choices given, cloud run is the only GCP product that gives the option to build applications (Whilst Dataproc is a bigdata tool; Pub/Sub is a messaging service; Baremetal as the name suggests)

Tip:
if App Engine was to be one of the options the question gets trickier.
remember App Engine does not have as many programming language support as CloudRun. CloudRun is build based on the experience from AppEngine.
upvoted 4 times
- bda92b3** 1 year, 4 months ago

Correct answer A
Cloud Run is a managed compute platform that lets you run containers directly on top of Google's scalable infrastructure.
upvoted 1 times
- MultiCloudIronMan** 1 year, 5 months ago

Selected Answer: A

You can develop on different programming languages on Cloud Run its serverless
upvoted 1 times
- MultiCloudIronMan** 1 year, 5 months ago

All the same, Bare Metal is a strong answer if you consider that you can bring your own tools as well.
upvoted 1 times

An organization has created a medical fitness app and wants to store a very large amount of data about exercise times, activity, and calories burned for later analysis.

Which data management solution should the organization use?

- A. Data lake
- B. Data warehouse
- C. Data field
- D. Database

Correct Answer: B

Community vote distribution

B (71%)

A (29%)

- Emma412

Highly Voted

10 months ago

Selected Answer: B

I strongly think that B is the correct answer

Reasoning:
I read an article from Qlik.com that compares DataLake and DataWarehouse.
In that i noticed 2 factors that bring out the difference.
1. DataLake better at handling unstructured Data whereas Datawarehouse is better at structured data
2. DataLake shines are real-time analytics and nothing is said about real-time analytics for Datawarehouse.
in the given question the data is primarily structured and explicitly mentions that data is required for analysis "Later". hence B should be the correct match/answer.
upvoted 6 times
- gyenesge

Most Recent

7 months, 1 week ago

Selected Answer: B

Datawarehouse is good for analysis. Datalake is more for storing semistructured data, or data in raw format.
upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: B

B. Data warehouse is correct. Because the data about exercise times, activity, and calories burned etc. falls under either structured data or semi-structured data which data warhouse (like BigQuery) can handle really well. And it is mentioned in question that data is needed for later analysis. Data warehouse (like BigQuery) can help in this analysis.
upvoted 4 times
- chai_gpt

7 months, 3 weeks ago

B. Data warehouse is correct. Because the data about exercise times, activity, and calories burned etc. falls under either structured data or semi-structured data which data warhouse (like BigQuery) can handle really well. And it is mentioned in question that data is needed for later analysis. Data warehouse (like BigQuery) can help in this analysis.
upvoted 2 times
- __rajan__

8 months ago

Selected Answer: A

A. Data lake

A data lake is a centralized repository that stores all types of data, including structured, semi-structured, and unstructured data. Data lakes are often used to store large amounts of data for later analysis.
upvoted 1 times
- Kaycee101

9 months, 2 weeks ago

B is the correct answer here
upvoted 2 times
- Dinraj

11 months ago

Selected Answer: B

I would go for B
upvoted 2 times
- inner4444

1 year ago

Selected Answer: A

Based on AI's answer it's A - data lake
upvoted 2 times

  **inner4444** 11 months, 4 weeks ago
because the data comes from the fitness app, it's most likely a structured data, I will change my answer to B
upvoted 1 times

  **andr31** 1 year, 1 month ago
Selected Answer: B
The data that the fitness company wants to store is structured and/or semi-structured and there's no mention of BLOBs or anything like that. The answer is B.
upvoted 2 times

  **hita97** 1 year, 1 month ago
Selected Answer: A
Data warehouse stores transformed data. For storing purpose Data Lake can be used
upvoted 1 times

  **Flexip** 1 year, 2 months ago
Selected Answer: A
An organization that has created a medical fitness app and wants to store a very large amount of data about exercise times, activity, and calories burned for later analysis should use a data lake (A). A data lake is a centralized repository that allows organizations to store all their structured and unstructured data at any scale. Data can be stored in its raw form and can be analyzed later using big data processing tools.
upvoted 4 times

  **BTAB** 1 year, 4 months ago
Selected Answer: B
Agree on B. Since there is no requirement for semi-structured and unstructured data in raw form, and the data mentioned in the question large amounts of structured data, I would go with Data Warehouse.
upvoted 2 times

  **CarloGiovanni** 1 year, 1 month ago
however, all data is structured so a warehouse fits better
upvoted 1 times

  **guilhermemls** 1 year, 4 months ago
Selected Answer: B
I would choose Data Warehouse as it is a solution to store Big Data and serve it for analytics services. For example, BigQuery has an ML solution embedded on it.
upvoted 2 times

  **DPonly** 1 year, 4 months ago
option A , as current requirement is storing data
upvoted 1 times

An organization is looking for a business intelligence solution that allows individual employees and end users to analyze business data and generate insights.

Which Google Cloud product or service should the organization use?

- A. Cloud Spanner
- B. BigQuery
- C. Looker
- D. Dataflow

Correct Answer: C

Community vote distribution

C (100%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: C

C. Looker

Looker is a cloud-based business intelligence (BI) platform that enables users to explore, analyze, and visualize data. It provides a variety of features that make it easy for users to create and share insights

upvoted 1 times

  **Emma412** 10 months ago

Selected Answer: C

Answer = C

Explanation: Merely by the definition Looker is a business intelligence

Tip: While Bigquery is also used for analyzing data, this is NOT specific to business insights

upvoted 1 times

  **tsk9921** 11 months, 2 weeks ago

C. Looker is the correct choice here

upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: C

C. Looker

upvoted 1 times

  **guilhermemls** 1 year, 4 months ago

Selected Answer: C

C, as Looker is a Business Intelligence platform.

upvoted 1 times

  **Haplo46** 1 year, 4 months ago

Looker, as it is a business intelligence platform that allows individual employees and end users to explore, analyze, and visualize data

upvoted 3 times

What is out-of-the-box observability within the context of cloud operations?

- A. Monitoring project development timelines
- B. Monitoring IT infrastructure and applications
- C. Monitoring customer satisfaction and feedback
- D. Monitoring cloud subscription costs

Correct Answer: B

Community vote distribution

B (87%)

13%

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
-   **Emma412** 10 months ago

Selected Answer: B

Answer = B

Explanation: Refer the documentation
<https://cloud.google.com/products/operations?hl=en> (goto first use-case)
"Cloud Logging and Cloud Monitoring provide your IT Ops/SRE/DevOps teams with out-of-the box observability needed to monitor your infrastructure and applications."
upvoted 2 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: B

Out-of-the-box observability within the context of cloud operations means having preconfigured monitoring and analysis tools that provide visibility into various aspects of an organization's IT infrastructure and applications without requiring extensive configuration or customization. This includes monitoring metrics such as system performance, availability, and error rates, as well as logs, traces, and other data that can be used to diagnose and troubleshoot issues. The goal is to provide quick and easy access to relevant information that can help IT teams understand the health and performance of their systems and applications.
upvoted 4 times

  **adz07** 1 year, 3 months ago

Selected Answer: B

It seems B. based on this

<https://cloud.google.com/blog/products/management-tools/cloud-operations-suite-gets-21-new-features>
upvoted 1 times

  **gviolant** 1 year, 4 months ago

Selected Answer: B

What is observability in Cloud?
In IT and cloud computing, observability also refers to software tools and practices for aggregating, correlating and analyzing a steady stream of performance data from a distributed application along with the hardware and network it runs on, in order to more effectively monitor, troubleshoot and debug the application
upvoted 4 times

  **rrgonzalez1992** 1 year, 4 months ago

Selected Answer: D

I think D is the correct answer
upvoted 2 times

An organization wants a serverless execution environment for building and connecting cloud services. They can't predict how many resources they'll need.

Which Google Cloud solution should they use?

- A. Compute Engine
- B. Google Kubernetes Engine
- C. Cloud Functions
- D. Bare Metal Solution for SAP

Correct Answer: A

Community vote distribution

C (100%)

-  **Kandankery** 5 months ago

C correct answer
upvoted 1 times
-  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times
-  **__rajan__** 8 months ago

Selected Answer: C

C. Cloud Functions.
As the question is asking for a serverless environment.
upvoted 2 times
-  **Emma412** 10 months ago

Selected Answer: C

Answer = C
It's easy to narrow down to "Cloud Functions"
Occasionally the phrase "can't predict how many.." may confuse if this is indirectly referring to Auto Spawning of Container nodes. But remember that GKE is running off of ComputeEngine. Refer to this link for precise explanation on how GKE runs.
https://cloud.google.com/kubernetes-engine/docs/concepts/kubernetes-engine-overview#how_works
upvoted 1 times

 **Nguyen25183** 1 year, 3 months ago

Same with AWS Lambda
upvoted 1 times

 **gviolant** 1 year, 4 months ago

Selected Answer: C

cloud functions
upvoted 2 times

 **lauraperdomo04** 1 year, 4 months ago

Selected Answer: C

cloud functions
upvoted 1 times

 **SANJANA_ALVA** 1 year, 4 months ago

C - cloud functions
upvoted 1 times

 **neerajRathi** 1 year, 4 months ago

Selected Answer: C

Cloud Functions
upvoted 1 times

 **Haplo46** 1 year, 4 months ago

Cloud Functions
Run code in the cloud with no servers or containers. Answer must be C
upvoted 1 times

  **DeNel** 1 year, 4 months ago

Selected Answer: C

The Google Cloud solution that the organization should use is Cloud Functions. Cloud Functions is a serverless execution environment for building and connecting cloud services. It allows you to run code without having to manage the underlying infrastructure, and it can be used to automate a variety of tasks. Cloud Functions is a good choice for the organization because it can be used to build and connect cloud services without having to predict how many resources they'll need.
upvoted 4 times

  **cloudmind** 1 year, 5 months ago

C - Cloud function is serverless
upvoted 2 times

  **Swati511** 1 year, 5 months ago

Compute is not serverless, it should be cloud functions
upvoted 1 times

  **Haplo46** 1 year, 5 months ago

C- Cloud Functions is a serverless execution environment for building and connecting cloud services
upvoted 2 times

  **Oksa** 1 year, 5 months ago

C - Cloud Functions
upvoted 2 times

  **FreddyBrainy** 1 year, 5 months ago

It should be C Cloud functions.
upvoted 2 times

  **kakaomg** 1 year, 5 months ago

C, Should not it be Cloud function ?
upvoted 2 times

An organization is struggling to ensure that their legacy systems meet rapidly changing IT needs.

Why should the organization use application programming interfaces (APIs)?

- A. To expose all underlying data so third parties can make amends
- B. To completely re-architect legacy applications
- C. To achieve long-term operational flexibility
- D. To operate on a pay-as-you go model so IT can scale flexibly

Correct Answer: C

Community vote distribution

C (80%)B (20%)

- stevee954** 7 months, 2 weeks ago

Wouldn't D accomplish C? not a great question.

upvoted 2 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C. To achieve long-term operational flexibility

upvoted 1 times
- Emma412** 10 months ago

Selected Answer: C

Answer = C

Explanation:

while logically C is the best answer, from exam standpoint we need to be 100% how this is derived.

By definition API gateway "API Gateway gives you high performance, scalability... Consumption-based and tiered pricing means you can better manage cost."

All of these are directly terms from "Operations Management". it also states the benefits of API gateways as "Take advantage of all the operational benefits such as....."

Hence this answer is a direct extract from google cloud documentation.

upvoted 2 times
- MNotABot** 11 months, 3 weeks ago

Answer is: "To achieve long-term operational flexibility".

You can "re-host" (Lift and Shift) your app on cloud without complete re-architecture of the whole APP.

upvoted 1 times
- tbolick6** 1 year, 3 months ago

Selected Answer: C

C for me

upvoted 1 times
- guilhermemls** 1 year, 4 months ago

Selected Answer: C

Answer C according to question 80

upvoted 3 times
- Swati511** 1 year, 5 months ago

It should be C, APIs are not used for change the underlying architecture of legacy systems, rather they help with integrations in the most cost effective way

upvoted 3 times
- FreddyBrainy** 1 year, 5 months ago

Making amends here I think I made a mistake:

In question 80: A hotel wants to modernize their legacy systems so that customers can make reservations through a mobile app. What's the benefit of using an application programming interface (API) to do this?

Answer: They do not have to re-write their legacy system.

If there's no need to re-write, there's should also be no need to re-architect it? Thus in reference to the answer on that question 80. Answer B is wrong.

So answer remains C.
upvoted 4 times

  **FreddyBrainy** 1 year, 5 months ago
Answer B.
upvoted 1 times

  **gviolant** 1 year, 5 months ago
Selected Answer: B
For me is B
upvoted 2 times

Question #181

Topic 1

An organization is looking for a solution that will allow them to build new applications, modernize old ones, and run applications across different cloud providers.

What solution should the organization use?

- A. Apigee Hybrid
- B. Anthos
- C. Cloud Run
- D. Compute Engine

Correct Answer: B

Community vote distribution

B (100%)

  **Leen8199** 2 months, 1 week ago
key word different cloud providers= anthos
upvoted 1 times

  **chai_gpt** 7 months, 3 weeks ago
Selected Answer: B
B is correct
upvoted 1 times

  **__rajan__** 8 months ago
Selected Answer: B
B. Anthos

Anthos is a managed platform that enables organizations to build, deploy, and manage modern applications anywhere. It provides a consistent set of tools and services for developing, deploying, and managing applications across multiple cloud providers and on-premises environments.
upvoted 1 times

  **SoftSami** 1 year, 3 months ago
Selected Answer: B
The organization should use Anthos. Anthos is a hybrid and multi-cloud application platform that allows organizations to build, deploy, and manage applications across different cloud providers, on-premises, and edge locations. It offers a consistent development and deployment experience across hybrid and multi-cloud environments, and provides a unified management interface for applications running on Kubernetes.
upvoted 2 times

  **SoftSami** 1 year, 3 months ago
Selected Answer: B
The solution that the organization should use to build new applications, modernize old ones, and run applications across different cloud providers is Anthos.
upvoted 2 times

An organization cannot afford to modernize their infrastructure, but they want to process data from their legacy system in a modern platform hosted by a business partner.

What solution should the organization choose to make their data accessible?

- A. Compute Engine
- B. Anthos
- C. An application programming interface
- D. Google Kubernetes Engine

Correct Answer: C

Community vote distribution

C (88%)

13%

- chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: C

C. An application programming interface

They can expose an API that can be used by their business partner.

upvoted 1 times
- tsk9921** 11 months, 2 weeks ago

C. API for me

upvoted 1 times
- tbolick6** 1 year, 2 months ago

Selected Answer: C

C - API to unlock value from legacy system.

upvoted 1 times
- guilhermemls** 1 year, 4 months ago

Selected Answer: C

Answer is C because of API serving as a bridge between legacy and modern systems, just like Haplo46 said.

upvoted 2 times
- Haplo46** 1 year, 5 months ago

For me C is correct, the API can act as a bridge between the legacy system and the modern platform, allowing the organization to process data from their legacy system without having to modernize the infrastructure.

upvoted 2 times
- Haplo46** 1 year, 5 months ago

Anthos lets you build and deploy modern apps in your data center, retail edge location, factory, stadium. For me B is the answer

upvoted 1 times
- Haplo46** 1 year, 5 months ago

I think I was wrong writing this comment, the right answer is C

upvoted 1 times
- Atish** 1 year, 5 months ago

Selected Answer: C

they want to process data from their legacy system in a modern platform and the way is to use API

upvoted 2 times

  **gviolant** 1 year, 5 months ago

Selected Answer: B

For me is B - Anthos

upvoted 1 times

Question #183

Topic 1

An international bank is looking for a serverless warehouse solution that lets them perform smart analytics.

Which Google Cloud product or service should the bank use?

- A. Cloud Spanner
- B. Compute Engine
- C. BigQuery
- D. Dataflow

Correct Answer: C

Community vote distribution

C (100%)

  **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

  **__rajan__** 8 months ago

Selected Answer: C

C is correct.

upvoted 1 times

  **SoftSami** 1 year, 3 months ago

Selected Answer: C

The international bank should use Google Cloud's BigQuery service, which is a fully managed, serverless data warehouse that allows for high-speed analysis of large datasets. It provides a range of built-in functions for analytics and can easily integrate with other Google Cloud services.

upvoted 2 times

  **bda92b3** 1 year, 4 months ago

C is correct

Google BigQuery's cloud-based data warehouse and analytics platform uses a built-in query engine and a highly scalable serverless computing model to process terabytes of data in seconds and petabytes in minutes.

upvoted 2 times

How does switching from on-premises to the cloud help organizations gain value over time?

- A. They can expand their internal application hosting infrastructure
- B. They can focus their efforts on solution development
- C. They can increase development of data recovery systems
- D. They can relax their on-premises data security protocols

Correct Answer: B

Community vote distribution

B (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: B

B is correct.

upvoted 1 times
-   **tsk9921** 11 months, 2 weeks ago

B. Moving to cloud helps them not worry about infrastructure anymore but focus on business development

upvoted 2 times
-   **WalterDeWalt** 1 year, 2 months ago

Selected Answer: B

B is correct

upvoted 1 times
-   **SoftSami** 1 year, 3 months ago

Selected Answer: B

B is correct

upvoted 2 times
-   **bda92b3** 1 year, 4 months ago

B is correct

upvoted 1 times

An organization is struggling to meet user demand for change and wants to modernize their legacy applications by moving the applications to the cloud.

Why would this help the organization satisfy user expectations?

- A. Customer data can be used to offer tailored content
- B. Toil automation helps make automatic updates
- C. Updates can be pushed out more quickly to repair bugs
- D. DevOps requires that industry trends be measured and tracked

Correct Answer: C

Community vote distribution

C (59%)A (41%)

- chpann

5 months, 3 weeks ago

Selected Answer: A

A caters to the customer's priority of change
upvoted 2 times
- sinh

7 months ago

Selected Answer: C

Other options can be realized without moving to the cloud.
upvoted 3 times
- gyenesge

7 months, 1 week ago

Selected Answer: A

I think A is correct. The question was not about bugs or quality issues, so C is not relevant. Though C is also an advantage for cloud solution.
upvoted 1 times
- stevee954

7 months, 2 weeks ago

The only one that adds user value directly is A. C: update can just be a fix.
upvoted 2 times
- chai_gpt

7 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times

__rajan__

8 months ago

Selected Answer: C

C. Updates can be pushed out more quickly to repair bugs.

Moving legacy applications to the cloud can help organizations satisfy user expectations by enabling them to push out updates more quickly to repair bugs.
upvoted 1 times

Kaycee101

9 months, 1 week ago

Selected Answer: A

Ans is A
upvoted 1 times

andr31

1 year, 1 month ago

Selected Answer: C

For me it's C because the company is trying to meet customer's demand for change and C covers that by mentioning "updates".
upvoted 2 times

AKhaled

1 year, 2 months ago

Selected Answer: C

for me C
upvoted 1 times

[Removed]

1 year, 3 months ago

Selected Answer: C

I don't like this question. The word "change" is confusing and should be "improvement". Either way, I'm going with C
upvoted 2 times

  **Ronniejsantos** 1 year, 3 months ago

Selected Answer: A

A for me
upvoted 1 times

  **bda92b3** 1 year, 4 months ago

Correct answer C
upvoted 1 times

  **gviolant** 1 year, 4 months ago

Selected Answer: A

for me A
upvoted 2 times

An organization is looking for a storage solution that will help them serve content to users worldwide. They need a solution that offers a high level of availability.

What feature of Cloud Storage would they benefit from?

- A. Multi-regional storage
- B. Global metadata
- C. Data encryption
- D. Object versioning

Correct Answer: A

Community vote distribution

A (100%)

- 

 **taro123** 1 week ago

Selected Answer: A

A is correct.
upvoted 1 times
- 

 **madcloud32** 6 months ago

A is correct.
upvoted 1 times
- 

 **gyenesge** 7 months, 1 week ago

Selected Answer: A

A seems to be the right one.
upvoted 1 times
- 

 **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- 

 **__rajan__** 8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
- 

 **NoCrapEva** 9 months, 4 weeks ago

Selected Answer: A

<https://cloud.google.com/blog/products/storage-data-transfer/choose-between-regional-dual-region-and-multi-region-cloud-storage>
upvoted 1 times
- 

 **Murtuza** 10 months, 3 weeks ago

Multiple Google Cloud services are managed by Google to be redundant and distributed within and across regions. These services optimize availability, performance, and resource efficiency.
upvoted 3 times
- 

 **Murtuza** 10 months, 3 weeks ago

An organization is struggling to ensure that their legacy systems meet rapidly changing IT needs.

Why should the organization use application programming interfaces (APIs)?
upvoted 1 times

What kind of architecture does a hybrid cloud offer organizations?

- A. A combination of different virtualized servers
- B. A combination of serverless computing and public cloud services
- C. A combination of private cloud infrastructure and public cloud services
- D. A combination of multiple public cloud provider services

Correct Answer: C

Community vote distribution

C (100%)

-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C. A combination of private cloud infrastructure and public cloud services

A hybrid cloud architecture combines a private cloud infrastructure with public cloud services. This allows organizations to take advantage of the benefits of both private and public clouds.

upvoted 1 times
-   **tsk9921** 11 months, 2 weeks ago

C. Hybrid definition is combination of on-premises (private) + public cloud

upvoted 3 times
-   **Jkzz** 1 year ago

Answer is C

upvoted 1 times
-   **WalterDeWalt** 1 year, 2 months ago

Selected Answer: C

too easy

upvoted 1 times
-   **[Removed]** 1 year, 3 months ago

Selected Answer: C

C. Nice n easy.

upvoted 1 times
-   **PHD_CHENG** 1 year, 4 months ago

Selected Answer: C

C is correct

upvoted 1 times

An international law firm needs their HR application to be able to share and access varying amounts of sensitive data across their branch offices.

Why would using Cloud Storage be beneficial for them?

- A. It stores geographically dispersed copies of data to prevent loss.
- B. It offers an affordable upgrade that enables data encryption.
- C. It randomly allocates data to different storage locations to protect against threats.
- D. It has a low minimum storage amount requirement.

Correct Answer: B

Community vote distribution

A (53%)	B (32%)	C (16%)
---------	---------	---------

-   **ManishSinha** Highly Voted 1 year, 3 months ago

I think A is better option. One of the primary benefits of using Google Cloud Storage is that it stores geographically dispersed copies of data to prevent data loss. This means that even if one copy of the data is lost or corrupted, other copies stored in different locations can be used to recover the data.

upvoted 6 times
-   **sinh** Most Recent 7 months ago

Selected Answer: B

B is correct.
The question text only mentions the sharing of sensitive data and does not mention the loss described in option A.

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: A

A. It stores geographically dispersed copies of data to prevent loss.

One of the key benefits of Cloud Storage is that it stores geographically dispersed copies of data to prevent loss. This is important for an international law firm that needs to be able to share and access varying amounts of sensitive data across their branch offices.

upvoted 1 times
-   **NoCrapEva** 9 months, 4 weeks ago

Selected Answer: A

By default - Cloud Storage encrypts data at no additional cost (Cloud Storage always encrypts your data on the server side, before it is written to disk, at no additional charge. - <https://cloud.google.com/storage/docs/encryption/>). This sorts out the sensitive data requirement of the HR company.
Data Loss is minimised by having geographically dispersed copies.
Answer B is possible but not required unless the customer wants key storage (HSM units) to secure/manage their cryptographic keys for an additional fee....

upvoted 3 times
-   **Emma412** 10 months ago

Selected Answer: B

Verified Answer = B
Explanation:
1. Data encryption: Cloud Storage uses server-side encryption to encrypt your data by default. You can also use supplemental data encryption options.
Refer- https://cloud.google.com/storage/docs/introduction#securing_your_data
2. Cloud Storage always encrypts your data on the server side, before it is written to disk, at no additional charge. Besides this there are additional ways to encrypt your data when using Cloud Storage.
Refer-<https://cloud.google.com/storage/docs/encryption>
3. If you need more control over key operations than what Google-managed encryption keys allows, you can use customer-managed encryption keys.

upvoted 1 times
-   **Emma412** 10 months ago

3. If you need more control over key operations than what Google-managed encryption keys allows, you can use customer-managed encryption keys. These keys are created and managed using Cloud Key Management Service (Cloud KMS), and you store the keys as software keys, in an

HSM cluster, or externally.
Refer-<https://cloud.google.com/storage/docs/encryption/customer-managed-keys>
upvoted 1 times

  **ddogg** 11 months ago
Question Asks about sensitive data. It must be C
upvoted 2 times

  **ddogg** 11 months ago
I mean B sorry
upvoted 1 times

  **inner4444** 1 year ago
Selected Answer: A
A - the correct one
upvoted 2 times

  **Jkzz** 1 year ago
Selected Answer: A
A - redundancy
upvoted 2 times

  **telosd** 1 year, 1 month ago
Selected Answer: A
A is the better option.
upvoted 1 times

  **Deep9548** 1 year, 2 months ago
Yes, better option is A here.
upvoted 1 times

  **Nguyen25183** 1 year, 3 months ago
someone is going to find the right answer for me pls. The exam is coming soon :((((
upvoted 1 times

  **[Removed]** 1 year, 3 months ago
Selected Answer: C
I don't like B as the answer here. Data is always encrypted. Encryption is not part of an "affordable upgrade". Only other answer that made sense to me is C.
upvoted 3 times

  **guilhermemls** 1 year, 4 months ago
Selected Answer: B
Answer B.
upvoted 2 times

  **Naree** 1 year, 5 months ago
Selected Answer: B
Cloud Storage always encrypts your data on the server side, before it is written to disk, at no additional charge. Besides this standard, Google-managed behavior, there are additional ways to encrypt your data when using Cloud Storage.
upvoted 2 times

A large retail organization uses traditional technology for their ecommerce website. During peaks in traffic, resources are often underutilized or overprovisioned. They have decided to migrate to cloud technology.

What aspect of cloud technology will benefit their ecommerce business?

- A. Agile infrastructure means that they only pay for what they need, when they need it.
- B. Operational expenditure means that their total cost of ownership is more predictable.
- C. Unlimited storage means that their website will never experience downtime.
- D. Shared responsibility means that the cloud provider brings increased visibility during peaks in traffic.

Correct Answer: D

Community vote distribution

A (100%)

- kakaomg**

Highly Voted

1 year, 5 months ago

A, Pay as you go
upvoted 6 times
- chai_gpt**

Most Recent

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__**

8 months ago

Selected Answer: A

A is correct.
upvoted 1 times
- Emma412**

10 months ago

Selected Answer: A

Suggested Answer D is incorrect on many counts
1. Creating visibility is first of all not cloud providers responsibility.
2. Secondly their responsibility is to secure the underlying infrastructure even as per definition of "shared responsibility"

The correct Answer must be = A
upvoted 1 times
- reachmalhotras**

11 months, 2 weeks ago

Selected Answer: A

Trick question, Moderator please check / correct the answer.
upvoted 1 times
- Jkzz**

1 year ago

A - pay for what you use
upvoted 2 times
- telosd**

1 year, 1 month ago

Selected Answer: A

A (pay as you go model????!!!!)
upvoted 1 times
- Sozan**

1 year, 2 months ago

A, PAYG
upvoted 1 times
- tbolick6**

1 year, 3 months ago

Selected Answer: A

Pay as you go.
upvoted 1 times
- lauraperdomo04**

1 year, 4 months ago

Selected Answer: A

Its A, PAYG
upvoted 1 times

  **WzWz** 1 year, 4 months ago

Selected Answer: A

In this case, it's A
upvoted 1 times

  **AK10000** 1 year, 4 months ago

Its A.
upvoted 1 times

  **guilhermemls** 1 year, 4 months ago

Selected Answer: A

A according to Pay as you go model.
upvoted 1 times

  **DeNel** 1 year, 4 months ago

Selected Answer: A

Pay as you go
upvoted 1 times

  **DPonly** 1 year, 5 months ago

its option A
upvoted 1 times

  **Haplo46** 1 year, 5 months ago

"auto-scaling" allows the organization to optimize their usage of resources and costs. Pay as you go is the right answer for me, option A
upvoted 3 times

An organization is looking for a hybrid-cloud management solution that will help them build and run applications on-premises and in the cloud.

What solution should the organization use?

- A. Anthos
- B. Cloud Functions
- C. App Engine
- D. Compute Engine

Correct Answer: A

Community vote distribution

A (100%)

- sivakarthick16** 5 months ago

Selected Answer: A

Anthos is a platform offered by Google Cloud that enables organizations to build and run applications both on-premises and in the cloud. It provides a consistent and unified experience across different environments, allowing for seamless application development and deployment. With Anthos, organizations can leverage the benefits of both on-premises infrastructure and cloud services, enabling flexibility and scalability in their application deployment strategies.

upvoted 1 times
- madcloud32** 6 months ago

A is correct. Anthos is for hybrid cloud.

upvoted 1 times
- chai_gpt** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times
- __rajan__** 8 months ago

Selected Answer: A

A is correct.

upvoted 1 times
- WalterDeWalt** 1 year, 2 months ago

Selected Answer: A

Anthos

upvoted 2 times
- Sozan** 1 year, 2 months ago

a is correct

upvoted 1 times
- bda92b3** 1 year, 4 months ago

Correct

upvoted 1 times

An organization with a commissions-based business model is struggling to predict cloud costs due to fluctuating revenue.

How can the organization optimize the use of their cloud resources?

- A. By applying intelligent recommendations
- B. By decentralizing financial decision making
- C. By sharing project ownership across all teams
- D. By increasing their capital expenditure

Correct Answer: A

Community vote distribution

A (67%)

B (33%)

- chai_gpt

7 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times
- __rajan__

8 months ago

Selected Answer: A

A. By applying intelligent recommendations

Intelligent recommendations can help organizations to optimize the use of their cloud resources by identifying and suggesting ways to reduce costs.
upvoted 1 times
- Emma412

10 months ago

Selected Answer: A

Answer = A
Explanation:
Refer the section "recommendations" under cost management
"View intelligent recommendations for optimizing your costs and usage. Easily apply these changes for immediate cost savings and greater efficiency."
<https://cloud.google.com/cost-management>
upvoted 1 times
- aws_munich

11 months, 3 weeks ago

Answer is A. Use GCP Recommender - <https://cloud.google.com/blog/products/management-tools/using-recommenders-keep-your-cloud-running-optimally>
upvoted 1 times
- CarloGiovanni

1 year, 1 month ago

Answer is A
upvoted 1 times
- telosd

1 year, 1 month ago

Selected Answer: A

apply recommendations
upvoted 2 times
- yizza

1 year, 2 months ago

Selected Answer: A

The organization can optimize the use of their cloud resources by applying intelligent recommendations. This means using tools and services that can analyze their cloud usage and identify areas where they can save money. For example, they could use a tool that can automatically shut down unused resources or that can move workloads to cheaper regions.

Decentralizing financial decision making, sharing project ownership across all teams, and increasing their capital expenditure would not necessarily help the organization optimize the use of their cloud resources. In fact, these options could make it more difficult to track and manage cloud costs
upvoted 3 times
- tbolick6

1 year, 2 months ago

Selected Answer: B

Decentralize

upvoted 1 times

  **Sozan** 1 year, 2 months ago

Selected Answer: B

B is the correct answer

upvoted 1 times

  **[Removed]** 1 year, 3 months ago

Selected Answer: B

Answer is B

upvoted 1 times

  **tbolick6** 1 year, 3 months ago

Selected Answer: B

Decentralize

upvoted 1 times

What DevOps practice should an organization use when developing their application to help minimize disruption caused by bugs?

- A. Pause production until all bugs have been eliminated
- B. Prioritize fixing large bugs during production because they are easier to review
- C. Implement small changes incrementally to reduce recovery time when bugs appear
- D. Implement large changes together to make rolling back easier when bugs appear

Correct Answer: C

Community vote distribution

C (100%)

-   **taro123** 1 week ago

Selected Answer: C

C is correct

upvoted 1 times
-   **madcloud32** 6 months ago

Selected Answer: C

C is correct

upvoted 1 times
-   **chai_gpt** 7 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times
-   **__rajan__** 8 months ago

Selected Answer: C

C. Implement small changes incrementally to reduce recovery time when bugs appear

One of the key principles of DevOps is to release changes frequently and in small batches. This helps to reduce the risk of disruption caused by bugs. If a bug is introduced in a small change, it is easier to identify and fix the bug without affecting a large number of users.

upvoted 1 times
-   **oriolman** 10 months, 2 weeks ago

Selected Answer: C

C, no doubt

upvoted 1 times
-   **Murtuza** 10 months, 3 weeks ago

Deploying changes to a small subset of users before rolling out to production (also known as a canary deployment) can help catch bugs before they impact the entire user subset

upvoted 1 times

An organization has been struggling to make operations more efficient.

What site reliability engineering (SRE) best practice should the organization follow to increase efficiency?

- A. Decrease over-reliance on data to make decisions.
- B. Assign exclusive production ownership to developers.
- C. Spend less time measuring outage impact.
- D. Focus on increasing toil automation.

Correct Answer: A

Community vote distribution

D (100%)

-   **Anna3456789** 2 weeks, 4 days ago

The answer is D. Toil refers to repetitive, manual operational work that does not provide long-term value and can be automated. By focusing on increasing toil automation, the organization can streamline operations, reduce manual intervention, and free up resources to focus on more strategic tasks. This approach improves efficiency by automating repetitive tasks, reducing human error, and allowing teams to focus on higher-value activities.

upvoted 2 times
-   **bad5fad** 1 month, 1 week ago

Selected Answer: D

The Answer is D

upvoted 3 times

An organization wants to ingest custom log data from GKE environments, virtual machines, and Google Cloud services.

Which Google Cloud tool should the organization use?

- A. Dialogflow
- B. Cloud Logging
- C. Cloud SDK
- D. Data Catalog

Correct Answer: B

What makes Google Kubernetes Engine (GKE) an effective solution for developers working to resolve application errors?

- A. It reduces the time needed to iterate on various solutions.
- B. It provides intelligent recommendations to optimize application development.
- C. It removes the need for containerizing applications.
- D. It stores application code on multiple servers to facilitate collaboration across silos.

Correct Answer: B

Community vote distribution

A (100%)

-   **bad5fad** 1 month, 1 week ago

Selected Answer: A

A is the correct answer

upvoted 2 times
-   **Vivek007** 1 month, 3 weeks ago

A: GKE's features like automated rollouts, rollbacks, and its managed environment can significantly reduce the time developers need to iterate on and deploy various solutions. The integration with CI/CD pipelines and tools also supports rapid iterations.

upvoted 4 times

What is an organization responsible for when migrating from on-premises to the cloud?

- A. Covering the cost of cloud service downtime
- B. Managing underlying network infrastructure
- C. Adapting to a pay-as-you-go cloud expenditure model
- D. Optimizing cloud storage systems

Correct Answer: D

Community vote distribution

C (100%)

-   **76ffe88** 1 month, 2 weeks ago

Selected Answer: C

it should be c pay as you go model

upvoted 2 times
-   **Vivek007** 1 month, 3 weeks ago

C. Adapting to a pay-as-you-go cloud expenditure model

Correct. Moving to the cloud often involves changing how IT expenditures are accounted for from a capital expenditure model (CapEx) to an operational expenditure model (OpEx). In the cloud, organizations typically pay for services as they use them, which requires adapting budgeting and financial management practices to a pay-as-you-go model.

upvoted 3 times

An organization wants to duplicate critical system components to enhance reliability and mitigate single points of failure. Which design consideration should the organization use?

- A. Redundancy
- B. Backups
- C. Latency
- D. Security

Correct Answer: A

An organization stores backup files in Cloud Storage. The files will be accessed annually to test the disaster recovery plans. Which storage class is the most cost-effective?

- A. Coldline class
- B. Standard class
- C. Nearline class
- D. Archive class

Correct Answer: D

  **Anna3456789** 2 weeks, 4 days ago

Can someone help? I'm between C. Nearline and D. Archive.
upvoted 1 times

An organization is running SQL Server on-premises and is struggling with capacity and management overhead. They want to modernize this database quickly by using Google products or services. What should the organization do?

- A. Refactor applications to use a cloud first database like Firestore.
- B. Perform a managed database migration to Cloud SQL.
- C. Export old tables from SQL Server to Cloud Storage.
- D. Migrate all SQL Server data to BigQuery.

Correct Answer: B

  **Vivek007** 1 month, 3 weeks ago

B: Cloud SQL is Google Cloud's fully managed database service that supports SQL Server, among other databases like MySQL and PostgreSQL. By migrating to Cloud SQL, the organization can leverage a managed service that handles much of the database administration tasks such as patching, updates, and setup, thereby reducing management overhead.

Cloud SQL also offers automatic scaling capabilities, addressing the concern about capacity that the organization is facing with their on-premise solution.

Migrating to Cloud SQL would involve minimal changes to the existing applications since it supports SQL Server, allowing for a smoother transition without needing extensive application refactoring or redevelopment.

upvoted 1 times

What is a benefit of the OpEx model for cloud security?

- A. Organizations do not need to make upfront capital investments in cloud security.
- B. Organizations can deploy custom security hardware.
- C. Organizations do not need to configure any security settings for cloud resources.
- D. The cloud provider guarantees security.

Correct Answer: A

What is the benefit of using a serverless data processing pipeline service?

- A. Full control over compute resources is provided.
- B. Processed data will not require analysis.
- C. Pipeline infrastructure is fully managed and scalable.
- D. Processed data is guaranteed to be free of errors.

Correct Answer: C

  **Vivek007** 1 month, 3 weeks ago

C: Managed and Scalable: The primary advantage of serverless data processing services is that they abstract the underlying infrastructure management away from the user. This means the cloud provider manages the scaling, maintenance, patching, and availability of the servers. The service automatically scales up or down based on the workload, which can significantly reduce operational complexity and cost.

upvoted 1 times

What does the shift toward cloud computing represent for an organization's transformation?

- A. An opportunity that is limited to large enterprises
- B. An opportunity to redefine existing business processes and services
- C. An opportunity that is only relevant to the IT department
- D. An opportunity to continue business as usual with new cost savings

Correct Answer: B

An organization is concerned about the unlikely event that Google Cloud infrastructure is physically accessed by someone with malicious intent. How is data protected in Google Cloud?

- A. Data is immediately deleted whenever an intrusion is detected.
- B. Data is stored on quantum computers with unbreakable encryption.
- C. Data is stored using robust encryption.
- D. Data is stored in random locations around the world to prevent it being found.

Correct Answer: C

  **Vivek007** 1 month, 3 weeks ago
C: Robust Encryption: Google Cloud uses multiple layers of encryption to protect data at rest and in transit. This means that even if someone were to physically access the hardware, the data stored would still
upvoted 2 times

An organization needs a flexible and scalable NoSQL database with strong web and mobile application support. Which Google Cloud product or service should the organization use?

- A. Cloud Spanner
- B. Cloud Storage
- C. BigQuery
- D. Firestore

Correct Answer: D

  **Vivek007** 1 month, 3 weeks ago
D: Firestore (formerly known as Cloud Firestore) is a highly scalable, serverless, NoSQL document database built for automatic scaling, high performance, and ease of application development. It provides powerful querying and real-time data synchronization capabilities, making it ideal for web and mobile applications that require a responsive and dynamic user experience. Firestore also integrates well with other Firebase features, enhancing its suitability for developing web and mobile applications.
upvoted 2 times

An organization is using three cloud vendors to maximize their available deployment locations worldwide. They are using GKE Enterprise to deploy Kubernetes applications across different clouds. What type is this deployment?

- A. On-premises
- B. Multi-cloud
- C. Hybrid-cloud
- D. Private cloud

Correct Answer: B

An organization wants to transfer some of its data from Google Cloud. Which of these statements is true?

- A. Customer data may not be transferred out of Google Cloud.
- B. Customers have full control of their data and may transfer it at any time.
- C. Outgoing data transfer must be enabled in the Google Cloud console.
- D. A technical support ticket must be raised with the correct department.

Correct Answer: B

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

B: Customer Data Control: One of the core principles of cloud services like Google Cloud is that customers retain ownership and control of their data. This means that customers can export or transfer their data whenever they need to, without any restrictions from the cloud provider. Google Cloud provides various tools and services that facilitate data export and transfer, supporting the organization's autonomy over their data.
upvoted 1 times

A vacation home rental organization wants to predict the popularity of properties in their upcoming busy season. They do not have a data science team, and want to use their in-house database administration skills to create a machine learning model. What should the organization do?

- A. Use custom training in Vertex AI with TensorFlow.
- B. Integrate pre-trained APIs into their application.
- C. Use BigQuery ML and create models using SQL.
- D. Build a model in AutoML using labeled data.

Correct Answer: C

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

C: BigQuery ML enables users with SQL skills to create and execute machine learning models directly within Google BigQuery. This approach is particularly suitable for organizations like the one described, which may not have specialized machine learning expertise but does have database administration skills. BigQuery ML allows these users to leverage their existing SQL knowledge to build and train predictive models, making it an accessible option for generating insights from their data.
upvoted 1 times

An organization needs to rapidly scale its use of computing resources and honor its commitment to environmental sustainability.

What should the organization do?

- A. Refactor application software to use less energy.
- B. Use a public cloud provider with energy-efficient data centers.
- C. Use a carbon-neutral energy provider for an existing on-premises data center.
- D. Purchase energy-efficient servers for an existing on-premises data center.

Correct Answer: *B*

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

B: Energy-Efficient Data Centers: Major public cloud providers invest heavily in optimizing their data centers for energy efficiency and often use renewable energy sources. These companies leverage advanced technologies to reduce power consumption and have the capital to invest in large scale sustainable energy projects. By using a public cloud provider known for energy-efficient practices, the organization can scale computing resources more flexibly and sustainably compared to maintaining their own data center.

upvoted 1 times

What is a defining feature of a non-relational database?

- A. Reporting across multiple data sources
- B. A strictly enforced schema
- C. A flexible data model
- D. Queries that join multiple tables

Correct Answer: *C*

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

C: Flexible Data Model: One of the key characteristics of non-relational databases (also known as NoSQL databases) is their flexible schema. Unlike relational databases, which require data to conform to a strict schema (defined data types and structures), non-relational databases allow data to be stored without predefined schema. This flexibility makes it easier to accommodate changes in data structure and the variety of data types typically found in big data and real-time web applications.

upvoted 1 times

What is an example of unstructured data?

- A. Historical stock inventory
- B. Product ratings
- C. Customer orders
- D. Call center transcripts

Correct Answer: D

  **Vivek007** 1 month, 3 weeks ago
D: Call Center Transcripts: This type of data is considered unstructured because it consists of free-form text that does not follow a predictable data model. Call center transcripts are typically generated from spoken language and recorded during customer service calls, containing varied vocabulary, sentence structures, and topics that do not fit neatly into pre-defined database fields.
upvoted 1 times

An organization has migrated all workloads to the cloud and is reviewing their cloud security posture. Who is now responsible for securing the physical infrastructure of the data centers?

- A. The organization and the cloud service provider
- B. Third-party security service providers
- C. The cloud service provider
- D. The organization

Correct Answer: C

An organization wants to lease the resources they need for their customized servers from a cloud provider on a pay-as-you-go basis, instead of paying one time for hardware.

Which service model should they use?

- A. Hybrid cloud
- B. Software as a service
- C. Infrastructure as a service
- D. Platform as a service

Correct Answer: C

An organization has petabytes of data gathered from a wide range of sources. They want to use the data for strategic analysis and to guide business decisions. What type of service should they use?

- A. Multi cloud environment
- B. Virtual machine environment
- C. Hybrid cloud environment
- D. Container environment

Correct Answer: C

-   **Vivek007** 1 month, 3 weeks ago

D: Container Environment: While containers themselves aren't specifically for handling data analytics, they are part of an infrastructure that can efficiently manage and scale data-intensive applications across various environments. Container orchestration platforms like Kubernetes can manage containers that run big data tools and applications, providing flexibility, scalability, and efficiency. This allows for the deployment of complex applications and data processing tasks that can handle large volumes of data effectively.

upvoted 3 times
-   **Hussiengamal** 6 days, 15 hours ago

So it will be C ?

upvoted 1 times

An organization wants a purpose-built AI solution to increase efficiency and provide personalized interactions for their customer care team. Which Google Cloud AI solution should they use?

- A. Text-to-Speech API
- B. Cloud Talent Solution
- C. Document AI
- D. Contact Center AI

Correct Answer: D

An organization wants to use an open source library with a flexible ecosystem of tools to create and train its own machine learning models. Which product or solution should the organization use?

- A. Cloud Functions
- B. Apache Beam
- C. Dataflow
- D. TensorFlow

Correct Answer: D

  **Vivek007** 1 month, 3 weeks ago

D: TensorFlow is an open-source machine learning library developed by Google. It provides a comprehensive, flexible ecosystem of tools, libraries, and community resources that let researchers push the state-of-the-art in ML, and developers easily build and deploy ML-powered applications. TensorFlow supports deep learning along with a suite of tools for model building, training, and deployment, making it ideal for a wide variety of machine learning applications across different platforms and languages.

upvoted 2 times

An organization wants to leverage cloud technologies but is concerned about vendor lock-in.

What would mitigate this concern?

- A. Open standards
- B. Database services
- C. Service level agreements
- D. Scalable infrastructure

Correct Answer: A

  **Vivek007** 1 month, 3 weeks ago

A: Use Open Standards and Technologies

Adopting open standards and technologies ensures that the organization's systems and data can be easily ported to different platforms. This includes using open-source software, open APIs, and standard protocols that are universally supported across cloud providers.

upvoted 1 times

An organization must identify and fix security vulnerabilities in its cloud infrastructure and applications. Which Google Cloud service should they use?

- A. Security Command Center
- B. Google Cloud Armor
- C. Cloud Storage
- D. VPC networks

Correct Answer: A

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

A: Security Command Center is Google Cloud's comprehensive security and risk management tool. It provides security analytics, vulnerability scanning, and threat detection to help organizations secure their cloud infrastructure, applications, and data. The Security Command Center enables visibility across the organization's assets in the cloud, identifies vulnerabilities and misconfigurations, and recommends corrective actions to enhance security. This tool is particularly valuable for continuous monitoring, compliance reporting, and gaining insights into potential security threats.

upvoted 1 times

A cinema company wants to build a model to predict customer visit patterns for the coming year. They have three years of customer visit data across 300 theaters; however, the data has been stored in different formats by different theaters. They must train the ML model. What should they do?

- A. Use the last year of data so there are fewer inconsistencies for the model to handle.
- B. Transform the data into a consistent format.
- C. Group different format types and train a different model for each group.
- D. Choose an ML model type that can process different formats of input data.

Correct Answer: B

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

B: Data Consistency: Before any effective machine learning modeling can occur, it is crucial that the data used is consistent and standardized. This means transforming all the data collected from different theaters into a uniform format. This process, often called data cleaning or preprocessing, ensures that the machine learning model has reliable and comparable input data across all samples, leading to more accurate and meaningful predictions.

upvoted 2 times

Within Google's Site Reliability Engineering framework, which concept measures how well a system is performing?

- A. Service-level agreement
- B. Service-level indicator
- C. Error reporting
- D. Service-level objective

Correct Answer: B

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

B: Service-Level Indicator (SLI): An SLI is a carefully defined quantitative measure of some aspect of the level of service that is provided. Most often an SLI is a metric or combination of metrics that directly reflect the performance, reliability, and overall health of a system. Common examples include uptime, response time, latency, error rates, and throughput. SLIs are used to objectively measure how effectively a service is meeting desired performance standards.

upvoted 1 times

What is the typical cloud spending behavior of most organizations?

- A. Decentralized and variable
- B. Centralized and variable
- C. Decentralized and fixed
- D. Centralized and fixed

Correct Answer: A

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

A: Decentralized and Variable: Cloud spending in many organizations tends to be decentralized and variable due to the nature of cloud services. Decentralized because different departments or teams within an organization often have the autonomy to procure and manage their own cloud resources based on their specific needs. Variable because cloud costs are typically based on usage (pay-as-you-go pricing models), which means spending fluctuates based on how much compute, storage, and network resources are consumed over time.

upvoted 1 times

An organization is developing a new container-based application. They do not know how popular the application will be when launched and they do not want to pay for idle infrastructure resources. Which benefit of serverless computing will address this concern?

- A. Disaster recovery
- B. Reduced development costs
- C. Scalability
- D. Built-in security

Correct Answer: C

  **jjkjjkkk** 1 month, 3 weeks ago
Scalability is correct
upvoted 1 times

How does the resource hierarchy in Google Cloud help organizations implement security policies?

- A. Policies can be applied at the folder level and are inherited by projects inside the folder.
- B. Projects in the resource hierarchy are not affected by security policies.
- C. Policies can only be applied at the organization level and affect all projects.
- D. Policies can only be applied to individual projects.

Correct Answer: A

Community vote distribution

A (100%)

  **shanwford** 3 weeks, 3 days ago

Selected Answer: A

Take care of the word "only" in C & D - so A is fine.

upvoted 1 times

  **Vivek007** 1 month, 3 weeks ago

A: Inheritance in the Resource Hierarchy: In Google Cloud's resource hierarchy, security policies can be applied at various levels, including the organization, folder, and project levels. A key aspect of this hierarchy is the inheritance of policies. When a policy is applied at the folder level, all projects and resources within that folder automatically inherit the policy. This feature simplifies security management by allowing organizations to apply broad policies at higher levels (like folders) and have them cascade down to all contained resources and projects. This hierarchical policy management ensures that all underlying resources comply with the organization's security protocols without the need to configure each project individually.

upvoted 1 times

An organization is considering the use of managed services when migrating to the cloud. Which routine tasks are typically provided automatically by a managed services platform?

- A. Managing user access
- B. Patching and upgrades
- C. Data archiving
- D. Cost optimization

Correct Answer: B

An organization has recently completed a migration from on-premises to Google Cloud.

How has cost management been affected?

- A. Costs will primarily shift from CapEx to OpEx.
- B. Costs will primarily shift from OpEx to CapEx.
- C. Cost management will stay the same, but the total cost of ownership (TCO) will be lower.
- D. Cost management will stay the same, but the total cost of ownership (TCO) will be higher.

Correct Answer: A

An organization wants to access a software application from a cloud vendor without the need to manage their own servers or write their own code.

Which service model does this represent?

- A. Infrastructure as a service
- B. Platform as a service
- C. Software as a service
- D. Functions as a service

Correct Answer: C

An organization wants to build a data pipeline to transform its data so it can be reconciled in a data warehouse. The solution must be scalable and require little or no management. Which Google product or service should the organization choose?

- A. Cloud Bigtable
- B. Cloud Storage
- C. Pub/Sub
- D. Dataflow

Correct Answer: *D*

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

D: Google Cloud Dataflow is a fully-managed service designed to simplify the tasks of processing large amounts of data in both batch and stream modes. Dataflow is especially suited for data integration and ETL (extract, transform, load) tasks, making it ideal for preparing data for a data warehouse. It is built on Apache Beam, providing a unified programming model to define data processing pipelines. The scalable nature of Dataflow means it can adjust resource allocation dynamically based on the workload, and its managed service ensures that the overhead of managing server infrastructure is minimal.

upvoted 2 times

When is data automatically encrypted in Google Cloud?

- A. When it is in transit only,
- B. When it is at rest and in transit.
- C. Data is not automatically encrypted by default.
- D. When it is at rest only.

Correct Answer: *B*

🗨️ 👤 **Vivek007** 1 month, 3 weeks ago

B: Automatic Encryption in Google Cloud: Google Cloud automatically encrypts data both when it is at rest and in transit. This means that:
At Rest: Data stored on disk in Google Cloud services, including but not limited to Google Cloud Storage, Compute Engine, and BigQuery, is encrypted before it is written to disk and decrypted when accessed.

In Transit: Data moving between user devices and Google data centers, as well as data moving within Google's network between data centers, is encrypted. This encryption protects data as it travels over the internet, ensuring that sensitive information is securely transmitted.

upvoted 1 times

An organization must verify the identity of a user seeking access to a system. Which aspect of cloud identity management does this action describe?

- A. Authorization
- B. Authentication
- C. Auditing
- D. Encrypting

Correct Answer: B

An organization is running Kubernetes applications across multiple cloud environments. They want a consistent and centralized management platform. Which service should they choose?

- A. Cloud Functions
- B. GKE Enterprise
- C. Cloud Run
- D. Compute Engine

Correct Answer: B

  **Vivek007** 1 month, 3 weeks ago
B: GKE Enterprise, also known as Anthos GKE or just Anthos, is a platform designed to enable the deployment and management of applications across various cloud environments, including both Google Cloud and other public or private clouds. Anthos provides a unified management interface that simplifies the operation of Kubernetes clusters across different environments, ensuring consistency, compliance, and security are maintained across the board.
upvoted 1 times

An organization wants to collect and store all logs generated by applications running in Google Cloud. Which service should they use?

- A. Cloud Profiler
- B. Cloud Monitoring
- C. Cloud Trace
- D. Cloud Logging

Correct Answer: D

  **Vivek007** 1 month, 3 weeks ago
D Cloud Logging (formerly known as Stackdriver Logging) is part of Google Cloud's operations suite and is specifically designed to handle the aggregation, analysis, and storage of log data and events from Google Cloud services and virtual machine instances. It provides real-time log management and analysis, allowing organizations to easily view, search, and alert on log data across all their Google Cloud resources.
upvoted 1 times

What is the benefit of using a unified cloud data solution?

- A. Data will always cost less to store.
- B. Data can be automatically secured from external threats.
- C. Data will always be backed up and cannot be lost or deleted.
- D. Data can enable innovation because it's no longer siloed.

Correct Answer: *D*

An organization stores its important industry data in a relational database. They want to create a new revenue stream by enabling third parties to use that data in their applications. Which cloud first approach should the organization choose?

- A. Add third-party users to their database.
- B. Offer chargeable downloads of archived data.
- C. Expose data through a chargeable API.
- D. Transfer data into a non-relational database.

Correct Answer: *C*

Which scenario is a good use case for machine learning?

- A. Classifying data with no prior examples
- B. Tasks that require human experience and intuition
- C. Creating customer recommendations
- D. Solving ethical dilemmas

Correct Answer: C

Community vote distribution

C (50%)B (50%)

-   **Anna3456789** 2 weeks, 3 days ago

Selected Answer: C

C. Machine learning is well-suited for tasks that involve pattern recognition, prediction, and decision-making based on data. Creating customer recommendations is a classic example of a use case where machine learning algorithms can be effectively applied.

It is not B because machine learning can augment human decision-making and automate repetitive tasks, but it may not fully replace tasks that require human experience, intuition, and subjective judgment, such as creative problem-solving, critical thinking, and ethical decision-making.

upvoted 1 times
-   **Biswada** 2 weeks, 6 days ago

Selected Answer: B

I think its B

upvoted 1 times

An organization needs to increase the speed at which they can train machine learning models. Which domain-specific hardware is designed for this task?

- A. Bare Metal Solution
- B. Preemptible or Spot VMs
- C. Cloud TPUs
- D. Containers

Correct Answer: C

Community vote distribution

C (100%)

-   **Anna3456789** 2 weeks, 3 days ago

Selected Answer: C

C. Cloud TPUs are domain-specific hardware designed specifically for accelerating machine learning model training tasks.

upvoted 1 times

An organization has migrated several large databases to the cloud. It wants to increase the value of its data, improve cost controls and strengthen regulatory compliance. What should the organization do?

- A. Delete data stored for over a year.
- B. Export relational data to modern NoSQL databases.
- C. Create monthly reports on data access and uses.
- D. Establish an effective data governance program.

Correct Answer: D

An organization has hired a team of data scientists and developers. They want to create unique value in their business by coding an advanced machine learning model in Vertex AI Workbench. Which service should the organization use to train the model?

- A. Custom training
- B. AutoML
- C. Compute Engine
- D. Prebuilt APIs

Correct Answer: A

  **Vivek007** 1 month, 3 weeks ago

A: Custom training allows organizations to train machine learning models using their own custom data and algorithms. This approach is suitable for cases where organizations have specific requirements, unique datasets, or specialized machine learning tasks that cannot be addressed by pre-built solutions. With custom training, organizations have full control over the model architecture, hyperparameters, and training process, enabling them to create highly tailored and optimized machine learning models to address their specific business needs.

upvoted 2 times

An organization is concerned that one of their applications takes too long to return a result. According to Google’s “Four Golden Signals,” which signal measures this aspect of the application’s performance?

- A. Saturation
- B. Errors
- C. Latency
- D. Traffic

Correct Answer: C

An organization is transforming their raw data into a format that can be used to derive business insights. Which step of the data value chain does this action represent?

- A. Data analysis
- B. Data collection
- C. Data storage
- D. Data processing

Correct Answer: D

An organization has a small development team that has created a web application which runs in a single container. They need a simple, serverless, and scalable way to host their container. Which Google service should the organization use?

- A. Compute Engine
- B. App Engine
- C. Kubernetes Engine
- D. Cloud Run

Correct Answer: D

  **Vivek007** 1 month, 3 weeks ago

D: Cloud Run is a fully managed compute platform that enables organizations to run stateless containers in a serverless environment. With Cloud Run, developers can deploy containerized applications quickly and easily without worrying about managing infrastructure. Cloud Run automatically scales the application based on incoming traffic, ensuring that resources are allocated efficiently to handle demand. This makes it an ideal choice for hosting containerized web applications, providing simplicity, scalability, and cost-effectiveness.

C. Kubernetes Engine

Kubernetes Engine is a managed Kubernetes service that enables organizations to deploy, manage, and scale containerized applications using Kubernetes. While it offers scalability and management features, it may be more complex and require more configuration than Cloud Run for a small development team seeking a simpler solution.

upvoted 1 times

An organization wants a centralized view of their cloud infrastructure in a fully managed system that includes uptime checks. Which Google Cloud service should they use?

- A. Cloud Trace
- B. Cloud Monitoring
- C. Cloud Logging
- D. Cloud Profiler

Correct Answer: B

An organization wants to migrate a workload to the cloud without changing the application code or architecture. Which migration path describes this approach?

- A. Rehosted
- B. Reimagined
- C. Replatformed
- D. Refactored

Correct Answer: A

Which scenario uses machine learning to unlock business value from unstructured data?

- A. Analyzing tabular records of product defects to predict future maintenance cycles.
- B. Recommending new products based on previous purchases.
- C. Monitoring financial transactions to identify potential fraud and risk.
- D. Determining customer sentiments from call center voice recordings.

Correct Answer: D

Community vote distribution

D (100%)

  **jancsn** 1 month, 3 weeks ago

Selected Answer: D

D. Determining customer sentiments from call center voice recordings.

This scenario uses machine learning to analyze unstructured data (voice recordings) to extract insights (customer sentiments). By understanding customer sentiments, businesses can make informed decisions to improve their products or services, enhance customer satisfaction, and ultimately drive business value.

upvoted 1 times

  **Vivek007** 1 month, 3 weeks ago

D: This scenario involves processing and analyzing voice recordings, which are a form of unstructured data. Machine learning techniques, particularly those related to natural language processing (NLP) and speech analysis, are used to analyze the voice data to extract insights about customer sentiments. This application of machine learning helps businesses understand their customers better, improve customer service, and tailor their strategies accordingly.

upvoted 1 times

An organization wants to build custom machine learning models. They require a managed platform that provides services to gather data, build models, and then deploy and monitor those models. Which service should they use?

- A. Document AI
- B. Natural Language API
- C. Kubernetes Engine
- D. Vertex AI

Correct Answer: D

In Google's cloud security model, how does availability contribute to a robust security posture for data?

- A. By restricting data access to authorized users
- B. By ensuring data meets industry standards
- C. By checking that data is accurate and trustworthy
- D. By ensuring data is reliable and accessible

Correct Answer: D

How does the legal concept of data sovereignty affect data?

- A. Data is subject to the laws and regulations of the country where it resides.
- B. A country has the right to access the data generated within its borders.
- C. An individual has the right to control their personal data.
- D. Data must always be encrypted in transit and at rest.

Correct Answer: A

An organization is concerned about their cloud costs. They want to be informed when their spending exceeds a specific threshold, rather than waiting to see their bill at the end of the month. What should the organization do?

- A. Pause virtual machines during non-business hours.
- B. Configure a budget threshold rule and alert.
- C. Adjust project resource quota policies.
- D. Use historical cost data to predict future overspend.

Correct Answer: *B*

An organization is running critical workloads in production and requires a Google Cloud support service with fast response times and a dedicated Technical Account Manager. Which customer care service level should the organization choose?

- A. Enhanced
- B. Basic
- C. Premium
- D. Standard

Correct Answer: *C*

A real estate organization processes photos of properties uploaded by their agents to an internal application. They must manually check and remove photos that contain personally identifiable information like passports and credit cards. They want a solution that their developer can use to integrate with their app as quickly as possible. Which solution should they use?

- A. Pre-trained APIs
- B. AutoML
- C. Custom training
- D. BigQuery ML

Correct Answer: *A*

What is generative AI?

- A. A machine learning model for language translation
- B. A specific field of AI that can be used to create content
- C. An umbrella term that includes all types of machine learning algorithms
- D. An AI model for labeling images

Correct Answer: B

What is a defining function of streaming analytics?

- A. Processing data records continuously
- B. Processing data records in batches
- C. Processing a one-off data backfill
- D. Accessing data with high latency

Correct Answer: A

An organization is deciding on the layout of their resource hierarchy in Google Cloud. They have several projects within a folder. What will happen when user access policies are applied to the folder?

- A. The policy will be inherited by the projects in the folder but will not affect their resources.
- B. The policy will be inherited by the projects and their resources within the folder.
- C. The policy will be applied to all folders within the organization.
- D. The policy applies to the folder only, and will not be inherited by any projects.

Correct Answer: B

An organization has an ecommerce application that requires a cost-effective, transactional database. The application will only serve customers in a single region. Which service should they use?

- A. BigQuery
- B. Cloud Bigtable
- C. Cloud Spanner
- D. Cloud SQL

Correct Answer: *D*

An organization wants to control what types of network traffic are allowed to enter and leave its network and access its applications. Which security measure should the organization use?

- A. 2-Step Verification
- B. Privileged access
- C. Encryption keys
- D. Firewall rules

Correct Answer: *D*

An organization has a large archive of unstructured data, including video and audio files. Which storage solution should the organization use?

- A. Cloud Spanner
- B. Cloud Storage
- C. Cloud Bigtable
- D. Cloud SQL

Correct Answer: *B*

A retail organization is training a model to recommend products to customers for an ecommerce website. The model was trained on previous purchases, but did not include demographic information on each buyer. What dimension of the data is responsible for the model's poor performance?

- A. Validity
- B. Accuracy
- C. Timeliness
- D. Completeness

Correct Answer: D

An organization has collected petabytes of historical data. They need an advanced analysis solution that is fast, scalable, and fully managed. Which Google product or service should the organization use?

- A. BigQuery
- B. Cloud SQL
- C. Firestore
- D. Cloud Storage

Correct Answer: A

What is a benefit of Google's purpose-built servers compared to standard servers?

- A. They are optimized for specific tasks making them more efficient.
- B. They are backward compatible with legacy disk drives.
- C. They are cheaper to build than standard servers.
- D. They run software that cannot be deployed on standard servers.

Correct Answer: A

What is the Site Reliability Engineering (SRE) term for an organization's desired level of reliability and performance?

- A. Service-level objective
- B. Enhanced support
- C. Scalable infrastructure
- D. Service-level indicator

Correct Answer: A

An organization is operating multiple workloads in containers and requires full control of how the workloads are configured. Which Google Cloud service should the organization use?

- A. Cloud Run
- B. Compute Engine
- C. Kubernetes Engine
- D. Cloud Functions

Correct Answer: C

When customer data is uploaded to Google Cloud, who owns the data?

- A. A third party
- B. The customer and Google share ownership
- C. Google
- D. The customer

Correct Answer: D

An organization is deploying their servers to the cloud using the infrastructure as a service model.

In the shared responsibility model, what is the cloud provider responsible for?

- A. Data access policies
- B. Security of the operating system
- C. Security of the software
- D. Physical security

Correct Answer: D

An organization runs a batch data analysis workload on a virtual machine (VM). The workload can be easily restarted without losing work, and is not time critical. Organizations must choose the lowest cost option to run the workload. What option should they choose?

- A. A Preemptible or Spot VM on Compute Engine
- B. A custom VM in a pay-as-you-go model on Compute Engine
- C. A standard VM in a pay-as-you-go model on Compute Engine
- D. A Cloud Function with a small memory limit

Correct Answer: A

Customers are reporting very high latencies when accessing an application from the United States. The application is currently running in a single region in Europe.

What should the organization do?

- A. Set up a new billing account in the United States.
- B. Run the application in additional zones in the European region.
- C. Run the application in additional regions in Europe.
- D. Run a replica of the application in a region in the United States.

Correct Answer: D

An organization needs protection against distributed denial-of-service (DDoS) attacks. Which Google Cloud service should the organization use?

- A. Google Cloud Armor
- B. Cloud Build
- C. Cloud VPN
- D. Security Command Center

Correct Answer: A

An organization needs to store daily transactional data such as customer records and purchase history. The data follows a consistent schema and is cross-referenced. Which type of service should the organization use?

- A. Non-relational database
- B. Data warehouse
- C. Relational database
- D. Data lake

Correct Answer: C

An organization is hosting an application in Europe, and customers in Asia are reporting slow response times despite their fast internet connection.

What is the problem?

- A. Not enough application servers
- B. Network bandwidth
- C. Network latency
- D. Misconfigured application servers

Correct Answer: C

A manufacturing organization has a large collection of images labeled as intact or defective parts. They want to use this data to build a simple solution to detect faulty parts on their production line. They have no data science expertise. Which solution should they use?

- A. Pre-trained APIs
- B. Document AI
- C. AutoML
- D. Discovery AI for Retail

Correct Answer: C

An organization has a large dataset that contains text transcripts of conversations between their customers and service representatives. They want an automated solution to identify the topics their customers care most about. Which service should the organization use?

- A. Vision API
- B. Speech-to-Text API
- C. Cloud Translation API
- D. Natural Language API

Correct Answer: D

An organization wants to analyze data in a data warehouse. How should they proceed?

- A. Import data into a semi-structured time-series database.
- B. Choose a system to store structured and semi-structured data that supports ad-hoc analysis and custom reporting.
- C. Copy unstructured data into a single large object store.
- D. Ensure data is stored in structured tables and rows to support transactional queries and relationships.

Correct Answer: B

An organization wants to ensure that they grant users only the permissions they require to perform their job responsibilities. Which security principle describes this approach?

- A. Cyber resilience
- B. Zero-trust
- C. Least privilege
- D. Security by default

Correct Answer: C

An organization is concerned about the risk of data loss that may occur due to hardware failures or cyber attacks. They want to restore their systems to a previous state if such an event occurs. What should the organization do?

- A. Set service level objectives (SLOs).
- B. Use Cloud Monitoring.
- C. Back up data regularly.
- D. Enable autoscaling.

Correct Answer: C

An organization processes batch sales data at the end of every month to analyze sales trends and derive business insights. They want to improve accuracy and make near real-time decisions. What should the organization do?

- A. Change from a relational database to a NoSQL database.
- B. Filter the data so reports are generated faster.
- C. Switch from batch processing to stream processing.
- D. Process batch reports weekly instead of monthly.

Correct Answer: C

An information security organization must ensure that its service providers have ISO 27001 certification. They must also access supporting documentation. Which Google Cloud tool should they use?

- A. Compliance Reports Manager
- B. Security Command Center
- C. Network Intelligence Center
- D. Cloud Monitoring

Correct Answer: A

An organization wants to adopt the advanced machine learning capabilities of the Google Cloud. However, regulations require data to be stored in an on-premises data center. Which approach should the organization use?

- A. A hybrid-cloud approach
- B. A multi-cloud approach
- C. A public-cloud approach
- D. A private-cloud approach

Correct Answer: A

An organization wants to run their custom application in the cloud in a flexible and scalable way without managing any infrastructure. Which service model should they use?

- A. Infrastructure as a service
- B. Platform as a service
- C. Software as a service
- D. Network as a service

Correct Answer: B

An organization is using Compute Engine and wants to receive sustained-use discounts. What should the organization do?

- A. Choose preemptible or spot instances when creating virtual machines.
- B. Choose low-powered virtual machines.
- C. Nothing. Sustained use discounts are automatically applied.
- D. Commit to virtual machine usage for a one or three-year period.

Correct Answer: C

An organization is evaluating its defenses against cyber security threats and is concerned about the risks of social engineering by cyber criminals. How might these attacks happen?

- A. SQL injection attacks
- B. Physical damage to hardware
- C. Phishing emails
- D. Distributed denial-of-service attacks

Correct Answer: C

An organization has a large VMWare environment that they want to migrate to the cloud. They want to retain existing operational processes and tools. Which Google Cloud service should the organization use?

- A. Compute Engine
- B. Google Cloud VMware Engine
- C. VMWare vSphere
- D. Bare Metal Solution

Correct Answer: B

An organization is building advanced machine learning models in Google Cloud by using TensorFlow. They want to develop their models faster with purpose-built hardware. Which solution should the organization use?

- A. TPUs
- B. CPUs
- C. CPUs
- D. GPUs

Correct Answer: A

A financial organization has many customers who close their accounts every year. The organization wants to use data and AI to identify at-risk customers, so they can retain customers by offering discounts and improved services. What should the organization do?

- A. Create a dashboard of previous customers that have exited, and look for obvious correlations in the visualization.
- B. Create a ML model based on the demographics and activities of previous customers that exited.
- C. Create a survey for all customers to identify their current level of satisfaction.
- D. Create a report based on last year's customer feedback.

Correct Answer: B

What is the purpose of an application programming interface (API)?

- A. To provide a set of instructions that allow computer programs to communicate with each other
- B. To provide cloud plugins for integrated development environments
- C. To connect networks from different cloud providers
- D. To manage multiple containerized workloads

Correct Answer: A

An organization supplies electric car chargers in a nationwide network. They store customer and charging data in BigQuery. They want to efficiently upgrade and maintain their network to reduce waste and achieve their sustainability goals. How can Google Cloud services help the organization?

- A. Secure data access by using Identity and Access Management.
- B. Provide access to third-party developers by creating an API using Apigee.
- C. Create a model in BigQuery to predict future maintenance schedules.
- D. Move data from BigQuery to Cloud Storage to reduce storage costs.

Correct Answer: C

An organization is planning to deploy a new workload to Google Cloud. They need an accurate estimate of the likely costs of running the workload. How should the organization create this estimate?

- A. Use the Google Cloud Pricing Calculator.
- B. Deploy workload to test environment to observe costs.
- C. Use historic costs of an existing similar workload.
- D. Refer to pricing information and manually calculate an estimate.

Correct Answer: A

An organization wants to refactor their application by using a microservices architecture when migrating to the cloud. Which benefit would this action provide?

- A. The refactored application is more efficient and scalable.
- B. No code changes will be required.
- C. This migration pattern provides the fastest path to the cloud.
- D. The application will become PCI-DSS compliant by default.

Correct Answer: A

An organization is using new technologies to change its business processes and culture to develop new customer experiences and adapt to market dynamics.

What is the name of this approach?

- A. Digital transformation
- B. Data center migration
- C. Operations optimization
- D. User personalization

Correct Answer: A

An organization is developing applications by using Kubernetes. They want their teams to spend more time developing rather than managing clusters. Which Google Cloud service should the organization choose?

- A. GKE Autopilot
- B. Compute Engine
- C. Cloud Run
- D. GKE Standard

Correct Answer: A

Get IT Certification

Unlock free, top-quality video courses on ExamTopics with a simple registration. Elevate your learning journey with our expertly curated content. Register now to access a diverse range of educational resources designed for your success. Start learning today with ExamTopics!

Start Learning for free