

1. What is the primary purpose of the Google Cloud Generative AI Leader training and exam?

- a) To train individuals to become machine learning engineers specializing in deep learning
- b) To certify professionals in advanced AI model development and infrastructure management
- c) To help professionals understand how generative AI can be applied within a business context using Google Cloud tools and services.
- d) To enable individuals to create custom foundation models from scratch

Answer:

c) To help professionals understand how generative AI can be applied within a business context using Google Cloud tools and services

This aligns with the guide's description that the exam is for professionals who need to understand how gen AI can be applied in business using Google Cloud

✗ Incorrect Options:

- a) Focuses on deep learning engineers, which is too technical for this leadership-level exam.
 - b) Certification in advanced model development is beyond the scope of this exam.
 - d) Creating custom foundation models is a technical task, not the goal of this leadership certification.
-

2. Which of the following is NOT listed as a key area assessed in the Generative AI Leader certification exam?

- a) Fundamentals of generative AI
- b) Google Cloud's generative AI offerings
- c) Advanced neural network architectures and design
- d) Business strategies for a successful gen AI solution

Answer:

c) Advanced neural network architectures and design

The exam focuses on fundamentals, Google Cloud offerings, model output techniques, and business strategies, and not deep technical design.

✗ Incorrect Options:

- a), b), d) are all explicitly listed as exam focus areas in the guide.
-

3. Approximately what percentage of the Generative AI Leader exam focuses on Google Cloud's generative AI offerings?

- a) ~15%
- b) ~20%
- c) ~30%
- d) ~35%

Answer:

d) ~35%

The exam focuses on fundamentals, Google Cloud offerings, model output techniques, and business strategies, and not deep technical design.

✗ Incorrect Options:

- a), b), d) are all explicitly listed as exam focus areas in the guide.
-

4. What are Foundation models defined as in the sources?

- a) Small AI models trained on specific datasets for narrow tasks
- b) Large AI models trained on massive datasets, allowing them to be adapted to many tasks
- c) Machine learning models designed exclusively for supervised learning
- d) AI models used only for generating text content

Answer:

b) Large AI models trained on massive datasets, allowing them to be adapted to many tasks

This is the definition provided in the guide.

✗ Incorrect Options:

- a) Describes narrow AI, not foundation models.
 - c) Supervised learning is a method, not a model type.
 - d) Foundation models are not limited to text generation.
-

5. Which of the following is a key feature of foundation models?

- a) They are designed for a very narrow range of use cases
- b) They are trained exclusively on labeled data
- c) They are flexible to a wide range of use cases
- d) They cannot be adapted to specialized domains

Answer:

c) They are flexible to a wide range of use cases

Flexibility is a core strength of foundation models.

✗ Incorrect Options:

- a) Opposite of flexibility.
 - b) They are trained on both labeled and unlabeled data.
 - d) They can be adapted to specialized domains.
-

6. What is the "art and science of creating effective inputs, known as prompts, for generative AI models to maximize their value"?

- a) Prompting
- b) Prompt engineering
- c) Model training
- d) Deep learning

Answer:

b) Prompt engineering

Defined in the guide as crafting effective prompts to guide model output.

✗ Incorrect Options:

- a) Prompting is the act, not the craft.
 - c) Model training is unrelated.
 - d) Deep learning is a broader ML technique.
-

7. According to the sources, which of the following is a capability of generative AI?

- a) Only summarizing existing information
- b) Creating new content
- c) Performing tasks that exclusively require human intelligence
- d) Analysing only numerical data

Answer:

b) Creating new content

Generative AI is specifically designed to generate new content such as text, images, code, and more.

✗ Incorrect Options:

- a) Summarizing is a capability, but not the only one—this option is too narrow.
 - c) Gen AI can assist with tasks that require human-like intelligence, but it doesn't exclusively perform such tasks.
 - d) It can analyze numerical data, but that's not its defining feature.
-

8. Which term refers to an application of Machine Learning that focuses on creating new content?

- a) Artificial intelligence (AI)
- b) Machine learning (ML)
- c) Generative AI
- d) Deep learning

Answer:

c) Generative AI

Generative AI is a subfield of ML focused on content creation.

✗ Incorrect Options:

- a) AI is broader and includes ML, robotics, etc.
 - b) ML is the broader category under which generative AI falls.
 - d) Deep learning is a technique used within ML, not a content-focused application.
-

9. What are Large Language Models (LLMs) described as?

- a) A type of foundation model designed to understand and generate human language
- b) Machine learning models that only process images
- c) Models that are trained on small, labeled datasets
- d) A subfield of Artificial Intelligence focusing on robotics [Not specified]

Answer:

a) A type of foundation model designed to understand and generate human language

LLMs are foundation models trained on large text datasets to perform language tasks.

✗ Incorrect Options:

- b) Incorrect—LLMs are not image processors.
 - c) They are trained on large, often unlabeled datasets.
 - d) Robotics is unrelated to LLMs.
-

10. What is Supervised learning primarily characterized by?

- a) Learning through interaction and feedback to maximize rewards
- b) Using unlabeled data to find natural groupings and patterns
- c) Training models on labeled data to predict outputs for new inputs
- d) Extracting complex patterns from data using many layers

Answer:

c) Training models on labeled data to predict outputs for new inputs

Supervised learning uses labeled data to train models for prediction.

✗ Incorrect Options:

- a) Describes reinforcement learning.
 - b) Describes unsupervised learning.
 - d) Describes deep learning, which is a technique, not a learning type.
-

11. Which type of data is defined as "raw, unprocessed information that hasn't been tagged and lacks meaning by itself"?

- a) Labeled data
- b) Structured data
- c) Unstructured data
- d) Unlabeled data

Answer:

d) Unlabelled data

Unlabeled data lacks tags or structure and is often used in unsupervised learning.

✗ Incorrect Options:

- a) Labeled data has tags.
 - b) Structured data is organized and searchable.
 - c) Unstructured data lacks format but may still have meaning.
-

12. Which of the following is the first step in the Machine Learning lifecycle mentioned?

- a) Model training
- b) Model deployment
- c) Data ingestion and preparation
- d) Model management

Answer:

c) Data ingestion and preparation

This is the foundational step before training a model.

✗ Incorrect Options:

- a) Model training comes after data preparation.
 - b) Deployment is a later stage.
 - d) Model management is ongoing after deployment.
-

13. What layer in the Generative AI landscape allows users to interact with and leverage the capabilities of AI?

- a) Infrastructure
- b) Model
- c) Gen-AI-powered application
- d) Agent

Answer:

c) Gen-AI-powered application

This is the user-facing layer that enables interaction with AI.

✗ Incorrect Options:

- a) Infrastructure supports the backend.
 - b) Model is the core algorithm.
 - d) Agent is the decision-making component.
-

14. What provides the core computing resources, including physical hardware like GPUs and TPUs, for generative AI?

- a) Agent
- b) Platform
- c) Infrastructure
- d) Model

Answer:

c) Infrastructure

Infrastructure includes the hardware and software needed to run AI models.

✗ Incorrect Options:

- a) Agents use infrastructure but don't provide it.
 - b) Platform offers APIs and tools, not raw compute.
 - d) Model is the algorithm, not the hardware.
-

15. Which Google generative AI model is described as supporting multimodal understanding, advanced conversational AI, content creation, and question answering?

- a) Gemma
- b) Imagen
- c) Veo
- d) Gemini

Answer:

d) Gemini

Gemini is Google's flagship gen AI model designed for multimodal tasks - text, images, code, and more.

✗ Incorrect Options:

- a) Gemma is for local, lightweight deployments.
 - b) Imagen is for text-to-image generation.
 - c) Veo is for video generation.
-

16. Which Google generative AI model is a text-to-image diffusion model that generates high-quality images from textual descriptions?

- a) Gemini
- b) Gemma
- c) Imagen
- d) Veo

Answer:

c) Imagen

Imagen generates high-quality images from text descriptions.

✗ Incorrect Options:

- a) Gemini is multimodal but not image specific.
 - b) Gemma is not for image generation.
 - d) Veo is for video, not images.
-

17. What is the primary function of Veo?

- a) Generating text summaries
- b) Generating video content based on text descriptions or still images
- c) Converting speech to text
- d) Analysing image content

Answer:

b) Generating video content based on text descriptions or still images

Veo is designed to create video content from prompts.

✗ Incorrect Options:

- a) Text summaries are not Veo's function.
 - c) Speech-to-text is handled by the Speech-to-Text API.
 - d) Image analysis is done by Cloud Vision API.
-

18. What is Gemini for Google Workspace designed to do?

- a) Help write and debug code in Google Cloud
- b) Integrate gen AI into familiar Workspace apps like Gmail and Slides
- c) Provide enterprise-grade contact center solutions
- d) Act as a research assistant for uploaded files

Answer:

b) Integrate gen AI into familiar Workspace apps like Gmail and Slides

Gemini enhances productivity tools like Gmail, Docs, and Slides with gen AI.

✗ Incorrect Options:

- a) Code assistance is a feature of Gemini for Google Cloud.
 - c) Contact center solutions are part of the Customer Engagement Suite.
 - d) Research assistant functionality is offered by NotebookLM.
-

19. Which tool allows you to upload files and then acts as a research assistant, summarizing key points, answering questions, and generating ideas grounded in your source material?

- a) Gemini App
- b) Gemini for Google Cloud
- c) NotebookLM
- d) Vertex AI

Answer:

c) NotebookLM

NotebookLM summarizes, answers questions, and generates ideas from uploaded documents.

✗ Incorrect Options:

- a) Gemini App is a chatbot, not a document assistant.
 - b) Gemini for Google Cloud is for cloud development tasks.
 - d) Vertex AI is for building and deploying ML models.
-

20. What is Google Cloud's unified ML platform that empowers you to build, train, and deploy ML applications?

- a) Google AI Studio
- b) Vertex AI
- c) Google Agentspace
- d) NotebookLM

Answer:

b) Vertex AI

Vertex AI is the central platform for building, training, and deploying ML models.

✗ Incorrect Options:

- a) Google AI Studio is for quick prototyping.
 - c) Google Agentspace is for building internal agents.
 - d) NotebookLM is for document-based research.
-

21. What is the key difference between Google AI Studio and Vertex AI Studio?

- a) Google AI Studio is for production-ready applications, while Vertex AI Studio is for quick prototyping
- b) Google AI Studio is free and for quick AI prototyping, while Vertex AI Studio is for building and deploying production-ready AI applications at scale
- c) Google AI Studio focuses on text generation, while Vertex AI Studio focuses on image generation
- d) Google AI Studio integrates with Google Workspace, while Vertex AI Studio integrates with Google Cloud

Answer:

b) Google AI Studio is free and for quick AI prototyping, while Vertex AI Studio is for building and deploying production-ready AI applications at scale

Google AI Studio is intended for experimentation, while Vertex AI Studio supports enterprise-grade deployment.

✗ Incorrect Options:

- a) Reverses the actual roles of the two platforms.
 - c) Both support text and image generation; this is not the key distinction.
 - d) Integration is not the main differentiator.
-

22. Which Customer Engagement Suite tool acts as an effective chatbot for customers?

- a) Agent Assist
- b) Conversational Agents
- c) Conversational Insights
- d) Google Agentspace

Answer:

b) Conversational Agents

These are designed to interact with customers directly, simulating human-like conversations.

✗ Incorrect Options:

- a) Agent Assist supports human agents, not customers directly.
 - c) Conversational Insights provides analytics, not interaction.
 - d) Google Agentspace is for internal use, not customer-facing.
-

23. What are "Tools" defined as within the context of an AI agent's components?

- a) The brains of the AI system that learn patterns from data
- b) An iterative process where the agent observes, interprets, reasons, and acts
- c) Functionalities that allow the agent to interact with its environment
- d) Predefined paths and actions for deterministic agents

Answer:

c) Functionalities that allow the agent to interact with its environment

Tools enable agents to perform actions like accessing data or triggering APIs.

✗ Incorrect Options:

- a) Describes the model, not tools.
 - b) Describes the reasoning loop.
 - d) Refers to deterministic agent behavior, not tools.
-

24. Which type of agent is defined with natural language using LLMs to give a real conversational feel to a chatbot?

- a) Deterministic agents
- b) Hybrid Agents
- c) Generative agents
- d) Platform agents

Answer:

c) Generative agents

These agents use LLMs to simulate natural, human-like dialogue.

✗ Incorrect Options:

- a) Deterministic agents follow fixed rules.
 - b) Hybrid agents combine deterministic and generative features.
 - d) Platform agents is not a defined category.
-

25. Which Google tool helps developers deploy AI models on edge devices?

- a) Vertex AI Pipelines
- b) Lite Runtime (LiteRT)
- c) Gemini Nano
- d) Vertex AI Model Registry

Answer:

b) Lite Runtime (LiteRT)

LiteRT is designed for lightweight, edge-based AI deployment.

✗ Incorrect Options:

- a) Pipelines automate ML workflows, not edge deployment.
 - c) Gemini Nano is a compact model, not a deployment tool.
 - d) Model Registry tracks versions, not deployment.
-

26. What is the primary function of the Speech-to-Text API?

- a) Converts text to natural-sounding speech
- b) Translates text documents
- c) Converts speech into text and transcribes audio/video content
- d) Extracts data from varied document formats

Answer:

c) Converts speech into text and transcribes audio/video content

This API is used for transcription and voice recognition.

✗ Incorrect Options:

- a) Describes Text-to-Speech API.
 - b) Translation is handled by the Translation API.
 - d) Document AI handles document extraction.
-

27. Which API is used for analysing image content, tagging images based on detected objects and text, and identifying faces and landmarks?

- a) Document AI API
- b) Cloud Video Intelligence API
- c) Cloud Vision API
- d) Natural Language API

Answer:

c) Cloud Vision API

This API is designed for image analysis and tagging.

✗ Incorrect Options:

- a) Document AI is for document processing.
 - b) Cloud Video Intelligence is for video, not images.
 - d) Natural Language API is for text analysis.
-

28. What is the Natural Language API designed to do?

- a) Convert speech to text
- b) Generate video content
- c) Help derive insights from unstructured text, understand sentiment, and classify content
- d) Translate formatted documents while keeping the original layout

Answer:

c) Help derive insights from unstructured text, understand sentiment, and classify content

This API processes and analyses text data.

✗ Incorrect Options:

- a) Speech-to-Text API handles this.
 - b) Video generation is unrelated.
 - d) Translation API handles document translation.
-

29. What are "Gems" within Gemini described as?

- a) Templates for re-using prompts
- b) Personalized AI assistants that provide tailored responses and streamline workflows
- c) Saved information for the model to use consistently
- d) External services connected via APIs

Answer:

b) Personalized AI assistants that provide tailored responses and streamline workflows

Gems are customizable assistants that enhance productivity.

✗ Incorrect Options:

- a) Templates are not the same as Gems.
 - c) Saved info is a feature, not a definition.
 - d) External services are handled via extensions or plugins.
-

30. Which prompting technique involves providing the model with one example to learn from?

- a) Zero-shot
- b) One-shot
- c) Few-shot
- d) Role

Answer:

b) One-shot

One-shot prompting gives the model a single example to guide its response.

✗ Incorrect Options:

- a) Zero-shot provides no examples.
 - c) Few-shot provides multiple examples.
 - d) Role prompting assigns a persona, not examples.
-

31. What is the purpose of the "Role" prompting technique?

- a) To provide multiple examples for the model to learn from
- b) To engage in a back-and-forth conversation with the AI
- c) To assign a persona to the model to influence its style, tone, and focus
- d) To connect the AI's output to verifiable sources of information

Answer:

c) To assign a persona to the model to influence its style, tone, and focus

Role prompting helps guide the model's behavior by giving it a specific identity or tone (e.g., "Act as a customer support agent").

✗ Incorrect Options:

- a) Describes few-shot prompting.
 - b) Describes prompt chaining.
 - d) Describes grounding, not role prompting.
-

32. What does RAG stand for in the context of improving generative AI model output?

- a) Reason and Act Generation
- b) Retrieval-augmented generation
- c) Recursive Agent Guidance
- d) Resource Allocation Grid

Answer:

b) Retrieval-augmented generation

RAG enhances model responses by retrieving relevant external data before generating output.

✗ Incorrect Options:

- a) Not a real term in this context.
 - c) Fabricated term.
 - d) Refers to resource planning, not AI.
-

33. Which step in the RAG process involves the LLM retrieving relevant information from external sources using tooling?

- a) Augmentation
- b) Generation
- c) Retrieval
- d) Iteration

Answer:

c) Retrieval

This is the first step in the RAG pipeline—gathering relevant data.

✗ Incorrect Options:

- a) Augmentation is the second step.
 - b) Generation is the third step.
 - d) Iteration is optional and comes later.
-

34. What prompting technique allows the LLM to reason and take action on a user query?

- a) CoT (chain-of-thought)
- b) Metaprompting
- c) ReAct (reason and act)
- d) Few-shot

Answer:

c) ReAct (reason and act)

ReAct combines reasoning and action-taking in a single prompt structure.

✗ Incorrect Options:

- a) CoT is for step-by-step reasoning, not action.
 - b) Metaprompting is about generating prompts.
 - d) Few-shot is about giving multiple examples.
-

35. Which sampling parameter controls the "creativity" or randomness of the model's word choices during text generation?

- a) Token count
- b) Temperature
- c) Top-p (nucleus sampling)
- d) Output length

Answer:

b) Temperature

Higher temperature = more randomness; lower = more deterministic output.

✗ Incorrect Options:

- a) Token count refers to input/output length.
 - c) Top-p controls randomness differently but not as directly.
 - d) Output length limits response size, not creativity.
-

36. Which of the following is listed as a limitation of foundation models?

- a) Unlimited knowledge beyond their training data
- b) Complete absence of bias
- c) Hallucinations
- d) Consistent performance across all edge cases

Answer:

c) Hallucinations

Hallucinations are when models generate plausible but incorrect or fabricated information.

✗ Incorrect Options:

- a) Models don't have unlimited knowledge.
 - b) Bias is a known issue, not absent.
 - d) They may fail on edge cases.
-

37. What does "Knowledge cutoff" refer to regarding AI models?

- a) The point at which a model stops learning
- b) The maximum amount of data a model can be trained on
- c) The specific date up to which an AI model is trained, meaning it might lack information about events after that point
- d) A method to filter out harmful content from model outputs

Answer:

c) The specific date up to which an AI model is trained, meaning it might lack information about events after that point

Models don't learn after their training data ends.

✗ Incorrect Options:

- a) Models don't learn continuously.
 - b) Training data size is unrelated.
 - d) Filtering harmful content is a separate concept.
-

38. In which scenario does Humans in the Loop (HITL) provide critical oversight, ensuring accuracy and reducing risks from automated systems?

- a) Automated data backup procedures
- b) High-stakes decisions in fields like healthcare and finance
- c) General content creation without sensitive information
- d) Simple data entry tasks

Answer:

b) High-stakes decisions in fields like healthcare and finance

HITL ensures human review in sensitive or high-risk applications.

✗ Incorrect Options:

- a), c), d) are low-risk or routine tasks where HITL is not essential.
-

39. What is "Fine-tuning" a technique used for?

- a) To increase the context window of a model
- b) To enhance a pre-trained or foundation model's performance for specific tasks or domains
- c) To manage and maintain models over time
- d) To filter out potentially harmful content

Answer:

b) To enhance a pre-trained or foundation model's performance for specific tasks or domains

Fine-tuning adapts general models to specific use cases.

✗ Incorrect Options:

- a) Context window is unrelated.
 - c) Model management is broader.
 - d) Filtering content is a safety feature, not fine-tuning.
-

40. Which Google Cloud tool is used to keep track of different versions of a model?

- a) Vertex AI Model Monitoring
- b) Vertex AI Feature Store
- c) Vertex AI Model Registry
- d) Vertex AI Pipelines

Answer:

c) Vertex AI Model Registry

This tool helps manage model versioning and lifecycle.

✗ Incorrect Options:

- a) Model Monitoring tracks performance, not versions.
 - b) Feature Store manages data features.
 - d) Pipelines automate workflows, not versioning.
-

41. Which of the following is NOT a factor to consider when planning your Gen AI strategy?

- a) Establishing a clear vision
- b) Ignoring responsible AI considerations
- c) Prioritizing use cases
- d) Measuring value

Answer:

b) Ignoring responsible AI considerations

Responsible AI is a critical component of any AI strategy. Ignoring it would be a major oversight and is clearly not recommended in the study guide.

✗ Incorrect Options:

- a) Establishing a clear vision is essential for strategic alignment.
 - c) Prioritizing use cases ensures focus and value delivery.
 - d) Measuring value helps track ROI and effectiveness.
-

42. When considering the "Needs" before starting a Gen AI project, which of the following is a key question?

- a) What is your connectivity?
- b) What's your budget?
- c) Do you have access to AI expertise?
- d) What are your project timelines?

Answer:

a) What is your connectivity?

Connectivity is a foundational requirement for deploying and accessing AI services, especially in cloud environments.

✗ Incorrect Options:

- b), c), and d) are also important but fall under “Resources” rather than “Needs” as categorized in the study guide.
-

43. What does "Responsible AI" primarily focus on?

- a) Maximizing profit margins through AI deployment
- b) Ensuring AI applications don't cause harm and are used ethically
- c) Developing AI models with the largest possible context window
- d) Automating all human decision-making processes

Answer:

b) Ensuring AI applications don't cause harm and are used ethically

Responsible AI is about fairness, transparency, and minimizing harm.

✗ Incorrect Options:

- a) Profit is a business goal, not the focus of Responsible AI.
 - c) Context window is a technical parameter.
 - d) Full automation is not the goal of Responsible AI.
-

44. What is one of the key factors to consider when choosing a model for your use case, related to input and output data types?

- a) Performance
- b) Availability
- c) Modality
- d) Latency

Answer:

c) Modality

Modality refers to the type of data (text, image, audio, video) the model handles.

✗ Incorrect Options:

- a) Performance is important but not specific to data types.
 - b) Availability is about uptime and access.
 - d) Latency refers to response time, not data type compatibility.
-

45. What framework helps organizations manage AI/ML model risks and ensure security?

- a) Responsible AI Framework
- b) Secure AI Framework (SAIF)
- c) ML Lifecycle Framework
- d) Human-in-the-Loop Framework

Answer:

b) Secure AI Framework (SAIF)

SAIF is Google's framework for managing AI risks and ensuring secure deployment.

✗ Incorrect Options:

- a) Responsible AI Framework focuses on ethics, not security.
 - c) ML Lifecycle Framework is about development stages.
 - d) HITL Framework is about human oversight, not security.
-

46. What is the first step when using NotebookLM to create your own study guide?

- a) Create a new notebook and upload a source
- b) Ask targeted questions to build study guide sections
- c) Navigate to the NotebookLM website and sign in
- d) Copy and paste information into a separate document

Answer:

c) Navigate to the NotebookLM website and sign in

Before creating or uploading anything, you must access the platform.

✗ Incorrect Options:

- a), b), and d) are subsequent steps after signing in.
-

47. When using NotebookLM to build a study guide, what is a recommended action after NotebookLM provides initial answers?

- a) Stop and move to the next topic
- b) Rephrase and summarize the information in your own words
- c) Discard the answers and find new sources
- d) Only copy-paste the answers verbatim

Answer:

b) Rephrase and summarize the information in your own words

This helps with retention and understanding.

✗ Incorrect Options:

- a) Stopping early limits learning.
 - c) Discarding answers wastes useful content.
 - d) Copy-pasting verbatim doesn't support deep learning.
-

48. Which of the following is a way to streamline prompting workflows in Gemini?

- a) Only using zero-shot prompting
- b) Randomly generating new prompts for each task
- c) Storing specific information for the model to use consistently
- d) Avoiding back-and-forth conversations with the AI

Answer:

c) Storing specific information for the model to use consistently

This allows for more efficient and context-aware interactions.

✗ Incorrect Options:

- a) Zero-shot is not always optimal.
 - b) Random prompts reduce consistency.
 - d) Back-and-forth conversations are often helpful.
-

49. What is the "Context window" in the context of AI models?

- a) The maximum number of outputs a model can generate
- b) The amount of text the model can consider
- c) The duration for which a model remains active
- d) The speed at which a model processes information

Answer:

b) The amount of text the model can consider

This defines how much input the model can process at once.

✗ Incorrect Options:

- a) Refers to output, not input.
 - c) Duration is unrelated.
 - d) Speed is about performance, not context.
-

50. What is "Drift monitoring" used for in Vertex AI Model Monitoring?

- a) To manage the data features a model uses
- b) To keep track of different versions of a model
- c) To watch for changes in the model's accuracy over time
- d) To automate machine learning tasks

Answer:

c) To watch for changes in the model's accuracy over time

Drift monitoring detects when a model's performance degrades due to changes in data.

✗ Incorrect Options:

- a) Feature management is handled by Feature Store.
 - b) Versioning is done via Model Registry.
 - d) Pipelines automate workflows, not monitor drift.
-